



Building Fire-Resilience Credit Model

Poisson Initiation with an 8-State Bayesian Containment Chain

MKM Research Labs

Version 0.30 — 08-June-2026

David K Kelly

CONFIDENTIAL — PROPRIETARY

Legal Notice

PROPRIETARY AND CONFIDENTIAL

This document, including all algorithms, models, methodology, formulas, calculations, and intellectual concepts contained herein, constitutes the exclusive intellectual property and confidential trade secrets of MKM Research Labs (“MKM”).

Any usage, reproduction, distribution, modification, or disclosure of this document or any portion thereof, without the express prior written authorisation from MKM Research Labs is strictly prohibited and constitutes an infringement of intellectual property rights.

The document is provided on an “as is” basis. MKM Research Labs makes no representations or warranties of any kind, express or implied, regarding its accuracy, reliability, or suitability for any particular purpose.

All rights reserved. © 2021–2026 MKM Research Labs.

Governed by the laws of the United Kingdom.

Contents

1	Executive Summary	3
2	Model Purpose and Scope	3
2.1	Purpose	3
2.2	Scope	4
2.3	Out of Scope	4
3	Mathematical Framework	4
3.1	Initiation (Model A)	4
3.2	Response Effectiveness (Model C)	5
3.3	Intensity Track (Model B)	5
3.4	Point of No Return	6
3.5	Containment Chain (Model D)	6
3.6	Resilience Credit	6
4	Input Parameters and Calibration	7
5	Implementation Details	7
5.1	Module Layout	7
5.2	Randomness	7
6	Validation and Backtesting	8
6.1	Automated Test Results	8
7	Sensitivity Analysis	8
7.1	Interpretation and known model weaknesses	10
8	Model Limitations and Known Weaknesses	11
9	Governance Validation Questions	11
9.1	Q1 — Model purpose and use	11
9.2	Q2 — Data quality and lineage	12
9.3	Q3 — Conceptual soundness	12
9.4	Q4 — Calibration and parameters	12
9.5	Q5 — Implementation verification	12
9.6	Q6 — Sensitivity and stability	12
9.7	Q7 — Limitations and weaknesses	12
9.8	Q8 — Ongoing monitoring	13
9.9	Q9 — Governance and approval	13
10	Change History	13
A	Initiation Parameters (Model A)	14

B	Initiation-Class Priors	15
C	Progression Dynamics (Models B & C)	15
D	Containment Transition Matrices (Model D)	16

1 Executive Summary

The Building Fire-Resilience Credit Model (MKM-FIRE-001) prices the fire resilience of a commercial asset so that fire loss aggregates into the same resilience-credit currency as the existing 10,000-storm Monte Carlo engine. The credit is consumed by the Commercial PRS Pricer as an independent *full-conflagration* leg, combined with the flood and wind legs through a root-sum-of-squares all-in coupon. The model is built from four governance-separated component models:

- **Model A — Initiation.** A Poisson ignition frequency λ_i per asset, composed from a base annual rate by commercial type and multiplicative modifiers for occupancy, process load, condition, protection and fire history.
- **Model B — Intensity Track.** A deterministic 15-minute latent intensity I_t whose growth is damped by passive (structural) defence and cut once active suppression bites.
- **Model C — Response Effectiveness.** A deterministic mapping from the asset's CDM resilience levels to the dynamic quantities Models B and D consume (detection time, suppression-bite time, growth rate, critical threshold, controllability, and the fire-service-reach gate).
- **Model D — Containment, Loss & Resilience Credit.** An 8-state discrete-time Bayesian chain marched in 15-minute steps to an absorbing terminal state, aggregated into the credit currency.

The central concept is the *point of no return* (PNR): the first step at which latent intensity I_t exceeds the critical threshold I_{crit} *before* suppression has bitten — or, equivalently, the first step of any fire in a building where suppression is structurally unreachable. Once PNR latches, the containment chain switches to an absorbing matrix whose Contained column is zero: the fire can no longer be contained and absorbs into {partial loss, total loss}. The PNR frequency, $n_{\text{PNR}}/n_{\text{sim}}$, multiplied by a 100% loss-given-event, is the bps the PRS pricer charges for the fire leg.

Status: Development (Tier 2, RAG Amber). All four component models and the end-to-end orchestrator are delivered with unit tests, including the fire-service-reach interaction. Every numeric parameter is a first-pass engineering-judgement seed (`config/fire_matrices.json`, `version "seed-0"`); the results in this document are therefore *directional, not calibrated*.

Key risk: The model uses placeholder calibration throughout. It is not approved for production until historical calibration completes and the Model Risk Committee (MRC) issues sign-off (remediation items FIRE-R1–R6).

2 Model Purpose and Scope

2.1 Purpose

To assign each commercial asset a fire-resilience credit — expressed as an annual unconditional loss frequency, a conditional containment rate, and a full-conflagration (PNR) frequency — that can be priced *pari-passu* with the flood and wind hazards already in the portfolio engine. The credit rewards genuine, evidenced fire resilience (detection, suppression, compartmentation, tested procedures) and

penalises the configurations that historically drive total losses, most notably tall buildings without adequate internal suppression.

2.2 Scope

The model covers commercial assets described by the v2 commercial-asset and resilience sections of the Common Data Model (CDM). It reads only the derived **AssetFireFeatures** bundle — never the raw CDM record — so that the feature contract is explicit, testable and stable against CDM schema churn. The in-scope drivers are: commercial type, occupancy status, business-rates category, property condition, the nine resilience-level fields (automatic detection, suppression systems, emergency procedures, structural fire resistance, compartments, fire stopping, external materials, access route, business continuity), fire-damage history, and number of storeys. The model produces a per-asset resilience credit and a portfolio roll-up; both are reproducible from a single random seed.

2.3 Out of Scope

- Physical (CFD) fire-spread and thermal-plume modelling.
- Smoke / life-safety egress and casualty modelling.
- Wildland–urban interface and external conflagration spread.
- Business-interruption quantum (only the resilience credit is produced; the loss-given-event is fixed at 100% for the conflagration leg).
- Explicit floor-level geometry — building height enters only through **NumberOfStoreys**.

3 Mathematical Framework

3.1 Initiation (Model A)

The per-asset annual ignition rate is the base rate for the commercial type scaled by five independent multipliers:

$$\lambda_i = \lambda_{0,c} m_{\text{occ}} m_{\text{proc}} m_{\text{cond}} m_{\text{protection}} m_{\text{history}},$$

where $\lambda_{0,c}$ is the base annual frequency by commercial type c ; m_{occ} from occupancy status; m_{proc} from the process load implied by type (raised, never lowered, by a business-rates override such as *Restaurant*); m_{cond} from property condition; $m_{\text{protection}}$ from the mean resilience level across the protection fields; and m_{history} from the severity and recency of any prior fire. Unknown CDM options fall back to a neutral multiplier of 1.0, so a partial record still yields a usable λ_i .

Over a run horizon of h years the effective rate is $\lambda_{\text{eff}} = \lambda_i h$. For each of n_{sim} draws a Poisson count $N \sim \text{Poisson}(\lambda_{\text{eff}})$ is sampled, and a fire instantiates for that draw when $N \geq 1$ — the direct analogue of storm-event generation, where only a subset of draws produce an event. Each instantiated fire is tagged with an **InitiationClass** drawn from the per-type categorical prior (Appendix B).

3.2 Response Effectiveness (Model C)

Each resilience level is mapped to an index $k \in \{0, \dots, 4\}$ in the CDM `RESILIENCE_LEVELS` vocabulary and to a fraction $f = k/4 \in [0, 1]$. A field-group effectiveness is the mean over its fields of the field's configured range linearly interpolated by f :

$$e_{\text{group}} = \frac{1}{|G|} \sum_{j \in G} [\ell_j + f_j (u_j - \ell_j)],$$

giving the aggregate passive effectiveness e_{passive} (structural resistance, compartments, fire stopping, external materials) and response effectiveness e_{response} (access route, procedures, continuity).

Detection time is $\tau_{\text{det}} = \tau_{\text{det},0} \delta_\ell$ where δ_ℓ is the detection multiplier for the asset's detection level. Suppression-bite time applies a gentle per-storey height penalty $p \in [p_{\min}, p_{\max}]$ interpolated over storeys *and* a suppression-level multiplier b_{supp} (a better system engages sooner), but is governed first by the **fire-service-reach gate**: external attack reaches only up to R storeys, and above R effective suppression requires internal sprinklers, i.e. a suppression level index $\geq \theta_{\text{int}}$. Writing s for the storey count and k_{supp} for the suppression-level index,

$$\text{reachable} = (k_{\text{supp}} \geq \theta_{\text{int}}) \vee (s \leq R),$$

$$\tau_{\text{supp}} = \begin{cases} \tau_{\text{supp},0} p b_{\text{supp}} & \text{if reachable,} \\ \tau_{\infty} (\gg \text{step budget}) & \text{otherwise.} \end{cases}$$

Compartmentation quality $q = \frac{1}{2}(k_{\text{comp}} + k_{\text{stop}})$ selects the well- vs. poorly-compartmented growth rate; the suppression level selects both the post-bite growth multiplier μ_{supp} and the critical intensity I_{crit} , the latter graded across all five levels (not a binary strong/weak step) so every suppression level shifts the controllability race. Finally the controllability score is the weighted blend

$$C = w_p e_{\text{passive}} + w_r e_{\text{response}} + w_s f_{\text{supp}},$$

(renormalised by $w_p + w_r + w_s$), which is remapped to a matrix-blend weight

$$w = \text{clamp}\left(\frac{C - C_{\text{floor}}}{C_{\text{ceil}} - C_{\text{floor}}}, 0, 1\right) \in [0, 1]$$

that mixes the pre-PNR transition matrices (Model D). A continuous blend, rather than a hard threshold, is what lets the well-graded score C actually grade containment: a hard switch saturates because passive and response defences alone already clear it, leaving suppression unable to move the matrix.

3.3 Intensity Track (Model B)

The latent intensity starts at $I_0 = 0$ and grows each 15-minute step by

$$g_t = g_{\text{base}} (1 - w_{\text{damp}} e_{\text{passive}}), \quad g_t \leftarrow g_t \cdot \mu_{\text{supp}} \text{ once suppression is active,}$$

$$I_{t+1} = I_t + g_t, \quad g_t \geq 0,$$

where g_{base} is the compartmentation-selected growth rate and suppression becomes active at the step $\lceil \tau_{\text{supp}} \rceil$. The *mean* profile fixes g_{base} and τ_{supp} , but each fire draws its own race realisation: per-fire lognormal jitter is applied once to the growth rate ($g_{\text{base}} \cdot e^{\sigma_g Z}$) and to the bite time ($\tau_{\text{supp}} \cdot e^{\sigma_\tau Z}$) before the march. Without this the PNR verdict would be a deterministic property of the asset, forcing the conflagration leg to be all-or-nothing (0 or the full ignition rate); the jitter makes it a per-fire Bernoulli outcome so the conflagration probability grades smoothly. Setting $\sigma_g = \sigma_\tau = 0$ recovers the deterministic track.

3.4 Point of No Return

PNR latches at the first step t for which suppression is not yet active and

$$I_t > I_{\text{crit}} \quad \vee \quad \neg \text{reachable},$$

and the latch is sticky (once set, it stays set). The second disjunct is the operational heart of the model: a building taller than the fire-service reach with no adequate internal suppression is uncontrollable *from ignition*, so PNR latches at step 1. This is what turns building height into a live driver of total loss.

3.5 Containment Chain (Model D)

Each fire is a discrete-time chain over the eight states

$$\begin{aligned} &S_0 \text{ Ignition, } S_1 \text{ Growth, } S_2 \text{ Detected, } S_3 \text{ Internal response, } S_4 \text{ External response,} \\ &S_5 \text{ Contained, } S_6 \text{ Partial loss, } S_7 \text{ Total loss,} \end{aligned}$$

with $\{S_5, S_6, S_7\}$ absorbing. At each step the intensity track advances and the active 8×8 row-stochastic matrix is selected: while PNR is unlatched the pre-PNR matrix is the controllability-weighted blend

$$M = w M_{\text{AA}} + (1 - w) M_{\text{weak}},$$

(row-stochastic for any $w \in [0, 1]$), interpolating continuously from the weak matrix at $w = 0$ to the controllable matrix at $w = 1$; once PNR latches, every transient state uses the absorbing point-of-no-return matrix, whose S_5 (Contained) column is identically zero. The next state is sampled from the current state's row (with defensive renormalisation) until an absorbing state is reached or the step budget is exhausted. The three base matrices are given in Appendix D.

3.6 Resilience Credit

Aggregating over the n_{sim} draws (of which n_{fires} ignite), with terminal counts $n_{\text{contained}}$, n_{partial} , n_{total} and n_{PNR} fires that reached the point of no return, the credit currency is

$$\begin{aligned} \text{loss frequency} &= \frac{n_{\text{partial}} + n_{\text{total}}}{n_{\text{sim}}}, & \text{containment rate} &= \frac{n_{\text{contained}}}{n_{\text{fires}}}, \\ \text{fire spread (bps)} &= \frac{n_{\text{PNR}}}{n_{\text{sim}}} \times 1.0 \times 10,000, \end{aligned}$$

the last being the full-conflagration leg the Commercial PRS Pricer consumes (frequency $\times$ 100% loss-given-event). It is combined with the flood/wind legs into an all-in coupon by root-sum-of-squares, $\text{all-in} = \sqrt{\sum_{\text{perils}} \text{leg}^2}$, treating the perils as independent.

4 Input Parameters and Calibration

All numeric seeds live in `config/fire_matrices.json` (version "seed-0") and are loaded by `config/fire.py`; no number is embedded in the model code. The resilience-level vocabulary is imported from `port.cdm.asset.resi` so the level-to-index-to-fraction mapping stays single-source with the rest of the CDM. The full parameter set — initiation base rates and the five modifier tables, initiation-class priors, the progression dynamics (growth, critical intensity, detection/suppression multipliers, passive/response ranges, height penalty, fire-service reach, timing, controllability weights) and the three transition matrices — is reproduced verbatim in the Appendices.

Calibration status. Every value is a first-pass engineering-judgement seed chosen for plausible *direction* and *ordering*, not magnitude: e.g. hotels and mixed-use ignite more often than offices; worse condition and worse protection raise frequency; better resilience levels shorten detection and suppression times and raise the critical intensity. Calibration against historical commercial fire-loss data (FIRE-R1) and an independent benchmarking exercise (FIRE-R2) are open remediation items; until they close, the outputs must be read as relative resilience signals, not absolute loss probabilities.

5 Implementation Details

5.1 Module Layout

- `config/fire.py` — parameter schema (enums, dataclasses, loader).
- `config/fire_matrices.json` — seed numeric parameters.
- `src/models/fire/data_structures.py` — runtime types (`AssetFireFeatures`, `ResponseProfile`, `IntensityTrack`, `ContainmentOutcome`, `AssetFireResult`).
- `src/models/fire/initiation.py` — Model A (Poisson initiation).
- `src/models/fire/response_effectiveness.py` — Model C.
- `src/models/fire/intensity_track.py` — Model B.
- `src/models/fire/containment.py` — Model D and the end-to-end orchestrator (`simulate_asset_fire`, `simulate_portfolio_fire`).
- `src/routes/commercial/hazard.py` — read-time join that attaches the fire leg (`fire_spread_bps`) to the asset hazard payload.

5.2 Randomness

All randomness flows through a single caller-owned `numpy.random.Generator`, so an entire portfolio run is reproducible from one seed — matching the storm and typhoon model convention. Model C is

deterministic, and Model B is deterministic *given a profile*; the per-asset mean profile is drawn into a per-fire realisation inside Model D, which jitters the bite time and growth rate once per fire and then samples the chain march. Every draw from the generator therefore still happens in Model D, and the whole run stays bit-for-bit reproducible.

6 Validation and Backtesting

Historical calibration and backtesting are open remediation items (FIRE-R1, FIRE-R2) — to be completed before production sign-off. Current validation is by directional unit tests, which confirm the model behaves as designed: a strong asset contains far more often and loses far less than a weak one; the point-of-no-return gate is reachable for genuinely weak or tall-and-unsprinklered assets but not for strong ones; the fire-service-reach interaction flips a building from fully containable to a guaranteed conflagration exactly at the reach height when internal suppression is inadequate; and a whole-portfolio run is bit-for-bit reproducible from a single seed. The sensitivity tables in the next section are produced by running the *actual* end-to-end model.

6.1 Automated Test Results

Automated test-results table to be generated. Stage 1 + Stage 2 unit tests live under `tests/models/fire/` (`test_initiation.py`, `test_response_effectiveness.py`, `test_intensity_track.py`, `test_containment.py`).

7 Sensitivity Analysis

The tables below are generated by `docs/models/sensitivities/fire_resilience/` running the real `simulate_asset_fire` against a fixed baseline (a fully-occupied 3-storey *Office* in *Fair* condition with every resilience field at *Meets minimum*), varying one input at a time over 20,000 Monte Carlo draws per cell.

The results show a clear **division of labour** across the three levers, which is how the credit is designed to respond.

Suppression drives containment and loss frequency. On the suppression sweep the containment rate climbs monotonically from $\approx 47\%$ at *Not assessed* to $\approx 88\%$ at *Verified*, and the annual loss frequency falls roughly five-fold ($1.14\% \rightarrow 0.24\%$). This is the continuous matrix blend at work: a better-suppressed asset has a higher controllability score C , a blend weight w nearer the controllable matrix, and so a higher chance of reaching S_5 (Contained). This is the economically meaningful credit the suppression system buys.

Compartmentation drives conflagration. Passive compartmentation selects the growth rate, so a poorly-compartmented asset grows fast enough to cross I_{crit} before the chain absorbs and shows a small but non-zero fire-spread (≈ 9 bps), whereas a well-compartmented one stays at zero. Commercial type and condition move loss frequency in the expected order (hotels, mixed-use and poor condition worst).

The fire-service reach drives the catastrophic tail. The height sweep (held at sub-threshold *Partial* internal suppression) is the headline non-linearity: at or below the 4-storey reach the credit is ordinary ($\approx 56\%$ containment, zero spread), but at *five storeys and above* suppression can no longer engage and the asset collapses to 0% containment and ≈ 203 bps of fire spread — a **cliff, not a gradient**.

The height $\times$ suppression grid confirms the cliff appears precisely when internal suppression is below the sprinkler threshold; *Meets minimum* or better keeps the asset reachable at any height.

Table 1: Fire-resilience credit by suppression-systems level (3-storey office baseline)

	Containment (%)	Loss freq (%)	Fire spread (bps)
Suppression: Not assessed	47.2	1.14	6.00
Suppression: Partial	55.7	0.9000	0.0000
Suppression: Meets minimum	62.1	0.7700	0.0000
Suppression: Enhanced	78.3	0.4400	0.0000
Suppression: Verified	87.8	0.2400	0.0000

Table 2: Fire-resilience credit by building height — the fire-service reach cliff (reach = 4 storeys; sub-threshold 'Partial' internal suppression)

	Containment (%)	Loss freq (%)	Fire spread (bps)
1 storeys	55.7	0.9000	0.0000
3 storeys	55.7	0.9000	0.0000
4 storeys	55.7	0.9000	0.0000
5 storeys	0.0000	2.03	203.0
8 storeys	0.0000	2.03	203.0
12 storeys	0.0000	2.03	203.0
20 storeys	0.0000	2.03	203.0

Table 3: Fire-resilience credit by compartmentation level

	Containment (%)	Loss freq (%)	Fire spread (bps)
Compartments: Not assessed	58.9	0.8350	9.50
Compartments: Partial	59.1	0.8300	9.00
Compartments: Meets minimum	62.1	0.7700	0.0000
Compartments: Enhanced	64.0	0.7300	0.0000
Compartments: Verified	65.8	0.6950	0.0000

Table 4: Fire-resilience credit by property condition

	Containment (%)	Loss freq (%)	Fire spread (bps)
Excellent	62.8	0.6150	0.0000
Good	63.6	0.6900	0.0000
Fair	62.1	0.7700	0.0000
Poor	63.3	0.8650	0.0000
Very poor	61.5	1.05	0.0000

Table 5: Fire-resilience credit by commercial type

	Containment (%)	Loss freq (%)	Fire spread (bps)
Office	62.1	0.7700	0.0000
Retail	64.0	1.21	0.0000
Hotel	64.8	1.51	0.0000
Leisure	63.8	1.35	0.0000
Healthcare	61.5	1.08	0.0000
MultiFamily	63.7	1.17	0.0000
MixedUse	64.8	1.59	0.0000
Other	64.8	1.59	0.0000

Table 6: Fire spread (bps) by height $\times$ suppression level — the fire-service-reach interaction

	Partial	Meets minimum	Enhanced	Verified
3 storeys	0.0000	0.0000	0.0000	0.0000
5 storeys	203.0	0.0000	0.0000	0.0000
8 storeys	203.0	0.0000	0.0000	0.0000
20 storeys	203.0	0.0000	0.0000	0.0000

7.1 Interpretation and known model weaknesses

The sensitivity surface above is directionally correct, but two structural features of the seed model deserve to be read alongside it:

- **The conflagration leg is muted for reachable buildings (by construction).** The chain march absorbs in ≈ 3 steps on average, while latent intensity climbs only ≈ 0.8 per step, so a *reachable* fire is contained or lost long before the slow point-of-no-return race can cross I_{crit} . The PNR gate therefore fires almost only for the structural reach cliff (tall, under-sprinklered) plus a sliver of no-suppression cases. This is *semantically* correct — a fire contained in three steps did not reach a point of no return — but it means fire-spread (bps) is driven chiefly by compartmentation and the reach cliff rather than by the active suppression level. Making spread respond strongly to suppression as well would require retuning the intensity-race timescale (a lower I_{crit} , of order 2–4) so the race competes with the three-step march; that is deferred as an enhancement because it is arguably less physical than letting compartmentation own conflagration.
- **The height sweep only shows the cliff at sub-threshold internal suppression.** A building with *Meets minimum* or better internal sprinklers is reachable at any height, so its height sweep is flat by design; the cliff table here deliberately holds suppression at *Partial* to expose the reach behaviour. A future enhancement would model fire-service reach as a brigade-specific, possibly height-graded capability rather than a single global 4-storey threshold, and would let internal-suppression adequacy degrade with height rather than flip at a hard index.

The exposed calibration knobs make the surface tunable without code changes: $C_{\text{floor}}/C_{\text{ceil}}$ position and stretch the containment band; `suppression_bite_multiplier_by_level` and `i_crit_by_level`

set the gradient steepness; and the `race_jitter` sigmas (σ_τ, σ_g) control how smoothly the conflagration probability grades.

8 Model Limitations and Known Weaknesses

- **No historical calibration (FIRE-L1).** Every parameter is a first-pass placeholder; outputs are relative signals, not absolute loss probabilities.
- **Height proxy (FIRE-L2).** Building height is proxied by `NumberOfStoreys`; the commercial schema carries no explicit floor-level or floor-area metric, and the fire-service reach is a single global threshold rather than a brigade-specific capability.
- **Muted conflagration leg for reachable buildings (FIRE-L3).** The chain march absorbs faster (≈ 3 steps) than the latent intensity race reaches I_{crit} , so the point-of-no-return / fire-spread leg responds chiefly to compartmentation and the reach cliff rather than to the active suppression level (see Section 7, “Interpretation and known model weaknesses”). Suppression instead expresses itself through containment rate and loss frequency. Retuning the intensity-race timescale to make spread suppression-sensitive is an open enhancement (FIRE-R5).
- **Single global reach threshold (FIRE-L4).** The fire-service reach is one global 4-storey cut-off and internal-suppression adequacy flips at a hard index, so the height cliff only appears at sub-threshold internal suppression. A brigade-specific, height-graded reach and a graded internal-suppression adequacy are an open enhancement (FIRE-R6).
- **Initiation class is cosmetic.** The initiation-class draw seeds an entry-point label but does not yet change the chain start state or dynamics.
- **Fixed loss-given-event.** The conflagration leg assumes a 100% loss given a point-of-no-return event; partial-loss quantum is not modelled.
- **Independence assumption.** The root-sum-of-squares all-in coupon treats fire, flood and wind as independent perils.

9 Governance Validation Questions

The following nine questions are mandated by Chapter 5 of the *Model Risk Governance for Vendors* handbook and must be explicitly addressed for every model in the inventory. Response status is tracked in the Model Governance Dashboard (MRC portal).

9.1 Q1 — Model purpose and use

“What is the model used for, and is that use consistent with its design?”

The model produces a fire-resilience credit (loss frequency, containment rate and a full-conflagration bps leg) consumed by the Commercial PRS Pricer as an independent peril. Use is consistent with design; it is *not* a life-safety or absolute-loss model and must not be used as one.

9.2 Q2 — Data quality and lineage

“Are the model inputs of sufficient quality, and is their lineage documented?”

Inputs are the CDM-derived `AssetFireFeatures`; lineage runs from the v2 CDM resilience fields through Model C. Missing options degrade gracefully to neutral multipliers. Data *quality* for calibration (historical fire-loss records) is not yet established — open under FIRE-R1.

9.3 Q3 — Conceptual soundness

“Is the methodology theoretically sound and appropriate?”

The four-model decomposition (Poisson initiation, per-fire-stochastic intensity track, resilience→dynamics mapping, controllability-blended absorbing Bayesian chain) is a standard and defensible structure. The point-of-no-return gate is the one bespoke mechanism and is documented in Section 3.4. Its known weakness — that the gate is muted for reachable buildings because the chain absorbs faster than the intensity race — is documented in Sections 7 and 8 (FIRE-L3).

9.4 Q4 — Calibration and parameters

“How were parameters estimated, and are they appropriate?”

All parameters are engineering-judgement seeds (`version "seed-0"`) chosen for direction and ordering, not magnitude. Statistical calibration is an open remediation item (FIRE-R1, FIRE-R2). This is the dominant model limitation.

9.5 Q5 — Implementation verification

“Has the implementation been verified against the specification?”

Yes, by unit tests covering each component model, the orchestrator and the fire-service-reach interaction, plus a reproducibility test. The sensitivity harness exercises the full end-to-end path.

9.6 Q6 — Sensitivity and stability

“How sensitive are outputs to inputs and assumptions?”

See Section 7. Outputs respond in the expected order to type, condition and resilience levels, and exhibit the designed height×suppression cliff. Stability under draw count is adequate at $n_{\text{sim}} = 20,000$; formal stability bands are an open item.

9.7 Q7 — Limitations and weaknesses

“Are the model’s limitations understood and documented?”

Yes — see Section 8 (FIRE-L1–L4 plus three further documented assumptions) and the “Interpretation and known model weaknesses” discussion in Section 7. The placeholder calibration and the height proxy are the most material; the muted conflagration leg (FIRE-L3) and single global reach threshold (FIRE-L4) are documented enhancements.

9.8 Q8 — Ongoing monitoring

“How will the model be monitored in production?”

Not yet in production. A monitoring plan (drift in input distributions, backtest of realised vs. predicted fire frequency) will be defined as part of FIRE-R3 before sign-off.

9.9 Q9 — Governance and approval

“What is the approval status and who owns the model?”

Owner: David K Kelly. Status: Development, Tier 2, RAG Amber. Not approved for production; MRC sign-off is contingent on closing FIRE-R1–R6.

10 Change History

Version	Date	Description
0.01	05-June-2026	Skeleton document. Stage 1 (Model A) and Stage 2 (Models B/C/D + orchestrator) delivered with unit tests; placeholder calibration.
0.20	06-June-2026	Completed all narrative sections; added the full mathematical framework (Models A–D, PNR, credit currency), governance validation questions, the matrices appendix, and the end-to-end sensitivity analysis.
0.30	08-June-2026	Recalibrated the resilience response: replaced the saturated hard controllability switch with a continuous matrix blend ($M = wM_{AA} + (1 - w)M_{weak}$); graded I_{crit} across all five suppression levels; added a per-level suppression-bite multiplier and per-fire race jitter (σ_τ, σ_g) so the conflagration probability is no longer all-or-nothing. Containment now grades $47\% \rightarrow 88\%$ across suppression levels (previously flat $\approx 88.7\%$). Reworked the height sweep to expose the 4-storey reach cliff and documented the muted-conflagration and single-reach-threshold weaknesses (FIRE-L3/L4, FIRE-R5/R6).

A Initiation Parameters (Model A)

Table 7: Base annual ignition frequency $\lambda_{0,c}$ by commercial type.

Commercial type	$\lambda_{0,c}$
Office	0.020
Retail	0.030
Hotel	0.035
Leisure	0.030
Healthcare	0.025
MultiFamily	0.030
MixedUse	0.040
Other	0.040

Table 8: Frequency modifiers. Process load is by type with a business-rates override (*Restaurant* = 1.25) that only ever raises the value.

Modifier	Option	Value
m_{occ} (occupancy)	Fully occupied	1.00
	Partially vacant	0.90
	Vacant	0.85
m_{proc} (process)	Office	1.00
	Retail	1.10
	Hotel	1.25
	Leisure	1.25
	Healthcare	1.10
	MultiFamily	1.05
	MixedUse	1.15
	Other	1.15
m_{cond} (condition)	Excellent	0.85
	Good	0.95
	Fair	1.05
	Poor	1.20
	Very poor	1.40
$m_{\text{protection}}$ (level)	Not assessed	1.10
	Partial	1.00
	Meets minimum	0.95
	Enhanced	0.90
	Verified	0.85
m_{history}	No prior	1.00
	Old minor	1.05
	Recent minor	1.10
	Recent moderate/severe	1.40

A fire counts as *recent* if it occurred within `recent_years` = 5 years.

B Initiation-Class Priors

Table 9: Categorical priors over initiation class by commercial type (IE = InternalElectrical, KO = KitchenOccupant, PT = PlantTechnical, PR = ProcessRelated, BS = BasementService, EE = ExternalExposure, Oth = Other). Rows sum to 1.

Type	IE	KO	PT	PR	BS	EE	Oth
Office	0.35	0.10	0.15	0.05	0.15	0.10	0.10
Retail	0.30	0.30	0.10	0.05	0.10	0.10	0.05
Hotel	0.28	0.25	0.15	0.05	0.12	0.10	0.05
Leisure	0.30	0.22	0.13	0.05	0.10	0.10	0.10
Healthcare	0.32	0.12	0.20	0.06	0.12	0.08	0.10
MultiFamily	0.35	0.30	0.08	0.02	0.10	0.08	0.07
MixedUse	0.28	0.12	0.15	0.12	0.10	0.08	0.15
Other	0.28	0.12	0.15	0.12	0.10	0.08	0.15

C Progression Dynamics (Models B & C)

Table 10: Growth, critical intensity, and per-level detection / suppression multipliers.

Block	Key	Value
Growth per step	well_compartmented	1.5
	poorly_compartmented	4.0
I_{crit} by suppression level	Not assessed	6
	Partial	10
	Meets minimum	16
	Enhanced	30
	Verified	55
Detection mult. δ_ℓ	Not assessed	1.00
	Partial	0.90
	Meets minimum	0.80
	Enhanced	0.72
	Verified	0.65
Suppression growth mult. μ_{supp}	Not assessed	1.00
	Partial	0.85
	Meets minimum	0.75
	Enhanced	0.62
	Verified	0.50
Suppression bite mult. b_{supp} (lower = sooner)	Not assessed	3.00
	Partial	2.00
	Meets minimum	1.30
	Enhanced	0.85
	Verified	0.55

Table 11: Passive and response effectiveness ranges $[\ell, u]$, interpolated by the field’s resilience-level fraction.

Group	Field	ℓ	u
Passive	StructuralFireResistanceAdequate	0.60	0.95
	CompartmentsProvided	0.55	0.90
	FireStoppingAtPenetrations	0.70	0.95
	ExternalMaterialsFireResistant	0.65	0.95
Response	AccessRouteResilient	0.70	0.95
	EmergencyProceduresTested	0.80	0.97
	BusinessContinuityPlanInPlace	0.85	0.98

Table 12: Height penalty, fire-service reach, timing, intensity dynamics and controllability.

Block	Key	Value
Height penalty	range	[1.05, 1.60]
	max_storeys	20
Fire-service reach	reach_storeys R	4
	internal_suppression_threshold θ_{int}	2
	no_reach_bite_steps τ_{∞}	999.0
Timing	detection_base_steps $\tau_{\text{det},0}$	2.0
	suppression_bite_base_steps $\tau_{\text{supp},0}$	3.0
Intensity dynamics	passive_damp_weight w_{damp}	0.60
	compartmentation_well_threshold	2.0
Race jitter	timing_sigma σ_{τ}	0.50
	growth_sigma σ_g	0.40
Controllability	weight (passive) w_p	0.35
	weight (response) w_r	0.35
	weight (suppression) w_s	0.30
	blend_floor C_{floor}	0.18
	blend_ceiling C_{ceil}	0.88

D Containment Transition Matrices (Model D)

Each is an 8×8 row-stochastic matrix over $\{S_0, \dots, S_7\}$; the step length is 15 minutes. The point-of-no-return matrix has a zero S_5 (Contained) column — once latched, containment is unreachable.

Table 13: **AA_controllable** — pre-PNR controllable endpoint; the active matrix at blend weight $w = 1$ (controllability $C \geq C_{\text{ceil}}$).

From\To	S_0	S_1	S_2	S_3	S_4	S_5	S_6	S_7
S_0	0.05	0.25	0.70	0	0	0	0	0
S_1	0	0.10	0.85	0	0	0	0.05	0
S_2	0	0	0	0.90	0.08	0	0.02	0
S_3	0	0	0	0	0.10	0.85	0.05	0
S_4	0	0	0	0	0	0.75	0.20	0.05
S_5	0	0	0	0	0	1.00	0	0
S_6	0	0	0	0	0	0	0.85	0.15
S_7	0	0	0	0	0	0	0	1.00

Table 14: **weak_controllable** — pre-PNR weak endpoint; the active matrix at blend weight $w = 0$ (controllability $C \leq C_{\text{floor}}$). For intermediate C the active matrix is $w M_{\text{AA}} + (1 - w) M_{\text{weak}}$.

From\To	S_0	S_1	S_2	S_3	S_4	S_5	S_6	S_7
S_0	0.05	0.70	0.25	0	0	0	0	0
S_1	0	0.10	0.45	0	0	0	0.45	0
S_2	0	0	0	0.40	0.40	0	0.20	0
S_3	0	0	0	0	0.15	0.15	0.70	0
S_4	0	0	0	0	0	0.25	0.50	0.25
S_5	0	0	0	0	0	1.00	0	0
S_6	0	0	0	0	0	0	0.50	0.50
S_7	0	0	0	0	0	0	0	1.00

Table 15: **point_of_no_return** — absorbing, post-PNR. The S_5 column is identically zero.

From\To	S_0	S_1	S_2	S_3	S_4	S_5	S_6	S_7
S_0	0	1.00	0	0	0	0	0	0
S_1	0	0	0	0	0	0	0.40	0.60
S_2	0	0	0	0	0	0	0.40	0.60
S_3	0	0	0	0	0	0	0.40	0.60
S_4	0	0	0	0	0	0	0.40	0.60
S_5	0	0	0	0	0	1.00	0	0
S_6	0	0	0	0	0	0	0.35	0.65
S_7	0	0	0	0	0	0	0	1.00