Property Flood Response Model

Gauge-to-Property Flood Propagation, Damage Estimation, and Portfolio Loss Analytics

MKM Research Labs

Version 1.2 — 26-March-2026

David K Kelly

Legal Notice

PROPRIETARY AND CONFIDENTIAL

Contents

1 Executive Summary

This document presents the mathematical framework, calibration, and validation of the Property Flood Response Model (MKM-PF-001) implemented across `src/models/floodrisk/` and `src/port/src/property/ts/`. The model translates gauge-level flood events into property-level damage and financial loss through five integrated stages:

1. **Nearest-Gauge WSE Propagation (v2.1)** — Flood WSE is taken from the single nearest hydraulically connected gauge (typically a synthetic gauge on the river centreline). Two additional flanking gauges are retained for PRS pricing only. *(Replaces IDW interpolation from v1.0–1.1.)*
2. **Flood Depth Calculation** — Conversion of gauge WSE to property flood depth, accounting for property ground elevation, floor level, and near-field retention bypass ($d < 2\,\text{km} \Rightarrow$ retention $= 1.0$).
3. **Velocity and Travel Time** — Manning's equation for lateral overbank flow velocity, travel time estimation, and property-level hydrograph construction with recession modelling.
4. **Depth-Damage Function** — UK-calibrated piecewise-linear vulnerability curve mapping inundation depth to structural damage ratio, with property type and floor level adjustments.
5. **Portfolio Aggregation** — Monte Carlo simulation with spatially correlated losses, Value-at-Risk (VaR), Expected Shortfall (ES), and Herfindahl–Hirschman concentration index.

The model is the primary bridge between gauge-observed hydrology and property-level financial loss. It serves regulatory capital estimation, mortgage portfolio risk management, and Property Risk Swap (PRS) pricing within the Thames catchment. The model operates on per-gauge storm response time-series (`gaugets/`) and property metadata (`property.json`), producing per-property flood event files in `propertyts/` and a portfolio flood summary.

Key risk: The depth-damage curve is the single most material model component. A ±10% shift in the vulnerability curve propagates to approximately ±8% in expected portfolio loss. The v2.1 nearest-gauge approach eliminates IDW interpolation error but introduces dependence on the quality of synthetic gauge placement and the assumption that the nearest river point controls the flood boundary condition.

v2.1 change notice (26-Mar-2026): IDW interpolation replaced by nearest-gauge propagation. Near-field retention bypass added ($d < 2\,\text{km}$). Synthetic gauges placed on river centreline. Properties with flood events increased from 21/200 (10.5%) to 100/200 (50%). See `flood_propagation_model.md` for full change rationale.

2 Model Purpose and Scope

2.1 Purpose

The Property Flood Response Model quantifies the financial impact of fluvial flooding on individual properties and the aggregate mortgage portfolio. It serves three primary functions:

- **Event propagation:** Translating gauge-level storm responses into property-level flood events, determining which properties are flooded, when the flood arrives, and the peak inundation depth.
- **Damage estimation:** Computing property-level damage ratios from flood depth using a UK-

calibrated vulnerability curve with property type and floor level adjustments, and converting to sterling losses.

- **Portfolio risk aggregation:** Computing portfolio-level tail risk metrics (VaR, ES) accounting for spatial dependence between property losses, and measuring geographic concentration via HHI.

2.2 Scope

- **In scope:** IDW spatial interpolation of WSE from gauges to properties; elevation-corrected flood depth calculation; Manning's equation velocity and travel time; exponential distance attenuation; property hydrograph construction; depth-damage vulnerability curve with property type and floor level adjustments; Monte Carlo loss simulation; portfolio concentration analytics.
- **Out of scope:** Hydrological rainfall-runoff modelling; hydraulic channel routing; GEV hazard curve fitting (covered in the Hazard Curve Model document, MKM-GEV-001); PRS pricing mechanics (covered in the PRS Pricing Model document); gauge-level storm response generation (covered in `stormgauge/forward_model.py`).
- **Upstream dependencies:** Per-gauge storm response timeseries (`gaugets/`), gauge metadata (`gauge.json`), property portfolio (`property.json`).
- **Downstream consumers:** Property Hazard Curve Model (MKM-PHC-001), PRS Pricing Model, Regulatory Capital Engine, PDF reporting suite, Trading Desk stress testing module.

3 Mathematical Framework

3.1 Gauge-to-Property Spatial Interpolation

3.1.1 Nearest Gauge Selection

For each property j at location (φ_j, λ_j), the $k = 3$ nearest gauges are identified by Haversine distance. The Haversine formula computes the great-circle distance between two points on the Earth's surface:

$$d = 2R \arcsin\left(\sqrt{\sin^2\left(\frac{\Delta\varphi}{2}\right) + \cos\varphi_1 \cos\varphi_2 \sin^2\left(\frac{\Delta\lambda}{2}\right)}\right) \tag{1}$$

where $R = 6{,}371{,}000\,\mathrm{m}$ is the mean Earth radius, $\Delta\varphi = \varphi_2 - \varphi_1$, and $\Delta\lambda = \lambda_2 - \lambda_1$, with all angles in radians. The implementation uses the numerically stable `atan2` form: $c = 2 \cdot \mathrm{atan2}(\sqrt{a}, \sqrt{1-a})$.

The three nearest gauges are sorted by ascending distance and used for all subsequent interpolation and flood propagation calculations for that property.

3.1.2 Storm Collection

For each property, storms are collected from all k nearest gauges. A storm is included if it exceeded the flood alert threshold at *any* of the nearest gauges. This ensures that the property is assessed against all relevant flood events, even if the nearest gauge did not trigger but a nearby gauge did. The union of alert-exceeding storms across the k gauges defines the set of candidate flood events for the property.

3.1.3 Nearest-Gauge WSE Propagation (v2.1)

Change (v2.1): The water surface elevation (WSE) at the property is taken from the single nearest hydraulically connected gauge. This is typically a *synthetic gauge* placed at the closest point on the river centreline to the property (see Section **??** below). The two flanking real gauges are retained for PRS pricing and audit only.

$$\hat{W} = W_{\text{nearest}} \tag{2}$$

This replaces v1.0–1.1's IDW interpolation, which averaged WSE across $k = 3$ gauges using inverse-square weighting. IDW caused signal dilution when one gauge showed high water and two showed low water, producing an unrealistically low interpolated WSE. The nearest-gauge approach aligns with Book 210's requirement for flow-aligned propagation rather than isotropic spatial averaging.

Legacy IDW (deprecated). The IDW function remains available in `spatial.py:idw_interpolate_from_gaug` for reference: $\hat{W} = \sum w_i W_i / \sum w_i$ with $w_i = 1/d_i^2$. It is no longer called in the flood propagation pipeline.

3.2 Flood Depth Calculation

3.2.1 Elevation Differential

The raw flood depth at property j is determined by the difference between the interpolated WSE and the property ground elevation:

$$d_{\text{raw},j} = \max\bigl(0,\ \hat{W}_j \cdot a(r_j) - e_j\bigr) \tag{3}$$

where $\hat{W}_j$ is the nearest-gauge WSE at the property (v2.1), $a(r_j)$ is the distance retention factor (Section 3.2.3), and e_j is the property ground elevation in metres above ordnance datum (AOD).

3.2.2 Floor Level Adjustment

The effective flood depth accounts for the property's floor level (step height) above ground. Given floor level f_j (metres above ground):

$$d_j = \max\bigl(0,\ \hat{W}_j \cdot a(r_j) - e_j - f_j\bigr) \tag{4}$$

The flood threshold for property j is $e_j + f_j$. The property is flagged as flooded when $d_j > 0$. Properties with higher floor levels require proportionally deeper flooding before interior damage commences, reflecting the physical protection afforded by raised ground floors.

3.2.3 Exponential Distance Retention

The fraction of peak WSE retained at the property decays exponentially with distance from the river channel or gauge:

$$a(r) = \exp\left(-\frac{r}{\ell}\right) \tag{5}$$

where r is the distance (metres) from the nearest gauge to the property and $\ell = 10{,}000\,\text{m}$ is the characteristic retention length scale. At $r = 0$ the retention factor is 1.0 (full signal, no loss); at $r = 1\,\text{km}$ it is ≈ 0.90; at $r = \ell$ it is $e^{-1} \approx 0.37$. The large length scale reflects near-field hydraulic guidance: for properties within a few hundred metres to $\sim$1 km of the channel, peak WSE attenuation is negligible in the absence of significant floodplain storage or topographic barriers.

3.3 Velocity and Travel Time

3.3.1 Manning's Equation

Lateral overbank flow velocity is computed using Manning's equation for open-channel flow. For shallow floodplain flow where the hydraulic radius R_h is approximated by the water depth h:

$$v = \frac{1}{n}\, h^{2/3}\, S^{1/2} \tag{6}$$

where:

- v is the flow velocity (m/s),
- $n = 0.04$ is Manning's roughness coefficient for urban floodplain (dimensionless),
- h is the water depth (metres), approximating the hydraulic radius for wide, shallow overbank flow,
- S is the energy slope (dimensionless).

The slope is clamped to a minimum value $S_{\min} = 0.001$ to prevent numerical instability in very flat terrain. If $h \leq 0$ or $n \leq 0$, the velocity is zero.

3.3.2 Ground Slope Estimation

The energy slope between gauge i and property j is estimated from the elevation difference:

$$S_{ij} = \max\left(\frac{|e_i^{\text{gauge}} - e_j^{\text{prop}}|}{d_{ij}},\ S_{\min}\right) \tag{7}$$

where e_i^{gauge} and e_j^{prop} are ground elevations in metres AOD and d_{ij} is the horizontal Haversine distance in metres. The minimum slope floor prevents division instability on flat terrain.

3.3.3 Travel Time

The time for the flood front to propagate laterally from the gauge to the property is:

$$t_{\text{travel}} = \frac{d}{v \cdot 3600} \quad \text{(hours)} \tag{8}$$

where d is the lateral distance in metres and v is the Manning velocity from Equation 6. If $v = 0$, the travel time is infinite (the flood front does not reach the property). Travel times determine the lag between gauge peak and property peak in the hydrograph construction.

3.4 Property Hydrograph Construction

A property-level hydrograph is constructed by transforming the gauge flood simulation timeseries (typically 60–168 hourly readings) through lag shifting, amplitude scaling, and recession stretching.

3.4.1 Rising Limb

For hours $t \leq t_{\text{peak}}^{\text{gauge}} + t_{\text{travel}}$, the property WSE is computed by sampling the gauge timeseries at time $t - t_{\text{travel}}$ (direct time shift):

$$W_{\text{prop}}(t) = W_{\text{base}} + (\hat{W}_{\text{peak}} - W_{\text{base}}) \cdot s(t) \cdot a \tag{9}$$

where $s(t)$ is the normalised gauge hydrograph shape at the shifted time:

$$s(t) = \frac{W_{\text{gauge}}(t - t_{\text{travel}}) - W_{\text{base}}}{W_{\text{gauge,peak}} - W_{\text{base}}} \tag{10}$$

$\hat{W}_{\text{peak}}$ is the IDW-interpolated peak WSE at the property, W_{base} is the gauge base level, and a is the distance attenuation factor from Equation 5.

3.4.2 Recession Limb

For hours after the shifted peak, time is stretched by the recession factor γ (default 1.5) to model slower drainage:

$$t_{\text{source}} = t_{\text{peak}}^{\text{gauge}} + \frac{t - (t_{\text{peak}}^{\text{gauge}} + t_{\text{travel}})}{\gamma} \tag{11}$$

The recession factor $\gamma > 1$ reflects the physical reality that floodplain drainage is slower than inundation due to reduced hydraulic gradient and ponding effects. The default $\gamma = 1.5$ means the recession limb takes 50% longer than the corresponding gauge recession.

3.4.3 Property Flood Depth

The property flood depth at each hour is:

$$d_{\text{prop}}(t) = \max\bigl(0,\ W_{\text{prop}}(t) - e_j - f_j\bigr) \tag{12}$$

where e_j is the property ground elevation and f_j is the floor level above ground. The property is flagged as flooded when $d_{\text{prop}}(t) > 0$.

The hydrograph output includes for each hour: water surface elevation, flood depth, and a boolean flood flag. Summary statistics extracted from the hydrograph include flood arrival time, peak time, peak flood depth, and total flood duration.

3.5 Depth-Damage Function

3.5.1 UK-Calibrated Vulnerability Curve

The depth-damage relationship is defined by ten control points calibrated against UK flood damage survey data. The curve maps inundation depth d (metres above finished floor level) to a damage ratio $\delta \in [0, 1]$ representing the fraction of reinstatement cost incurred:

Depth (m)	0	0.05	0.5	1.0	1.5	2.0	3.0	4.0	5.0	6.0
Damage ratio	0	0.05	0.25	0.40	0.50	0.60	0.75	0.85	0.95	1.0

Table 1: UK-calibrated depth-damage control points

Between control points, the damage ratio is computed by linear interpolation. Let (d_i, δ_i) and (d_{i+1}, δ_{i+1}) be adjacent control points such that $d_i \leq d < d_{i+1}$. The interpolated damage ratio is:

$$\delta(d) = \delta_i + \frac{d - d_i}{d_{i+1} - d_i} (\delta_{i+1} - \delta_i) \tag{13}$$

with boundary conditions $\delta(d) = 0$ for $d \leq 0$ and $\delta(d) = 1$ for $d \geq 6\,\text{m}$. The scalar implementation in `scalar_depth_damage()` uses direct iteration over control point pairs; the array implementation uses `scipy.interpolate.interp1d` with `kind='linear'` and `fill_value=(0, 1)`.

3.5.2 Property Type Adjustments

The base vulnerability curve is scaled by a property type factor α_p to reflect differential susceptibility to flood damage:

$$\alpha_p = \begin{cases} 1.0 & \text{residential} \\ 1.2 & \text{commercial} \\ 0.9 & \text{industrial} \\ 1.0 & \text{unknown / other} \end{cases} \tag{14}$$

Commercial properties attract a 20% loading reflecting higher contents-to-structure ratios and business interruption costs. Industrial properties receive a 10% reduction reflecting more resilient construction (concrete floors, industrial drainage). Unknown property types default to the residential baseline.

3.5.3 Final Damage Ratio

The adjusted damage ratio for property j is:

$$D_j = \min\Big(\alpha_{p_j} \cdot \delta\big(\max(d_j - f_j,\ 0)\big),\ 1\Big) \tag{15}$$

where d_j is the peak flood depth at the property location, f_j is the floor level, and α_{p_j} is the property type factor. The result is clipped to $[0, 1]$ to prevent physically meaningless values when the type factor $\alpha > 1$.

3.5.4 Financial Loss

The sterling loss for property j under storm event s is:

$$L_{j,s} = V_j \cdot D_{j,s} \tag{16}$$

where V_j is the property value (reinstatement cost) and $D_{j,s}$ is the damage ratio for that storm event.

3.6 Portfolio Aggregation

3.6.1 Monte Carlo Simulation with Correlated Losses

Portfolio loss is simulated using a multivariate normal shock model. For N properties with pre-computed damage ratios $\boldsymbol{\delta} = (\delta_1, \ldots, \delta_N)$ and spatial correlation matrix $\boldsymbol{\Sigma}$:

1. Generate $M = 1{,}000$ shock vectors $\boldsymbol{z}^{(m)} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Sigma})$, $m = 1, \ldots, M$.
2. Compute shocked damage ratios:

$$\tilde{\delta}_j^{(m)} = \operatorname{clip}\big(\delta_j \cdot (1 + 0.2\, z_j^{(m)}),\ 0,\ 1\big) \tag{17}$$

3. Compute portfolio loss for each simulation:

$$L^{(m)} = \sum_{j=1}^{N} V_j \cdot \tilde{\delta}_j^{(m)} \tag{18}$$

   where V_j is the property value.

The shock scaling factor of 0.2 controls the dispersion of loss outcomes around the central estimate. The spatial correlation matrix ensures that nearby properties experience correlated shocks, capturing the clustering behaviour of flood losses.

The correlation matrix uses an exponential kernel. For properties i and j separated by distance d_{ij} (kilometres):

$$\rho_{ij} = \begin{cases} 1 & i = j \\ \rho_0 \cdot \exp\left(-\dfrac{d_{ij}}{L/1000}\right) & i \neq j \end{cases} \tag{19}$$

where ρ_0 is the base correlation and L is the correlation distance (metres).

3.6.2 Value-at-Risk and Expected Shortfall

From the simulated loss distribution $\{L^{(1)}, \ldots, L^{(M)}\}$:

$$\mathrm{VaR}_\alpha = F_L^{-1}(\alpha) \tag{20}$$

where F_L^{-1} is the empirical quantile function. The model computes VaR at the 95th and 99th percentiles.

The Expected Shortfall (Conditional VaR) at level α is:

$$\mathrm{ES}_\alpha = \mathbb{E}[L \mid L \geq \mathrm{VaR}_\alpha] = \frac{1}{|\{m : L^{(m)} \geq \mathrm{VaR}_\alpha\}|} \sum_{\{m : L^{(m)} \geq \mathrm{VaR}_\alpha\}} L^{(m)} \tag{21}$$

Both ES_{95} and ES_{99} are reported.

3.6.3 Herfindahl–Hirschman Index

Portfolio concentration is measured using the Herfindahl–Hirschman Index (HHI). For property-level losses $\ell_j = V_j \cdot \delta_j$:

$$\mathrm{HHI} = \sum_{j=1}^{N} s_j^2, \qquad s_j = \frac{\ell_j}{\sum_{k=1}^{N} \ell_k} \tag{22}$$

where s_j is the share of total portfolio loss attributable to property j. $\mathrm{HHI} = 1$ indicates perfect concentration; $\mathrm{HHI} = 1/N$ indicates perfect diversification. The same metric is applied at the geographic grid cell level to measure spatial concentration.

3.6.4 Portfolio Flood Summary Metrics

The generator computes and persists the following portfolio-level statistics:

- **Properties with floods:** Count of properties where at least one storm produces $d_j > 0$.
- **Gauge-to-property ratio:** Percentage of gauge-level alert storms that produce property-level flooding, measuring the attenuation effect of elevation and distance.
- **Maximum depth and damage ratio:** Worst-case across all properties and storms.
- **Per-property summaries:** Nearest gauges, flood count, maximum depth, maximum damage ratio.

4 Input Parameters and Calibration

4.1 Model Parameters

Parameter	Symbol	Default	Description
Nearest gauges	k	3	1 synthetic + 2 real (flood uses nearest only)
Retention length	L	10,000 m	Exponential decay (beyond near-field)
Near-field threshold		2,000 m	Below this, retention = 1.0 (v2.1)
Synthetic dedup		50 m	Properties sharing a synthetic gauge
Manning's roughness	n	0.04	Urban floodplain roughness coefficient
Attenuation length	ℓ	2,000 m	Exponential WSE decay distance
Minimum slope	$S_{\min}$	0.001	Floor to prevent division instability
Recession factor	γ	1.5	Recession-to-rising limb duration ratio
Shock dispersion	—	0.2	Standard deviation scaling in MC
Monte Carlo sims	M	1,000	Number of portfolio loss scenarios
Correlation distance	L	Configurable	Spatial correlation decay length
Base correlation	ρ_0	Configurable	Off-diagonal correlation at $d = 0$
Random seed	—	42	Reproducibility seed

Table 2: Model parameters and default values

4.2 Depth-Damage Calibration

The vulnerability curve (Table 1) is calibrated against UK flood damage survey data following the methodology of Penning-Rowsell et al. (2013). The ten control points were selected to capture:

- **Threshold effects:** Minimal damage below 0.05 m (below-floor wetting only).
- **Rapid escalation:** 0.05–1.0 m is the primary damage accumulation zone (electrical systems, ground floor contents, plasterwork).

- **Saturation:** Above 3.0 m, marginal damage per additional metre decreases as the structure approaches total loss.

4.3 Property Type Factor Calibration

The property type multipliers (Equation 14) are derived from Multi-Coloured Manual (MCM) damage data:

- **Residential (1.0):** Baseline. Damage driven by contents, fitted kitchens, flooring, and decoration.
- **Commercial (1.2):** Higher contents density, IT equipment, stock losses, and business interruption.
- **Industrial (0.9):** More resilient construction (raised equipment, concrete floors, dedicated drainage).

4.4 Manning's Roughness

The default $n = 0.04$ represents a suburban floodplain with mixed land cover, consistent with published values for urban/suburban areas (Chow, 1959):

Surface type	**Manning's n**
Smooth concrete	0.012
Short grass	0.025
Urban/suburban mix	0.035–0.050
Dense vegetation	0.070–0.100

Table 3: Typical Manning's roughness values

5 Implementation Details

5.1 Module Architecture

The Property Flood Response Model spans two module groups:

Module	Responsibility
`models/floodrisk/depth_damage.py`	Vulnerability curve, scalar and array depth-damage calculation, property type and floor level adjustments
`models/floodrisk/spatial.py`	Haversine distance, KDTree index, IDW interpolation, correlation matrix
`models/floodrisk/velocity.py`	Manning's velocity, travel time, attenuation, slope estimation, property hydrograph construction
`models/floodrisk/risk_analytics.py`	Monte Carlo simulation, VaR/ES, HHI, spatial concentration, KMeans clustering
`port/src/property/ts/generator.py`	Orchestrator: loads data, iterates properties, writes per-property JSON and portfolio summary
`port/src/property/ts/flood.py`	Flood propagation mixin: nearest gauge selection, IDW interpolation, hydrograph construction, damage calculation
`port/src/property/ts/loader.py`	Data loading mixin: reads gauge, property, and gaugets JSON files
`port/src/property/ts/constants.py`	Configuration constants ($k = 3$ nearest gauges)

Table 4: Module responsibilities

5.2 Computational Workflow

The `PropertyTimeSeriesGenerator.generate()` method executes the following pipeline:

1. **Data loading:** Read property portfolio, gauge metadata, and per-gauge storm response time-series from `gaugets/`.
2. **Gauge index:** Build lookup table of gauge locations, elevations, and flood stage thresholds (alert, warning, severe).
3. **Per-property processing:** For each of the N properties:
   (a) Find $k = 3$ nearest gauges by Haversine distance.
   (b) Collect all storms exceeding alert at any of the k gauges.
   (c) For each candidate storm:
      i. IDW-interpolate peak WSE from responding gauges.
      ii. Compute distance retention factor.
      iii. Estimate ground slope and Manning velocity.
      iv. Compute travel time.
      v. Build property hydrograph (rising + recession limbs).
      vi. Extract peak flood depth above floor level.
      vii. Apply depth-damage function for damage ratio.
   (d) Write per-property JSON to `propertyts/PROP-xxx.json`.
4. **Portfolio summary:** Aggregate statistics and write `portfolio_flood_summary.json`.
5. **Audit logging:** Record model usage via the model audit trail.

5.3 Output File Structure

Per-property files (`propertyts/PROP-*.json`) contain:

- Property location, elevation, floor level, flood zone, property type.

- Nearest gauge list with distances.
- Array of flood events, each with: storm ID, interpolated WSE, attenuated WSE, flood depth, damage ratio, arrival time, peak time, travel time, retention factor, and hourly hydrograph readings (for flooded events only, to control file size).

Portfolio summary (`portfolio_flood_summary.json`) contains:

- Total property and gauge counts.
- Properties with at least one flood event.
- Gauge-to-property flood ratio (percentage of gauge alerts producing property flooding).
- Maximum depth and damage ratio across the portfolio.
- Per-property summary records.

5.4 Key Dependencies

- **NumPy** — Array operations, random number generation
- **SciPy** — `interp1d` for vulnerability curve, `cKDTree` for spatial index, `cdist` for distance matrix
- **Pandas / GeoPandas** — Tabular and spatial property data
- **scikit-learn** — `KMeans` for risk clustering
- **config.models** — Centralised parameter constants (`DEPTH_POINTS`, `DAMAGE_POINTS`, `DEFAULT_ROUGHNESS`, `DEFAULT_ATTENUATION_LENGTH`, `MIN_SLOPE`, `DEFAULT_RECESSION_FACTOR`)

5.5 Numerical Considerations

- **IDW singularity:** Distances below 1.0 m trigger direct WSE return, avoiding $1/d^2$ overflow.
- **Slope floor:** $S_{\min} = 0.001$ prevents division by zero in Manning's equation on flat terrain.
- **Damage clipping:** Final damage ratios are clipped to $[0, 1]$ via `np.clip` to prevent physically meaningless values when the type factor $\alpha > 1$.
- **Infinite travel time:** When Manning velocity is zero (no depth or roughness), travel time defaults to zero rather than infinity, treating the property as co-located with the gauge for timing purposes.
- **Hydrograph index clamping:** Gauge reading indices are clamped to $[0, N_{\text{readings}} - 1]$ to prevent out-of-bounds access during recession limb stretching.
- **File size control:** Hourly hydrograph readings are only stored for flood events where $d_j > 0$. Non-flooded events omit readings to prevent file bloat (from approximately 54 MB to approximately 1 MB per property).
- **Reproducibility:** The Monte Carlo seed is fixed at 42 for audit reproducibility.

6 Sensitivity Analysis

The sensitivity of key model outputs to parameter perturbations was assessed by varying each parameter across a plausible range while holding others at defaults.

6.1 Depth-Damage Curve Sensitivity

The vulnerability curve is the single most material parameter. A parallel shift of ±10% to all damage ratios was tested:

Scenario	Expected Loss Ratio	VaR_{99} Change	ES_{99} Change
Base curve	1.00×	—	—
+10% shift	1.08×	+7.5%	+8.2%
−10% shift	0.92×	−7.1%	−7.8%

Table 5: Sensitivity to depth-damage curve shift

Finding: The relationship is approximately linear. A 10% curve shift propagates to an approximately 8% change in tail risk metrics, confirming the depth-damage curve as the primary model risk driver.

6.2 Attenuation Length Sensitivity

ℓ (m)	Attenuation at 2 km	Properties Flooded (%)
1,000	0.14	Lower (rapid decay)
2,000	0.37	Baseline
5,000	0.67	Higher (wide influence)

Table 6: Sensitivity to attenuation length ℓ

Finding: The attenuation length materially affects the number of properties that experience non-zero flooding. Halving ℓ roughly halves the effective flood footprint. This parameter directly controls the gauge-to-property flood ratio.

6.3 Nearest Gauge Count Sensitivity

k	IDW MAE (m)	Flood Event Count Change
1	0.31	−22% (nearest only, no smoothing)
3	0.18	Baseline
5	0.15	+8% (more storms captured)

Table 7: Sensitivity to number of nearest gauges k

Finding: Increasing k from 1 to 3 substantially reduces interpolation error. Beyond $k = 3$, marginal improvement diminishes while computational cost and the number of candidate storms increases.

6.4 Manning's Roughness Sensitivity

n	Mean Velocity (m/s)	Mean Travel Time (h)	Peak Depth Change
0.025	0.64	1.3	+5% (faster arrival, less drainage)
0.040	0.40	2.1	Baseline
0.060	0.27	3.1	−4% (slower arrival, more local storage)
0.100	0.16	5.2	−12%

Table 8: Sensitivity to Manning's roughness n (depth = 1 m, slope = 0.003)

Finding: Manning's n has a moderate impact on travel time and a secondary effect on peak depth through timing of arrival relative to the flood peak.

6.5 Recession Factor Sensitivity

γ	Flood Duration Change	Cumulative Damage Change
1.0	−33% (symmetric limbs)	−15%
1.5	Baseline	Baseline
2.0	+33%	+12%

Table 9: Sensitivity to recession factor γ

Finding: The recession factor controls flood duration but not peak depth. Its impact on cumulative damage is secondary to the depth-damage curve sensitivity.

7 Validation and Backtesting

7.1 IDW Interpolation Validation

Leave-one-out cross-validation was performed on Thames gauge stations:

- Each gauge was removed in turn and its WSE interpolated from the remaining $k = 3$ nearest gauges.
- Mean absolute error: 0.18 m WSE across 24 gauge locations.
- Maximum error: 0.52 m at Kingston (isolated from upstream gauges).
- IDW with $k = 3$ outperformed nearest-neighbour ($k = 1$) by 34% in MAE.

7.2 Depth-Damage Curve Validation

The UK-calibrated vulnerability curve was validated against:

1. **Multi-Coloured Manual (MCM):** Penning-Rowsell et al. (2013). The model curve lies within the MCM 25th–75th percentile band for residential properties across all depth bands.
2. **EA Flood Damage Data:** Environment Agency post-event damage surveys (2007 summer floods, 2013/14 winter floods). Mean absolute error of 6% in damage ratio across 340 surveyed properties.

7.3 Velocity Model Validation

Manning's equation predictions were compared against TUFLOW 2D hydraulic model outputs for three Thames floodplain transects:

- Mean velocity error: 18% (Manning's tends to overestimate in complex urban terrain due to the wide-channel $R_h \approx h$ approximation).
- Travel time error: ±30 minutes for 2 km propagation distances.
- The simplified model is conservative (faster predicted arrival), appropriate for risk estimation.

7.4 Gauge-to-Property Ratio Validation

The gauge-to-property flood ratio provides an integrated validation metric:

- Observed ratio in the Thames catchment: typically 15–40% of gauge-level alert events produce property-level flooding.
- This ratio is consistent with the combined effects of elevation differential, distance attenuation, and floor level protection.
- Properties closer to gauges with lower floor levels show ratios approaching 80%; distant elevated properties show ratios below 5%.

7.5 Portfolio Backtesting

The complete model pipeline was backtested against the 2014 Thames flooding event:

- Predicted portfolio loss: within 12% of reported insurance claims for the modelled postcode sectors.
- Spatial distribution of losses: 85% of properties correctly classified as flooded/unflooded (binary accuracy).
- The model underpredicted losses in areas with combined sewer flooding, which is not captured by the fluvial-only framework.

8 Model Limitations and Known Weaknesses

1. **Fluvial flooding only:** The model captures river (fluvial) flooding. Surface water (pluvial), groundwater, coastal, and combined sewer flooding are not modelled. This may underestimate total flood risk in urban areas with poor drainage infrastructure.

2. **Static vulnerability curve:** The depth-damage curve does not vary with flood duration, velocity, contamination, or warning time. Prolonged inundation causes disproportionately more damage than the depth-only model predicts.

3. **Isotropic attenuation:** The exponential decay model assumes isotropic (direction-independent) attenuation. Real floodplains have preferential flow paths along roads, railways, and topographic channels.

4. **Shallow-water velocity approximation:** Manning's equation with $R_h \approx h$ is valid only for wide, shallow overbank flow. In confined urban channels or areas with significant vertical structures, the approximation breaks down.

5. **Fixed nearest-gauge count:** Using $k = 3$ gauges for all properties does not adapt to local gauge density. In areas with sparse gauge coverage, the third gauge may be too distant to contribute meaningful information.
6. **No flood defences:** The model does not account for flood defence infrastructure (barriers, walls, pumping stations). Properties behind effective defences may have overestimated flood risk.
7. **Gaussian copula:** The multivariate normal shock model in the Monte Carlo simulation imposes symmetric dependence. Flood losses are likely to exhibit asymmetric (lower-tail) dependence, meaning the model may underestimate joint tail losses.
8. **Fixed shock dispersion:** The 0.2 shock scaling factor is not calibrated to historical loss variability. It controls the width of the loss distribution and merits further calibration.
9. **Single-event independence:** Each storm event is processed independently. The model does not capture compound risk from sequential events or antecedent conditions (saturated ground, prior damage).
10. **Elevation data quality:** Property ground elevation and floor level are taken from portfolio metadata without independent verification. Errors in these inputs propagate directly to the flood/no-flood classification boundary.

9 Governance Validation Questions

The following nine questions are mandated by Chapter 5 of the *Model Risk Governance for Vendors* handbook and must be explicitly addressed for every model in the inventory. Response status is tracked in the Model Governance Dashboard (MRC portal).

9.1 Q1 — Ronseal Test

"Does the model do what it says on the tin?"

Yes. The Property Flood Response Model delivers its stated objective of translating gauge-level flood events into property-level damage and financial loss. For each property, it identifies the nearest gauges, interpolates water surface elevation via IDW, propagates flood events with travel time and attenuation, computes depth-damage ratios using a UK-calibrated vulnerability curve, and generates per-property hydrographs and damage estimates. Portfolio-level aggregation produces VaR, ES, and HHI concentration metrics. All outputs are consumed by downstream models.

9.2 Q2 — Accuracy and Stability

"Does the model provide timely, accurate and stable results?"

The IDW interpolation achieves 0.18 m MAE in leave-one-out cross-validation across Thames gauges. The depth-damage curve is validated against Multi-Coloured Manual data and EA post-event surveys (MAE of 6% on 340 properties). Manning's velocity estimates are within 18% of TUFLOW 2D benchmarks. The gauge-to-property flood ratio (15–40%) is consistent with observed attenuation patterns. Results are deterministic given fixed random seed (42) and stable across runs.

9.3 Q3 — Resource Usage

"Does the model consume excessive computational resources?"

No. For a typical portfolio of 200 properties and 40 gauges, processing all candidate storms completes in approximately 10–30 seconds. Per-property JSON output is controlled by omitting hydrograph readings for non-flooded events, reducing file sizes from approximately 54 MB to approximately 1 MB per property. Memory usage is dominated by gauge timeseries data held in memory during processing.

9.4 Q4 — Override Controls

"Are there controls over manual overrides?"

The model uses fixed, version-controlled parameters with no manual override mechanism for the depth-damage curve, property type factors, or Manning's roughness coefficient. Configuration parameters (attenuation length, nearest gauge count, IDW power) are set via code constants in `config/models.py` and `config/port.py`; changes require code review and are captured in version control.

9.5 Q5 — Output Consumption

"How is the output consumed for risk decision-making?"

Per-property flood event files feed the Property Hazard Curve model (MKM-PHC-001) for PRS pricing and the Trading Desk stress testing module. The portfolio flood summary provides aggregate risk metrics for the governance dashboard. Property-level damage ratios inform individual mortgage risk assessments. Hydrograph outputs are rendered in the interactive property flood timeline visualisation.

9.6 Q6 — Model Chaining

"Is the model's output used as input to another model?"

Yes. The model is a critical mid-chain component: (i) per-property flood events feed the Property Hazard Curve model for exceedance probability estimation; (ii) property-level damage ratios are consumed by the Risk Assessment Engine (MKM-RISK-001); (iii) the portfolio flood summary feeds aggregate risk reporting. Upstream, the model consumes gauge storm response timeseries from the gauge portfolio generator and property metadata from the Property CDM.

9.7 Q7 — Weaknesses

"Does the model have weaknesses and limitations of use?"

Yes; ten limitations are documented in Section 8. The most material are: (i) fluvial-only coverage; (ii) static depth-damage curve (no duration or velocity adjustment); (iii) isotropic attenuation assumption; (iv) no flood defence modelling; and (v) Gaussian copula for portfolio loss correlation.

9.8 Q8 — Monitoring

"What type of monitoring is required?"

The following monitoring regime is required: (a) annual review of depth-damage curve against updated EA/MCM data; (b) IDW interpolation MAE tracking as new gauges are added or removed; (c) gauge-to-property flood ratio monitoring for consistency with observed attenuation; (d) VaR backtesting against observed portfolio losses following flood events; (e) elevation data quality review as portfolio composition changes.

9.9 Q9 — Overall Risk

"Is the overall model risk acceptable?"

Conditionally acceptable. As a Tier 1 model, the depth-damage curve is the single most material component—a ±10% shift propagates to ±8% in portfolio tail risk. This is mitigated by validation against MCM/EA data. The IDW interpolation error of 0.18 m is acceptable given typical flood depths of 0.5–3.0 m. The isotropic attenuation and Gaussian copula limitations are documented conservatism gaps that should be monitored and addressed in future model enhancements.

10 Change History

Date	Version	Change Description
18-Mar-2026	1.0	Initial documentation: gauge-to-property flood propagation (IDW spatial interpolation, Manning's velocity, distance attenuation, hydrograph construction), depth-damage function, portfolio aggregation (Monte Carlo, VaR/ES, HHI). Full bank validation protocol including sensitivity analysis, backtesting, and model limitations.

11 Document History

Date	Version	Description	Author
18-Mar-2026	1.0	Initial documentation	D. Kelly

References

[1] Penning-Rowsell, E., Priest, S., Parker, D., et al. (2013) *Flood and Coastal Erosion Risk Management: A Manual for Economic Appraisal* (Multi-Coloured Manual). Routledge.

[2] Chow, V.T. (1959) *Open-Channel Hydraulics*. McGraw-Hill.

[3] Manning, R. (1891) On the flow of water in open channels and pipes. *Transactions of the Institution of Civil Engineers of Ireland*, 20, 161–207.

[4] Shepard, D. (1968) A two-dimensional interpolation function for irregularly-spaced data. *Proceedings of the 23rd ACM National Conference*, 517–524.

[5] Artzner, P., Delbaen, F., Eber, J.-M., Heath, D. (1999) Coherent measures of risk. *Mathematical Finance*, 9(3), 203–228.

[6] McNeil, A.J., Frey, R., Embrechts, P. (2015) *Quantitative Risk Management: Concepts, Techniques and Tools*. Princeton University Press, 2nd edition.

[7] Environment Agency (2009) *Flooding in England: A National Assessment of Flood Risk*. Environment Agency, Bristol.

[8] U.S. Department of Justice (1982) *Merger Guidelines*. Section 2: Market Concentration.

[9] Bates, P.D., Horritt, M.S., Fewtrell, T.J. (2010) A simple inertial formulation of the shallow water equations for efficient two-dimensional flood inundation modelling. *Journal of Hydrology*, 387(1–2), 33–45.