Storm Sequence Generator

Multi-Storm Sequence Model v2.0 (MKM-SS-001)

MKM Research Labs

Version 2.0 — 11-March-2026

David K Kelly

Legal Notice

Contents

1 Executive Summary

This document presents the mathematical framework, implementation, and validation of the Storm Sequence Generator (MKM-SS-001), a significant upgrade from single-event storm simulation to **multi-storm sequence modelling** as the primary method of synthetic event generation within the Physical Risk platform.

Version 2.0 replaces the single-storm paradigm (one intensity sample → one water-level response) with the *insurance hours clause* model: a group of 1–5 related precipitation events, collectively referred to as a *StormSequence*, all fitting within a **168-hour (7-day) event window**. All precipitation is required to end by hour 156, leaving a minimum 12-hour drainage observation window. This design directly mirrors the *loss occurrence period* clause used in catastrophe reinsurance and parameterised cat-bond structures.

Four sequence types are defined, ordered by inter-storm coupling:

1. **Isolated** — a single precipitation event (1 storm).
2. **Doublet** — two events with a medium recovery gap (24–72 h).
3. **Cluster** — 3–4 events with medium gaps; typical of active frontal passages.
4. **Persistent** — 4–5 events with short gaps (6–36 h); modelling sustained blocking-pattern precipitation.

The probability of a multi-storm sequence increases with intensity category: 5% (minimal) to 85% (catastrophic), reflecting the meteorological tendency for the most severe atmospheric systems to produce extended sequences of precipitation events.

Key design principle: All sequences are guaranteed to fit within the event window without rejection sampling or fallback logic. Storm durations are budgeted as $d_{\max} = 156/n$ per storm, and any sampled gaps that would overflow the remaining window are scaled proportionally. This guarantees correctness in $O(1)$ and produces zero "could not fit" warnings at any scale.

Downstream impact: The multi-storm framework trains a significantly more capable flood probability classifier (MKM-FC-001 v2.0). With 10,000 sequences providing richer multi-peak hydrograph shapes and **ten times the training data** (1,680,000 vectors per gauge vs. 168,000 in v1.x), the stressm classifier achieves a **mean AUC-ROC of ≈0.993** across all Thames gauges, a further 4% improvement over the already-exceptional v1.3 result.

2 Model Purpose and Scope

2.1 Purpose

The Storm Sequence Generator serves three principal purposes within the Physical Risk platform:

1. **Synthetic sequence generation.** Produce a catalogue of 10,000 StormSequence objects covering the full intensity spectrum (moderate through catastrophic), each containing 1–5 temporally structured precipitation events.
2. **Hydrograph composition.** Provide temporal structure (storm start times, durations, inter-storm gaps, peak positions) so the Storm-Gauge Forward Model can compose multi-peak hydrographs showing catchment recovery between events.
3. **Classifier training data.** Feed the stressm flood probability classifier (MKM-FC-001 v2.0)

with training sequences that include antecedent-wetness dynamics: a gauge that partially recovered from Storm 1 is at elevated risk from Storm 2.

2.2 Scope

- **In scope:** Sequence type selection; storm count and duration sampling; inter-storm gap sampling and budget-scaling; intensity correlation within sequences; precipitation calculation; event window validation; batch generation across all intensity categories.
- **Out of scope:** Spatial footprint and track geometry (handled by MKM-SG-001); GEV hazard curve fitting (MKM-GH-001); PRS pricing (MKM-PR-001); climate change projections; tidal interactions.
- **Upstream dependencies:** Intensity category weights and base intensity parameters (configured in `batch_generator.py`).
- **Downstream consumers:** stressm flood classifier training pipeline (`src/port/src/stressm.py`); portfolio stress test backend (`src/routes/trading/port_stress.py`).

2.3 Model Chain Position

The sequence generator sits between intensity parameterisation and gauge-level response:

Intensity Params $\xrightarrow{\text{batch\_generator}}$ **StormSequence (MKM-SS-001)** $\xrightarrow{\text{stressm.py}}$ Multi-peak hydrograph $\rightarrow$ Classifi

2.4 Regulatory Context

This model is classified as a *Tier 2* component. It does not directly produce a regulated output but is the primary input to the stressm flood classifier, which feeds the Portfolio Stress P&L calculation displayed to trading desk risk managers.

3 Mathematical Framework

3.1 Event Window and Timing Constraints

The insurance *hours clause* defines a 168-hour (7-day) event window W:

$$W = 168\,\text{h} \tag{1}$$

All precipitation within a sequence must end by hour W_P:

$$W_P = W - W_{\text{drain}} = 168 - 12 = 156\,\text{h} \tag{2}$$

where $W_{\text{drain}} = 12\,\text{h}$ is the minimum drainage observation window required by the contract specification. The drainage window for a specific sequence is the time from the end of the last storm to the end of the event window:

$$W_{\text{drain}}^{(\text{seq})} = W - (t_{\text{start},n} + d_n) \tag{3}$$

where $t_{\text{start},n}$ and d_n are the start time and duration of the last (index n) storm. Validity requires $W_{\text{drain}}^{(\text{seq})} \geq W_{\text{drain}} = 12\,\text{h}$.

3.2 Sequence Type Taxonomy

Four sequence types are defined in Table 1. The storm count for `cluster` and `persistent` types is drawn uniformly from their ranges at generation time.

Table 1: Sequence type definitions.

Type	Storm Count	Gap Type	Physical Analogue
`isolated`	1	—	Single frontal system
`doublet`	2	MEDIUM	Atlantic low followed by secondary low
`cluster`	3–4	MEDIUM	Active frontal passage (2–4 days)
`persistent`	4–5	SHORT	Blocking pattern; near-continuous rain

3.3 Sequence Type Selection

For each event, the generator first decides whether to produce a multi-storm sequence at all, with probability depending on the intensity category:

$$P(\text{sequence}) = \begin{cases} 0.05 & \text{minimal} \\ 0.15 & \text{baseline} \\ 0.30 & \text{moderate} \\ 0.50 & \text{severe} \\ 0.70 & \text{extreme} \\ 0.85 & \text{catastrophic} \end{cases} \tag{4}$$

If a sequence is generated, the type is drawn from a weighted categorical distribution with weights that shift toward more coupled types as intensity increases (Table 2).

Table 2: Sequence type weights [$P(\text{doublet})$, $P(\text{cluster})$, $P(\text{persistent})$] by intensity category. Isolated events are a separate Bernoulli trial (Equation 4).

Category	Doublet	Cluster	Persistent
minimal / baseline	0.60	0.30	0.10
moderate	0.50	0.35	0.15
severe	0.30	0.50	0.20
extreme	0.20	0.45	0.35
catastrophic	0.15	0.45	0.40

3.4 Storm Duration Sampling

Each storm's duration is sampled from a triangular distribution with parameters (a, c, b) = (min, mode, max) determined by the intensity category (Table 3).

$$D \sim \text{Triangular}(a_{\text{cat}},\ c_{\text{cat}},\ b_{\text{cat}}) \tag{5}$$

Table 3: Triangular duration parameters (hours) by intensity category.

Category	**Min**	**Mode**	**Max**	**Mode duration**
minimal	2	6	12	6 h
baseline	6	12	24	12 h (half-day)
moderate	12	30	60	30 h (1.25 days)
severe	24	48	96	48 h (2 days)
extreme	36	72	144	72 h (3 days)
catastrophic	48	120	240	120 h (5 days)

3.4.1 Budget Cap

To guarantee that all storms in a sequence fit within $W_P = 156\,\text{h}$, each sampled duration is capped to a per-storm budget:

$$d_i = \min\left(D_i,\ \frac{W_P}{n}\right), \quad i = 1, \ldots, n \tag{6}$$

where n is the number of storms in the sequence. This is the key invariant of v2.0: **storm durations are bounded before gap sampling**, so the remaining budget for gaps is always non-negative.
The per-storm cap as a function of sequence size:

- Isolated ($n = 1$): cap = 156 h — effectively unconstrained for all categories.
- Doublet ($n = 2$): cap = 78 h.
- Cluster ($n = 3$): cap = 52 h; ($n = 4$): cap = 39 h.
- Persistent ($n = 4$): cap = 39 h; ($n = 5$): cap = 31.2 h.

Note that for catastrophic sequences (mode = 120 h, max = 240 h), the cap is active for all cluster/persistent configurations, clamping what would otherwise be multi-day storms into the available budget. This is physically appropriate: extreme storms in rapid succession are each necessarily shorter-lived.

3.5 Inter-Storm Gap Sampling

Gaps are sampled from triangular distributions whose type depends on the sequence type (Table 4). SHORT gaps model active blocking patterns; MEDIUM gaps model partial catchment recovery between frontal passages.

Table 4: Inter-storm gap parameters (hours) by gap type.

Type	**Used by**	**Min**	**Mode**	**Max**	**Physical interpretation**
SHORT	persistent	6	18	36	Minimal drainage; high compounding
MEDIUM	doublet, cluster	24	48	72	Partial drainage (30–60% recovery)

3.5.1 Gap Budget Scaling

Let $R = W_P - \sum_{i=1}^{n} d_i$ denote the remaining time budget after placing all n storms, and let $G_1, \ldots, G_{n-1}$ be raw gap samples from Equation (5). If the total raw gap exceeds the budget:

$$\sum_{j=1}^{n-1} G_j > R \tag{7}$$

then all gaps are scaled proportionally:

$$g_j = G_j \cdot \frac{R}{\sum_{k=1}^{n-1} G_k}, \quad j = 1, \ldots, n-1 \tag{8}$$

This preserves the relative spacing between storms (doublet gaps remain larger than persistent gaps) while guaranteeing the timing constraint. If $\sum G_j \leq R$, the raw samples are used unmodified. The combined guarantee from Equations (6) and (8) ensures that $\sum d_i + \sum g_j \leq W_P$ holds in all cases.

3.6 Storm Placement

Storm start times are computed sequentially from the placed durations and gaps:

$$t_{\text{start},1} = 0 \tag{9}$$

$$t_{\text{start},i} = t_{\text{start},i-1} + d_{i-1} + g_{i-1}, \quad i = 2, \ldots, n \tag{10}$$

The last storm ends at $t_{\text{start},n} + d_n \leq W_P$, verified by the post-generation validator (Section 5.4).

3.7 Intensity Sampling Within Sequences

The base intensity F_0 for the sequence is drawn from a Normal distribution per intensity category (Table 5), floored at 0.1:

$$F_0 \sim \max\big(0.1,\ \mathcal{N}(\mu_{\text{cat}},\ \sigma^2_{\text{cat}})\big) \tag{11}$$

Table 5: Base intensity distribution parameters by category.

Category	μ (mean)	σ (std)
minimal	0.30	0.10
baseline	0.60	0.15
moderate	1.00	0.20
severe	1.80	0.30
extreme	3.00	0.50
catastrophic	5.00	0.80

The individual storm intensities are drawn as correlated perturbations around F_0 using the following asymmetric scheme. Let $\delta \in [-0.2, +0.2]$ be the variation range ($\pm 20\%$). For the first storm ($i = 0$):

$$F_0^{(i=0)} = \max(0.1,\ F_0 \cdot (1 + \delta_0))\,, \quad \delta_0 \sim \text{Uniform}(-0.2,\ 0) \tag{12}$$

The first storm is always drawn with a downward variation (0 to -20%) from the base, reflecting that the triggering event may be somewhat weaker than the sequence peak.

For subsequent storms ($i \geq 1$):

$$F_0^{(i)} = \max(0.1,\ F_0 \cdot (1 + \delta_i))\,, \quad \delta_i \sim \begin{cases} \text{Uniform}(0,\ +0.2) & \text{with probability 0.70} \\ \text{Uniform}(-0.2,\ 0) & \text{with probability 0.30} \end{cases} \tag{13}$$

The 70% probability of maintained or increasing intensity models the antecedent soil moisture mechanism: catchments partially saturated by Storm $i-1$ respond more severely to Storm i even if raw precipitation is similar.

3.8 Precipitation Calculation

The precipitation (mm) for storm i is:

$$P_i = P_{\text{base}} \cdot F_0^{(i)} \tag{14}$$

where P_{base} is a catchment-specific baseline precipitation. For the Thames catchment, $P_{\text{base}} = 35\,\text{mm}$. The peak position within the storm (dimensionless, $[0, 1]$) is drawn from a symmetric Beta distribution:

$$\text{peak\_position} \sim \text{Beta}(2, 2) \tag{15}$$

This produces a bell-shaped distribution centred at 0.5, favouring mid-storm peaks while allowing early and late peaks.

3.9 Aggregate Sequence Statistics

The following aggregate statistics are computed for each StormSequence:

$$P_{\text{total}} = \sum_{i=1}^{n} P_i \tag{16}$$

$$F_{\max} = \max_i F_0^{(i)} \tag{17}$$

$$\bar{F} = \frac{1}{n} \sum_{i=1}^{n} F_0^{(i)} \tag{18}$$

$$F_{\text{cum}} = \sum_{i=1}^{n} F_0^{(i)} \tag{19}$$

$$D_{\text{total}} = \sum_{i=1}^{n} d_i + \sum_{j=1}^{n-1} g_j \tag{20}$$

F_{cum} is particularly important for portfolio stress testing: it captures the compounding effect of multiple storms rather than just the single peak.

4 Input Parameters and Calibration

4.1 Batch Generation Parameters

The primary entry point is `generate_event_set()` in `batch_generator.py`. The default intensity weights follow spec Section 4.3:

Table 6: Default intensity category weights for batch generation.

Category	Weight
moderate	0.40 (40%)
severe	0.35 (35%)
extreme	0.20 (20%)
catastrophic	0.05 (5%)

Default batch size: 10,000 sequences. The weights are normalised internally so that custom weights need not sum to exactly 1.0.

4.2 Catchment-Specific Parameters

Table 7: Catchment base precipitation calibration.

Catchment	P_{base} (mm)
thames	35.0
rhine	40.0
danube	45.0
mississippi	55.0

4.3 Fixed Constants

Table 8: Fixed model constants.

Constant	Symbol	Value	Description
Event window	W	168 h	Hours clause event window
Minimum drainage	W_{drain}	12 h	Post-precipitation observation
Max precipitation window	W_P	156 h	$W - W_{\text{drain}}$
Intensity variation	δ	±20%	Per-storm variation around base
Correlation prob.	p_{corr}	0.70	Prob. of subsequent storm $\geq$ first
Beta shape (peak)	(α, β)	(2, 2)	Peak position distribution
Min intensity floor		0.1	Floor on intensity factor

5 Implementation Details

5.1 Software Architecture

The model is implemented in the `src/port/src/storm_multi/` package:

- `core/data_structures.py` — `SequenceStorm`, `StormSequence` dataclasses; `SequenceType`, `GapType` enumerations; `make_storm_id()`, `make_sequence_id()` UUID generators.
- `generators/sequence_generator.py` — `SequenceGenerator` class; `generate()`, `_build_sequence()` methods.
- `generators/batch_generator.py` — `generate_event_set()` function; default intensity weights; base intensity parameters per category.
- `generators/duration_sampler.py` — `sample_duration()`, `DURATION_PARAMS` table.
- `generators/gap_sampler.py` — `sample_gap()`, `GAP_PARAMS` table.
- `generators/intensity_sampler.py` — `should_generate_sequence()`, `sample_sequence_type()`, `get_storm_count()`, `sample_storm_intensity()`.
- `utils/validation.py` — `validate_sequence()`, `validate_event_set()`, `fits_in_window()`.
- `utils/serialization.py` — JSON serialisation/deserialisation.

5.2 Generation Flow

The primary entry point `generate_event_set(count, catchment_id, seed)` executes the following for each of the N sequences:

1. Draw intensity category from the weighted categorical distribution (Table 6).
2. Draw base intensity $F_0 \sim \mathcal{N}(\mu_{\text{cat}}, \sigma^2_{\text{cat}})$, floored at 0.1.
3. Call `SequenceGenerator.generate(category, F_0)`.
4. Within `generate()`: determine sequence type via `should_generate_sequence()` and `sample_sequence_type` determine storm count via `get_storm_count()`.
5. Call `_build_sequence(seq_id, seq_type, n, category, F_0, now)`:
   (a) Compute $d_{\max} = W_P/n$.
   (b) For each storm: sample duration $D_i \sim$ Triangular, cap to $d_{\max}$; sample intensity factor F_i.
   (c) Compute remaining budget $R = W_P - \sum d_i$.
   (d) Sample raw gaps $G_1, \ldots, G_{n-1}$; scale if $\sum G_j > R$.
   (e) Compute start times from Equations (9)–(10).
   (f) Assemble `SequenceStorm` objects; compute aggregates.
6. Append `StormSequence` to output list.

5.3 Reproducibility

The batch generator uses two `numpy.random.RandomState` instances initialised from a single integer seed:

- Outer RNG: category assignment draws and base intensity draws.
- `SequenceGenerator` RNG: internal duration, gap, and intensity draws.

Both are seeded from the same `seed` parameter, ensuring full reproducibility.

5.4 Validation

Post-generation validation is performed by `validate_event_set()` across the full batch. Checks include:

- Sequence type is one of {isolated, doublet, cluster, persistent}.

- `num_storms` equals `len(storms)`.
- Storm count is in $[1, 5]$.
- Gap count equals `num_storms` $- 1$.
- Temporal ordering: each storm starts after the previous storm ends.
- Last storm ends by $W_P = 156\,\text{h}$ (with 10^{-9} tolerance for floating-point accumulation).
- Drainage window $\geq 12\,\text{h}$.
- All intensity categories are valid.
- `total_duration_hours` is consistent with sum of durations and gaps (within $0.5\,\text{h}$ tolerance).

The budget-cap and gap-scaling guarantees of Sections 3.4.1 and 3.5.1 ensure that timing violations are structurally impossible. Validation nonetheless runs as a defence-in-depth check.

5.5 CLI Integration

The model is invoked via:

```
python app.py port --stressm [--num-storms N] [--gauge-id GID]
                   [--train-classifier]
```

Output files are written to `data/output/stressm/`:

- `doublet_sequences.json` — serialised sequence catalogue.
- `sequences_summary.json` — batch-level statistics.
- `gauge_summary.json` — per-gauge response summary.
- `GAUGE-{id}.joblib` — trained GBM classifiers (with `--train-classifier`).

6 Sensitivity Analysis

6.1 Sequence Type Distribution

The expected distribution of sequence types across a batch can be derived from the probability tables. For the default intensity weights (Table 6) and using the moderate-to-catastrophic range:

Table 9: Expected sequence type breakdown for the default batch (10,000 sequences). Figures are approximate, based on weighted-average sequence probabilities.

Type	Expected count
isolated	≈5,800
doublet	≈2,000
cluster	≈1,600
persistent	≈600

6.2 Duration Cap Activation Rate

The cap $d_{\max} = W_P/n$ is active whenever the sampled triangular duration exceeds it. For catastrophic events in a 5-storm persistent sequence: $d_{\max} = 156/5 = 31.2\,\text{h}$, against a triangular mode of $120\,\text{h}$. Approximately 99% of catastrophic durations would be capped in this configuration. For isolated catastrophic events, $d_{\max} = 156\,\text{h}$ and the cap is never active (triangular max $= 240\,\text{h}$, but isolated sequences have only 1 storm).

For the default batch, the cap activation rate across all configurations is approximately 35% of all storm-duration draws.

6.3 Sensitivity to Gap Scaling

Gap scaling (Equation 8) is triggered when raw gaps overflow the remaining budget. The remaining budget R is highest for isolated events (effectively $\leq 156\,\mathrm{h}$) and smallest for persistent catastrophic sequences. In practice, for moderate-to-severe persistent sequences, gap scaling is triggered in approximately 20% of sequences; for catastrophic persistent sequences, in approximately 60%.

When scaling is triggered, the characteristic gap ratios are preserved: a doublet gap is always larger than a persistent gap by the same factor. The only change is a uniform reduction in all gaps. This is physically consistent: during an active blocking pattern with limited available time, all inter-storm intervals are compressed.

7 Validation and Backtesting

7.1 Structural Validation

Running `validate_event_set()` on a batch of 10,000 sequences consistently returns zero invalid sequences. All timing constraints, gap counts, storm counts, and intensity categories are satisfied by construction.

7.2 Type Distribution Verification

Monte Carlo verification across 100 independent seeds of 10,000 sequences each confirms that the empirical type distribution is within 2% of the theoretical expected values (Table 9).

7.3 Downstream Classifier Performance

The most direct validation of the sequence generator's quality is the performance of the downstream stressm flood classifier (MKM-FC-001 v2.0) trained on its output. Results across all 52 Thames gauges:

Table 10: Stressm classifier AUC-ROC by model version (all Thames gauges).

Training data	Mean AUC-ROC	Training vectors/gauge
v1.x single-storm (1,000 storms $\times$ 168 h)	0.9525	168,000
v2.0 multi-storm (10,000 sequences $\times$ 168 h)	$\approx$0.993	1,680,000

The 4% AUC improvement reflects two compounding effects: (i) **ten times more training data**, reducing variance in GBM tree estimates; (ii) **richer hydrograph shapes** from multi-peak sequences, exposing the classifier to the antecedent-wetness dynamics that are most diagnostically important for compound flood events.

7.4 Gauge-Level Validation Example

A representative individual gauge result (`GAUGE-8f1d368d`):

- **AUC-ROC:** 0.9927
- **Accuracy:** 0.9776
- **Flood rate (training):** 8.4%
- **Dominant features:** `log_h_s` (≈70%), `log_t_end` (≈24%)

This pattern — high accuracy, physically meaningful feature dominance, realistic flood rate — is consistent across all gauges trained to date.

8 Model Limitations and Known Weaknesses

1. **Budget cap distorts duration distribution.** The cap $d_{\max} = W_P/n$ systematically truncates the upper tail of the triangular duration distribution, particularly for high-intensity sequences with many storms. This understates the duration of the most extreme multi-storm events.
2. **Gap scaling is proportional, not physical.** When gaps are scaled down to fit the budget, all gaps are reduced uniformly. In reality, some inter-storm intervals are physically determined by atmospheric wave speeds and synoptic-scale patterns that would produce minimum gap durations even under severe budget constraint.
3. **Independent storm placement.** Storm durations and peak positions within a sequence are sampled independently. Real compound events exhibit serial correlation: a very long first storm changes the soil moisture state in a way that physically shortens the effective duration of the second event. This is not modelled.
4. **No spatial coupling in multi-storm sequences.** Each storm in a sequence inherits the same spatial footprint model. In reality, successive storms in a blocking pattern may have different tracks, different footprint widths, and different centres of activity.
5. **Catchment base precipitation is fixed.** P_{base} is a constant per catchment and does not vary seasonally or with antecedent conditions. Winter events in saturated catchments would produce greater runoff from the same precipitation total.
6. **Stationarity.** The intensity category weights and gap type assignments are fixed. Climate change scenarios (increased persistence, higher intensities) are not yet parametrised within the sequence framework.

9 Governance Validation Questions

The following nine questions are mandated by Chapter 5 of the *Model Risk Governance for Vendors* handbook and must be explicitly addressed for every model in the inventory. Response status is tracked in the Model Governance Dashboard (MRC portal).

9.1 Q1 — Ronseal Test

"Does the model do what it says on the tin?"

Yes. The model generates catalogues of StormSequence objects respecting the 168-hour insurance hours clause, with the four documented sequence types (isolated, doublet, cluster, persistent), calibrated duration and gap distributions, and intensity correlation within sequences. All sequences are validated post-generation and consistently pass zero invalid sequences against a comprehensive set of timing and structural checks. The downstream stressm classifier trained on these sequences achieves mean AUC-ROC ≈0.993, confirming that the generated hydrograph shapes are physically discriminative.

9.2 Q2 — Accuracy and Stability

"Does the model provide timely, accurate and stable results?"

Generation of 10,000 sequences completes in under 60 seconds. The budget-cap and gap-scaling mechanisms guarantee structural correctness without rejection sampling or stochastic retries, making execution time fully deterministic. With a fixed seed, output is exactly reproducible. The type distribution is stable across seeds (within 2% of theoretical). The downstream AUC improvement from v1.x to v2.0 is consistent across all 52 gauges.

9.3 Q3 — Resource Usage

"Does the model consume excessive computational resources?"

No. Generation of 10,000 sequences uses NumPy triangular and Normal draws in a tight Python loop and completes in under 1 minute on a single core. Memory usage is dominated by the resulting list of dataclasses (≈50 MB for 10,000 sequences with full storm arrays). Serialisation to JSON adds approximately 10 seconds for the full catalogue. Classifier training on 1.68M vectors per gauge requires approximately 2–5 minutes per gauge using `sklearn.GradientBoostingClassifier`; the full 52-gauge portfolio trains in approximately 2 hours.

9.4 Q4 — Override Controls

"Are there controls over manual overrides?"

Intensity category weights, gap type mappings, sequence probability table, and type weight tables are all defined as named constants in their respective Python modules (`batch_generator.py`, `gap_sampler.py`, `intensity_sampler.py`). No mechanism exists for runtime override of these tables. Changes require code modification with version control tracking. The `force_sequence_type` parameter provides a deterministic override for testing purposes (e.g. generating a pure doublet batch), subject to normal version control review.

9.5 Q5 — Output Consumption

"How is the output consumed for risk decision-making?"

Generated sequences are consumed in two ways: (i) the stressm training pipeline (`app.py port --stressm --train-classifier`) uses sequences to synthesise multi-peak hydrographs and train per-gauge GBM flood classifiers; (ii) the Portfolio Stress tab (Tab 8) on the Trading Desk presents sequence-level scenarios to traders and risk managers, showing P(flood) per gauge, stressed P&L per trade, and portfolio-level stress P&L in GBP.

9.6 Q6 — Model Chaining

"Is the model's output used as input to another model?"

Yes. StormSequence objects are consumed by the stressm hydrograph synthesiser, which uses per-storm parameters (start time, duration, intensity factor, peak position) to compose a multi-peak time series at each gauge. This time series is the training input for MKM-FC-001 v2.0. At inference time, the same hydrograph composition pipeline is used for the portfolio stress scenario runner. The model has no upstream model dependencies beyond intensity category parameterisation, which is set by the analyst.

9.7 Q7 — Weaknesses

"Does the model have weaknesses and limitations of use?"

Six limitations are documented in Section 8. The most material are: (i) the budget cap systematically truncates the most extreme multi-storm durations; (ii) gap scaling is proportional rather than physically motivated; (iii) independent storm placement within sequences ignores serial correlations in soil moisture response. These limitations are consistent with the model's Tier 2 classification as a training-data generator; they do not impair the downstream classifier's ability to learn discriminative flood-probability signals from realistic multi-peak hydrographs.

9.8 Q8 — Monitoring

"What type of monitoring is required?"

Annual monitoring should include: (a) review of the sequence type distribution from the most recent 10,000-sequence batch against the theoretical expected values; (b) comparison of the stressm classifier AUC across gauges against the v2.0 baseline (≈0.993), flagging any gauge where AUC falls below 0.95; (c) review of the duration cap activation rate, with a threshold of >50% indicating that the category weights or duration parameters may need recalibration; (d) review of gap scaling frequency, with >70% scaling in any sequence type indicating that the base distributions are inconsistent with the event window constraint.

9.9 Q9 — Overall Risk

"Is the overall model risk acceptable?"

Acceptable. The model is Tier 2 and serves solely as a training data and scenario generator. The guaranteed-fit property eliminates the most significant operational risk of the v1.x generator (rejection sampling fallbacks corrupting the type distribution). The dramatic improvement in downstream classifier performance (AUC 0.9525 → ≈0.993) demonstrates that the richer training data directly translates to better risk discrimination. The principal residual risk—budget-cap distortion of extreme multi-storm durations—is conservative in direction (understating the most severe events) and is mitigated by the 5% catastrophic weight in the default batch.

10 Change History

- **v2.0 (11-Mar-2026):** Initial documentation for the multi-storm sequence generator. This is a complete replacement of the single-storm paradigm (MKM-SI-001) for stress testing purposes. Introduces the four-type sequence taxonomy, budget-based window fitting, and correlated per-storm intensity sampling.

11 Document History

Date	Version	Description	Author
11-Mar-2026	2.0	Initial documentation — multi-storm sequence generator (MKM-SS-001)	D. Kelly