

Storm Generator v2.0

Implementation Specification

Thames Valley Physical Risk Swaps - Enhanced Multi-Storm Sequences

Project: PRS Platform Storm Generator v2.0

Date: March 5, 2026

Prepared for: David Kelly

Scope: Phase 1 (Extended Durations) + Phase 2 (Storm Sequences with Drainage Windows) + Spatial Correlation

Version: 2.0-multi-storm-spatial

Executive Summary

This specification provides detailed implementation guidance for a focused storm generation system targeting **single-scenario, two-storm (doublet) sequences** with explicit inter-storm drainage periods and **spatially correlated gauge responses** to model compound flooding events in the Thames Valley.

Key Innovation

Moving from independent storm sampling to **temporally-correlated multi-storm sequences with spatial coherence** where:

- Each scenario contains **1-5 precipitation events** (isolated, doublet, cluster, or persistent patterns)
- Clear **inter-storm gaps** (6–144 hours depending on type) allow observable river level decline
- All precipitation stops **≥12 hours before hour 168** to show final drainage
- Antecedent soil moisture from Storm 1 amplifies Storm 2 response

- **Spatially correlated precipitation fields** across 40 Thames gauges capture realistic storm footprints
- Hydraulic and hydrologic models work together to capture compounding effects

Success Criteria

- Generate 10,000 realistic two-storm flood scenarios
 - Respect 168-hour insurance event windows
 - Show clear water level decline between precipitation pulses
 - Model soil moisture persistence and compounding effects
 - Generate spatially coherent precipitation across 40 Thames gauges
 - Maintain market-calibrated intensity distribution (20-60 bps fair spreads)
 - File sizes remain manageable (<50 MB for 10K events)
 - Performance: <2 minutes to generate 10K doublet scenarios with correlated gauge responses
-

1. Project Context and Constraints

1.1 Document Purpose

This is a **specification document**, not implementation code. It describes:

- What needs to be built
- How components should interact
- Data structures and schemas
- Algorithms and logic flows
- Integration points with existing code
- Testing and validation criteria
- Mathematical foundations for hydrologic-hydraulic coupling
- Spatial correlation methodology for gauge networks

1.2 Project Constraints

Scenario structure: Four sequence types

(isolated/doublet/cluster/persistent) with 1-5 storms per sequence

Event count: 10,000 doublet scenarios maximum

Event window: 168 hours (7 days) per insurance industry standard

Drainage requirement: Final 12 hours (156–168h) must be precipitation-free

Gauge network: 40 Thames gauges from Richmond to Tilbury (defined in catchments/thames.py)

Spatial coherence: Precipitation must be realistically correlated across gauge network

Climate scenarios: NOT in scope (present-day conditions only)

Backward compatibility: Must maintain ability to load existing storm files

1.3 Known Limitations

Validation data: Limited historical gauge data availability for doublet sequences

Antecedent conditions: Initial implementation uses simplified state model; full state-space model (Section 6) is future enhancement

Tidal influence: Current gauge levels may show tidal undulations; tide schedule injection capability to be added if needed

Spatial correlation: Initial implementation uses isotropic exponential kernel; anisotropic models for future enhancement

2. Core Architecture: Multi-Storm Sequences

2.1 Overview

Each "scenario" is a **StormSequence** object containing 1-5 **StormScenario** objects (storms), where multiple storms are separated by inter-storm gaps that allow partial drainage. The entire sequence fits within a 168-hour event window, with all precipitation ending by hour 156 to allow final drainage observation.

2.1.1 Sequence Type Definitions

Type	Storm Count	Total Duration	Inter-storm Gaps	Use Case
isolated	1	6-240 hours	N/A	Single storm
doublet	2	48-156 hours	12-72 hours	Two storms, short gap
cluster	3-4	96-156 hours	24-96 hours	Multiple storms
persistent	4-5	96-156 hours	12-72 hours	Persistent wet pattern

Table 1: Sequence type definitions

2.2 Data Structures

2.2.1 StormScenario (Individual Storm)

```
@dataclass
```

```
class StormScenario:
```

```
    """
```

```
A single precipitation event within a sequence.
```

```
    """
```

```
# Identity
```

```
storm_id: str # Format: "STORM-{uuid4}"
```

```
scenario_id: str # Parent sequence ID
```

```
# Temporal
```

```
start_time: str # ISO format, hours from sequence start
```

```
duration_hours: int # Storm duration
```

```
# Intensity
```

```
intensity_category: str # 'moderate', 'severe', 'extreme', 'catastrophic'
```

```
intensity_factor: float # Multiplier for base intensity
```

```
precipitation_mm: float # Total precipitation
```

```
# Metadata
```

```
generated_at: str
catchment_id: str
```

2.2.2 StormSequence (Multi-Storm Container)

```
@dataclass
```

```
class StormSequence:
```

```
    """
```

A sequence of 1-5 related storm events treated as one flood event.
Insurance industry 'hours clause' implementation.

```
    """
```

```
# Identity
```

```
sequence_id: str # Format: "SEQ-{uuid4}"
```

```
sequence_type: str # 'isolated', 'doublet', 'cluster', 'persistent'
```

```
# Temporal structure
```

```
sequence_start: str      # ISO format
```

```
sequence_end: str        # Must be  $\leq$  156 hours from start
```

```
total_duration_hours: int # Actual duration of precip + gaps
```

```
event_window_hours: int   # Always 168
```

```
drainage_window_hours: int # Final dry period ( $\geq$ 12 hours)
```

```
# Storm composition
```

```
storms: List[StormScenario] # 1-5 storms
```

```
num_storms: int            # 1-5
```

```
# Temporal gaps
```

```
inter_storm_gaps_hours: List[int] # Gaps between consecutive storms
```

```
# Aggregate statistics
```

```
total_precipitation_mm: float
```

```
max_intensity_mm: float
```

```
avg_intensity_mm: float
```

```
cumulative_intensity_factor: float
```

```
# Antecedent conditions (for gauge response)
```

```
antecedent_soil_moisture: Optional[str] = None # 'dry', 'normal', 'wet'
```

```
antecedent_groundwater: Optional[str] = None # 'low', 'normal', 'high'
```

```
# Spatial correlation metadata
spatial_correlation_enabled: bool = False
spatial_correlation_range_km: Optional[float] = None

# Metadata
generated_at: str
catchment_id: str
```

Validation rules:

- `num_storms` must equal `len(storms)`
- `storms` list: minimum 1, maximum 5
- `sequence_type` in {'isolated', 'doublet', 'cluster', 'persistent'}
- `len(inter_storm_gaps_hours)` must equal `num_storms - 1`
- `sequence_end` $\leq$ 156 hours (leaves $\geq$ 12h drainage)
- `total_duration_hours` = `sum(storm durations)` + `sum(inter_storm_gaps_hours)`
- Temporal ordering: `storms[i].start_time` < `storms[i+1].start_time`

5. Spatial Correlation Across Thames Gauges

5.1 Objective

Introduce a **gauge-level spatial correlation model** so that hourly precipitation for each StormSequence is simulated consistently across the 40 Thames gauges, rather than treating gauges independently. This ensures:

- Realistic spatial coherence of storm footprints
- Correlated flood peaks across the catchment
- Accurate portfolio loss aggregation for insurance applications
- Proper representation of widespread vs localized storm events

5.2 Gauge Set and Geometry

We use the 40 ordered gauge points defined in `catchments/thames.py`:

- `thames.GAUGE_POINTS`: list of 40 tuples (lat, lon, elevation_meters), ordered west–east from Richmond Lock to Purfleet
- `thames.GAUGE_NAMES`: corresponding location names for each gauge

Geographic extent:

- Western boundary: Richmond Lock (-0.3072° , 51.4573°)
- Eastern boundary: Purfleet (0.2855° , 51.4581°)
- Linear river distance: ~ 65 km
- Catchment width: ~ 5 - 15 km (tidal Thames corridor)

From these gauge coordinates we derive:

- A 40×40 **distance matrix** D in km between all gauge pairs (great-circle distance)
- An **index ordering** consistent across all storm and gauge response modules

5.3 Correlation Kernel

Let $i, j \in \{0, \dots, 39\}$ index gauges in `GAUGE_POINTS` order.

Step 1: Compute pairwise distances

Great-circle distance d_{ij} (km) between gauges i and j using Haversine formula.

Step 2: Define exponential correlation function

$$C_{ij} = (1 - \text{nugget}) \cdot \exp\left(-\frac{d_{ij}}{r_h}\right) + \text{nugget} \cdot \mathbb{1}_{i=j} \quad (1)$$

where:

- r_h = horizontal correlation range (km), base parameter
- **nugget** = small-scale noise parameter (0–0.1), represents measurement error and micro-scale variability

- $\mathbb{1}_{i=j}$ = indicator function (1 if $i = j$, 0 otherwise)

Step 3: Intensity-dependent correlation range (optional)

For more intense or longer-duration storms, increase effective correlation range to represent more spatially coherent weather systems:

$$r_h^* = r_h \cdot [1 + \rho_{\text{intensity}} \cdot (I_{\text{seq}} - 1)] \quad (2)$$

where:

- I_{seq} = cumulative_intensity_factor from the StormSequence (normalized around 1.0)
- $\rho_{\text{intensity}}$ = sensitivity parameter (0–1), tuning coefficient

Physical interpretation:

- Moderate localized storms: shorter correlation range (15-30 km)
- Severe widespread systems: longer correlation range (40-80 km)
- Extreme events: maximum spatial coherence across entire catchment

5.4 Default Parameter Values (Thames-Calibrated)

Based on Thames catchment geometry and typical UK precipitation patterns:

Parameter	Default Value	Rationale
base_range_km	40.0	Half of Thames corridor length
nugget	0.05	5% micro-scale variability
rho_intensity	0.4	40% range increase for extreme events
sigma_lognormal	0.4	40% spatial CV in precipitation

Table 2: Default spatial correlation parameters for Thames

Effective ranges by intensity:

- Moderate ($I_{\text{seq}} = 0.8$): $r_h^* = 40 \times (1 + 0.4 \times -0.2) = 36.8$ km

- Severe ($I_{\text{seq}} = 1.2$): $r_h^* = 40 \times (1 + 0.4 \times 0.2) = 43.2$ km
- Extreme ($I_{\text{seq}} = 2.0$): $r_h^* = 40 \times (1 + 0.4 \times 1.0) = 56.0$ km

5.5 Generation Algorithm

We generate a **space–time precipitation field** conditioned on the 1D temporal profile of each StormSequence.

High-level flow:

1. Build 40×40 correlation matrix C using intensity-conditioned range
2. Compute Cholesky decomposition L such that $C = LL^T$
3. For each hour $t = 0, \dots, 167$:
 - a. Get catchment-mean precipitation P_t from StormSequence
 - b. If $P_t = 0$, set all gauge precipitation to zero
 - c. Otherwise:
 - Sample independent standard normals $\epsilon \sim N(0, I_{40})$
 - Generate correlated field $z = L\epsilon$
 - Transform to lognormal multipliers
 $m_i = \exp(\sigma z_i - \sigma^2/2)$
 - Normalize multipliers to preserve catchment mean
 - Assign gauge precipitation $P_{i,t} = P_t \cdot m_i$
4. Return dictionary {gauge_name: precip[0:168]} for all 40 gauges

Mathematical detail:

Lognormal transformation: Ensures non-negative precipitation multipliers with mean = 1.

$$M_i = \exp(\sigma Z_i - \sigma^2/2) \quad (3)$$

where $Z_i \sim N(0, 1)$ are correlated Gaussian variables and σ controls spatial variability.

Normalization: Preserve catchment-mean precipitation.

$$\tilde{M}_i = \frac{M_i}{\sum_j w_j M_j} \cdot \sum_j w_j \quad (4)$$

where w_j are gauge weights (equal weights for Thames: $w_j = 1/40$).

Gauge precipitation:

$$P_{i,t} = P_t \cdot \tilde{M}_i \quad (5)$$

5.6 Performance Considerations (40 Gauges)

Computational complexity:

- Correlation matrix construction: $O(N^2) = O(1600)$ operations per sequence
- Cholesky decomposition: $O(N^3/3) = O(21,000)$ operations per sequence
- Per-hour sampling: $O(N^2) = O(1600)$ operations $\times$ 168 hours
- Total per sequence: $\sim 300,000$ floating-point operations

Performance optimizations:

1. **Pre-compute distance matrix:** Calculate once at module load, not per sequence
2. **Cache Cholesky factors:** For identical correlation parameters, reuse decomposition
3. **Vectorized operations:** Use NumPy broadcasting for hourly sampling loop
4. **Parallel sequence generation:** Spatial correlation is independent across sequences

Expected performance:

- Single sequence with spatial correlation: < 0.01 seconds
- 10,000 sequences: < 2 minutes (meets success criteria)
- Memory footprint: ~ 3 MB per sequence (40 gauges $\times$ 168 hours $\times$ 8 bytes)

5.7 Integration with Gauge Response Model

Update `compute_sequence_gauge_response` to accept spatially correlated precipitation:

Option A: Pre-computed precipitation series

```
def compute_sequence_gauge_response(sequence: StormSequence,
gauge_id: str,
```

```
gauge_params: dict,
spatial_field: Dict[str, np.ndarray]
) -> np.ndarray:
"""
```

Use spatially correlated precipitation at gauge_id to compute hourly river levels.

Args:

```
sequence: StormSequence object
gauge_id: Gauge identifier (from thames.GAUGE_NAMES)
gauge_params: Hydrologic/hydraulic parameters for this gauge
spatial_field: Pre-computed precipitation for all gauges
```

Returns:

```
Array of length 168 (hourly river levels)
```

"""

```
precip_series = spatial_field[gauge_id]
```

```
# Initialize hydrologic state
```

```
soil_moisture = initial_soil_moisture()
```

```
groundwater = initial_groundwater()
```

```
river_level = baseline_level(gauge_id)
```

```
levels = np.zeros(168)
```

```
for hour in range(168):
```

```
    # Update hydrologic state
```

```
    soil_moisture, groundwater, quickflow = update_hydrology(
        soil_moisture, groundwater, precip_series[hour], hour
    )
```

```
    # Update river level (hydraulic routing)
```

```
    river_level = update_hydraulics(
        river_level, quickflow, gauge_params
    )
```

```
    levels[hour] = river_level
```

return levels

Option B: On-the-fly generation

Spatial correlation module can be called internally by gauge response model if precipitation field not pre-computed.

5.8 Calibration and Validation

Calibration data sources:

1. Historical Thames events: 2013/14 winter floods, 2000 autumn floods, 2003 floods, Storm Henk 2024
2. DEFRA gauge networks: Hourly precipitation and river level data
3. Met Office rainfall radar: 1km gridded precipitation data for spatial patterns

Parameter estimation:

- **Correlation range (r_h):** Fit empirical semivariogram from radar rainfall data
- **Nugget:** Estimate from gauge network vs radar discrepancies
- **Intensity sensitivity ($\rho_{\text{intensity}}$):** Compare storm footprint sizes for different intensity levels
- **Spatial variability (σ):** Fit from gauge precipitation ratios during historical events

Validation checks:

1. **Pairwise correlation:** Compare synthetic vs observed correlation of hourly precipitation at gauge pairs
2. **Event-total correlation:** Validate spatial correlation of cumulative storm precipitation
3. **Spatial patterns:** Visual comparison of synthetic storm footprints vs historical radar imagery
4. **Peak timing:** Verify downstream hydrograph correlation structure (peak alignment along main stem)
5. **Portfolio metrics:** Check that spatial correlation impacts aggregate loss distributions as expected

Statistical tests:

- Kolmogorov-Smirnov test: Synthetic vs observed distance-correlation curves ($\alpha = 0.05$)
- Correlation matrix comparison: Frobenius norm between synthetic and historical correlation matrices
- Spatial moment matching: Mean, variance, and range of spatial precipitation fields

5.9 Configuration and Versioning

New configuration file: input/thames/spatial_correlation.json

```
{
"spatial_correlation": {
"enabled": true,
"model_type": "exponential",
"base_range_km": 40.0,
"nugget": 0.05,
"rho_intensity": 0.4,
"sigma_lognormal": 0.4,
"num_gauges": 40,
"gauge_source": "catchments.thames.GAUGE_POINTS",
"calibration": {
"events": ["Thames_2013_14", "Thames_2000", "Storm_Henk_2024"],
"data_sources": ["DEFRA_gauges", "MetOffice_radar"],
"calibration_date": "2026-03-05",
"validation_metrics": {
"correlation_rmse": 0.08,
"ks_test_pvalue": 0.12
}
}
}
}
```

Generation metadata addition:

Each StormSequence object includes spatial correlation metadata in output JSON:

```
"spatial_correlation_metadata": {
  "enabled": true,
  "range_km": 40.0,
  "intensity_factor": 1.2,
  "effective_range_km": 43.2,
  "nugget": 0.05
}
```

Schema version bump: From 2.0-doublet to 2.1-doublet-spatial

Backward compatibility:

- If `spatial_correlation.enabled = false` or config missing, default to spatially uniform precipitation
- Legacy storm files (v2.0) can be loaded; spatial correlation status inferred from schema version
- Generator can produce both v2.0 (uniform) and v2.1 (spatial) outputs based on configuration flag

5.10 Output Data Structures

Enhanced gauge response output: `gauge_responses_spatial.json`

```
{
  "schema_version": "2.1-doublet-spatial",
  "generated_at": "2026-03-05T16:20:00Z",
  "num_sequences": 10000,
  "num_gauges": 40,
  "spatial_correlation": {
    "enabled": true,
    "range_km": 40.0
  },
  "gauge_responses": [
    {
      "sequence_id": "SEQ-abc123",
      "gauges": {
        "Richmond Lock": {
          "hourly_precipitation_mm": [0.0, 0.0, 2.1, ...], // 168 values
          "hourly_level_m": [2.4, 2.4, 2.6, ...] // 168 values
        },
        "Richmond Bridge": {
```

```

"hourly_precipitation_mm": [0.0, 0.0, 1.8, ...],
"hourly_level_m": [2.3, 2.3, 2.5, ...]
},
// ... 38 more gauges
},
"spatial_statistics": {
"mean_precipitation_mm": 3.2,
"spatial_cv": 0.38,
"max_gauge_ratio": 1.6,
"min_gauge_ratio": 0.4
}
},
// ... 9999 more sequences
]
}

```

File size estimate:

- Per sequence: 40 gauges × 168 hours × 2 variables × 8 bytes = 107 KB
- 10,000 sequences: ~1.0 GB uncompressed
- With JSON compression (gzip): ~250-350 MB (acceptable for distribution)

5.11 Testing and Validation Requirements

Unit tests:

1. **Distance matrix:** Verify great-circle distances match expected values for known gauge pairs
2. **Correlation matrix:** Check positive definiteness, symmetry, diagonal = 1
3. **Cholesky decomposition:** Verify $C = LL^T$ within numerical tolerance
4. **Lognormal multipliers:** Confirm $\text{mean}(M) \approx 1$ before normalization
5. **Normalization:** Verify weighted mean of normalized multipliers = 1
6. **Catchment conservation:** Spatial mean of gauge precipitation = catchment mean

Integration tests:

1. Generate 1000 doublet sequences with spatial correlation
2. Verify all gauge precipitation series are non-negative
3. Check that high-intensity sequences show stronger spatial coherence
4. Confirm gauge responses show realistic downstream propagation
5. Validate that spatial correlation reduces for distant gauge pairs

Performance tests:

1. Benchmark: 10,000 sequences with spatial correlation in <2 minutes
 2. Memory usage: <5 GB peak memory for batch generation
 3. Parallel efficiency: >90% scaling on 4-8 cores
-

6. Mathematical Foundations

6.1 State-Space Model for Antecedent Conditions (Future Enhancement)

For a more rigorous treatment of compounding effects, implement a full state-space model:

State evolution:

$$S_{t+1} = f_S(S_t, G_t, P_t, \Delta t) + \epsilon_S \quad (6)$$

$$G_{t+1} = f_G(S_t, G_t, \Delta t) + \epsilon_G \quad (7)$$

where:

- f_S, f_G = nonlinear state transition functions
- ϵ_S, ϵ_G = stochastic components (process noise)

Observation model:

$$L_t = h(R_t, S_t, G_t) + \nu_t \quad (8)$$

where:

- h = observation function (rating curve)

- ν_t = measurement noise

Calibration approach:

1. Use historical Thames doublet events (2013/14, 2000, 2003 floods)
 2. Fit state-space model using Extended Kalman Filter or Particle Filter
 3. Estimate amplification coefficients λ_S, λ_G via maximum likelihood
 4. Validate that synthetic doublets match observed peak ratios (Storm 2 / Storm 1)
-

12. Document Control

Version: 2.1-doublet-spatial

Status: SPECIFICATION - Ready for Implementation

Date: March 5, 2026

Author: Technical Specification Team

Approved By: David Kelly

Change Log:

- v1.0 (Feb 26, 2026) - Original multi-storm specification
- v2.0-doublet (Mar 5, 2026) - Focused on two-storm sequences, added drainage window constraint
- v2.1-doublet-spatial (Mar 5, 2026) - Added Section 5: Spatial Correlation Across Thames Gauges

Next Steps:

1. Review and approve spatial correlation specification
 2. Implement spatial correlation module
(models/spatial_correlation.py)
 3. Calibrate parameters using Thames historical data
 4. Validate spatial patterns against radar rainfall
 5. Integrate with gauge response model
 6. Generate test catalog with spatial correlation enabled
-

References

- [1] Original Storm Generator Implementation Spec v1.0. (2026, February 26). Thames Valley Physical Risk Swaps project documentation.
- [2] Salvadori, G., & De Michele, C. (2004). Frequency analysis via copulas: Theoretical aspects and applications to hydrological events. *Water Resources Research*, 40(12).
- [3] Serinaldi, F., & Grimaldi, S. (2023). A Framework of Dependence Modeling and Evaluation System for Compound Flooding. *Water Resources Research*, 59(7).
- [4] Environment Agency. (2024). Thames flooding January 2024 - Section 19 investigation report. *UK Government*.
- [5] Bárdossy, A., & Pegram, G. (2013). Interpolation of precipitation under topographic influence at different time scales. *Water Resources Research*, 49(8), 4545-4565.
- [6] Papalexiou, S. M., & Serinaldi, F. (2020). Random Fields Simplified: Preserving Marginal Distributions, Correlations, and Intermittency, With Applications From Rainfall to Humidity. *Water Resources Research*, 56(2).