

MKM RESEARCH LABS

Weather Patterns to Physical Risk Swaps

A Practitioner's Guide to Flood
Risk Management and Other
Hazards

David K Kelly

Edited and Foreword by Johnny Mattimore

First Edition: 14th April 2025

MKM Research Labs

Weather Patterns to Physical Risk Swaps

A Practitioner's Guide to Flood Risk
Management and Other Hazards

First Edition: 14th April 2025

By David K Kelly
Chief Science Officer
MKM Research Labs

Edited and Foreword by Johnny Mattimore
Chief Executive Officer
MKM Research Labs

ISBN: TBA

Copyright Statement

© 2025 MKM Research Labs. All rights reserved.

The author has asserted their moral right to be identified as the author of this work in accordance with the Copyright, Designs and Patents Act 1988.

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of MKM Research Labs, except in the case of brief quotations embodied in critical reviews and specific other noncommercial uses permitted in accordance with the provisions of the Copyright, Designs and Patents Act 1988, or under the terms of any licence permitting limited copying issued by the Copyright Licensing Agency, Saffron House, 6-10 Kirby Street, London EC1N 8TS. Any person who does any unauthorised act in relation to this publication may be liable to criminal prosecution and civil claims for damages.

For permission requests, contact:

admin@mkmresearchlabs.com

1 THE PAVEMENT
BUSHY PARK ROAD
TEDDINGTON
ENGLAND TW11 9JE

www.mkmresearchlabs.com

This page has been intentionally left blank.

About the Author



David K Kelly is a seasoned investment banking professional with 30 years of experience, having held senior leadership positions in the front office and risk departments at global Tier 1 banks.

His entrepreneurial ventures include his most recent role as co-founder and Chief Science Officer of MKM Research Labs, a research company focused on developing and donating open-source business solutions for finance, particularly in the physical risk space.

Throughout his career, David has demonstrated expertise in implementing advanced modelling solutions for capital requirements under new regulatory frameworks. He is a leading advocate for Model Risk Governance for vendors that provide data and analytics to Finance.

David's educational background includes a first-class degree in Mathematics from the University of Bristol and Honours in Advanced Studies (Part III) in Pure Mathematics from the University of Cambridge.

About the Editor

Johnny Mattimore has over 35 years in financial services. He originally trained as a mathematician and markets economist working on major trading floors of capital markets firms. He worked for the first half of his career in investment banking in analysis, trading and structuring.

He then spent the second half of his career in hedge funds and asset management across a range of roles running money, risk management, and trading systems design and integration. More recently, he has focussed his efforts on the integration of cutting-edge risk solutions into the financial ecosystem for banking, insurance and asset management.

He is a leading proponent of open-source standards and has a high level of expertise in the integration of data, methods and models into finance for lines of business and group functions. As well as having held mainstream roles in large organisations, he has a long history in developing new businesses, particularly in hedge funds and fintech.

Johnny graduated from the University of Bristol, UK, in 1988 with a BSc in Mathematics. He has written extensive research throughout his career, including major works for risk, emerging markets, structured debt and investment methodologies, ranging from discretionary to systematic.

About MKM Research Labs



MKM Research Labs

MKM stands at the forefront of open-source innovation in banking technology, mainly focused on addressing the banking sector's critical challenges of physical risk management. Founded in January 2025 by Johnny Mattimore, David Kelly, and Fearghal McGoveran, MKM consolidated years of prior open-source work into a cohesive entity dedicated to developing and donating open-source business solutions for finance.

"We are not in the weather prediction business. Rather, we quantify banking risk by modelling the probability distributions of extreme but plausible weather events that trigger floods. We focus on assigning accurate probabilities and banking impact severities to these tail events, enabling banks to properly assess their exposure to hazard-related risks." - MKM Exco.

The company's flagship initiative focuses on transforming flood risk management through open-source banking technology. Their approach applies capital markets technology to physical risk management, enabling flood risk to be packaged, priced, and traded in real-time, similar to interest rate and credit default swaps. Central to their work is developing a Common Domain Model (CDM) implementation that standardises flood and hazard risk data and models, including integrating fluid dynamics for flood modelling.

MKM leads crucial projects in Physical Risk & Resilience while spearheading the extension of the Common Domain Model to property and physical hazards. Their work extends beyond traditional risk management to include the OS-SFT (Open-Source Sustainable Finance Taxonomy) project, which aims to harmonise regulatory interpretation across jurisdictions. Through their open-source approach under the FINOS umbrella, MKM is building a new market infrastructure that promises to ensure the resilience of property finance for decades to come.

MKM actively participates in the Banking Services Open-Source movement under the auspices of ISDA and FINOS.

Obtaining a Copy of the Book for Free

The book will be available at www.mrmresearchlabs.com as a free download.

Acknowledgments

The Author would like to thank many people who have provided support, encouragement and technical advice. The author would like to highlight fellow MKM founders Johnny Mattimore and Fearghal Mcgoveran for their encouragement and proofreading, as well as Vanessa Balmora for her technical guidance in flood-related things. The author recognises the excellent collaboration between Sheroy Marker and Aditi Gautam from Nvidia during the discovery phase, and strategic discussions for further research with Aaron Erickson, Gabe Hege and Jochen Papenbrock.

For open-source guidance and leadership, the Author would like to acknowledge the support of the Linux Foundation, FINOS and OS-Climate organisations, particularly Gabriele Columbro, Jane Gavronski, Michael Tiemann, Matt Sandoe, Steven Tebbe, Richard Harmon, Vincent Valdeira, and the founder and former CEO of OS-Climate, Truman Semans.

For the Common Domain Model, the Author would like to thank Ian Sloyan, who has provided invaluable guidance on our initial designs for the Physical Risk CDM, which can be found at <https://github.com/MKM-Research-Labs/Physrisk-cdm>.

Finally, there are many organisations who have contributed to the thinking behind this book over recent years, including, but not limited to: ISDA (International Swap Dealers Association), ICMA (International Capital Markets Association), ISLA (International Securities Lending Association), BNP Paribas, Citi, Goldman Sachs, JP Morgan, Morgan Stanley, UBS, Fannie Mae, Red Hat, NVIDIA, JBA Risk Management, OASIS Loss Modelling Framework, and Imperial College London.

To my lovely wife, Victoria, and my wonderful children, David Henry, Georgina, Angus, and my most faithful dog, Otis, a Standard Wired Dachshund.

To Johnny Mattimore, whose infectious enthusiasm for this book and thoughtful critique has kept me (nearly) sane.

Table of Contents

Foreword	15
Prologue	20
Overview of Solution Evolution.....	27
Chapter 1 - Physical Risk in Finance.....	29
Chapter 2 - Data Framework	43
Chapter 3 - Weather Prediction	72
Chapter 4 - Hydrological Modelling	91
Chapter 5 - Flood Risk Assessment	131
Chapter 6 - Impact Quantification	162
Chapter 7 - Insurance Risk Assessment.....	183
Chapter 8 - Mortgage Risk Assessment.....	207
Chapter 9 - Physical Risk Swaps	216
Chapter 10 - Model Risk Governance	241
Chapter 11 - The Cutting-Edge	253
Glossary of Terms.....	274
Property CDM	297
Bibliography.....	306

Table of Figures, Tables and Formulae

- Figure 1: MKM Solution Overview
- Figure 2: Basel III Eligibility Framework
- Figure 3: Physical Risk CDM Overview
- Figure 4: Weather to Risk Transfer via CDM
- Figure 5: Property Adaption Measures
- Figure 6: MKM Reporting Framework
- Figure 7: Physical Risk Swap
- Figure 8: PRS and Flood Participants
- Figure 9: Components of a Model
- Figure 10: Model Risk Cycle
- Figure 11: Stages for PRS Development
- Figure 12: Connection of CorrDiff to LisFlood

- Table 1: Format Comparison
- Table 2: Options for a Physical Risk Service
- Table 3: MKM Framework and External Forcings
- Table 4: Flood Depth by Return Period
- Table 5: Insurance Premium Adjustment

- Formula 1: Stochastic Atmospheric State Evolution
- Formula 2: Hidden Markov Models
- Formula 3: Generalized Linear Models
- Formula 4: Nash Cascade Catchment Response
- Formula 5: Hydrologic Routing Differential
- Formula 6: Modified Puls Flow Rate
- Formula 7: Muskingum River Basin
- Formula 8: Muskingum-Cunge Discharge
- Formula 9: Saint-Venant Equations
- Formula 10: Hazard Curve Development
- Formula 11: Vulnerability Curve
- Formula 12: Belanger and Bourdeau-Brien
- Formula 13: SOP Fundamental
- Formula 14: Phase-specific SOP

- Formula 15: Hazard Reduction
- Formula 16: Return Period Adjustment
- Formula 17: Total Property Resilience
- Formula 18: Pure Premium
- Formula 19: Multivariate Premium Calculation
- Formula 20: Customer Lifetime Value
- Formula 21: Cost Frequency Bayesian
- Formula 22: Predictive Tree
- Formula 23: Discount Cash Flow
- Formula 24: Mortgage PD and LGD
- Formula 25: Stochastic House Price
- Formula 26: CDS Survival
- Formula 27: PRS Survival

This page has been intentionally left blank.

Foreword

David Kelly and I have come a long way together on the journey of integrating physical risk into the banking system. While this journey is quite advanced, it is far from over. The completion of the first edition of this book marks a meaningful milestone in our journey. It represents the first major publication from our new firm, MKM Research Labs, which was founded by the three of us - Johnny Mattimore, David Kelly, and Fearghal McGoveran - in January 2025 to consolidate our decades of experience in banking and finance.

“This book provides a visionary guide to integrating Physical Risk into the banking system, at high speed, large scale and reduced cost for all.”- Johnny Mattimore, MKM.

Introduction

The motivation for writing the book is identical to that for creating the firm: to combine cutting-edge artificial intelligence, data, models and deep domain expertise to solve some of the most challenging problems in finance. One of those problems is how to integrate physical risk into the banking ecosystem at high speed, at large scale and at reduced cost for all.

The book is crafted from many years of real-world experience and scholarly excellence, combined with the essential reflection and vision necessary to have a meaningful impact on the reader.

Context for the Book

Put simply, this book makes the subject of physical risk easy and accessible for bankers to understand. It provides bankers with what they need in order to understand physical risk without the need to become experts in every underlying skill. Uniquely, in my view, it does this as a single source, as Kelly has taken the best of current research and models from hundreds of sources and combined them with his individual skill and knowledge, into an accessible guide for bankers.

To achieve this simplicity, it has three common threads running throughout it, taken from great minds of history: Newton, Aristotle and Michelangelo. All of these threads are familiar to us and are identifiable in all works of great success for humankind.

Giants

"If I have seen further, it is by standing on the shoulders of giants."- Sir Isaac Newton, (1675).

In this phrase, Newton acknowledges that his discoveries were made possible by building on the work of great thinkers before him - such as Galileo Galilei, Johannes Kepler, and René Descartes.

Following this tradition, this book illustrates how a solution for the bankers' problem is only possible by building on the work of others. So, the book draws together many other areas of skill developed over many decades in the modern fields of physical risk, including: weather forecasting, hydrology, hydraulics, specialist common data models, artificial intelligence and the imperative of time series models for pricing risk in banks. Through these guiding steps, Kelly illuminates the pathway for solving the bankers' problem of physical risk integration into its existing core system with minimal friction.

Value

"The whole is greater than the sum of its parts"- Aristotle, circa (340 BC).

While the exact wording may not appear verbatim in Aristotle's surviving works, the idea reflects a key concept from his Metaphysics, where he discusses how complex systems and organisms have properties that cannot be explained solely by analysing their individual components.

Thus, the brilliance of this book lies in Kelly's ability to connect the many areas of expertise in an accessible way. Each component is demystified so that by the end, the reader feels confident that the formerly opaque and confusing becomes clear and lucid.

Moreover, it then becomes evident how these individual components alone do not solve the bankers' integration problem of estimating the financial risk in terms of real money. Instead, it shows that only by combining these components, with defined sequencing and coupling can one extract the required value. Thus, it delivers on the promise to the banker of proving that "the whole is greater than the sum of the parts"

Ambition

"The greater danger for most of us lies not in setting our aim too high and falling short, but in setting our aim too low and achieving our mark." - Michelangelo Buonarroti (1475-1564).

While possibly apocryphal, this phrase nevertheless captures the ambitious spirit of Michelangelo's life and work which gave rise to his monumentally daring projects. Today, it is still a powerful motivational idea associated with his legacy.

Rephrasing this for our era and for bankers, an appropriate rewrite might be: "In banking, as in life, the greater danger lies not in bold targets unmet, but in modest goals too easily achieved." So, one could play it safe and get modest results, but as any financier knows, it is bold goals that drive meaningful change to deliver outsized growth while maintaining stability and repeatability.

Kelly outlines the opportunity that I coined, and which he and I have jointly developed for a new tradeable asset to manage, hedge and transfer physical risk: the Physical Risk Swap (PRS).

We developed this idea for a capital markets PRS over the last five years. Our first publication was in November, 2023, as a LinkedIn post. This was where we first coined the term Physical Risk Swap (PRS), comparing it to a Credit Default Swap (CDS), first used in Collateralised Debt Obligations (CDO), and with a securitised form of a Parametric Insurance Contract.

Our invention of PRS is truly ambitious: it not only transforms the world of physical risk into a tradeable asset, but it is predicated on the ambition of creating an open source architectural framework for the entire banking ecosystem.

In this book, Kelly's explanation takes a huge step forward compared with other related literature detailing how to achieve the ambition of PRS across a stack of models. In my view, the creation of a market in PRS is the

pinnacle of our joint ambition and the ambition of our firm, MKM Research Labs.

Kelly explains this transformational leap forward to PRS by leveraging the achievements of others in different domains to deliver value faster; by combining the sum of the parts to deliver even greater total value; and, finally, by setting the ambition high . In our vision, PRS creates a brand new growth opportunity while enhancing systemic stability.

This combination of Giants, Value and Ambition is what makes this book a wonderful piece of craftsmanship and a truly outstanding contribution to banking and society.

Final Comments

Throughout reading this book, perhaps remember this one enduring fact: without a functioning, solvent banking system, almost all that we depend upon, and all that we take for granted is at risk of being taken away from us unless we protect ourselves from extreme physical risks. This book contributes substantially to the body of accessible knowledge to help us preserve all those things that we hold dear.

Enjoy the read and then the ride as you witness how managing physical risk in banking will change the world for the better, for the benefit of all and forever.

Johnny Mattimore
CEO and Co-founder of MKM Research Labs

This page has been intentionally left blank.

Prologue

The banking world stands at a critical inflection point. A threat more insidious than market volatility, more persistent than credit cycles, and more fundamental than technological disruption is reshaping the risk landscape. Once considered a future concern, physical risk has evolved into an immediate crisis that threatens the very foundations of our banking system.

Consider this stark reality: The US property market, valued at \$52.5 trillion and burdened by \$13.6 trillion in debt, faces an unprecedented challenge. In 2024, we observed a mass exodus of insurers from coastal markets, with 42% of US coastal insurers withdrawing from Florida and the Carolinas. This situation left \$4.2 trillion in mortgage collateral exposed and uninsured. The reduction in available insurance is not a gradual shift; flood insurance premiums skyrocketed 58% yearly, pricing out homeowners and destabilising property valuations across entire regions.

The scale of exposure to the UK property market is staggering. Along the Gulf and Atlantic Coasts alone, 7.3 million homes face the threat of storm surge flooding. Nationwide, 17% of all properties risk flood damage, with the highest concentrations in Florida, Louisiana, California, New York, and New Jersey. In California, 4.6 million properties stand in wildfire zones, while Texas faces approximately 137 tornadoes annually, affecting over 13 million people. These are not distant threats - they are present realities affecting properties that underpin mortgages, secure investments, and represent lifelong savings for millions of Americans.

We face a fundamental mismatch in time horizons that threatens the stability of our banking system. Banks write mortgages with 30-year terms, while insurers underwrite policies annually. As physical risks intensify, insurers can and do retreat from markets they deem too risky, leaving banks exposed to decades-long commitments on potentially worthless assets. This temporal disconnect echoes the opacity that preceded the 2008 banking crisis, but with a crucial difference - physical risk is not a complex banking instrument that can be traded away. Flood risk is tied to the immutable realities of geography, weather patterns and infrastructure, and the risk of banking loss is tied to individual property value, flood resilience, and outstanding mortgage debt.

More recently, banking regulators in various jurisdictions have been working on incorporating climate risk more explicitly into regulatory frameworks, but these are generally still developing and vary by region. The Basel Committee on Banking Supervision has published principles for

the effective management and supervision of climate-related banking risks, but these haven't yet translated into standardised capital requirements. However, there are concerns that individual regulators could use their discretionary powers under Pillar 2 to impose capital charges of up to 15% for unhedged physical climate risk exposures.

Some countries are taking more direct approaches to managing climate-related physical risks. For instance, the Italian government has implemented measures requiring all corporations to obtain specific flood-related insurance coverage, shifting some of this risk management responsibility directly to businesses rather than leaving it entirely within the banking regulatory framework.

Many lenders are considering such regulatory risk and exiting mortgage provisions backed by any property in a heightened flood-risk area.

The traditional bank risk management playbook offers no solutions. Our banking systems, sophisticated in handling market and credit risk, remain primitive in their approach to physical risk. The tools, models, and frameworks that served us well in managing traditional banking risks prove inadequate in the face of this new challenge.

Yet, within this crisis lies opportunity. The same forces driving insurance retreat could catalyse innovation in risk transfer, valuation, and resilience investment. When properly directed, the banking sector's capacity for innovation could transform physical risk from a threat to a frontier for new banking products, improved risk management, and enduring property finance.

This book presents the MKM Framework for Physical Risk, a transformational approach that bridges the gap between academic-driven scientific research and banking mathematics, property-level analysis and systemic risk management, and current challenges and future solutions. The stakes could not be higher. The integrity of our mortgage markets, the stability of our banking system, and the resilience of our communities all hang in the balance.

To provide some context, the Pew Research Centre outlines that home equity is a significant component of household wealth, accounting for 45% of net worth among UK homeowners in 2021. While the UK and EU have slightly lower concentrations, Chinese households, according to East-West Property, hold 65% of real estate.

The time for incremental solutions has passed. We need a fundamental reimaging of measuring, managing, and transferring physical risk. The future of finance depends on it. This book aims to forge a path from

forecasting weather patterns to live revaluation of a portfolio of banking assets.

The ultimate goal is to apply Capital Markets technology as a bridge into the insurance industry, triggering a recovery in insurance capacity and leading to lower and more available underwriting capacity.

The Hazard Landscape

It is noticeable in the space of Hazards that they cover. A multitude of events. At this stage in the book, it is worth outlining a classification.

Hydrological Risks (Water-Related)

- **Flooding:** River, coastal, or flash floods due to heavy rainfall or storm surges.
- **Drought:** Prolonged water shortages affecting agriculture and water supply.
- **Landslides:** Soil and rock movement triggered by heavy rain or erosion.

Meteorological Risks (Weather-Related)

- **Hurricanes/Typhoons/Cyclones:** Strong winds, heavy rain, storm surges.
- **Tornadoes:** Highly destructive rotating windstorms.
- **Extreme Heat:** high temperatures with health and infrastructure issues.
- **Extreme Cold & Winter Storms:** Freezing, heavy snowfall, ice storms.
- **Hailstorms:** Large hailstones causing damage to property and agriculture.

Geological Risks (Earth-Related)

- **Earthquakes:** Ground shaking causing structural damage and tsunamis.
- **Volcanic Eruptions:** Lava flows, ash clouds, and toxic gases.
- **Tsunamis:** ocean waves triggered by underwater earthquakes or landslides.

Biological Risks (Ecosystem-Related)

- **Wildfires:** Uncontrolled fires from dry conditions, lightning, or human activity.
- **Pest Infestations:** Insects or diseases affecting agriculture and forestry.

- **Pandemics & Epidemics:** diseases impacting health and economies.

Although this book reflects physical risk as a present danger, for completeness, the Intergovernmental Panel on Climate Change (IPCC) of the United Nations, in its sixth assessment report, classifies physical risks related to climate change into three main categories:

Chronic Physical Risks: are sudden, extreme weather events that cause immediate damage. Examples include:

- Hurricanes and typhoons.
- Floods and storm surges.
- Heatwaves.
- Wildfires.

Chronic Physical Risks: are long-term, gradual changes in climate patterns that progressively impact ecosystems, economies, and societies. Examples include:

- Rising sea levels.
- Increasing average temperatures.
- Changes in precipitation patterns.
- Ocean acidification.

Compound and Cascading Risks: occur when multiple climate hazards interact, intensifying their effects. Examples include:

- A heatwave exacerbating drought and wildfires.
- Flooding damages infrastructure, leading to economic and health crises.
- Sea level rise increasing storm surge damage.

About this book - A focus on flood risk

The world of finance, meteorology, fluid dynamics, and instrument pricing is awash with academic articles and dense mathematical formulae. The purpose of this book is not to dive into the maths but to help capital markets practitioners make sense of the problem of flood risk in finance and forge a pathway to implementation that will need to straddle multiple academic expertise.

At some point, we will introduce mathematical equations. The author believes they are helpful for explanation but not essential for an overall appreciation of the concept. Therefore, non-mathematicians can skim through with a knowing nod, leaving it to the quants to enjoy. The sections, including mathematical equations are highlighted in blue.

There are many hazards that groups like OS-Climate continue to lead with excellent research. Hazards such as wildfire, soil erosion, drought and pestilence. Maybe not the last. We are focused on floods arising from weather formations to set up the correct framework for risk transfer within a capital markets context. To that end, there is some overlap with wind and coastal damage as hazards arising from cyclones.

The author and the MKM Exco leave the modelling of potential climate changes for horizons out to 2100 to academic bodies under the IPCC's umbrella. MKM is dedicated to implementing a framework that complies with how banking market risk management is adopted by all banks that warehouse risk. The management of hazard risk follows the same four steps currently applied to instrument pricing: Identifying hazard risk, Measuring hazard risk, Monitoring exposure to hazards, and building Control tools to affect risk transfer.

Impact of Flood Risk on UK Properties - a Case in Point

While the focus here is on the UK property sector, the underlying issues apply straightforwardly to the EU, US, China, and Japan.

Flood risk has emerged as a critical factor reshaping property markets, regulatory frameworks, and banking stability across the UK, with profound implications for the broader banking sector. Recent comprehensive assessments reveal an escalating threat that demands urgent attention from banking institutions, policymakers, and property owners alike.

Current Risk Exposure and Projections

The latest National Assessment of Flood Risk (NaFRA) identifies 6.3 million properties in England currently classified as flood-prone, with 4.6 million at risk from surface water flooding and 2.4 million from rivers/sea flooding. This represents a substantial portion of the nation's housing stock, with climate change projections suggesting this figure could rise dramatically to 8 million properties (1 in 4) by mid-century, driven by increasingly extreme weather patterns, heavier rainfall, and accelerating sea-level rise.

Regional disparities in flood risk exposure are significant and economically consequential. East Anglia, the North West, and Yorkshire face the highest exposure rates, with 13-18% of properties in high-risk zones. Urban areas with ageing or inadequate drainage infrastructure, such as parts of London and Birmingham, are particularly vulnerable to

surface water flooding, which now accounts for 57% of all flood insurance claims.

Quantifiable Market Impacts

The Bayes Business School's "Residential Property Flood Risk UK 2023" research report reveals that flood risk creates measurable valuation penalties in the property market. Properties affected by flood risk sell at an 8.14% average discount compared to equivalent non-affected properties, escalating dramatically to 31.3% for very high-risk areas. This price differential correlates with flood probability, with each 1% increase in risk-reducing property values by 0.13-0.19%.

The market dynamics reflect rational economic behaviour: high-risk homes experience a 50% slower price growth than properties in safer areas. Critically, the average flood risk of sold properties between 2006 and 2021 was 8.01%, compared to an unsold property's 8.63%, indicating that flood exposure significantly impacts marketability. Analysis shows that zero exposure to flooding can increase a property's saleability from 63.3% to 65.6%.

Insurance Challenges and Banking Stability Risks

The UK's government-backed insurance scheme, Flood Re, in their Transition Plan Report 2023, currently protects approximately 200,000 households in high-risk areas. Still, its scheduled expiration in 2039 raises serious concerns about future insurance affordability and availability. With the Association of British Insurers pointing out that the average flood claims cost £32,000 per incident, insurers are increasingly factoring climate projections into their risk models and pricing.

This creates a potential systemic risk to banking stability through collateral devaluation. Properties in high-risk areas face growing loan defaults as declining property values outpace mortgage balances, potentially creating clusters of negative equity in vulnerable regions. JBA Risk Management estimates the annual cost of flood damage to UK properties at £527 million, a figure projected to rise substantially without significant adaptation measures.

Regulatory Response and Adaptation Measures

Regulatory frameworks have evolved to incorporate stricter building codes in response to these growing risks. New floodplain developments must now implement elevated foundations, flood barriers, and sustainable drainage systems (SuDS). The Sequential Test mandates avoiding high-risk zones for new construction unless no viable alternatives exist.

In March 2025, the UK Government and Environment Agency announced a record £2.65bn investment in flood defences, aimed at building or repairing over 1,000 flood schemes and protecting 66,500 properties by 2026, focusing on vulnerable regions like the Thames Estuary and Humber Basin. Despite these efforts, climate projections by the National Assessment of Flood and Coastal Erosion Risk in England 2024 indicate that 637,600 properties could face high river/sea flood risk without accelerated adaptation by 2070 (a 73% increase), and 1.8 million may be exposed to severe surface water flooding (a 66% increase).

Future Outlook and Banking Sector Implications

The banking sector is responding through market innovations, including mandatory climate risk disclosures for mortgages and the development of parametric insurance models that provide rapid payouts based on predetermined trigger events rather than assessed damage. Property flood resilience (PFR) measures are increasingly demanded in the valuation during property assessments, with buyers willing to pay premiums for homes with implemented flood mitigation features.

Physical Risk resilience has become a cornerstone of long-term investment decisions in the property sector, with valuation models increasingly incorporating sophisticated flood risk metrics. This evolution represents both a challenge and an opportunity for banking institutions to develop physical risk, dependent lending practices and innovative banking products that incentivise adaptation.

The comprehensive data indicates that flood risk is not merely an environmental concern but a material banking risk that requires integrated assessment within investment, lending, and insurance frameworks. As the impact of physical risk develops with increased urbanisation, the property sector's ability to price and mitigate these risks accurately will be critical to maintaining banking stability and market efficiency.

Overview of Solution Evolution

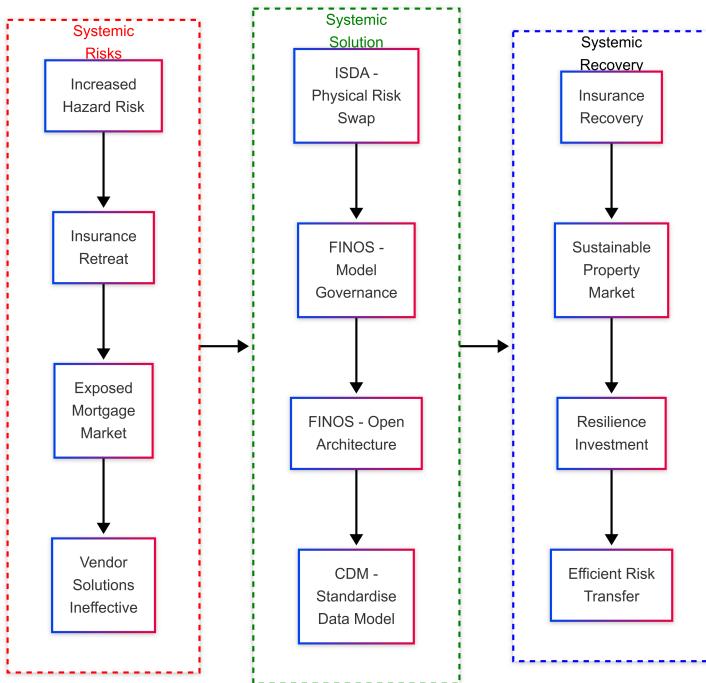


Figure 1: MKM Solution Overview

History and evolution of the solution

The approach proposed by MKM Research Labs addresses the double jeopardy of increased insurance payouts due to hazard risks and their impact on assets, notably properties and mortgages. The insurance industry's response is to retreat capacity and only focus on low-risk areas. The challenge for banks and mortgage providers is exacerbated by vendor solutions that don't really work for them.

The response is, therefore, to increase capacity for underwriting such risk by working from capital markets under the auspices of ISDA and FINOS, which provide a collective legal and data framework from upon which the industry can build models that accurately cost (as exemplified in

this book) flood risk and transfer risk through new derivative instruments such as physical risk swaps.

The efficiency of risk transfer will enable a higher degree of assurance around underwriting hazard (flood) risk and its impact on the current portfolio of assets. With capital markets technology that understands how to manage derivative-type exposure, the increase in capacity will lead to lower insurance premiums.

A critical societal side effect of improved hazard risk discipline thanks to ISDA and the standardisation of data through FINOS CDM is the drive for a more effective causality between resilience investment and lower mortgage and insurance costs.

“Wouldn’t it be nice if the mortgage provider, insurance underwriter, and local government spoke the same language based on common data? Crazy, I know, but we have to start somewhere, and all great journeys start with a single step.” - David Kelly, MKM and Author of that single step.

Chapter 1 - Physical Risk in Finance

The banking sector's challenge is to enhance risk measurement and develop new frameworks for understanding how physical risks propagate through banking networks. We require innovative ways to model the relationship between weather events and asset values, new data standards to capture physical risk exposure, and novel banking instruments to effectively transfer and manage these risks. This transformation necessitates bridging multiple disciplines, from weather prediction through hydrology and loss modelling to banking mathematics and machine learning.

This chapter investigates the key dimensions of physical risk in finance and examines why traditional approaches are inadequate. We analyse how emerging technologies and methodologies facilitate a more sophisticated understanding of physical risk and why this evolution is crucial for ensuring banking stability in an era of evolving physical risks. Most importantly, we delineate a pathway for transforming physical risk from an unquantified threat into a manageable, transferable element of banking risk management.

“The banking world's approach to physical risk is undergoing a fundamental transformation - from a future consideration to an immediate driver of value, from an unquantified threat to a tradable asset, from a peripheral concern to a core determinant of banking stability.” -
David Kelly, MKM.

Overview of Physical Risk Challenge

Physical risk in the banking sector has evolved from a peripheral concern to a central challenge in modern risk management. The landscape of physical risk has transformed dramatically, characterised by the increasing frequency and severity of chronic physical events on regional

and global scales. What was once considered idiosyncratic risks are now manifesting as systemic challenges, fundamentally altering how we approach risk assessment and management.

The Swiss Re Institute sigma No. 1/2024 report highlights that the number of medium-sized natural catastrophe events has grown by an average of 7.5% annually since 1994, with severe convective storms (SCS) showing the most significant increase in insured losses, rising by 9.7% annually in inflation-adjusted terms.

The Swiss Re Institute sigma No. 1/2023 report also notes a long-term upward trend in insured losses from natural catastrophes, driven by economic growth, urbanisation, and increased storm activity. Insured losses have grown 5-7% annually since 1992, with severe weather events being a key driver.

The complexity of physical risk assessment demands increasingly granular analysis at the asset level. Banking institutions must now evaluate thousands or millions of assets in their portfolios, considering precise physical locations and projecting expected losses across multiple physical hazards. This granular approach represents a significant departure from traditional portfolio-level risk assessment methods.

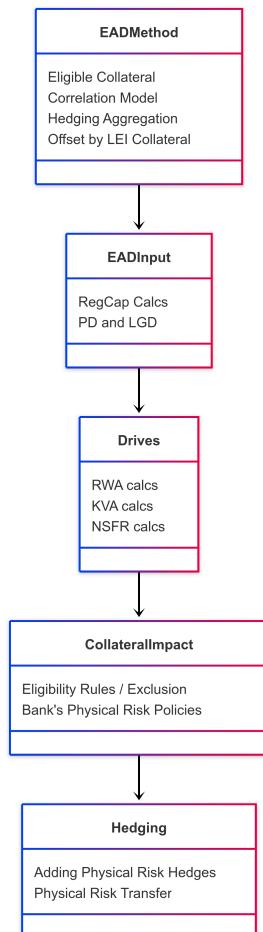


Figure 2: Basel III Eligibility Framework

The data requirements for such detailed analysis present substantial challenges. Banking institutions require robust data on asset locations, hazard projections, and potential damage functions. Significant uncertainties and knowledge gaps in weather projections and impact modelling complicate these data needs. Integrating diverse data sources, from high-resolution terrain models to real-time weather measurements, demands new data management and standardisation approaches.

The banking sector is exploring innovative risk transfer mechanisms in response to these challenges. New securities, derivatives, and risk transfer products are being developed, including specialised mortgage-backed securities (MBS) and physical risk swaps (PRS). These instruments aim to provide more effective means of managing and transferring physical risks analogous to specialised Collateralised Debt Obligations (CDO) and Credit Default Swaps (CDS).

The industry's response to these challenges requires a fundamental evolution in risk management approaches. Traditional frameworks, built on historical data and statistical relationships, struggle to capture the changing nature of physical risks. The non-linear characteristics of weather patterns, feedback loops, and fractal-esq interactions necessitate new modelling paradigms that can handle these complexities while remaining viable for practical implementation.

Modern risk management requires banks to connect the science of physical risk with banking risk assessment to be practical in daily situations. This involves integrating advanced statistical methods, including domain-specific AI and Bayesian approaches, to capture better the complexity of weather patterns and their implications for banking.

Developing standardised data frameworks and robust governance structures is essential to supporting consistent risk assessment across institutions.

Implementing these new approaches presents significant operational challenges. Banks are required to develop and validate complex models while managing vast amounts of data. The computational demands of physical risk assessment require scalable technological solutions that can handle increasingly sophisticated analysis requirements.

As we move forward, the industry needs a comprehensive approach to addressing these interrelated challenges. Given that physical risk presents a present issue for banks, a key aspect of any solution is to decide on the mix between the level of collaboration through groups such as FINOS, where the output is open source, and what is considered proprietary. The need for two counterparties to operate under the same legal, data and model framework encourages market liquidity, whereas proprietary intelligence provides a profitability edge.

This book offers an approach that progresses from fundamental weather prediction to banking impact assessment while prioritising practical implementation considerations. We start by exploring the elements of physical risk assessment and then create an integrated framework that merges these elements into a cohesive whole suitable for application in modern banking institutions.

Current Approaches and Limitations

The banking sector's traditional approach to physical risk assessment has evolved from catastrophe modelling developed for the insurance industry. This heritage reveals both the strengths and limitations of current methodologies. Insurance companies have invested decades in developing sophisticated models to price annual policies' flood, fire, and storm risks. However, they did not design these models to address the broader needs of banking institutions managing long-term exposures through mortgages, infrastructure investments, and structured products.

Current approaches typically follow a four-step process:

- Hazard identification.
- Exposure analysis.
- Vulnerability assessment.
- Loss calculation.

Hazard identification relies heavily on historical data. Exposure analysis focuses on asset location and characteristics, while vulnerability assessment examines the relationship between hazard intensity and potential damage. Finally, loss calculation combines these elements to estimate possible banking impacts. However, this seemingly straightforward process masks significant methodological weaknesses that limit its effectiveness for modern risk management.

A fundamental limitation lies in how these models handle risk quantification. Many existing models do not explicitly separate frequency and severity modelling, departing from established best practices in property insurance. Furthermore, they often produce deterministic expected loss values rather than complete probability distributions, severely limiting their utility for sophisticated risk management applications.

Specifically, the vulnerability models that map hazard characteristics to potential damage typically employ deterministic rather than probabilistic approaches, failing to capture the full spectrum of uncertainty inherent in physical risk assessment.

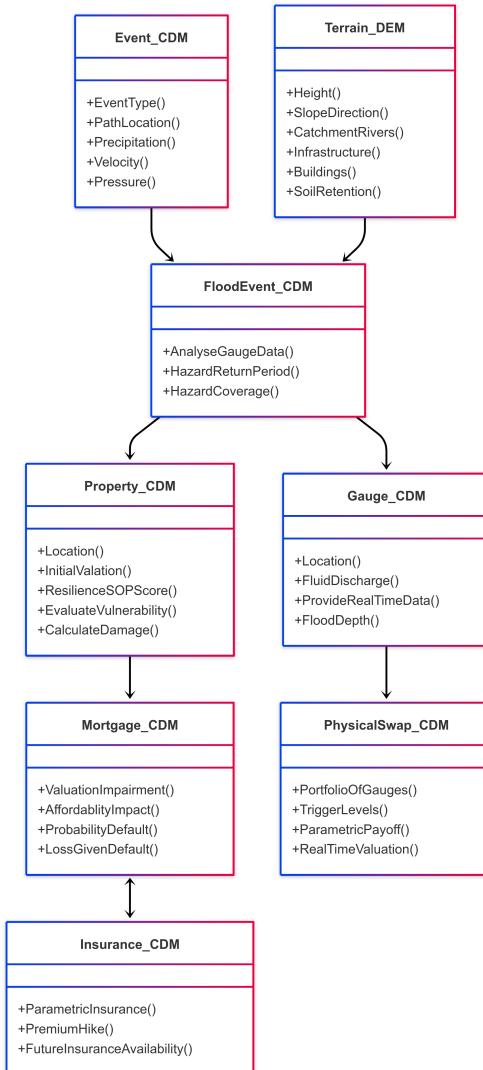


Figure 3: Physical Risk CDM Overview

This approach, while logical, faces significant limitations when applied to modern banking risk management. Relying on historical data becomes increasingly problematic as weather patterns' frequency and severity of

extreme events follow random paths. The assumption of stationarity—that past patterns reliably indicate future risks—no longer holds.

The current models struggle with the temporal mismatch between insurance and banking time horizons. Insurance models are calibrated for annual policy terms, making them ill-suited for assessing the decades-long exposure of a mortgage portfolio. The models cannot adequately capture how physical risks might evolve over a 30-year mortgage term, nor how these risks interact with property values, insurance availability, and mortgage affordability.

“Existing approaches fail to capture how physical events trigger cascading effects through banking networks. When insurers retreat from high-risk areas, they leave properties uninsured and impact property values, mortgage default probabilities, and, ultimately, the stability of mortgage-backed securities markets.”

This limitation is exacerbated by the lack of clear attribution between risk factors and their economic consequences, making it difficult to conduct meaningful stress testing or design optimal insurance solutions.” -

David Kelly, MKM.

The Challenge of Model Governance

The most frustrating aspect of the vendor landscape, regarding weather and physical risk, is the after-thought considerations for model governance. These considerations compound these issues and make bank adoption under the auspices of OCC SR 11-7 and other regulatory requirements problematic.

Quick Introduction to OCC SR 11-7

SR 11-7, formerly known as the “Supervisory Guidance on Model Risk Management,” was introduced in April 2011 by the U.S. Federal Reserve and the Office of the Comptroller of the Currency (OCC) to address the growing reliance on quantitative models in banking institutions.

SR 11-7 establishes comprehensive standards for managing risks from flawed model design, implementation errors, or misuse, which could lead to banking losses, poor decision-making, or reputational harm. It defines a model as a quantitative method that processes inputs into estimates using statistical, economic, or mathematical techniques, emphasising three core components:

- Input Data.
- Processing Logic.
- Output Reporting.

The guidance mandates rigorous model validation, governance policies, documentation standards, and ongoing monitoring to ensure models perform as intended while aligning with an institution's risk tolerance. By requiring "effective challenge" through independent review and structured accountability, SR 11-7 aims to balance innovation with risk mitigation across the model lifecycle.

Physical risk models shortfall against SR 11-7

Many current models in physical risk are academic in their nature and outlook and thus function as "black boxes" without clear documentation of their methodologies and assumptions. This lack of transparency makes thorough model validation nearly impossible. Such behaviour contradicts the banking industry, where the model engines are similar, and the intellectual property is embedded in their implementation. In all areas of physical risk, the underlying models originate from academia and are publicly accessible. For instance, fluid dynamics is covered in the first semester of a physics degree. The lack of transparency regarding applications or benchmarks for flood risk modelling further complicates the situation, hindering coherent comparison and systematic improvement of various modelling approaches.

Different vendors use varying assumptions, are very guarded about their methodologies, and use public and private data sources differently, making it difficult for banks to compare and aggregate risk assessments. This fragmentation hampers the development of market-based solutions for transferring and managing physical risk.

The limitations become particularly chronic when examining flood risk, which exemplifies the complexity of modern physical risk assessment. Flood models are required to integrate high-resolution terrain data, precipitation forecasts, infrastructure conditions, and property characteristics. To be remotely accurate, they must account for gradual changes in flood patterns and sudden shifts in insurance market

behaviour. The recent exodus of insurers from Florida and the Carolinas demonstrates how quickly risk assessments can become outdated when market conditions change.

Basel III/IV regulations now explicitly require banks to account for physical risk in their capital calculations, but the tools for doing so remain primitive. Current approaches often resort to simple overlays or adjustment factors rather than an integrated assessment of how physical risks affect the probability of default, loss-given default, and other key risk metrics.

A particularly problematic aspect of current approaches is their origin outside the insurance domain. Many existing models were initially developed for land-use planning or other purposes, making adapting to insurance pricing and risk management applications problematic. This misalignment between the original design and the current application creates fundamental model reliability and appropriateness challenges.

These limitations indicate the fundamental need to reimagine physical risk assessment in banking. We need frameworks that can:

- Bridge multiple time horizons, from weather events to mortgage terms.
- Integrate physical and banking risk metrics.
- Support standardisation in particular market risk data.
- Develop streamlined processes to support secondary market-making activities.
- Enable dynamic updating as conditions change.
- Facilitate risk transfer and management.

The following chapters present such a framework, building from weather prediction to banking impact assessment while addressing the limitations of current approaches. The goal is not to replace existing catastrophe models but to extend and adapt their capabilities to the broader needs of modern banking risk management.

Common Domain Model: The Cornerstone of Physical Risk Management

The Common Domain Model (CDM) originated as a collaborative effort by three significant banking trade associations: the International Swaps and Derivatives Association (ISDA), the International Capital Market

Association (ICMA), and the International Securities Lending Association (ISLA). These organisations developed the CDM to standardise the representation of banking products, lifecycle events, and processes, aiming to streamline operational efficiency, enhance data consistency, and reduce costs and fragmentation within the banking markets.

In September 2022, ISDA, ICMA, and ISLA selected the Fintech Open Source Foundation (FINOS) to host and manage the CDM as an open-source project, marking a critical shift towards fostering community-driven innovation and widespread industry adoption. The transition to FINOS included creating a centralised repository under the Community Specification Licence to allow banking institutions and other stakeholders to collaborate transparently and inclusively.

FINOS officially launched the CDM project in February 2023, integrating it fully into its ecosystem. This collaboration aligns with FINOS's broader mission to promote open-source standards for interoperability, regulatory alignment, and innovation in banking services. The CDM has since evolved into a key tool for automating trade processing, managing banking transactions, and supporting digital frameworks like smart contracts and regulatory reporting.

Core Components of the CDM

The CDM defines and organises the banking transaction lifecycle into the following hierarchical models to ensure efficiency, transparency, and standardisation:

- **Product Model:** Describes banking instruments like contracts used to transfer banking risk.
- **Event Model:** Represents lifecycle events such as trade executions and modifications.
- **Process Model:** Lays the groundwork for automating and standardising industry processes.
- **Reference Data Model:** Captures foundational details such as parties, legal entities, and rate indexes essential for modelling trades and agreements.

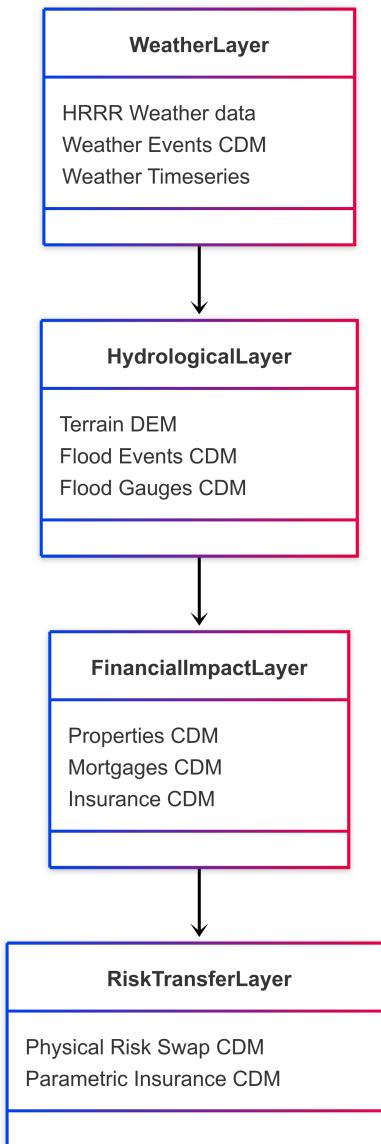


Figure 4: Weather to Risk Transfer via CDM

Purpose and Benefits

The introduction of the CDM addresses several challenges in the banking markets:

- **Interoperability and Automation:** By creating a common language, the CDM reduces the variations in how firms record trade lifecycle events, enabling straight-through processing and minimising operational inefficiencies.
- **Regulatory Compliance:** The standardised model promotes consistency and transparency in regulatory reporting, easing compliance and increasing alignment between banking institutions and regulators.
- **Innovation Acceleration:** The CDM provides a robust foundation for integrating emerging technologies like distributed ledger technology (DLT), smart contracts, and artificial intelligence into banking markets.

Scope of CDM Application

The development of CDM capability operates under Fino's governance, which is developed through a “community-driven governance model”; the CDM is openly accessible under the Community Specification Licence.

Working groups of industry practitioners and experts oversee its development and evolution, ensuring the model remains adaptable and relevant to industry needs.

The CDM is implemented across various use cases, including:

- Processing OTC derivatives, cash securities, commodities, and securities financing (e.g., repos and securities lending).
- Serving as the backbone for ISDA's Digital Regulatory Reporting (DRR), streamlining the creation of human-readable and machine-executable rules for regulatory compliance.
- Enhancing efficiency in post-trade processes such as collateral management and lifecycle management of banking products.

The CDM represents an industry-wide effort to establish a standardised operational framework that reduces manual processes while introducing clarity and efficiency in banking markets. This leads to solutions with the following critical positive features.

- **Transparency:** Open structure and standardisation ensure clarity in data representation.
- **Scalability:** Designed to handle vast transaction volumes efficiently.
- **Technology Agnostic:** Supports integration into various systems and platforms.

The MKM Framework and the CDM

The first layer handles weather pattern prediction. Whether using AI-enhanced forecasting models or traditional meteorological approaches, all weather events must conform to the Weather Events CDM structure. This standardisation enables consistent interpretation of weather predictions across different modelling platforms and institutions.

The second layer translates weather patterns into flood events through hydrological modelling. Here, the Flood Events and Flood Gauges CDM structures ensure that the outputs are consistently structured and interpretable regardless of the specific model used, such as open-source LISFLOOD or any proprietary alternatives. This standardisation is particularly crucial for calibration against historical gauge measurements.

The third layer maps physical events to banking impacts using:

- **Property** = property characteristics.
- **Mortgage** = collateralise loan terms.
- **Insurance** = property insurance.

This standardisation ensures that vulnerability assessments and damage calculations can be consistently applied across portfolios and institutions, even when using different underlying models for damage estimation.

The fourth layer enables risk transfer through the Physical Risk Swap CDM structure. This standardisation is essential for developing liquid markets in physical risk transfer instruments, ensuring all participants understand the underlying risk measures and transfer mechanisms, primarily through derivatives.

Derivatives serve as powerful banking instruments in capital markets, enabling the efficient transfer of risk between parties. By allowing businesses, investors, and institutions to hedge or speculate on price movements, interest rates, credit defaults, and other banking variables such as physical risk, derivatives redistribute risk to those more willing or capable of managing it.

The key benefits of derivatives, which are well-established in debt markets under CDO/CDS structures, make the transition to MBS/PPRS structures entirely natural for the banking market. The management of physical risk can follow this natural progression and reap the following benefits from the outset:

- **Risk Allocation:** Transfers risk to those most capable of managing it (e.g., insurers, hedge funds) at a competitive price. Only insurers cover physical risk.
- **Capital Relief:** Banks reduce regulatory capital requirements by transferring credit risk via derivatives.
- **Diversification:** Investors access exposures (e.g., the concentration of mortgages in flood-prone areas) that would otherwise be difficult to obtain.

Throughout all layers, model risk governance is paramount. Compliance with SR 11-7 principles ensures that models are:

- Properly validated and documented.
- Based on sound theoretical foundations.
- Regularly reviewed and updated.
- Subject to appropriate controls and oversight.
- Transparent in their assumptions and limitations

The MKM framework Key principles:

- **Model Independence:** While requiring compliance with CDM structures and governance standards, the framework remains model-agnostic. Institutions can choose or develop models that best suit their needs, provided they interface correctly with the CDM structures.
- **Data Standardisation:** The CDM defines all critical attributes across the seven key components. This standardisation enables consistent data collection, validation, and exchange across institutions.
- **Governance Integration:** Model risk governance is built into the framework rather than added as an afterthought. Compliance with standards like SR 11-7 is inherent in the framework's design; otherwise, adoption by banking firms is a non-starter.
- **Market Development:** The standardised CDM structures, particularly for Physical Risk Swaps, provide the foundation for developing liquid markets in physical risk transfer instruments.

- **Implementation Flexibility:** While the CDM structures are standardised, practitioners can tailor their implementation to specific institutional needs and regulatory requirements. This flexibility ensures the framework can adapt to different contexts while maintaining consistency.

This integrated approach transforms physical risk management from disparate models and approaches into a coherent, governed framework built on standard definitions and structures. This foundation enables the banking sector to develop more sophisticated approaches to measuring, managing, and transferring physical risk.

The following chapters examine each framework component in detail, from the specific attributes defined in each CDM structure to the practical implementation of compliant models. We focus on how standardisation and governance enable more effective risk management while supporting market development.

Chapter 2 - Data Framework

The banking industry's approach to physical risk assessment stands at a critical juncture where the need for standardised data models intersects with increasingly complex environmental challenges. The Common Domain Model (CDM) emerges as a crucial framework for addressing this complexity, particularly in the context of flood risk assessment and management - one of the most significant physical risks facing banking institutions today.

"The Common Domain Model represents more than a technical standard - it bridges weather time series creation and banking markets, physical reality and banking value, data and decision. Without this bridge, we cannot build the market infrastructure needed for managing physical risk in the decades to come." - David Kelly, MKM.

Traditional approaches to physical risk assessment have been fragmented, with different institutions developing proprietary models and data structures. This fragmentation has created inconsistent risk evaluations across institutions, made it challenging to aggregate and compare risk assessments, increased operational costs, and limited the ability to share and validate models across the industry. The CDM addresses these challenges by providing a standardised framework that enables consistent representation of physical risk data and models across the banking sector.

The CDM's extension into physical risk assessment builds upon established flood damage modelling frameworks, demonstrating the critical importance of standardised approaches. Recent research has shown that flood damage models require the integration of hazard parameters, exposure data, and vulnerability information - all of which align closely with CDM principles. By standardising the representation of environmental data, asset characteristics, and susceptibility to damage, the CDM creates a common language for physical risk assessment.

Modern risk assessment through the CDM framework offers significant advantages for banking institutions. It enables consistent modelling of physical risk parameters and facilitates integration with existing economic models. The framework supports regulatory reporting requirements while seamlessly integrating national and local-scale models.

Perhaps most importantly, it improves risk aggregation and analysis capabilities, supporting more accurate pricing of physical risk in banking instruments and facilitating the development of new risk transfer mechanisms.

The evolution of CDM in physical risk assessment will support integration with advanced AI and machine learning techniques, the development of standardised risk transfer instruments, and improved collaboration between banking institutions and climate science organisations. As physical risks continue to intensify and reshape banking markets, the CDM provides the foundation for a more resilient and adaptable approach to risk management.

The following sections explore these aspects in detail, providing practical implementations and case studies demonstrating the CDM's application in physical risk assessment and management. We will examine how banking institutions can leverage the CDM to bridge the gap between physical science and banking mathematics, between property-level analysis and systemic risk management, and between current challenges and future solutions.

A Quick History of ISDA

Given that ISDA will be referenced frequently, this is an opportune moment to provide an introduction to this industry body and its current form, including outreach.

The International Swaps and Derivatives Association (ISDA) has a rich history that dates back to its foundation in 1985. It was established to standardise and streamline the over-the-counter (OTC) derivatives market, addressing the risks and inefficiencies associated with bespoke and unregulated transactions.

ISDA was initially formed to provide a framework for the growing derivatives market, publishing its first industry document, the Code of Standard Wording, Assumptions, and Provisions for Swaps (SWAPS Code).

The organisation released its first standardised agreements in 1987, such as the Interest Rate and Currency Exchange Agreement and the

Interest Rate Swap Agreement. These documents established consistent terms for single-currency and multi-currency derivatives.

The 1992 ISDA Master Agreement significantly evolved the banking industry, creating a comprehensive framework for Over-The-Counter (OTC) derivatives transactions. This was followed by an updated 2002 ISDA Master Agreement incorporating lessons from market crises in the late 1990s.

ISDA played a central role in resolving risks during the Russian banking crisis, the Asia banking meltdown, and the collapse of firms like Lehman Brothers. Its documentation provisions, such as close-out netting and collateral management, helped stabilise the market during these crises.

Beyond contracts, ISDA introduced initiatives like the banking products markup language (FpML) for efficient data exchange in derivatives trading and protocols for handling specific market changes (e.g., credit events).

Today, ISDA serves over 1,000 member institutions across 76 countries, including banks, asset managers, corporations, and government entities. ISDA collaborates with regulators and policymakers to promote market transparency and risk reduction.

ISDA continues to adapt to global banking changes. For example, it introduced updated terminology with the 2014 ISDA Credit Derivatives Definitions and has been a proponent of centralised clearing to manage counterparty risk.

Origins of the Common Domain Model (CDM)

The Common Domain Model (CDM) represents one of modern banking technology's most significant collaborative efforts. Spearheaded by the International Swaps and Derivatives Association (ISDA) and developed in partnership with key industry bodies, including ICMA, ISLA, and FINOS, the CDM has emerged as a transformative force in banking data standardisation.

At its core, the CDM provides an open-source data model that creates standardised digital representations of events and processes throughout the lifecycle of banking products. What began as an initiative focused on derivatives has evolved into a comprehensive framework encompassing bonds, loans, and an expanding array of banking instruments. The 2022 release of CDM Version 5.0 marked a significant expansion of this coverage, incorporating sophisticated modelling capabilities for interest rate derivatives, equity derivatives, credit derivatives, bonds, repos, and securities lending.

The development of the CDM follows a unique collaborative approach. Through GitHub, ISDA members, market infrastructures, and third-party vendors work together to evolve and refine the model. This open-source methodology ensures that the CDM remains responsive to industry needs while fostering innovation and adaptation. The result is a living framework that grows with the industry rather than a static standard imposed from above.

The International Swaps and Derivatives Association (ISDA) initiated the Common Domain Model (CDM) project.

ISDA launched the CDM to standardise how banking products, notably derivatives, are represented and processed. The goal was to create a machine-readable, open-source model that ensures consistency across trading, risk management, and regulatory reporting.

The project was first introduced around 2017, and ISDA has been collaborating with technology firms, banking institutions, and regulators to refine and expand its use. The CDM is particularly relevant for smart contracts, distributed ledger technology (DLT), and automation in banking markets.

The International Capital Market Association (ICMA) and the International Securities Lending Association (ISLA) joined ISDA's Common Domain Model (CDM) initiative in 2020.

- **ICMA** joined in June 2020 to extend the CDM's coverage to repo and bond markets.
- **ISLA** joined in October 2020 to bring securities lending into the CDM framework.

Their participation helped expand the CDM beyond derivatives, creating a standardised digital representation of transactions across repo, securities lending, and bond markets.

The MKM founders recognised early on that the adoption of CDM from the outset was critical to facilitate the integration of these non-traditional datasets into the bank's risk architecture. For the first time since CDM's inception, we have data considerations (alongside model governance), including standardisation from the outset.

The Power of CDM

The CDM's true power lies in its ability to create a common language for banking institutions, market infrastructures, and regulators. By establishing standardised ways to represent trade and transactional lifecycle events and processes, the CDM enables seamless communication across different platforms and organisations. This standardisation drives operational efficiency, reduces data translation and reconciliation costs, and creates new opportunities for automation and innovation.

The adoption of CDM is voluntary, yet its value proposition has driven significant uptake among major banking institutions, including JP Morgan, Morgan Stanley, Goldman Sachs, and BNP Paribas. While implementation remains optional, CDM's influence across the industry continues to expand as organisations recognise its benefits: streamlined data management, enhanced automation capabilities, and simplified regulatory reporting. Success in traditional banking products has established a foundation for CDM's extension into emerging domains, including physical risk assessment.

Looking ahead, the CDM's role in physical risk modelling promises to be transformative. By providing a standardised framework for representing physical risk data, the CDM can help bridge the gap between climate academic research and banking markets, enabling more efficient risk transfer mechanisms and supporting the development of new banking products designed to manage physical risk exposure.

Risk Disclosure and the Headache for Banks

Risk Disclosure, as required by the regulators and supervisory bodies, has expanded in the years after the 2008 market disruption. The direction of travel is set, and no regulator will get promoted by reducing the burden of disclosure reporting. The cost of compliance is a significant drain on all banks and presents a barrier to entry for new asset classes such as banks. Dealing with the cost of regulation is another aspect of the MKM Framework for physical risk.

Banks' disclosure reporting costs encompass direct banking expenditures and indirect operational challenges driven by regulatory complexity, technological demands, and evolving compliance requirements. These costs arise from mandatory reporting standards, data management infrastructure, personnel expertise, and potential penalties for non-compliance. Key findings from industry studies and regulatory sources highlight several critical dimensions:

Direct Costs of Compliance

- Banks face rising costs due to increasingly granular reporting requirements. For example, the SEC's 2020 overhaul of bank disclosure rules introduced new demands for detailed metrics like loan maturity distributions and allowance for credit losses, requiring system upgrades and data validation processes. Similarly, the EU's IReF framework mandates extensive parallel reporting runs, doubling infrastructure costs during transitional phases.
- Modernising legacy systems to handle large-scale data processing (e.g., Basel III's standardised disclosure templates) incurs significant upfront costs. Investments in cloud infrastructure, automated workflows, and tools for data lineage tracking are critical but costly.
- Skilled compliance teams and IT specialists are essential for interpreting regulations and ensuring reporting accuracy. The FDIC estimates that preparing mandatory banking disclosures alone requires banks, as noted in Ruth Picker's " Applying IFRS Standards," to pass on an estimated 1.5% in higher loan spreads to cover the labour and technical costs of compliance. Additionally, third-party legal and audit services further amplify expenses.
- Disclosure mandates can distort lending behaviour. For instance, banks exempted from Community Reinvestment Act (CRA) reporting after 2005 reduced small-business lending in low-income areas, prioritising affluent markets to avoid reputational risks.
- Failure to meet standards like GDPR or Dodd-Frank triggers fines and legal fees. For example, Equifax's 2017 data breach resulted in \$700 million in settlements and lost customer trust.

CDM and Regulatory Disclosure

Integrating new datasets into existing bank systems has historically been complex and enormously costly, especially when the motivation is to satisfy supervisory bodies. Physical risk data presents particular challenges, requiring institutions to connect diverse and fragmented data sources, ranging from weather pattern models to property specifications.

However, the CDM's structured approach to data representation offers a solution that extends beyond mere standardisation—it creates a direct path to regulatory compliance through the FINOS Digital Disclosure Reporting framework.

The power of this approach lies in its end-to-end workflow design. The CDM enables a comprehensive data flow from initial event simulation to final risk valuation for physical risk assessment. The workflow begins with Event_CDM, which standardises the representation of extreme weather

event pathways, leveraging advanced AI platform components. Weather pattern parameters such as temperature, pressure and precipitation feed into terrain modelling through standardised digital elevation models, followed by flood event creation based on elevated gauge levels.

Standardising Critical Datasets

The CDM standardises damage and vulnerability curves at the property level, enabling consistent valuation impacts across immediate property valuation loss, mortgage valuation impairment, heightened insurance premiums, and triggering events for physical risk swap payments. This standardisation ensures that all risk-relevant data flows seamlessly into regulatory reporting frameworks.

The substantial efficiency gains come from how this standardised data structure aligns with existing regulatory regimes. By implementing the CDM correctly, institutions can automatically feed their physical risk assessments into:

- Basel III requirements.
- BCBS 239 Risk Principles.
- SA-CCR for derivatives and long-term settlement.
- BCBS 530 Physical/Climate Risk Principles.

This alignment means that once physical risk data is represented in CDM form, regulatory disclosures become a natural extension of existing processes rather than a separate compliance burden. The CDM acts as a bridge between new physical risk data and established regulatory frameworks, ensuring that data integration investments serve risk management and compliance needs.

Integrating with Credit Risk Flow Process

In the credit risk workflow that is notoriously complex, the CDM acts as a HUB to the spoke of 100s of legacy systems and thus decouples the integration of new, often incompatible Physical risk data from 100s of separate data integration headaches that paths the way to methodological integration of physical risk factors into core banking systems. This systematic approach allows banks to calculate the delta between traditional risk metrics and those incorporating physical risk data, providing transparent justification for risk adjustments that regulators increasingly demand.

“The FINOS and their open-source framework further amplify these benefits by allowing institutions to leverage shared effort, knowledge and interpretations. Rather than each bank developing its mapping between physical risk data and regulatory requirements, the industry can collaborate on standardised approaches, mutualising development costs and ensuring consistency in regulatory submissions.” - Johnny Mattimore, MKM.

This integration framework demonstrates why getting the CDM implementation right from the start is crucial. A well-structured CDM implementation solves today's data integration challenges and creates a foundation for efficient regulatory reporting that can evolve with changing requirements and expanding physical risk datasets.

Regulatory Framework and Climate Risk Disclosure

While physical risk represents an immediate challenge rather than a future climate consideration, banking institutions must navigate an evolving regulatory landscape focused on various disclosure drivers across multiple jurisdictions - from immediate hazard events to long-dated weather and climate scenarios. The CDM framework provides a structured approach to meeting these regulatory requirements while addressing current physical risk challenges.

Central Banks and Stress Testing

Central banks and banking regulators have responded to political pressure and have developed comprehensive frameworks for assessing and managing climate-related risks. The Bank of England exemplifies this approach and provides insight into how regulators consider climate risk and their demand on banks for further disclosure reporting.

The Author believes that Physical risk has created a present issue for banking and, therefore, does not need to participate in the climate discussions. The reality is that the narrative that the future will create more severe and frequent flood events prevails with politicians and their regulators, so all disclosure reporting for physical risk comes under climate risk.

The Climate Biennial Exploratory Scenario (CBES) is a stress-testing exercise conducted by the Bank of England to assess how climate change

might impact the UK banking system. First implemented in 2021, it evaluates banks' and insurers' resilience to transition risks (economic shifts toward a low-carbon economy) and physical risks (direct consequences of climate change like floods and fires).

The CBES examines three potential scenarios over a 30-year horizon: Early Action (immediate policy implementation limiting warming to 1.8°C by 2050), Late Action (delayed implementation causing disorderly transitions), and No Additional Action (leading to 3.3°C warming and severe physical consequences).

Key findings reveal that climate risks could reduce banking institutions' profits by 10-15% annually, with the No Additional Action scenario projecting cumulative losses of up to £334 billion by 2050. Under this worst-case scenario, approximately 7% of UK households could become uninsurable due to escalating climate risks.

While banking institutions are progressing in climate risk management, significant improvements in modelling capabilities and data collection are needed. The Bank of England plans to refine future CBES exercises to understand the banking system's vulnerabilities better and support the transition to a low-carbon economy.

Extension of CDM to Physical Risk and Physical Risk Swaps

Standardisation is the foundation of effective physical risk management in banking markets. The Common Domain Model (CDM) is core to this framework, which provides standardised definitions and structures for seven key components: Weather Events, Flood Events, Flood Gauges, Properties, Mortgages, Insurance, and Physical Risk Swaps. This standardisation is not merely a technical convenience—it is the bedrock upon which banking practitioners can build all subsequent risk measurement and management capabilities.

Consider a property, for example. The CDM defines the critical attributes for any property assessment beyond valuation: location and physical characteristics and their relationship to terrain, flood defence structures, and historical events. This standardised definition ensures that whether a bank uses the model for mortgage valuation or an insurer for policy pricing, they work with consistent, comparable data structures. This standardisation reaches the first stage of the ultimate risk transfer mechanism in the form of Physical Risk Swaps, which are recognised by all actors in the banking markets.

Based on this CDM foundation, the framework supports a variety of modelling approaches, both public and proprietary. For instance,

institutions might decide to implement the public Joint Research Centre (JRC) spatially distributed hydrological model, LISFLOOD, to simulate hydrological processes such as rainfall-runoff, flood forecasting, and water resource assessments over large transnational river catchments.

The crucial point is not which specific models are used but that they interface with the standardised CDM structures and comply with model risk governance requirements, particularly the principles outlined in SR 11-7.

- **Property-level:** risk assessment requirements demand a granular analysis of physical risk exposure. The CDM's Property structure standardises this assessment, ensuring consistent evaluation and reporting of risk factors such as elevation, flood defence characteristics, and historical exposure.
- **Regulatory Stress:** testing scenarios increasingly examine mortgage portfolio vulnerability. The CDM's Mortgage component enables institutions to model how physical risks affect property values and default probabilities, providing the detailed analysis regulators demand.
- **Insurance Availability:** and pricing are emerging as critical regulatory concerns. The CDM's Insurance structure standardises the representation of coverage, premiums, and market capacity, facilitating the reporting of insurance-related risks.

Developing Physical Risk Swaps through the CDM framework also anticipates regulatory needs for risk transfer mechanisms. By structuring these instruments similarly to Credit Default Swaps (CDS), the framework leverages existing regulatory frameworks while addressing new risk transfer requirements.

Technical Implementation and Core Components

The technical implementation of the CDM framework focuses on three foundational components that form the basis for physical risk assessment and management:

Property CDM: Key Utility Components

The CDM captures critical valuation factors through:

- Detailed property attributes (size, bedrooms, construction quality).
- Transaction history with actual sale prices and rental income.
- Property type and classification for comparative analysis.

- Income generation potential (rental, holiday let status).

Flood risk assessment:

- Geolocation data (latitude/longitude, elevation, BNG references).
- Environment Agency flood zone classification.
- Distance measurements to water bodies (rivers, lakes, coasts, canals).
- Historical flood events with dates and severity.
- Information around return period predictions (e.g. flood height for a one-in-100-year event).

Implemented resilience measures:

- Physical barriers (flood panels, airbrick covers, waterproof walls).
- Structural modifications (raised foundations, reinforced walls).
- Drainage systems and water diversion measures.
- Natural solutions (rain gardens, strategic planting, bioswales).

Construction specifications:

- Building materials and methods (brick/block, timber frame).
- Foundation types and floor construction.
- Wall construction and insulation status.
- Age and period of construction.
- Recent renovations and improvements.

Energy performance data:

- EPC and carbon ratings.
- Annual energy consumption by source (gas, electricity, renewables).
- Building fabric thermal performance (insulation, glazing).
- Heating systems and renewable installations.
- Energy costs and potential savings calculations.

Contextualises property risk:

- Urban/rural classification.
- Local authority and administrative boundaries.
- Ground conditions (soil type, subsidence risk, contamination).
- Environmental quality indicators (air, water, noise).
- Proximity to infrastructure and services.

This Property CDM structure creates a standardised, interoperable data framework that supports accurate risk assessment, valuation modelling, and resilience planning across the property finance ecosystem. It provides the foundational data layer for mortgage risk calculations,

insurance pricing, and physical risk transfer mechanisms, enabling more accurate climate adaptation investment decisions.

Mortgage CDM: Key Utility Components

Risk-adjusted Valuation Framework:

- Loan-to-value calculations that incorporate physical risk factors.
- Standardised property valuation methodologies accounting for resilience features.
- Temporal adjustments to reflect changing risk profiles over loan lifetime.
- Comparative valuation benchmarks across risk categories.

Default probability adjustments:

- Correlation modelling between climate events and payment delinquency.
- Stress testing frameworks for extreme weather scenarios.
- Risk segmentation methodologies for portfolio management.

Insurance availability:

- Premium escalation modelling and affordability thresholds.
- Coverage gap identification and risk quantification.
- Integration with Flood Re and other market-based insurance schemes.

Collateral Risk Evaluation:

- Repair cost estimation standardisation for common damage types.
- Residual value calculation methodologies post-event.
- Refinancing risk assessment due to insurance shifts.

System Integration Standards:

- Data exchange protocols with existing credit risk systems.
- Reporting frameworks aligned with regulatory requirements.
- Portfolio aggregation methods for institution-wide exposure assessment.
- Stress testing integration with existing banking stability frameworks.

Risk Transfer Enablement:

- Standardised categorisation of risk profiles supporting securitisation.
- Data requirements for creating physical risk swaps Definitions of trigger events and verification methodologies Standardization of settlement processes and documentation.

The Mortgage CDM structure creates a robust framework for incorporating physical risk into mortgage finance operations. It enables more accurate pricing of climate-related risks, supports the development of new risk transfer instruments, and facilitates integration with existing credit risk management systems. This standardisation helps banking institutions better understand their exposure to physical climate risks and develop appropriate mitigation strategies.

Parametric Flood Insurance CDM: Key Utility Components

Objective Trigger Mechanisms

- Standardised measurement units for consistent applications.
- Precise specification of trigger types across contracts.
- Historical threshold exceedance tracking for frequency analysis.

Transparent Payout Structure

- Tiered payout levels linked to specific trigger thresholds.
- Predefined payout amounts eliminate subjective assessment.
- Graduated response reflecting increasing severity of events.
- Historical comparison data to contextualise trigger points.

Data Source Standardisation

- Detailed gauge specifications and certification status.
- Consistent measurement frequency and methodology.
- Clear data transmission protocols and accessibility.
- Maintenance and operational status tracking.

Contract Clarity and Efficiency

- Standardised policy structures for rapid deployment.
- Clear identification of covered properties and limits.
- Explicit premium allocation for flood coverage.
- Defined contract terms and status tracking

Risk Transfer Transparency

- Standardised format enabling easy comparison between policies.
- Historical frequency data supporting accurate pricing.
- Sensor reliability metrics for risk assessment.
- Location-specific sensor data for precise risk evaluation.

Integration Capabilities

- Link to Property CDM through PropertyID reference.
- Connection to third-party data sources through standardised interfaces.
- Compatibility with existing insurance documentation.
- Support for automated claims processing.

The Parametric Flood Insurance CDM establishes a foundation for innovative insurance products that significantly enhance the efficiency of flood risk transfer. Eliminating the subjectivity and delays associated with traditional claims assessment facilitates quicker payouts and more accurate risk pricing. This standardisation promotes the growth of insurance markets in high-risk areas where traditional coverage may be unaffordable or unavailable.

The structure also facilitates integration with the broader CDM framework, enabling banking institutions to accurately incorporate insurance information into mortgage risk assessments and property valuations. This creates a complete picture of physical risk exposure across the banking system.

Physical Risk Swap CDM: Key Utility Components

Risk Transfer Framework

- Counterparty specifications and legal entity identifiers.
- Standardised documentation aligned with ISDA/FIA conventions.
- Consistent reference data for physical risk exposures.
- Compatible structure with existing derivatives infrastructure

Trigger “Default” Definitions

- Event classification taxonomy (flood, hurricane, wildfire, etc.).
- Severity threshold specifications mirroring parametric insurance.
- Duration and spatial extent parameters.
- Data source authentication protocols.

Settlement Mechanisms

- Cash settlement calculation methodologies.
- Physical delivery specifications for applicable scenarios.
- Collateral valuation adjustments for impaired assets.
- Netting and aggregation rules across multiple contracts.

Market Infrastructure Support

- Clearing compatibility specifications.
- Trade reporting requirements aligned with regulatory frameworks.
- Transaction lifecycle event processing.

- Position reconciliation methodologies.

Pricing Transparency

- Pricing models and inputs.
- Mark-to-market valuation methodologies.
- Risk premium calculation approaches.
- Historical data requirements for backtesting.

Regulatory Alignment

- Capital requirement calculation methods.
- Risk-weighted asset determination.
- Disclosure and reporting specifications.
- Stress testing integration protocols.

The Physical Risk Swap CDM creates a crucial bridge between traditional banking markets and physical risk management. Adopting familiar CDS-like structures reduces the implementation barrier for market participants and leverages existing trading infrastructure. This standardisation supports the development of liquid markets for physical risk transfer, enabling more efficient risk distribution across the banking system.

The MKM framework integrates seamlessly with the Property, Mortgage, and Parametric Insurance CDM components to create a comprehensive physical risk management ecosystem. Property data feeds into risk assessment, mortgage exposure determines hedging requirements, and parametric triggers align with swap settlement conditions.

These components interact through standardised interfaces, enabling the development of Physical Risk Swaps that mirror CDS structures. This approach allows institutions to:

- Price physical risk transfer accurately.
- Create standardised risk transfer documentation.
- Have legal certainty with ISDA-created.
- Establish clear trigger events and settlement processes.
- Enable market-making and trading.
- Support regulatory reporting requirements.

The implementation leverages existing banking market infrastructure while introducing new elements specific to physical risk. For example, the framework incorporates flood modelling outputs from systems like LISFLOOD while maintaining familiar banking product structures that facilitate market adoption.

High-Resolution Rapid Refresh (HRRR): Advanced Weather Prediction for Physical Risk Assessment

Assessing physical risks demands increasingly granular and timely weather data. Enter the High-Resolution Rapid Refresh (HRRR) model, a breakthrough in operational weather prediction that has transformed our ability to capture and analyse atmospheric conditions across the United States. As a cornerstone of modern weather forecasting infrastructure, HRRR represents a significant advancement in spatial resolution and update frequency, making it an invaluable tool for physical risk assessment.

Technical Foundation and Capabilities

At its core, HRRR is a convection-allowing numerical weather prediction model operated by the National Oceanic and Atmospheric Administration (NOAA). Its remarkable spatial resolution of 3 kilometres across the contiguous United States sets it apart, providing a level of detail that captures local weather phenomena crucial for risk assessment. This granularity allows risk managers to identify and evaluate location-specific weather patterns that might affect physical assets, from individual buildings to entire infrastructure networks.

The model's "rapid refresh" capability delivers frequently updated forecasts, enabling near-real-time risk assessment and decision-making. This combination of high-resolution and frequent updates makes HRRR particularly valuable for analysing chronic physical risks, such as severe weather events that could impact property portfolios or infrastructure investments.

Cloud-Native Architecture and Data Accessibility

The evolution of HRRR data architecture reflects the broader shift toward cloud-native computing in banking risk assessment. Through NOAA's Big Data Program, HRRR output is now available on Amazon Web Services (AWS) in two formats: the traditional GRIB2 format and the modern cloud-optimised Zarr format.

The evolution of the High-Resolution Rapid Refresh (HRRR) data architecture exemplifies the broader transition toward cloud-native computing in various industries, including banking risk assessment.

Through the NOAA Big Data Program (BDP), HRRR output is now available on Amazon Web Services (AWS) in the traditional GRIB2 and modern cloud-optimized Zarr formats.

Comparison of GRIB2 and Zarr Formats:

GRIB2 Format

- A long-standing format used for compressing and distributing large numerical weather model outputs as two-dimensional grids.
- While GRIB2 files are efficient for storage, they pose challenges for users needing frequent, high-throughput access. Files can be hundreds of megabytes and require substantial memory and processing resources to extract relevant data.
- GRIB2 remains essential for applications requiring complete datasets, such as initialising high-resolution simulations requiring all variables at multiple levels.

Zarr Format

- Created to address the limitations of GRIB2, Zarr is a cloud-native, chunked, and compressed storage format designed for high-speed access and ease of use with open-source software, such as xarray
- HRRR Zarr archives reformat GRIB2 output into smaller, optimised chunks (as small as 1 MB), reducing processing and access time by approximately 40 times compared to GRIB2 files. This primarily benefits machine learning, data analysis, or real-time forecasting applications.
- Zarr data is organised hierarchically in an AWS Simple Storage Service (S3) bucket, allowing users to access only the necessary subdomains or variables, significantly reducing data transfer times and costs.

Implications for Cloud-Native Computing

The NOAA/AWS HRRR initiative highlights the growing reliance on cloud-native architectures to improve accessibility, scalability, and efficiency in handling large datasets:

- **Efficiency Gains:** Cloud computing resources allow faster analysis and higher throughput, crucial in time-sensitive applications such as weather forecasting or banking modelling.

- **Scalability:** By leveraging AWS, the HRRR architecture accommodates growing data volumes—currently exceeding 145 terabytes and growing daily—with minimal operational overhead.
- **Broader Applications:** The HRRR Zarr format lends itself well to use cases requiring high-velocity data access, such as extreme weather event analysis and predictive modelling, which mirrors banking services' adoption of cloud-native architectures to handle complex calculations like credit risk analysis.

This dual-format data availability demonstrates a practical strategy for connecting legacy systems with modern cloud-native infrastructures, enabling various sectors to meet their computational and analytical needs better.

Efficient Graphical Subsetting

Efficient Geographic Sub-setting refers to the process of selecting a relevant subset of geographic data from a larger dataset in an optimised way.

It is commonly used in fields like geospatial analysis, geographic information systems (GIS), machine learning, and database management to extract only the necessary geographic data while minimising computational costs.

Key Concepts & Techniques

- **Spatial Indexing:** Using data structures like QuadTrees, R-Trees, or KD-Trees to quickly locate and extract relevant geographic regions without scanning the entire dataset.
- **Bounding Boxes:** Defining rectangular or polygonal boundaries around the area of interest to filter data efficiently.
- **Grid-Based Partitioning:** Dividing a geographic space into smaller cells and selecting only relevant ones, reducing the need to process unnecessary regions.
- **Geohashing:** Converting geographic coordinates into string representations that allow for quick lookup and comparison.
- **Sampling & Clustering:** Using techniques like k-means clustering or stratified sampling to extract representative geographic subsets.
- **Database Query Optimisation:** Leveraging spatial databases (PostGIS, Google BigQuery GIS, etc.) to run optimised queries that retrieve only the required geographic subset.

Applications

- **Mapping & Visualization:** Extracting only the relevant region for interactive maps.
- **Machine Learning on Geospatial Data:** Reducing data size before training models.
- **Satellite & Remote Sensing:** Processing only the necessary geographic area to save computational resources.
- **Logistics & Supply Chain:** Optimising delivery routes by selecting relevant city or regional subsets.

Processing of Specific Data Subsets.

Modern data science workflows involve multiple stages—data collection, modelling, visualisation, and deployment. Streamlined integration ensures that each step seamlessly connects, reducing inefficiencies, improving performance, and enabling automation.

Traditional data processing and analysis methods require manual data cleaning and transformation, significantly slowing the analytical process. Streamlined integration allows for automated ETL (Extract, Transform, Load) pipelines, reducing time spent on repetitive tasks. For instance, using Apache Spark or Pandas with cloud storage solutions like AWS S3 or Google BigQuery enables real-time data streaming, dramatically accelerating the initial phases of data science work.

Data science workflows frequently deal with big data, necessitating tools that efficiently query and subset information. Modern integration ensures compatibility between distributed storage systems such as Hadoop, Snowflake, and Delta Lake with machine learning frameworks like TensorFlow, PyTorch, and Scikit-Learn. A prime example is using SQL-based querying in Spark to preprocess large datasets before training machine learning models, allowing data scientists to work with volumes of data that would otherwise be unmanageable.

Data scientists, engineers, and analysts are required to collaborate effectively, and smooth integration enables these collaborative workflows. Integration with Git, DVC (Data Version Control), MLflow, and Kubernetes ensures reproducibility and precise experiment tracking across teams. This is exemplified when using MLflow to track hyperparameters and model performance across different training runs, enabling teams to understand which approaches yield the best results and ensuring that successful experiments can be reproduced consistently.

Businesses increasingly require real-time insights, as delayed processing often leads to missed opportunities. Modern workflows integrate streaming analytics platforms like Kafka, Flink, and Kinesis with AI models, enabling automated decision-making at scale. A compelling example is fraud detection in banking. AI models process transactions in real-time using event-driven architectures, allowing banking institutions to identify and prevent fraudulent activities before they complete rather than conducting post hoc analysis of suspicious transactions.

Scalable & Flexible Deployments streamlined integration allows for easy model deployment across cloud, edge, or hybrid environments. Frameworks like Docker, Kubernetes, FastAPI, and TensorFlow Serving significantly simplify the deployment process, enabling utilising the same models across diverse computing environments. For example, deploying a recommendation system via an API using FastAPI and serving models with TensorFlow Serving enables consistent, high-performance predictions regardless of where the system is hosted while maintaining the flexibility to scale resources based on demand patterns.

Feature	Spatial Dimension	Temporal Dimension
Represents	Location (where)	Time (when)
Measured in	Distance area volume	Seconds minutes days
Example Data	Latitude longitude depth	Timestamp event sequence
Common Uses	Maps 3D modelling GIS	Time series forecasting logs
Change Characteristics	Often fixed or slow-changing	Continuously evolving

Table 1: Format Comparison

Cost Optimisation and resource efficiency are other substantial benefits of integrated workflows, as they reduce unnecessary computation, cutting cloud computing costs substantially. Serverless architectures like AWS Lambda and Google Cloud Functions enable pay-as-you-go processing, eliminating the need for constantly running infrastructure that may sit idle during periods of low demand. A practical example is using serverless data pipelines instead of always-on virtual machines for preprocessing

tasks, which can dramatically reduce operational costs while maintaining or even improving performance for sporadic or unpredictable workloads.

Benefits

While quite generic, the benefits are worth listing. Cost Optimisation and resource efficiency are music to any head of IT at any bank.

- Saves time by automating repetitive tasks.
- Handles big data efficiently.
- Enables real-time AI-driven decisions.
- Improves collaboration across teams.
- Reduces costs by optimising cloud and compute resources.

Forecast Data Organisation and Structure

The HRRR-Zarr architecture employs a sophisticated organisational structure that balances analytical flexibility with computational efficiency. The data is systematically organised along multiple dimensions:

The primary organisation follows a hierarchical structure based on:

- **Model run Datetime:** Typically, this refers to the timestamp when executing a computational model. This timestamp is crucial for tracking, logging, and analysing model results over time. It is used in AI to track when a model was trained or made predictions.
- **Vertical Levels:** The HRRR model employs 51 vertical levels, consistent with the Rapid Refresh (RAP) model. These levels are hybrid, transitioning from a terrain-following sigma coordinate near the surface to isobaric levels in the mid and upper atmosphere. This approach reduces numerical noise in the upper levels, particularly mountainous regions.
- **Low-Level Heights:** The HRRR model's lowest level is approximately 8 metres above ground level (AGL), with the spacing between levels increasing from about 15-30 metres near the surface to 400 metres in the mid-troposphere and reaching about 700 metres near the tropopause. The model also provides outputs at specific heights, such as 10 metres and 80 metres AGL, which are commonly used in applications like wind energy analysis.
- **Forecast Adjustability:** HRRR data can be interpolated to desired heights above ground for specific applications (e.g., wind turbine hub heights), ensuring versatility for various sectors like renewable energy.
- **Analyses and Forecasts:** Represent the state of the atmosphere at the time of model initialisation. These analyses are derived using

advanced data assimilation techniques, integrating various observational datasets with model simulations to provide an accurate starting point. Forecasts are updated every hour and offer parameters such as precipitation, wind, temperature, and other meteorological variables.

- **Time-Lagged Ensemble (TLE) Approach:** A unique feature of HRRR operations is its use in producing Time-Lag Ensembles (TLEs). Sequential model runs are treated as ensemble members for a given valid time, allowing evaluation of forecast uncertainty and consistency. TLEs help forecasters assess trends and variability in predictions over time.

- **Physical Organization of Data as 2D Tiles for Analysis Files:** HRRR analysis files (denoted as FOO) are divided into 2D tiles, each covering a grid size of 150x150 points. This segmentation facilitates efficient data retrieval and minimises overhead for users seeking specific regions or parameters.

- **Physical Organization of Data as 3D Cubes for Forecast Files:** Forecast files are stored as 3D data cubes, where the three dimensions correspond to forecast lead time (XX), and spatial resolutions of 150x150 grid points for latitude and longitude.

- **Initialisation Times:** Times are set at 0000, 0600, 1200, and 1800 UTC, with a forecast duration extending up to 48 hours (XX=48). For other initialisation times, the duration is limited to 18 hours (XX=18). This organisation supports applications that require temporal and spatial forecasting across varying lead times.

Spatial vs. Temporal Dimensions in Data Analysis

Spatial and temporal dimensions represent fundamental frameworks for organising and analysing data across various scientific disciplines. While they serve distinct purposes, their integration often yields powerful insights that neither dimension could provide alone.

Understanding Spatial Dimensions

Spatial dimensions refer to physical space, typically represented through latitude, longitude, and altitude coordinates. These dimensions allow us to:

- Map geographic locations with precision.
- Model three-dimensional objects and environments.
- Analyse spatial relationships between different entities.

Geographic Information Systems (GIS), physics simulations, urban planning, and computer graphics are typical applications. For instance, a satellite image with precise coordinates allows researchers to track deforestation patterns across specific regions.

Understanding Temporal Dimensions

Temporal dimensions relate to time—when events occur, how long they last, and the sequence in which they unfold. These dimensions are crucial for:

- Tracking changes over time.
- Identifying patterns and cycles.
- Predicting future trends based on historical data.

Time series analysis, forecasting models, and historical tracking systems rely heavily on temporal dimensions. Stock market analysis, for example, depends on precise timestamps to identify trading patterns throughout the day.

Key Differences

The most potent analyses often emerge when spatial and temporal dimensions are combined, creating spatiotemporal frameworks that capture where and when phenomena occur. This integration enables:

- Climate scientists to track temperature changes across regions over multiple decades.
- Epidemiologists to monitor disease spread through populations and geographic areas.
- Transportation planners to optimise traffic flow based on historical patterns at specific locations.
- Environmental researchers to model pollution dispersion across both space and time.
- Self-driving vehicles represent a cutting-edge application of spatiotemporal analysis, continuously processing both location data and temporal patterns to navigate safely through changing environments.

By understanding how spatial and temporal dimensions interact, researchers can develop more comprehensive models that capture the full complexity of dynamic systems in our world.

Industry Standard Parameters and Variables

The HRRR model implements the World Meteorological Organization's (WMO) standardised parameter definitions, representing the foundational variables used throughout the weather industry as the global standard.

These parameters fall into several key categories:

Atmospheric State Variables

- Temperature (K) at multiple pressure levels and 2m above ground.
- Pressure (Pa) from surface to upper atmosphere.
- Relative Humidity (%) across vertical profile.
- Specific Humidity (kg/kg) for moisture content.
- Wind Components (m/s) - U and V vectors at multiple levels.
- Vertical Velocity (Pa/s) for atmospheric motion.

Precipitation and Hydrology

- Total Precipitation Rate (kg/m²/s).
- Accumulated Precipitation (kg/m²).
- Snow Water Equivalent (kg/m²).
- Frozen Precipitation Types (snow, sleet, freezing rain, graupel, hail).
- Precipitation Type Probabilities (%).
- Surface Runoff (kg/m²).

Radiative and Energy Parameters

- Downward Short-wave Radiation (W/m²).
- Upward Long-wave Radiation (W/m²).
- Surface Heat Fluxes (W/m²).
- Ground Heat Flux (W/m²).
- Cloud Radiative Properties.

Surface and Boundary Layer

- Surface Temperature (K).
- Soil Moisture Content (kg/m²).
- Surface Roughness (m).
- Planetary Boundary Layer Height (m).
- Surface Wind Components (m/s).
- Friction Velocity (m/s).

Derived Meteorological Products

- CAPE (Convective Available Potential Energy, J/kg).
- CIN (Convective Inhibition, J/kg).
- Storm Relative Helicity (m²/s²).
- Various Stability Indices.
- Echo Top Height (m).

- Radar Reflectivity (dBZ).

The Weather Industry's Universal Language

The HRRR model speaks the universal language of operational meteorology by implementing standardised parameters from the World Meteorological Organisation (WMO). These carefully defined variables represent decades of international collaboration and consensus in the weather industry, forming a comprehensive atmospheric measurement and prediction framework. This long-standing collaboration leading to standard metrics is where banks can stand on the shoulders of meteorological giants to advance their banking risk management.

State Variables

At the foundation of HRRR's capabilities lie the essential atmospheric state variables. The model tracks temperature variations across multiple pressure levels, from the surface to the upper atmosphere, while simultaneously monitoring pressure systems and humidity patterns. Wind patterns are captured through detailed U and V vector components, providing a complete picture of atmospheric motion at every level. These fundamental measurements form the backbone of all weather prediction and analysis.

The model's precipitation and hydrological parameters offer crucial insights for operational forecasting. Beyond simple rainfall measurements, HRRR provides a detailed analysis of precipitation rates, accumulation patterns, and precipitation type classifications. The system distinguishes between rain, snow, and frozen precipitation while tracking crucial hydrological factors like surface runoff and snow water equivalent. These parameters prove especially valuable for flood prediction and winter weather operations.

Energy Transfer

Energy transfer within the Earth's atmosphere is a dynamic process involving multiple pathways of radiation and surface-level exchanges. These mechanisms collectively govern Earth's energy balance and influence temperature regulation, particularly in urban and agricultural settings.

Shortwave radiation from the sun enters Earth's atmosphere, with about 30% reflected due to the albedo effect from clouds and surface features, while 70% penetrates further. Of this penetrating energy, approximately 20% is absorbed by atmospheric gases, and the remainder reaches Earth's surface as direct and diffuse radiation. This solar energy

varies regionally, with humid areas experiencing reduced solar heating due to cloud cover, while dry regions receive higher solar flux.

At Earth's surface, energy interacts through multiple processes: absorption converts solar radiation into heat energy; sensible heat flux transfers energy between the surface and atmosphere through conduction and convection; latent heat flux moves energy through evaporation and transpiration (especially in vegetated areas); and ground heat flux transfers energy into or out of the subsurface depending on temperature differentials between the surface and ground below.

Boundary Layer

The surface and boundary layer parameters provide critical information about the interface between Earth's surface and the atmosphere. HRRR measures surface temperature, wind patterns, soil moisture content, and surface roughness characteristics. Calculating the planetary boundary layer height offers crucial insights into pollution dispersion and convective potential, making these parameters especially valuable for air quality forecasting and aviation operations.

This parameter set embodies the weather industry's standardised approach to atmospheric science. The parameters maintain uniform definitions across different modelling systems worldwide, ensuring seamless integration with existing weather visualisation and analysis tools.

This standardisation extends beyond mere measurement to encompass computation methods and quality control procedures, making HRRR data immediately applicable within existing weather industry workflows. HRRR maintains backward compatibility with historical datasets through this standardisation while pushing the boundaries of weather prediction capabilities.

Integration with Risk Assessment Workflows

For banking institutions and risk managers, the HRRR data architecture offers several key advantages:

- **Efficient Data Access:** The ability to access data directly from specific geographic regions and time periods reduces storage requirements and processing overhead.
- **Scalable Analysis:** The cloud-native format enables parallel processing and distributed computing, essential for analysing large portfolios or conducting systematic risk assessments.

- **Real-time Integration:** The frequent update cycle of HRRR data supports near-real-time risk monitoring and assessment workflows.
- **Cost-effective Operations:** The ability to subset data and process it in parallel can significantly reduce computing costs and processing time.

The HRRR system represents more than just a weather model; it exemplifies the evolution of physical risk data architecture toward cloud-native, analysis-ready formats that support modern risk assessment requirements. As physical risk analysis advances, the accessibility and efficiency of HRRR data will play an increasingly important role in enabling sophisticated risk assessment capabilities.

The Case for Collaborative Physical Risk Data Infrastructure

Current Challenges for Individual Banks

Banking institutions face unprecedented challenges in assessing and managing physical risks from climate change. Unlike traditional banking risks, physical risks from events like floods, hurricanes, and wildfires present unique data challenges:

- **Geographic Data Complexity:** Physical risk assessment requires integrating diverse geospatial datasets (flood plains, elevation models, property boundaries) with banking exposure data. Most banks lack in-house GIS expertise to process these effectively.
- **Computational Intensity:** Running high-resolution weather, fluid dynamics models against mortgage portfolios demands significant computational resources
- **Data Acquisition Costs:** High-quality terrain data from specialised providers is expensive, often costing millions per year for comprehensive coverage, creating redundant spending across the sector.

Why Collaboration Makes Sense for Physical Risk

Physical risk data presents an ideal case for inter-bank collaboration because:

- **Pre-competitive Information:** Raw physical risk data (flood maps, temperature projections) is not a source of competitive advantage—

interpretation and integration with proprietary portfolios is where differentiation occurs.

- **Scale Economies:** The fixed costs of building and maintaining physical risk data infrastructure can be distributed across multiple institutions, making sophisticated analysis accessible to smaller lenders.
- **Standardisation Benefits:** A collaborative approach would foster standardised risk assessment methodologies, improving transparency and comparability for regulators and investors.
- **Reduced Duplication:** Currently, multiple banks independently purchase identical datasets from the same vendors, creating industry-wide inefficiency.

Practical Implementation Model

A collaborative physical risk data service could function as:

Component	Description	Key Benefit
Shared Data Lake	Centralized repository of climate models, historical event data, and geospatial information	Eliminates redundant data acquisition costs
Standardized APIs	Common interfaces for integrating physical risk data with proprietary systems	Simplifies technical integration
Federated Computation	Distributed processing allowing banks to run analyses against shared datasets without exposing portfolios	Preserves competitive information
Regulatory Alignment	Built-in compliance with emerging climate disclosure requirements (TCFD, ECB Guidelines)	Reduces compliance overhead

Table 2: Options for a Physical Risk Service

Evidence of Feasibility

Several precedents demonstrate this approach is viable:

- The Insurance Development Forum has already created a shared catastrophe modelling platform for insurers facing similar physical risk assessment challenges.

- The OS-Climate initiative brings together banking institutions to develop open-source climate risk tools, showing an appetite for pre-competitive collaboration.
- Banking regulators like the Bank of England and ECB are promoting standardised climate scenario analysis, creating natural incentives for shared infrastructure.

Addressing Potential Obstacles

While promising, several challenges must be addressed:

- **Governance Structure:** A neutral third-party or industry consortium model would be needed to ensure equitable access and management.
- **Cost Allocation:** Usage-based pricing models could ensure institutions pay proportionate to their benefit from the system.
- **Data Privacy:** Federated learning approaches can allow analysis without exposing sensitive portfolio details.
- **Integration with Legacy Systems:** Standardized APIs and data formats would be essential to overcome the technical fragmentation highlighted in your document.

Conclusion

The physical risks posed by climate change present a challenge that transcends individual institutions, which are often ill-prepared to tackle these issues on their own. Through collaboration on a unified data infrastructure for assessing physical risks, banks can conduct more sophisticated analyses at a lower cost while maintaining competitive advantages by applying these insights to their specific portfolios and business strategies.

The rising regulatory pressure around climate risk disclosure makes this collaboration increasingly urgent. Banks that delay may struggle to meet reporting requirements and accurately price climate-related risks in their mortgage and commercial real estate portfolios.

Chapter 3 - Weather Prediction

Weather pattern prediction lies at the heart of physical risk assessment for flooding events. However, bank's needs differ fundamentally from traditional weather forecasting. Rather than attempting to predict specific weather conditions at particular times, banks need to understand and characterise the distribution of possible weather patterns, including tails, that could lead to flooding events. Such a solution requires a sophisticated framework that combines classical weather pattern modelling with advanced artificial intelligence and statistical techniques.

"Weather prediction models for physical risk assessment must depart from traditional forecasting approaches.

Rather than seeking precise predictions of specific conditions, banks require a framework that maps the complete probability space of weather pattern evolution, with particular attention to those pathways that could lead to extreme precipitation events. This fundamental shift in perspective drives every aspect of the MKM Framework and modelling approach."- David Kelly, MKM.

This chapter presents our approach to weather pattern prediction, explicitly built for physical risk assessment in banking systems. We begin by examining the fundamental principles of weather pattern modelling, which are adapted and focused on the bank's particular needs in flood risk assessment. We then explore how artificial intelligence and Bayesian methods enhance predictive capabilities, allowing banks to extract more granular insights from weather data than traditional approaches allow.

The chapter continues with the MKM Framework for modelling weather pattern distributions, providing a complete characterisation of possible pattern evolutions with particular attention to those most relevant to flood risk. Finally, we examine the crucial process of

transforming weather patterns into precipitation time series that can drive subsequent hydrological modelling that then drives time series in gauge levels. It is essential to mention that time series of market data is at the heart of bank's risk systems.

Throughout the chapter, we carefully balance sophisticated mathematics and practical utility. While the underlying techniques may be complex, we focus on producing results to support meaningful risk assessment. Each component of the MKM Framework builds upon established physical principles while leveraging modern computational methods to extract additional insights relevant to flood risk assessment for banks.

The methods presented here provide the foundation for all subsequent banking risk analysis. We develop a robust and physically consistent approach to weather pattern prediction to create a solid base for meaningfully assessing physical risks in banking systems.

Weather Model Fundamentals

Understanding weather system dynamics is the foundation of weather pattern prediction. While traditional weather pattern models excel at capturing large-scale atmospheric and oceanic processes, MKM Framework requires a fundamentally different approach—one focused on the specific patterns and interactions that drive flood risk. This shift in perspective shapes how we construct and apply a new modelling framework to physical risk.

The approach outlined below builds upon the fundamental physical laws governing atmospheric behaviour—mass, energy, and momentum conservation. These principles form an immutable foundation, expressed through coupled partial differential equations describing the evolution of key atmospheric variables. However, where traditional models might prioritise global-scale predictions over decades or centuries, our framework focuses on the regional-scale patterns and interactions most relevant to flood risk assessment that aligns with weather patterning, so out to 5 days. This starkly contrasts climate models that do not predict changes in precipitation patterns.

The spatial resolution of our modelling framework reflects this specialised focus. We employ a nested grid structure that provides excellent resolution in regions of interest for flood risk assessment, similar to how risk managers drill down into risk datasets.

This approach allows us to capture crucial local-scale processes - orographic precipitation, convective storm development, and boundary layer interactions - while maintaining computational efficiency. The grid structure adapts dynamically, concentrating computational resources that prove most valuable for risk assessment.

Parameterisation

Parameterisation is one of the most challenging aspects of weather modelling. Many crucial processes, particularly those related to cloud formation and precipitation, occur at scales smaller than any practical model grid. Our framework addresses this challenge through physical parameterisations and machine-learning approaches that capture sub-grid-scale processes. This hybrid approach is valuable for representing the intense, localised precipitation events that often drive flooding risk.

The MKM framework demands careful attention to external forcings. Beyond the standard considerations of greenhouse gas concentrations and solar variations, we pay particular attention to factors influencing extreme precipitation events.

Land use changes, aerosol distributions, and soil moisture conditions receive explicit treatment, as these factors can significantly affect the development and intensity of extreme weather patterns.

Initialisation

Model initialisation presents unique challenges for our purposes. While traditional weather pattern models might initialise from long-term average conditions, our framework requires a precise representation of current atmospheric states. We employ sophisticated data assimilation techniques that combine observations from multiple sources - ground stations, satellites, and weather balloons - to construct initial conditions that capture the subtle atmospheric features that might influence weather pattern evolution.

The MKM Framework places particular importance on model risk governance and evaluation, a regulatory requirement of the bank. Beyond traditional model skill metrics, we employ specialised validation approaches focused on representing extreme but not biblical events. Validation includes careful assessment of precipitation intensity distributions, spatial coherence of weather patterns, and the realistic evolution of atmospheric features associated with flooding events.

The ultimate value of this model framework lies in its ability to support robust physical risk assessment in banking systems. While sophisticated in their treatment of global weather dynamics, traditional weather pattern models often prove insufficient for the specific needs of flood risk

assessment. The MKM Framework addresses this limitation by maintaining physical rigour while focusing on the weather patterns and processes most relevant to flooding events that impact a bank's asset and lending portfolio.

This focused approach yields several crucial advantages for risk assessment. By capturing the full range of possible weather pattern evolutions, from industry and academic configurations to rare but dangerous combinations of atmospheric conditions, we provide a solid foundation for quantifying flood risk probabilities. The framework's attention to local-scale processes and extreme event precursors proves valuable for identifying potentially dangerous weather patterns before they fully develop.

The MKM Model Framework seamlessly integrates with subsequent components of our bank risk assessment framework. It is physically consistent with weather patterns, providing ideal inputs for our AI-based pattern recognition and Bayesian analysis approaches.

“This integration of the MKM Model Framework ensures that our more sophisticated statistical analyses build upon meteorologically sound foundations rather than operating in isolation from physical constraints” - Johnny Mattimore, MKM.

Perhaps most importantly, the framework maintains transparency and interpretability throughout its operation. While the underlying mathematics may be complex, the model's predictions can be traced back to fundamental physical principles and verified against known atmospheric dynamics. This transparency and explainability proves crucial for bank risk managers and regulators who are required to understand and justify their risk assessments.

In later chapters, this physically grounded approach will prove invaluable as we examine specific applications in weather pattern prediction and flood risk assessment. Our framework provides a robust foundation for meaningful physical risk assessment by carefully balancing sophisticated mathematics with practical utility.

AI and Bayesian Approaches to Weather Pattern Prediction

Integrating artificial intelligence with Bayesian methods represents a transformative approach to weather pattern prediction. While traditional

meteorological models excel at physics-based forecasting, they often struggle to capture the fine-grained patterns and subtle precursors that can signal the development of extreme weather events. Our framework addresses this limitation by combining deep learning architectures with Bayesian inference to extract more granular insights from High-Resolution Rapid Refresh (HRRR) data than conventional approaches allow.

The cornerstone of our approach lies in treating weather pattern prediction not as a deterministic forecast but as a Bayesian inference problem. We begin with prior distributions informed by a physical understanding of atmospheric dynamics and climatological patterns. These priors encode fundamental constraints such as conservation laws and thermodynamic relationships, ensuring that our predictions remain physically plausible even as we push the boundaries of traditional forecasting.

Deep learning enters our framework through carefully designed neural network architectures that learn to identify and track weather pattern evolution. Convolutional neural networks (CNNs) prove particularly valuable for recognising spatial patterns in atmospheric fields, while recurrent architectures capture temporal dependencies in weather system evolution. However, we move beyond traditional deterministic neural networks by employing Bayesian neural networks that provide crucial uncertainty quantification at each step of the prediction process.

The Bayesian neural network approach offers several key advantages over conventional deep learning methods. Rather than producing single-valued predictions, these networks generate entire probability distributions over possible weather patterns.

“This probabilistic output aligns with our broader framework for physical risk assessment, providing a richer characterisation of potential weather evolution pathways.” David Kelly, MKM

The MKM Framework pays particular attention to the challenge of rare event prediction. Traditional machine learning methods often struggle with imbalanced datasets where extreme events are underrepresented. We address this by combining techniques that appropriately weigh extreme event precursors while maintaining the network's skill at predicting more common weather patterns.

The power of our AI framework lies in its ability to learn from vast amounts of historical weather data while maintaining physical consistency. Where traditional numerical weather prediction models rely solely on physical equations, our approach combines physics-based understanding with patterns learned from millions of historical weather

observations. This hybrid approach proves particularly valuable for identifying subtle precursors to extreme weather events that conventional methods might miss.

Data assimilation is one of the most crucial aspects of our framework. Weather prediction inherently depends on diverse data streams, including ground stations, weather balloons, satellite observations, and radar measurements. Our Bayesian approach provides a mathematically rigorous framework for combining these varied data sources, weighing their relative uncertainties, and accounting for spatial and temporal correlations in measurement errors.

Advanced Neural Network Architecture for Weather Prediction

This neural network architecture integrates three key innovations tailored for weather prediction challenges, addressing critical limitations in traditional approaches:

Attention Mechanisms for Dynamic Feature Prioritization.

The model employs sophisticated attention mechanisms to weigh atmospheric variables across spatial and temporal dimensions dynamically:

- Spatial Attention identifies critical regions, such as developing low-pressure systems or oceanic temperature anomalies that disproportionately influence future weather patterns.
- Temporal Attention isolates pivotal time steps, including rapid intensification phases in cyclones or diurnal heating cycles, enabling the model to focus on nonlinear transitions.

This approach mirrors human forecasters' ability to prioritise "weather makers" but implements this capability quantitatively using transformer-based architectures.

Physics-Aware Skip Connections

The architecture maintains multi-scale information essential for physical consistency through specialised skip connections:

- Variable Conservation for critical quantities like potential vorticity and moisture flux are preserved across layers via residual pathways, preventing degradation in deep networks.

- Multi-Resolution Processing: Shallow layers retain high-resolution cloud microphysics details, while deeper layers capture planetary-scale Rossby wave interactions.

These connections emulate the grid nesting used in numerical models like WRF but in a learned, adaptive manner.

Variational Inference for Uncertainty Quantification

Unlike deterministic NWP ensembles requiring 50+ computational runs, this framework implements:

- Learned Priors: Bayesian layers that parametrise distributions over key variables using historical extremes from CMIP6/ERA5 datasets.
- Stochasticity Injection: Monte Carlo dropout during inference that generates 100+ ensemble members at 1/50th the computational cost of dynamical models.

This approach effectively captures epistemic uncertainty (model limitations) and aleatoric uncertainty (inherent atmospheric chaos).

Operational Advantages Over Traditional NWP

Combining these elements, the architecture achieves ECMWF-level skill in 72-hour hurricane track forecasts while resolving convective-scale features (3 km vs. 9 km in IFS). However, like all ML approaches, it remains constrained by training data coverage—rare events like "bomb cyclones" still require hybrid modelling with physics constraints.

Local adaptation is vital to our approach. While weather patterns operate on large scales, their manifestation often depends heavily on local topography, land use patterns, and regional weather pattern characteristics. Our framework addresses this through a hierarchical structure that combines global pattern recognition with locally trained models that capture region-specific weather behaviour. This multi-scale approach proves particularly valuable for flood risk assessment, where local factors can significantly influence precipitation patterns.

The challenge of non-stationarity in weather patterns requires special attention. Traditional machine learning models often assume statistical stationarity in their training data—an assumption that becomes increasingly problematic if weather patterns shift. We address this through transfer learning techniques and continuous model updating, allowing our framework to adapt to evolving weather patterns while maintaining its foundation in historical data.

Extreme event prediction presents particular challenges for AI-based approaches. The relative rarity of such events in historical data can lead traditional machine-learning models to underestimate their probability. Our framework addresses this through several complementary strategies.

Synthetic data generation leverages physical principles and catastrophe models to enhance the realism and diversity of training datasets for rare events like extreme weather. By integrating physical models of phenomena such as storm surges or wind patterns with statistical extreme value distributions, synthetic data algorithms produce artificial event footprints that preserve the spatial-temporal relationships and intensity profiles observed in historical records.

Aspect	Traditional NWP	Proposed Framework
Compute Cost	Exascale HPC required	GPU-optimized <10 nodes
Lead Time	Hours for initialization	Real-time streaming
Uncertainty Metrics	Post-hoc calibration	Native probabilistic outputs

Table 3: MKM Framework and External Forcings

For example, catastrophe models combine deterministic simulations of known events with probabilistic extrapolations beyond historical observations, using techniques like Generative Adversarial Networks (GANs) to maintain statistical fidelity while expanding coverage to low-frequency, high-impact scenarios. This augmentation helps address the inherent scarcity of extreme event data in climate modelling, particularly for compound hazards like coastal flooding, where physical processes interact nonlinearly.

Our Bayesian framework provides a natural paradigm for quantifying uncertainties in rare event predictions by systematically incorporating prior scientific knowledge and updating probability estimates as new data emerges. Unlike frequentist methods that struggle with zero-event scenarios, Bayesian approaches use hierarchical models and empirical Bayes techniques to pool information across related hazards, enabling robust frequency estimates even for events with no historical precedents.

For tail-event weather forecasting that does not have to be extreme or biblical, this manifests as probabilistic predictions that account for both model structural uncertainties and climate variability, expressed through credible intervals rather than single-point estimates. By integrating physical constraints through informed priors and continuously assimilating observational data, Bayesian methods produce dynamically calibrated risk assessments essential for resilient infrastructure planning and adaptive climate policies.

The computational demands of this sophisticated approach require careful consideration. Where possible, we employ efficient approximations and model compression techniques that maintain prediction quality while reducing computational overhead. This efficiency proves crucial for operational risk assessment, where rapid analysis of emerging weather patterns can distinguish between effective and ineffective risk mitigation strategies.

Distribution Path Modelling

Modelling weather pattern distributions represents a fundamental shift from traditional meteorological approaches. Rather than attempting to predict specific outcomes, we aim to characterise the complete space of possible weather pattern evolutions, particularly those that could lead to extreme precipitation events. The challenge of mapping and analysing this full probability space of weather pattern evolution necessitates a sophisticated mathematical framework that combines stochastic differential equations, copula-based dependency structures, and advanced clustering techniques.

Why does Copula trump Gaussian?

Copula-based dependency structures offer significant advantages over Gaussian models in characterising the probability space of weather pattern evolutions, particularly for extreme precipitation events. Here's why:

Capturing Non-Gaussian Dependencies

- **Tail dependence:** Copulas explicitly model dependencies in extreme values (e.g., simultaneous heavy rainfall and specific atmospheric conditions), which Gaussian models often underestimate.
- **Nonlinear relationships:** Weather variables like temperature and rainfall exhibit complex, nonlinear interdependencies.

Flexibility in Marginal Distributions

- Copulas decouple marginal distributions (e.g., rainfall intensity, storm duration) from their dependence structure, allowing mixed distributions (e.g., gamma for precipitation, normal for temperature).
- For example, in flood risk analysis, coupling Gaussian mixture models with copulas improved the joint modelling of flood peak, volume, and duration.

Robustness for Extreme Events

- Gaussian copulas require high percentile thresholds (e.g., 0.8) to model extremes effectively, while copulas like Gumbel and Student's t perform well even at lower percentiles (0.2–0.8)[1]. This is critical for rare but high-impact events like extreme precipitation.
- Hierarchical Archimedean copulas enable asymmetric dependence modelling across different distribution parts, improving flood frequency projections under climate change.

State Space Representation

Our approach begins with a comprehensive state space representation of atmospheric conditions.

State space representation treats the atmosphere as a dynamical system governed by conservation laws for mass, momentum, energy, and moisture. Each grid point in the three-dimensional atmospheric grid (e.g., latitude-longitude-altitude) is assigned a state vector containing the variables we have observed in the HRRR data, such as:

- Temperature.
- Pressure.
- Zonal and meridional wind components
- Specific humidity
- Ice/water content

We model the atmosphere as a high-dimensional system where each point in state space captures the parameters above across a three-dimensional grid.

“This high-dimensional system creates a substantial computational challenge with millions to a billion degrees of freedom in modern models. We then couple this evolution with differential equations derived from the Navier-Stokes equations, thermodynamic principles, and parameterised sub-grid processes (e.g., convection, radiation, turbulence).” David Kelly, MKM

This representation allows us to track the spatial relationships between variables and their temporal evolution through the atmosphere's complex dynamics.

These atmospheric states' evolution follows deterministic physical laws and inherently stochastic processes, which we capture through carefully constructed stochastic differential equations (SDEs).

Formula 1: Stochastic Atmospheric State Evolution

These take the form:

$$dX(t) = a(X(t), t)dt + b(X(t), t)dW(t)$$

Where

- $X(t)$ represents our state vector of atmospheric variables.
- The drift term $a(X,t)$ encodes the deterministic physics - thermal gradients driving heat transfer, pressure differences generating winds, moisture transport mechanisms, and the crucial interactions between terrain and atmospheric flow.

Meanwhile, the diffusion term $b(X,t)dW(t)$ acknowledges and quantifies the uncertainties inherent in weather pattern evolution, from measurement errors in initial conditions to the fundamental limits of predictability in chaotic atmospheric systems.

The intricate dependencies between atmospheric variables demand more sophisticated modelling than traditional correlation measures can provide. We employ a hierarchical copula framework that captures direct relationships between variables and their conditional dependencies. Through a vine copula structure, we decompose the high-dimensional dependency network into a cascade of bivariate relationships.

Vine copulas balance flexibility with computational feasibility, making them a cornerstone of modern multivariate dependence modelling. Using a vine copula approach addresses the limitations of traditional multivariate copulas (e.g., Gaussian or Archimedean), which often impose restrictive symmetry or tail dependence assumptions. Vine copulas organise dependencies using a sequence of trees (called a “regular vine” or “R-vine”), where each tree level represents conditional relationships between variables. The structure includes:

- C-vines Feature a central "root node" at each tree level, applicable when one variable dominates. We can choose between temperature, precipitation, or atmospheric pressure for weather prediction.
- D-vines: Use a path-like structure, ideal for sequential dependencies (e.g., time-series data)

This approach allows us to match each relationship with the most appropriate copula family - using Gaussian copulas where dependencies are approximately linear, Student-t copulas where we see strong tail dependence, and Gumbel copulas for asymmetric relationships in extreme values.

These copula relationships themselves evolve dynamically as atmospheric conditions change. A stable summer weather pattern exhibits different variable dependencies than an intense developing storm system. Our framework captures these shifting relationships by allowing copula parameters to vary with the underlying weather regime, providing crucial flexibility when modelling the transition into extreme event scenarios.

Weather Pattern Distributions

Weather pattern distributions leveraging Monte Carlo frameworks enhanced by importance sampling and adaptive strategies are designed to capture rare meteorological events while efficiently minimising computational costs. These methods address the challenges of traditional approaches, which often require prohibitively large sample sizes to resolve low-probability phenomena.

In capital markets, we recognise this as Skew plus Kurtosis, so adjusting the distribution towards tail events is industry standard practice. In the insurance industry, they take it a step further by adopting the modified Ornstein-Uhlenbeck process, which biases temperature simulations toward rare extremes, allowing for a non-zero estimation of events like 100-year heatwaves with far fewer simulations than the brute-force Monte Carlo method. The problem for capital markets is that this approach looks too skewed, but it is worth considering if the Monte-Carlo is used for capital purposes.

The resulting collection of weather pattern evolutions provides rich material for analysis through our clustering framework. Moving beyond simple distance-based clustering, we employ a modified DBSCAN (Density-Based Spatial Clustering of Applications with Noise) algorithm that considers weather patterns' spatial structure and temporal evolution.

Meteorological data often includes transient phenomena (e.g., isolated thunderstorms) or measurement errors. DBSCAN distinguishes these as noise, ensuring clusters represent meaningful weather patterns; in particular, isolated weather stations with anomalous readings are flagged as outliers and sporadic events such as localised flooding are separated from systematic trends.

The clustering operates in a feature space, combining direct state variables with derived quantities like vorticity, precipitation potential, and pattern similarity measures grounded in atmospheric physics.

This clustering process serves multiple crucial functions in our framework. Beyond identifying typical weather pattern trajectories, it reveals potential extreme event pathways that more traditional analysis methods might miss. The clusters provide natural categories for probability assignment through our Bayesian framework, which weighs historical frequency, physical plausibility, and risk relevance to assign meaningful probabilities to different evolution pathways.

The integration of these modelling components requires careful attention to feedback and validation. Each generated weather pattern must satisfy fundamental physical constraints - conservation of mass, energy balance, and fundamental thermodynamic relationships. We employ continuous validation against historical extreme events to verify our model's ability to reproduce known patterns and to ensure we capture the essential physical mechanisms that drive extreme weather evolution.

The probability assignment process represents one of the most delicate aspects of our framework. Traditional approaches often rely too heavily on historical frequencies, which can underestimate the probability of extreme events. Our Bayesian framework instead combines multiple streams of evidence - historical observations, physical constraints and pattern similarity metrics - to assign probabilities that better reflect current and emerging risks.

“Energy constraints play a vital role in weather pattern evolution and guide much of our modelling approach.

Atmospheric systems cannot arbitrarily transition between states; they must follow physically realistic paths that respect energy conservation and thermodynamic principles.” - Johnny Mattimore, MKM.

Our stochastic differential equations explicitly incorporate these constraints, ensuring that even extreme scenarios maintain physical plausibility.

Pattern similarity plays a crucial role in both clustering and probability assignment. Rather than relying on simple Euclidean distances, we employ similarity measures that capture meteorologically significant features—the structure of pressure gradients, the organisation of moisture fields, and the coherence of wind patterns. These physically motivated similarity metrics help ensure our clustering results reflect meaningful weather pattern families rather than arbitrary mathematical groupings.

Handling Extreme Events

Our framework's handling of extreme events deserves particular attention. Extreme weather patterns often emerge from subtle interactions between multiple atmospheric features - interactions that more straightforward modelling approaches might miss. Our copula-based dependency modelling proves particularly valuable here, capturing the complex relationships that can align to produce severe precipitation events. Adaptive sampling ensures we explore these crucial edge cases thoroughly rather than treating them as mere statistical outliers.

The output of this distribution path modelling feeds directly into our precipitation time series generation, but its value extends beyond simple inputs for the following modelling stage. The clustered patterns and their assigned probabilities provide insight into the mechanisms of extreme weather evolution, helping identify potentially dangerous atmospheric configurations before they manifest as flood events. This understanding proves invaluable for long-term risk assessment and adaptation planning.

Our framework carefully balances complexity and interpretability. While the underlying mathematics may be sophisticated, the resulting weather pattern distributions and their evolution pathways must remain interpretable to meteorologists and risk managers. We achieve this through careful visualisation techniques and physically meaningful

pattern classifications that bridge the gap between mathematical abstraction and practical risk assessment.

Our framework's success depends heavily on efficiently exploring the vast state space of possible weather patterns. While robust, traditional Monte Carlo methods leave significant coverage gaps that might miss critical pattern evolution pathways. To address this limitation, we employ sophisticated quasi-Monte Carlo (QMC) techniques that systematically explore the state space through carefully constructed low-discrepancy sequences.

The advantage of QMC methods is that they achieve more uniform coverage of the sampling space than conventional random sampling. We primarily employ Sobol's sequences, chosen for their superior properties in the high-dimensional spaces characteristic of atmospheric systems. These sequences ensure that our sampling systematically explores different combinations of atmospheric variables, leaving no significant gaps in our coverage of possible weather pattern evolutions.

Atmospheric systems present particular challenges due to their high dimensionality and discontinuous behaviour near weather fronts and other sharp transitions. We address these challenges through a hybrid approach that combines QMC sampling with targeted importance sampling in regions of particular interest for flood risk.

Targeted importance sampling combines statistical efficiency with domain knowledge to better capture meaningful tail events in flood modelling without wasting computational resources on unlikely scenarios.

Traditional Monte Carlo methods distribute samples uniformly throughout the parameter space, which is inefficient for flood modelling since certain regions of the parameter space (such as specific rainfall-soil moisture combinations) disproportionately impact outcomes. Targeted importance sampling addresses this by:

- Adjusting the sampling distribution to focus on regions of higher relevance, such as urban drainage systems and complex topography
- Applying a higher sampling density in areas with greater uncertainty or sensitivity, such as rainfall intensity thresholds that trigger nonlinear responses or soil saturation levels near critical points.
- Applying temporal targeting such as seasonal patterns when antecedent conditions create higher risk or known weather system types that correlate with flooding

“The art of target sampling is not to “do a Cat insurance” on the tail of the distribution and seek to maintain some statistical validity with just a touch of skew.” - David Kelly, MKM.

The integration of QMC methods with our copula-based dependency structures requires careful attention to the mapping between uniform QMC points and the actual atmospheric state space. We employ dimension-reduction techniques that identify the most influential combinations of atmospheric variables, allowing us to focus our QMC sampling on these critical subspaces while handling less crucial dimensions through conventional methods.

Synthesising all the elements

Our framework concludes by synthesising all these elements - stochastic differential equations, copula-based dependencies, clustering analysis, and QMC sampling - into a coherent system for generating and analysing weather pattern distributions. The resulting output provides a set of possible weather patterns and a structured understanding of how these patterns might evolve and interact to produce conditions conducive to flooding events.

This sophisticated approach to distribution path modelling forms the foundation for our subsequent precipitation analysis and flood risk assessment. By maintaining physical consistency while efficiently exploring the full range of possible weather pattern evolutions, we create a robust basis for quantifying and managing physical risk in banking systems with the benefit of complete transparency for model governance that banking regulators require.

Our framework's ultimate validation comes through its practical application in risk assessment. The weather pattern distributions it generates must strike a careful balance between being comprehensive enough to capture the full range of possible outcomes and focused enough to provide actionable insights for risk management. Our framework achieves this balance through careful calibration and continuous refinement, providing a solid foundation for the precipitation analysis in the next section.

Time Series Analysis of Precipitation Patterns

One of the most challenging aspects of physical risk assessment is transforming weather pattern predictions into usable precipitation time series. Traditional approaches to precipitation modelling rely heavily on a

direct statistical analysis of historical data, but this proves insufficient for our purposes.

We require a more sophisticated framework that can:

- Capture the full distribution of possible precipitation patterns.
- Include the critical tail events that are bad but not catastrophic.
- Pay particular attention to extreme events that drive flooding risk.

The goal is to simulate realistic precipitation scenarios, including unusual but not catastrophic weather events, with a special focus on the types of extreme precipitation that lead to flooding hazards. The Bayesian paradigm provides the ideal foundation for this analysis.

"Rather than treating precipitation as a deterministic outcome, we view it as an uncertain process about which we can make probabilistic inferences. This approach allows us to combine multiple sources of information - from high-resolution radar data to ground-based measurements - in a mathematically rigorous way."

Johnny Mattimore, MKM.

At its core, our Bayesian framework treats precipitation patterns as the outcome of underlying atmospheric processes about which we have incomplete information. We start with prior distributions based on a physical understanding of atmospheric dynamics and regional climatology. These priors are then updated with observational data through the likelihood function, leading to posterior distributions representing our updated beliefs about precipitation patterns.

The power of this approach becomes particularly evident when dealing with extreme events. Traditional frequency-based statistics often struggle with rare events due to limited historical data. The Bayesian framework allows us to incorporate physical constraints and expert knowledge about extreme precipitation mechanisms, leading to more robust estimates of tail probabilities.

Spatial correlation presents another crucial challenge in precipitation modelling. Rain doesn't fall uniformly across a catchment area - topography, prevailing winds, and other factors create complex spatial patterns. Our framework employs spatial correlation structures that adapt to changing weather patterns while respecting physical constraints. This spatial correlation ensures that our synthetic time series exhibit realistic spatial coherence across catchment areas.

The temporal evolution of precipitation patterns requires equal attention. Weather systems don't evolve randomly - they follow physical laws and exhibit various forms of persistence. Our Bayesian model captures these temporal dependencies through state-space representations that explicitly model the evolution of precipitation systems. This includes short-term persistence during storm events and longer-term patterns related to seasonal and climatic factors.

"The final output of our analysis consists of synthetic precipitation time series that serve as inputs to hydrological models. These aren't simple point forecasts - they represent complete probabilistic descriptions of possible precipitation patterns. Each synthetic series maintains consistency with known physical constraints while incorporating the uncertainties inherent in weather prediction."- David Kelly, MKM.

This comprehensive approach to time series analysis bridges the gap between weather pattern prediction and flood risk assessment. By maintaining physical consistency while properly characterising uncertainty, we provide the robust inputs needed for subsequent risk quantification. The Bayesian framework ensures that our models can adapt as new data becomes available, making them particularly valuable for long-term risk assessment.

"It is this high-speed adaptation that is of huge value to banking, where intraday updates make meaningful differences to potential asset valuation and impairment decisions."- Johnny Mattimore, MKM.

The MKM Framework transforms weather prediction from a meteorological exercise into a critical banking risk management tool. While traditional banks may view weather modelling as outside their domain, our approach bridges this gap by directly connecting atmospheric science to financial outcomes. Banks can dynamically adjust collateral valuations in response to emerging weather patterns before physical damage occurs, accurately provide for expected losses in loan portfolios with exposure to flood-vulnerable regions, and optimise capital reserves by replacing overly conservative buffer approaches with precision-targeted allocations based on scientific weather patterns distributions.

Our framework provides banks with transparent model governance that satisfies regulatory requirements for explainability and physical basis.

The adaptive risk recalibration capabilities respond to rapid changes in weather conditions, similar to how market risk systems adjust to volatility shifts. Banks gain competitive differentiation through superior physical risk assessment that manifests in pricing advantages and reduced losses across their portfolios.

The output of this weather prediction framework—probabilistic precipitation time series—integrates seamlessly with banks' existing risk infrastructure. The time series format mirrors the market data inputs already used in banks' core systems. Uncertainty quantification through Bayesian methods aligns with modern market risk management approaches. The distributed sampling approach enables stress testing and scenario analysis familiar to risk managers, allowing them to work with physical risk in the same paradigm they apply to market and credit risk.

By transforming complex meteorological science into usable financial decision inputs, the MKM Framework enables banks to incorporate physical risk into their core operations without requiring wholesale system redesigns or specialised meteorological expertise. Weather patterns become another dimension in the risk assessment matrix that responds to the same rigorous quantitative approaches banks already employ for their traditional risk factors.

In the following chapters, we will demonstrate how these precipitation time series drive hydrological models that ultimately translate into metrics directly relevant to banking: asset valuation impacts, expected loss adjustments, and capital requirement calculations that reflect the full spectrum of physical risks facing modern financial institutions. The bridge from atmosphere to balance sheet is now complete, allowing for sophisticated risk management that accounts for our physical world's increasingly volatile reality.

Chapter 4 - Hydrological Modelling

Assessing physical risk in water-related hazards fundamentally depends on our ability to understand and model how water moves through the environment. Hydrological modelling is the critical bridge between weather events and their impacts on the flood risk assessment. It translates meteorological data into actionable insights about flood risks, water scarcity, and other hazards. This chapter explores the complex world of hydrological modelling, examining how different approaches can simulate water movement through natural and built environments.

"All hydrological models, regardless of their complexity or simplicity, are ultimately expressions of the same fundamental physical laws established by Newton. We may dress them differently, adapt their form to various scales and purposes, but beneath these surface variations lie the immutable principles of mass and energy conservation." - David Kelly, MKM.

Hydrological modelling is grounded in fundamental physical principles - primarily Newton's laws of motion and thermodynamics. These universal physical laws govern how water moves and changes state throughout its journey, from the moment it falls as precipitation until it reaches its ultimate destination, whether that's the ocean, groundwater, or the atmosphere through evaporation.

While models may vary significantly in complexity and approach, they all ultimately derive from these same physical principles. This process involves multiple interconnected components: precipitation patterns, surface runoff, infiltration into the soil, groundwater movement, and channel flow dynamics. The accuracy of physical risk assessment depends critically on how well we can represent these processes in our models.

Hydrological modelling has evolved significantly with advances in computational power and data availability, though all approaches - from the simplest to the most complex - remain fundamentally rooted in classical physics. Whether a model is lumped, distributed, empirical, or physically based, we can easily trace their underlying equations to the same core physical principles of mass, energy, and momentum conservation.

Tracing back to Newton is a crucial aspect of such modelling. The intellectual property of those who model such outcomes for their insurance clients does not stem from their original creation of the models, as these originate from academia and are, therefore, in the public domain. The intellectual property should align with that of banks and their instrument pricing models, which is based on their implementation rather than solely on the accuracy and granularity of their results and scalability.

Early models were necessarily simple, treating entire watersheds as single units with uniform characteristics. Modern approaches can incorporate detailed spatial variations in terrain, land use, and soil properties, leading to more accurate water movement and accumulation predictions.

Hydrological Model Overview

Hydrological models can be classified along several dimensions, each offering different advantages for specific applications:

Spatial Representation

- **Lumped Models:** These models treat the watershed as a uniform unit with uniform characteristics. While computationally efficient and requiring minimal data, they sacrifice spatial detail in favour of simplicity. They remain valuable for rapid assessments and preliminary analyses.
- **Semi-Distributed Models:** Representing a middle ground, these models divide watersheds into sub-basins, each with its characteristics. This approach balances computational efficiency with the need to represent spatial variability. It makes them particularly useful for medium-to large watersheds where some spatial detail is essential but complete distribution isn't necessary.
- **Fully Distributed Models:** These models represent the ultimate spatial detail, dividing the study area into a fine grid where each cell can have unique properties. While computationally intensive, they provide the most detailed representation of spatial processes and are essential for applications requiring high spatial resolution.

Process Representation

The way models represent physical processes varies significantly, reflecting different philosophical approaches to hydrological modelling:

- **Empirical Models:** Built on statistical relationships derived from observed data, these models excel in stability and simplicity but may struggle when conditions differ from their training data. They serve well in operational settings where quick, reliable results are needed for familiar conditions.
- **Conceptual Models:** These models balance empirical and physical approaches, representing key processes through simplified equations. They maintain physical meaning while remaining computationally manageable, making them popular for many practical applications.
- **Physical Models:** These models offer the most robust theoretical foundation. They demand data and computation based on fundamental water movement and conservation equations. Still, they provide the best framework for understanding process changes under novel conditions.

When integrating backtesting into hydrological model applications for physical risk assessment, it becomes critical to address the following enhanced considerations:

- **Robustness to non-stationarity:** Backtesting against historical extreme floods helps validate model performance under shifting climatic regimes. Traditional backtesting assumes stationarity, so newer frameworks combine historical validation with stress-testing against synthetic extremes beyond observational records.
- **Uncertainty quantification:** Backtesting enables empirical estimation of prediction intervals by analysing model predictive capability across diverse historical scenarios. This directly informs risk pricing through metrics like reliability, sharpness, and skill scores.
- **Computational efficiency:** While backtesting requires intensive scenario runs, reduced-complexity models can maintain usability if they demonstrate consistent skill across validation periods, stability in parameter transferability, and scalable performance when applied to large portfolios of risk locations.
- **Implementation challenges:** Long-term observational datasets are a prerequisite for meaningful backtesting, non-stationarity may render older events less representative of future hazards, and trade-offs exist

between model complexity and computational demands for ensemble backtesting.

Precipitation Modelling

Precipitation modelling is one of the most challenging aspects of weather prediction, owing to the complex interplay of multiple physical processes across various spatial and temporal scales. The fundamental challenge is that many critical precipitation processes occur at scales significantly smaller than typical model grid cells, necessitating sophisticated parameter schemes to represent their collective effects.

Modern precipitation modelling approaches encompass a rich spectrum of methodologies, ranging from purely statistical frameworks that leverage historical data patterns to entirely mechanistic models that simulate fundamental physical processes. At one end of this spectrum, statistical approaches treat precipitation as a stochastic process, using probability theory and historical observations to capture temporal and spatial patterns without explicitly modelling the underlying atmospheric physics.

These methods include traditional time series analysis, which can identify cyclical patterns and trends, and more sophisticated machine learning techniques, which detect complex, nonlinear relationships in precipitation data. As we move along the spectrum, semi-empirical models combine statistical relationships with basic physical constraints, balancing computational efficiency and physical realism. At the entirely mechanistic end, models directly simulate atmospheric dynamics, thermodynamics, and microphysics, requiring substantial computational resources but providing explicit representations of the physical processes driving precipitation formation.

Among these approaches, Hidden Markov Models (HMMs) have emerged as particularly sophisticated tools for capturing the complex temporal dynamics of precipitation patterns. HMMs conceptualise precipitation as a system that transitions between hidden weather states, each associated with distinct precipitation characteristics. These hidden states represent different atmospheric conditions or weather regimes that aren't directly observable but manifest in measurable precipitation patterns.

The power of HMMs lies in their probabilistic framework, which comprises transition probabilities between weather states and emission probabilities that govern the likelihood of specific precipitation amounts given each state. This dual-layer structure allows HMMs to capture the temporal persistence of weather patterns and the variability in precipitation intensity. The Markov property—where the current state

depends only on the previous state—provides a computationally tractable way to model the temporal evolution of weather systems.

HMMs employ specialised algorithms, such as the Viterbi algorithm for precipitation modelling, which identifies the most likely sequence of weather states given observed precipitation patterns, and the Baum-Welch algorithm for parameter estimation from historical data. These algorithms enable the model to learn the underlying weather state dynamics and their relationship to precipitation from observational data. This learning capability makes HMMs particularly valuable for regions with complex precipitation patterns.

While HMMs provide sophisticated statistical frameworks for precipitation modelling, other approaches, such as Generalized Linear Models (GLMs) and mechanistic models based on radar data and physical principles, offer complementary capabilities. Each approach offers distinct advantages and limitations, with the choice often depending on the specific application and available computational resources.

The Art of Hydrology and Hydraulic Maintenance

"Hydrologic modelling shows us the journey of every raindrop across watersheds—predicting how much water arrives and when. Hydraulic modelling reveals how water surges through channels, spills across floodplains, and exerts a force on everything in its path. Together, they form the essential narrative of water's movement through our world: one tells us what's coming, the other shows us what happens when it gets here."

The hydrologic and hydraulic routing field embodies the fundamental tension in environmental modelling between theoretical rigour and practical utility. From the most straightforward Muskingum calculations to the most sophisticated two-dimensional Saint-Venant implementations, each approach balances physical realism, data requirements, computational feasibility, and decision-making needs.

This balance is not merely a technical consideration but a philosophical one—it recognises that models serve as tools for understanding and decision-making rather than perfect representations of reality. The ongoing challenge of finding appropriate complexity for specific applications continues to drive innovation in the field, from

methodological advances to computational techniques and uncertainty characterisation.

As weather patterns alter hydrological regimes and urbanisation transforms watersheds, the importance of robust routing methods only grows. Integrating traditional process understanding with modern computational techniques and data science approaches promises a future where we can better understand, predict, and manage water's journey through our environment – an endeavour with profound implications for infrastructure resilience, ecological sustainability, and human well-being.

Three distinct yet interconnected approaches form the foundation of our understanding of water's movement through landscapes: hydrological modelling, hydrologic routing, and hydraulic modelling. Each represents a different scale of analysis, a different set of simplifications, and a different perspective on water's journey.

Hydrological modelling embraces the entire water cycle within a watershed. It begins with rain falling from the sky. As it is intercepted by vegetation, it follows that water infiltrates into soils, percolates to groundwater, evaporates back to the atmosphere, or becomes surface runoff. This holistic approach asks: Of all the water that falls as precipitation, how much becomes runoff, and when does it reach the stream network?

The watershed resembles a canvas showcasing details of land use, soil types, topographic features, and vegetation patterns. Each element affects the fate of water–urban areas shed runoff quickly, forests capture and slowly release moisture, and agricultural fields change with the seasons. Hydrological models integrate these complex interactions through a series of mathematical relationships, ranging from simple empirical equations to sophisticated physics-based formulations.

Hydrology modelling connects the atmospheric drivers (precipitation, temperature, solar radiation) to the terrestrial response, creating the crucial link between meteorology and hydrologic routing.

As water converges into channels and reservoirs, hydrologic routing takes centre stage. This approach focuses on how the shape and timing of flood waves change as they move downstream—a process fundamental to flood forecasting and reservoir operations.

Hydrologic routing embraces the principle of continuity: water cannot be created or destroyed, only stored or released. The elegant simplicity of the continuity equation—where the difference between inflow and outflow equals the rate of change in storage—betrays the complex interplay between channel characteristics and flow dynamics.

Methods like the Muskingum approach represent channels as a series of natural reservoirs, using empirical relationships to describe how flood waves attenuate and lag as they travel downstream. Reservoir routing similarly tracks the balance between inflows, outflows, and changes in storage, enabling operators to manage releases for flood control, water supply, hydropower generation, and environmental flows.

These approaches sacrifice the detailed mechanics of fluid flow for computational efficiency and practical applicability. They recognise that for many applications, we need not resolve the intricacies of hydraulics to predict the essential characteristics of downstream flow—the peak, timing, and volume that drive decision-making.

Hydraulic modelling is a crucial next step when we need to understand not only the volume of water passing a point and the timing but also its depth, velocity, energy, and spatial distribution. Based on the principles of fluid mechanics—particularly the conservation of mass, momentum, and occasionally energy—these models clarify the detailed behaviour of water in channels, floodplains, and hydraulic structures.

Hydraulic models address whether water will reach a gauge's alert level, resulting in flooding. How effective is this proposed flood control structure? Which areas will be inundated during a 100-year flood event? Which bridge designs can withstand the force of floodwaters? The answers influence billions in infrastructure investments and numerous decisions impacting public safety and property.

The MKM framework effectively blurs the distinctions, creating integrated platforms where rainfall-runoff processes feed seamlessly into routing methods, which in turn drive detailed hydraulic analyses. This integration reflects the reality of the water cycle itself—a continuous system where neat categorisations inevitably break down.

This tension between the mathematical precision of our models and the chaotic reality of natural systems demands a particular mindset—one that balances confidence in scientific principles with humility about their limitations. The mindful modeler understands that models are not truth, but tools for approaching truth; not reality itself, but lenses through which we perceive reality.

The Zen of water modelling lies in finding harmony between theoretical elegance and practical utility, between excessive simplification and needless complexity.

As weather patterns alter hydrological regimes and urbanisation transforms watersheds, the importance of robust routing methods only grows. Integrating traditional process understanding with modern computational techniques and data science approaches promises a future

where we can better understand, predict, and manage water's journey through our environment - an endeavour with profound implications for infrastructure resilience, ecological sustainability, and human well-being.

The following chapters will explore each element of this trinity in detail, examining the theoretical foundations, practical applications, and emerging frontiers of hydrological modelling, hydrologic routing, and hydraulic modelling. Like water itself, our exploration will flow from the broad watershed scale to the detailed mechanics of channel flow, always seeking to understand the essence of water's movement through our world.

Hidden Markov Models

Formula 2: Hidden Markov Models

The mathematical foundations of Hidden Markov Models in precipitation modelling can be expressed through several key equations. The model is characterised by:

$$P(St|St-1) = A[i,j]$$

Where $A[i,j]$ represents the transition probability matrix between weather states i and j , the emission probabilities for precipitation amounts are typically modelled using a mixed distribution:

$$P(Rt|St) = wO(St)\delta_0 + (1-wO(St))\gamma(\alpha(St), \beta(St))$$

Where $wO(St)$ is the probability of a dry day in state St , δ_0 is the Dirac delta function at zero (representing no rain), and $\gamma(\alpha, \beta)$ is the gamma distribution with shape parameter α and scale parameter β specific to each state.

Formula 3: Generalized Linear Models

For Generalized Linear Models (GLMs), the precipitation process is often separated into occurrence and intensity components:

$$\text{logit}(P(Rt > 0)) = X\beta + \epsilon$$

Where X represents predictor variables and β their coefficients. For non-zero precipitation:

$$\log(Rt|Rt > 0) = Z\gamma + \eta$$

Where Z may include both atmospheric predictors and temporal dependencies.

In physically-based models, the fundamental equations include the conservation of water vapour:

$$\partial q / \partial t + v \cdot \nabla q = S - C + D$$

Where q is specific humidity, v is wind velocity, S represents sources (evaporation), C represents sinks (condensation), and D represents diffusion.

Hidden Markov Model Approach

These are the stages the Hidden Markov Model is doing in practice:

- **Identify weather states:** Define a set of hidden weather states (typically 2-5 states) that might represent different atmospheric patterns, such as "dry," "light precipitation," or "heavy precipitation."
- **Train the model:** Use historical precipitation data to estimate two sets of probabilities, notably transition probabilities (how likely the weather is to change from one state to another) and emission probabilities (How much rain typically falls when in each state).
- **Handle zero rainfall:** Create a special case for days without rainfall since precipitation data contains many zeros.
- **Model non-zero rainfall:** For days when it does rain, use a continuous distribution to model the amount of rainfall
- **Make predictions:** Once trained, use the model to predict near-future precipitation patterns by estimating the most likely sequence of hidden states.

The model outputs from precipitation modelling systems provide essential inputs for downstream risk assessment:

Primary Outputs:

- Precipitation intensity (mm/hr or mm/day).
- Precipitation duration.
- Precipitation frequency.
- Spatial distribution of rainfall.
- Storm movement and evolution.

Derived Statistics:

- Annual precipitation totals.
- Seasonal patterns and cycles.
- Intensity-Duration-Frequency (IDF) curves.
- Return period precipitation depths.
- Drought indices.

Risk Assessment Metrics:

- Extreme event probabilities.
- Precipitation deficits.
- Spatial correlation structures.
- Seasonal forecasting indicators.

Model Diagnostic Parameters:

- State transition probabilities (for HMMs).
- Parameter distributions.
- Uncertainty bounds.

As discussed in the following section, these outputs form crucial inputs for runoff modelling, creating a seamless link between precipitation processes and their hydrological consequences. The temporal and spatial resolution of outputs varies by model type:

- **Global Weather Models:** Typically daily or sub-daily, at grid scales of 50-100km
- **Regional Weather Models:** Hourly to daily, at 10-50km resolution
- **Statistical Models:** Can be configured for various temporal scales, often focused on point locations
- **Radar-based Models:** Sub-hourly, at kilometre or sub-kilometre resolution

The integration of these outputs with runoff models requires careful consideration of:

- Scale compatibility.
- Uncertainty propagation.
- Temporal aggregation effects.
- Spatial interpolation requirements.

The combination of precipitation and runoff modelling creates a comprehensive framework for physical risk assessment, where uncertainties and biases from both components must be carefully managed to provide reliable risk metrics for decision-making.

A persistent challenge in global weather pattern models has been their tendency to underestimate the intensity of heavy precipitation events. This systematic bias largely stems from their coarse spatial resolution, limiting their ability to resolve localised intense precipitation events. Recent advances in computing power have enabled the development of higher-resolution models, particularly at regional scales, where improved representation of topographic effects has led to more accurate precipitation patterns.

Modelling convective precipitation over land presents particular challenges, as it involves complex interactions between surface heating, atmospheric stability, and moisture transport. The daily precipitation cycle, especially in tropical regions, has historically been challenging to capture in models accurately. This difficulty extends to simulating seasonal monsoon patterns, which require an accurate representation of large-scale atmospheric circulation patterns and their interaction with local topography and surface conditions.

An innovative development in precipitation modelling has incorporated cloud-resolving components within larger-scale models. This approach simulates convective processes rather than relying on parameterisation schemes, improving precipitation intensity and temporal distribution representation. However, such approaches' computational cost often necessitates careful consideration of the resolution and domain size trade-offs.

The modelling process typically separates precipitation into two distinct components: occurrence and intensity. The occurrence model determines whether precipitation occurs on a given day, while the intensity model simulates the amount of precipitation on wet days. This two-step approach allows for a more accurate representation of precipitation patterns, particularly in regions with distinct wet and dry seasons.

Precipitation outputs from weather models often require additional processing for practical applications in hydrology and risk assessment. Statistical downscaling and bias correction techniques bridge the gap between model resolution and relevant spatial scales for impact assessment.

Model evaluation remains a critical aspect of precipitation modelling, requiring careful comparison of simulated precipitation statistics against observations across multiple spatial and temporal scales. Key metrics include the mean precipitation and measures of variability, extreme events, and spatial patterns. The availability of high-quality observational datasets, including satellite-based precipitation estimates, has dramatically enhanced our ability to evaluate and improve model performance.

Current research focuses on enhancing the representation of fundamental processes such as convection, cloud microphysics, and land-atmosphere interactions. These improvements and increasing computational capabilities gradually reduce systematic biases in precipitation modelling. However, significant challenges remain, particularly in simulating extreme precipitation events and capturing regional-scale precipitation patterns accurately.

Integrating precipitation modelling with physical risk assessment requires careful consideration of model uncertainties and limitations. Risk managers must understand the capabilities and constraints of different modelling approaches to make informed decisions about their application in specific contexts. This understanding becomes particularly crucial when using precipitation models to assess future distributions, where the interaction between changing precipitation patterns and other variables can significantly impact risk profiles.

Runoff Modelling

The transformation of precipitation into runoff represents a critical link in the hydrological cycle, with profound implications for physical risk assessment. While conceptually straightforward, this process embodies remarkable complexity in practice, encompassing multiple interacting pathways through which water moves through a catchment system. Understanding these pathways and their relative importance under different conditions forms the foundation of modern runoff modelling approaches.

The evolution of runoff modelling reflects a continuous tension between complexity and practicality, leading to diverse approaches ranging from purely empirical through conceptual to fully physically based frameworks.

Empirical approaches draw purely from statistical relationships between rainfall inputs and runoff outputs. They offer computational efficiency but sacrifice explicit representation of physical processes. While limited in their transferability to conditions outside their calibration domain, these models continue to find practical application in risk assessment scenarios where rapid computation is essential.

Moving along the spectrum of complexity, conceptual models attempt to bridge the gap between empirical and physical approaches by representing key hydrological processes through simplified mathematical frameworks. These models typically conceptualise catchments as collections of interconnected storage elements, each representing different components of the hydrological system. The popularity of

conceptual models in practical applications stems from their ability to maintain physical interpretability while remaining computationally tractable.

At the most sophisticated end of the spectrum, physically-based models attempt to represent the mechanics of water movement through the catchment system. These models solve fundamental mass, momentum, and energy conservation equations, providing detailed representations of infiltration, subsurface flow, and channel routing processes. While theoretically more robust, their practical application often faces challenges related to data requirements and computational demands.

The spatial representation of catchment processes introduces another dimension of complexity in runoff modelling. Lumped models, treating entire catchments as single units, offer simplicity but may miss critical spatial variations in catchment response.

Semi-distributed models partition catchments into hydrologically similar units, providing some predictive capability and a pragmatic balance between spatial detail and computational efficiency.

While fully distributed models offer the most detailed spatial representation, they also demand extensive data and computational resources that improvements in predictive capability may not always justify.

The runoff generation process encompasses multiple mechanisms operating at different temporal and spatial scales. Surface runoff, generated when rainfall intensity exceeds soil infiltration capacity or when soil becomes fully saturated, represents the most rapid response component. Subsurface flow through soil layers provides a slower response mechanism, while groundwater contribution to streamflow represents the slowest component. The relative importance of these mechanisms varies with catchment characteristics and storm properties, necessitating careful consideration in model development.

Unit Hydrograph Concept

Within this framework, the unit hydrograph concept has emerged as a foundational approach to understanding and predicting catchment response to rainfall.

A hydrograph serves as the fundamental graphical representation of a watershed's response to rainfall, depicting how streamflow changes over time following a precipitation event. Developed by Sherman in 1932, the unit hydrograph models the direct runoff response to a standardised rainfall input (typically 1 inch or 1 cm) distributed uniformly across a

watershed during a specified duration. This elegant concept transforms the complex rainfall-runoff relationship into a predictable, reproducible pattern based on several key principles: time invariance (consistent watershed response over time), linearity (proportional relationship between rainfall and runoff), and superposition (ability to combine responses from sequential rainfall events).

The beauty of the unit hydrograph lies in its practicality—it establishes a watershed's unique "signature" or response pattern that can be scaled and applied to predict runoff from storms of varying intensities and durations. This makes it an invaluable tool for flood forecasting, water resource management, and infrastructure design.

Despite its simplifying assumptions—uniform rainfall distribution and consistent watershed conditions—the unit hydrograph concept remains one of hydrology's most enduring and valuable tools for translating the complex interactions between precipitation and watershed processes into practical, quantifiable predictions.

This method, while simplified, provides crucial insights into how catchments transform excess rainfall into runoff. The conventional approach derives unit hydrographs from observed rainfall and runoff data for single storm events, carefully separating baseflow from the total hydrograph to isolate the direct runoff component. This empirical foundation, however, faces practical limitations in ungauged catchments, leading to the development of synthetic unit hydrograph methods.

SCS Dimensionless, Snyder and GUH

Hydrologists derive unit hydrographs through various methods, including direct analysis of observed rainfall-runoff data, synthetic approaches like Snyder's method that relate hydrograph characteristics to watershed properties, or the dimensionless hydrograph that standardises the shape based on peak flow timing.

These synthetic approaches, including the widely adopted SCS Dimensionless Unit Hydrograph and Snyder's Method, enable hydrograph estimation without observed data. More sophisticated techniques, such as the Geomorphological Unit Hydrograph (GUH), forge explicit connections between hydrograph shape and catchment physical characteristics, incorporating stream order statistics and network width functions. The time-area method offers another perspective, conceptualising catchment response through isochrones representing zones of equal travel time to the outlet.

Synthetic Unit Hydrographs: Practical Solutions for Ungauged Catchments

The challenge of predicting hydrological responses in watersheds without gauging stations has driven the development of ingenious synthetic approaches. These methods represent different philosophical approaches to the same fundamental problem—how to translate watershed characteristics into expected runoff patterns without direct observations.

The Soil Conservation Service (now Natural Resources Conservation Service) approach embodies the power of standardisation in hydrological science. Rather than deriving unique solutions for each watershed, the SCS method offers a standardised template derived from analysing numerous catchments across diverse landscapes. This dimensionless curvilinear shape can be scaled to fit specific watersheds using two key parameters: time to peak and peak discharge.

The SCS method's enduring popularity stems from its pragmatic balance between simplicity and effectiveness. Assuming a consistent hydrograph shape across watersheds allows practitioners to focus on determining the appropriate scaling factors. The method relies on a Peak Rate Factor (PRF)—typically set at 484 in traditional units—which can be adjusted based on terrain characteristics, with higher values for steeper watersheds and lower values for flatter terrain.

While often approximated as a triangular shape for calculation simplicity, the complete SCS dimensionless hydrograph captures the characteristic asymmetry of natural runoff events—a rapid rise followed by a more gradual recession. This standardised approach makes the SCS method accessible to practitioners with limited data or hydrological expertise, explaining its widespread adoption in engineering practice and regulatory frameworks worldwide.

Despite its empirical nature and generalised assumptions, the SCS method continues to provide reasonable estimates for many practical applications. Its limitations primarily emerge in watersheds with unusual characteristics or in situations requiring more nuanced representations of catchment behaviour. In these cases, modifications to the standard PRF value or more sophisticated methods may be necessary.

Snyder's Method: Regional Empiricism in Hydrograph Synthesis

Snyder's approach pioneered in the late 1930s, represents an early attempt to codify the relationship between watershed characteristics and hydrograph shape through regional empirical equations. Unlike the standardised shape of the SCS method, Snyder's approach directly estimates key hydrograph parameters—specifically basin lag time and peak discharge—based on physical watershed characteristics.

The cornerstone of Snyder's method is the relationship between basin lag time and watershed geometry, particularly length and centroid distance. By incorporating these physical features, Snyder recognised that watershed shape fundamentally influences the timing and magnitude of runoff. The regional coefficients in Snyder's equations acknowledge that watersheds in different geographic regions respond differently to rainfall, even when their physical dimensions are similar.

This regional sensitivity makes Snyder's method adaptable across diverse landscapes, from the original Appalachian watersheds where it was developed to mountainous regions and coastal plains. However, this adaptability comes with the challenge of determining appropriate regional coefficients, which typically requires calibration against observed data from similar watersheds in the region.

Snyder's approach bridges purely empirical and physically based methods. While not directly derived from physical principles, it acknowledges that watershed geometry meaningfully influences hydrological response, creating an intuitive connection between watershed characteristics and hydrograph parameters. The method has proven particularly valuable for medium to large watersheds and continues to influence modern hydrological engineering, including integration into contemporary software like HEC-HMS.

The practical implementation of Snyder's method involves not only calculating peak discharge and timing but also constructing the complete shape of the hydrograph. Parameters such as hydrograph width at 50% and 75% of peak flow help define this shape, enabling engineers to estimate not only when and how high a flood will peak but also how long elevated flows will persist—critical information for infrastructure design and emergency planning.

Geomorphological Unit Hydrograph: Catchment Form as Hydrological Function

The Geomorphological Unit Hydrograph (GUH) represents a departure from earlier synthetic approaches. Rather than relying primarily on empirical relationships or standardised shapes, GUH explicitly connects hydrograph characteristics to the structure of drainage networks within watersheds. This approach, pioneered by Rodriguez-Iturbe and Valdes in 1979, recognises that river networks encode information about how watersheds process rainfall.

The GUH approach draws on Horton's stream ordering laws, which quantify river networks' hierarchical structure. When analysed through the lens of Horton-Strahler ordering, these networks reveal remarkable regularities across diverse landscapes—patterns in how streams branch,

how their lengths increase with order, and how drainage areas expand. The GUH approach posits that these network characteristics fundamentally determine how rainfall is transformed into runoff.

Unlike other methods that treat watersheds as black boxes or simplified geometric shapes, GUH acknowledges the watershed as a complex system where water flows through a network of pathways with varying travel times. The approach models runoff as a probabilistic process where water particles move through different stream orders before reaching the watershed outlet, with each transition governed by geomorphological ratios derived from the stream network.

This statistical translation of watershed structure into runoff dynamics makes GUH particularly valuable for ungauged basins where traditional calibration data are unavailable. Hydrologists can develop reasonably accurate hydrograph estimates without historical flow records by extracting network characteristics from maps or digital elevation models—a significant advantage in data-sparse regions.

However, the GUH approach introduces complexity in both its conception and application. The original formulations involve probability theory and differential equations, which can be challenging to implement in practice. This complexity has led to various simplifications and approximations, such as the use of equivalent Nash cascade models that preserve key GUH properties while enhancing computational tractability.

The GUH approach also raises interesting theoretical questions about the linearity assumption inherent in unit hydrograph theory. Because GUH parameters like time to peak depend on flow velocity, which may vary during storm events, the method contains implicit non-linearities that challenge traditional unit hydrograph assumptions. Some hydrologists have suggested that "geomorphological response function" might be a more accurate term than "unit hydrograph" for this approach.

Despite these complexities, GUH represents an essential bridge between traditional hydrology and the emerging field of hydrogeomorphology, recognising that watershed form and hydrological function are intimately connected. Modern GIS capabilities have made GUH implementation more accessible, allowing extraction of network parameters from digital elevation models rather than laborious manual map analysis.

In practice, these synthetic unit hydrograph approaches are not competing alternatives but complementary tools in the hydrologist's toolkit, each with distinct strengths and limitations. The SCS method offers accessibility and standardisation, Snyder's approach provides regional sensitivity and geometric consideration, while GUH establishes explicit connections to watershed structure and network properties.

Modern hydrological practice often involves hybrid approaches that combine elements of different methods. For example, regionalised SCS parameters might be developed through GUH analysis, or Snyder's coefficients might be adjusted based on stream network characteristics. Digital tools and geographic information systems have revolutionised all these methods, enabling rapid parameter extraction and reducing the historical barriers to implementation.

The evolution of synthetic unit hydrograph methods reflects a broader trend in hydrology toward physically based approaches that capture the underlying processes governing watershed response. While perfectly representing these complex processes remains challenging, each generation of methods has moved closer to linking hydrograph characteristics to fundamental watershed properties.

Despite sophisticated computational models now available, these synthetic unit hydrograph approaches remain valuable for their balance of physical intuition and practical applicability. They offer conceptual frameworks that help practitioners understand watershed behaviour, not just predict it. This conceptual understanding is particularly valuable in ungauged basins where data limitations preclude more data-intensive approaches.

As land use changes and watershed characteristics modify, these synthetic methods evolve. Regional coefficients require periodic recalibration, and the fundamental relationships between watershed form and hydrological function may shift in response to changing conditions. Yet the core principles—connecting physical watershed characteristics to runoff patterns through standardised shapes, empirical relationships, or network properties—remain as relevant as ever in modern hydrological practice.

Mathematical representations have further enriched the hydrograph modelling toolkit, providing rigorous frameworks for understanding catchment response.

Nash Cascade: Visualising Catchment Response

Formula 4: Nash Cascade Catchment Response

The Nash cascade model conceptualises the catchment as a series of n linear reservoirs, each with storage coefficient K , yielding an impulse response function:

$$h(t) = (1/K(n-1)!) * (t/K)^{(n-1)} * e^{(-t/K)}$$

Where $h(t)$ represents the instantaneous unit hydrograph ordinate at time t , this elegant formulation captures the essential features of catchment response while maintaining mathematical tractability.

The convolution integral forms the mathematical backbone of unit hydrograph application:

$$Q(t) = \int[0 \text{ to } t] I(\tau) * h(t-\tau) d\tau$$

Where $Q(t)$ is the direct runoff hydrograph, $I(\tau)$ is the excess rainfall intensity, and $h(t-\tau)$ is the unit hydrograph ordinate. In discrete form, this becomes:

$$Q(t) = \sum[j=1 \text{ to } m] P(j) * U(t-j+1)$$

where $P(j)$ represents excess rainfall in period j , and $U(t-j+1)$ is the unit hydrograph ordinate.

The Nash cascade model offers an elegant conceptual framework for understanding how watersheds transform rainfall into streamflow. Imagine a series of connected reservoirs—like a sequence of small ponds—each feeding into the next. This is essentially how the Nash model represents a catchment, with water flowing through multiple storage elements before reaching the outlet.

In this model, the number of reservoirs (n) and their storage behaviour (K) together determine how quickly or slowly the watershed responds to rainfall. More reservoirs create a smoother, more delayed response, much like how a complex watershed with many storage zones tends to delay and attenuate flood peaks. Similarly, larger storage values (K) represent a watershed that holds water longer before releasing it downstream.

When rain falls on the watershed, this model traces how that water pulse moves through the system. At first, little water reaches the outlet as the reservoirs begin filling. Gradually, flow increases as water cascades through the system. Eventually, the flow peaks and then recedes as the reservoirs progressively empty. This rising, peaking, and falling pattern creates the bell-shaped hydrograph observed in natural streams.

The beauty of the Nash model lies in its ability to capture complex watershed behaviour with just two parameters. A watershed with many storage zones (wetlands, lakes, soil layers) might have more reservoirs, while a flashy urban watershed might use fewer. Similarly, a watershed with high infiltration capacity and substantial groundwater storage would have a more significant storage coefficient than an impervious, steep catchment.

The Convolution Integral: Translating Rainfall Patterns to Streamflow

The convolution integral represents the fundamental mathematical operation connecting rainfall to streamflow using a unit hydrograph. Given its an integral, it accounts for how every bit of rainfall contributes to streamflow over time.

If we consider rainfall as a sequence of individual pulses, each pulse generates its own small streamflow response, resembling the unit hydrograph but scaled by the amount of rainfall. The convolution process effectively stacks these individual responses together, correctly offset in time, to yield the total streamflow.

When an intense rainstorm occurs over several hours, the initial rainfall begins generating streamflow immediately, while subsequent rainfall enhances this primary response. As time goes on, the impact of earlier rainfall decreases, while the effect of later rainfall increases. The convolution integral accurately tracks and adds these overlapping contributions together.

Hydrologists typically use a discretised version of this process, dividing rainfall into short time intervals (perhaps hourly or sub-hourly). For each rainfall interval, they calculate the resulting streamflow contribution across all future time steps. Then, at any specific time, they aggregate the contributions from all previous rainfall to determine the total streamflow.

Our mathematical framework enables hydrologists to convert complex, variable rainfall patterns into anticipated streamflow responses. It underpins nearly all rainfall-runoff modelling applications, from flood forecasting to water resource planning. It provides the crucial link between what falls from the sky and what flows in our rivers.

Modern deconvolution techniques enable the extraction of unit hydrographs from complex observed hydrographs through Fourier analysis or matrix operations, while data-driven approaches leverage machine learning to bypass traditional hydrograph separation requirements. Each method presents distinct advantages and limitations, with selection typically guided by data availability, catchment characteristics, and specific modelling objectives.

The outputs of runoff models provide essential information for risk assessment:

Hydrograph Characteristics:

- Peak flow rate (Q_p) in m^3/s .
- Time to peak (t_p) in hours.
- Hydrograph volume in m^3 .
- Flow duration at specific thresholds.

Derived Metrics:

- Annual maximum series.
- Flow duration curves.
- Flood frequency distributions.
- Base flow separation indices.

Risk Assessment Parameters:

- Inundation extent predictions.
- Alert Exceedance probabilities.
- Above alert duration analysis.

These outputs form the foundation for various risk assessments, from flood plain mapping to infrastructure design and insurance pricing. Depending on the application and catchment characteristics, the temporal resolution of outputs typically ranges from sub-hourly to daily.

Model development follows a structured yet iterative process. It begins with formulating a perceptual model that captures the modeller's understanding of key catchment processes. This conceptual framework is then translated into mathematical expressions, leading to decisions about model structure and parameter relationships. The crucial calibration and validation steps follow, where model parameters are optimised against observed data, and performance is verified using independent datasets.

Runoff models are used in physical risk assessment in multiple domains, from insurance and reinsurance to infrastructure investment and property valuation. For infrastructure investment, they support design flood estimation and adaptation planning. In property valuation, runoff models contribute to flood risk zonation, influencing investment decisions.

The challenge of uncertainty pervades all aspects of runoff modelling, manifesting in parameter estimation, input data quality and model structure. Managing these uncertainties requires careful attention to model selection criteria, considering the purpose of analysis, data availability, required accuracy, and resource constraints. The communication of these uncertainties to stakeholders remains a critical aspect of model application in risk assessment contexts.

Recent Developments in Runoff Modelling

Recent advancements in runoff modelling have been significantly enhanced by integrating machine learning approaches. Long Short-Term Memory (LSTM) networks have demonstrated superior performance in capturing temporal dependencies compared to traditional methods, particularly during high-flow periods.

Hybrid approaches combining LSTM with signal decomposition techniques like Variational Mode Decomposition (VMD) have improved Nash-Sutcliffe Efficiency and reduced root-mean-square error across diverse conditions. Similarly, Random Forest models have outperformed traditional conceptual frameworks by leveraging ensemble techniques to produce more stable and precise runoff estimates, especially when incorporating antecedent runoff data to improve accuracy across seasonal variations.

Data assimilation techniques have transformed how hydrological models integrate real-time observations. The Ensemble Kalman Filter and variational techniques have been widely employed to incorporate near-real-time data such as streamflow and remotely sensed soil moisture, significantly enhancing prediction accuracy.

The development of frameworks like the Parallel Data Assimilation Framework (PDAF) has enabled better integration of high-resolution land surface and meteorological observations across large spatial domains. Despite these advances, challenges persist in addressing non-stationarity in hydro-climatic conditions and managing limited high-quality field data, though machine learning models have shown promise in effectively utilizing sparse datasets to partially mitigate these constraints.

Practical Implementation

The practical implementation of runoff models in risk assessment requires careful attention to best practices in model selection, data management, and result communication. Success depends not only on technical proficiency but also on a clear understanding of model limitations and uncertainties. Regular model reviews and updates are required to ensure continued relevance and accuracy in risk assessment applications.

The dynamic nature of weather and human development patterns ensures that runoff modelling remains an active area of research and development. As our understanding of hydrological processes deepens and computational capabilities expand, the challenge is to balance model sophistication with practical utility in risk assessment applications. This

balance becomes particularly crucial as we face increasing uncertainty in future weather conditions and their implications for water-related risks.

Hydrologically Tracing Water's Journey

Hydrologic routing addresses the specific challenge of tracking how water moves through landscape features once it enters channels, rivers, and reservoirs. This approach emphasises pragmatic simplification while maintaining sufficient physical realism for practical applications.

Hydrologic routing rests upon fundamental principles of water movement. At its core lies the continuity equation - the mathematical expression of mass conservation that states the rate of change in storage equals the difference between inflow and outflow:

Formula 5: Hydrologic Routing Differential

$$dS/dt = I(t) - O(t)$$

Where:

- S represents storage volume
 - I(t) captures the inflow rate at time t
 - O(t) describes the outflow rate at time t
-

This elegant relationship forms the mathematical backbone for tracking water's movement through the landscape. However, a solvable system requires complementary relationships between storage and discharge.

Model Parameters and Implementation

Effective implementation of hydrologic routing involves careful consideration of several parameter types:

Physical Parameters:

- Channel roughness coefficients.
- Cross-sectional geometry parameters.
- Channel bed slope.
- Expansion/contraction coefficients.

Numerical Parameters:

- Time step size (Δt).

- Space step size (Δx).
- Numerical scheme coefficients.
- Convergence tolerances.

Boundary Condition Parameters:

- Upstream flow conditions.
- Downstream stage relationships.
- Lateral inflow coefficients.
- Structure operation rules

The practical implementation involves careful attention to data requirements. Channel geometry data, typically obtained through field surveys or remote sensing, must capture relevant features while maintaining computational efficiency. Boundary conditions require careful consideration to ensure physical realism and numerical stability.

Hydrologic routing methods have evolved distinct approaches for different landscape features:

Water bodies with relatively horizontal surfaces - lakes, reservoirs, and wide floodplains - invite a simplification known as level pool routing. Here, the water surface remains horizontal throughout, allowing storage to be calculated directly from the water level.

The Modified Puls Method and Flow Routing

Formula 6: Modified Puls Flow Rate

The Modified Puls method operationalises this approach through a finite difference approximation:

$$S(j+1) + (\Delta t/2)O(j+1) = S(j) + (\Delta t/2)O(j) + (\Delta t/2)(I(j) + I(j+1))$$

Where:

- $S(j)$ is the storage volume at time step j .
- $S(j+1)$ is the storage volume at the next time step $j+1$.
- $O(j)$ is the outflow rate at time step j .
- $O(j+1)$ is the outflow rate at the next time step $j+1$.
- $I(j)$ is the inflow rate at time step j .

- $I(j+1)$ is the inflow rate at the next time step $j+1$.
 - Δt is the time step interval.
-

The Modified Puls method provides a straightforward way to track water movement through reservoirs, channel reaches, or storage areas. At its core, this method balances what comes in, goes out, and stays stored in the system.

This method is particularly useful because it handles the relationship between storage and outflow. In natural systems, outflow typically depends on storage—more water stored means higher pressure and faster outflow. The Modified Puls method accounts for this by using a weighted average of outflows (both current and future) rather than just a single value.

In practice, engineers first establish a relationship between storage and outflow for the specific system (like a reservoir or river reach) through direct measurement or hydraulic calculations. Then, for each time step in a storm event, they solve for the unknown future outflow using known values of current storage, current outflow, and inflows (both current and future).

The method is especially valuable because it preserves mass conservation while being computationally simple. It's widely used in flood routing applications, reservoir operations, and stormwater management because it balances accuracy with practicality. Modern software packages often implement Modified Puls routing because it provides reliable results without demanding excessive computational resources or complex data inputs.

Channel Routing: Muskingum

Rivers and streams present a different challenge, as water moves longitudinally with varying speeds and depths. The Muskingum method, developed in the 1930s for flood control studies in the Muskingum River basin, addresses this through a conceptual model relating storage to weighted combinations of inflow and outflow:

Formula 7: Muskingum River Basin

$$S = K[xI + (1-x)O]$$

Where:

- S is the storage volume in the river reach (m^3).
- I is the inflow rate to the reach (m^3/s).
- O is the outflow rate from the reach (m^3/s).
- K represents a storage constant approximating travel time through the reach (seconds).
- x serves as a weighting factor (between 0 and 0.5), balancing the influence of inflow and outflow
 - $x = 0$ represents pure translation (reservoir-type storage)
 - $x = 0.5$ represents pure translation (kinematic wave)
 - Typical values range from 0.1 to 0.3 for natural channels

This conceptualisation leads to an efficient routing equation:

$$O(j+1) = C_0 \cdot I(j+1) + C_1 \cdot I(j) + C_2 \cdot O(j)$$

Where:

- $O(j+1)$ is the outflow at the next time step $j+1$.
- $I(j+1)$ is the inflow at the next time step $j+1$.
- $I(j)$ is the inflow at the current time step j .
- $O(j)$ is the outflow at the current time step j .
- C_0 , C_1 , and C_2 are routing coefficients derived from K , x , and Δt (time step)

The routing coefficients are calculated as:

$$C_0 = (-Kx + 0.5\Delta t)/(K - Kx + 0.5\Delta t) \quad C_1 = (Kx + 0.5\Delta t)/(K - Kx + 0.5\Delta t) \quad C_2 = (K - Kx - 0.5\Delta t)/(K - Kx + 0.5\Delta t)$$

Routing coefficients (C_0 , C_1 , and C_2) are derived from the physical parameters K and x and the computational time step. This approach has become a workhorse for operational hydrology, offering accuracy when adequately calibrated.

Considering channels as having both "prism storage" (regular channel volume) and "wedge storage" (additional volume from flood waves), the Muskingum method captures how flood waves change shape as they move downstream. The method's conceptual foundation lies in relating the water stored in a river reach to both inflow and outflow rather than to either value alone.

The storage constant approximates how long water typically takes to travel through the river reach, representing the natural lag in the system. The weighting factor x indicates how much the storage in the reach depends on inflow versus outflow. When x equals zero, storage depends entirely on outflow (like a level pool), while higher values of x (though typically not exceeding 0.5) indicate more significant influence of the inflow.

The Muskingum method transforms these physical concepts into a straightforward computational procedure. For each time step, the outflow is calculated using a weighted combination of current and previous inflows plus the previous outflow. This approach effectively tracks how flood waves attenuate (reduce in peak) and lag (shift in timing) as they move downstream.

River engineers and hydrologists value the Muskingum method for several reasons. It preserves the conservation of mass, runs efficiently even with limited computational resources, requires minimal data inputs, and conceptually aligns with physical river processes. Properly calibrating using observed inflow-outflow data provides reliable results for many natural river systems. These benefits have established the Muskingum method as a standard tool in flood forecasting, water resource planning, and reservoir operation studies.

Muskingum Calculations

To perform Muskingum calculations, you need the following data:

- Streamflow data is usually measured in cubic feet per second (cfs).
- A roughness coefficient, C , characterises the stream's resistance to flow.
- A time lag, T , represents the delay between changes in upstream conditions and downstream responses.

Once you have this data, you can calculate the Muskingum parameters:

- Instantaneous discharge (Q) at each time step.
- Area-averaged velocity (V) at each time step.

- Storage (S) at each time step.
- Time of travel (Tt) between upstream and downstream locations.

Once you have these coefficients, you can use them to estimate the discharge at an ungauged location using the Muskingum-Cunge method. This method assumes that the relationship between discharge and upstream and downstream conditions is linear, which may not always be accurate but is often a good approximation.

You can calculate the Muskingum parameters using the following equations:

Instantaneous discharge (Q): $Q = C \cdot A \cdot V^2 / (2 \cdot g)$, where A is the cross-sectional area, g is the acceleration due to gravity, and V is the average velocity.

Storage (S): $S = A \cdot (h_{down} - h_{up}) / 1000$, where h_{down} and h_{up} are the water surface elevations downstream and upstream of a given cross-section.

Time of travel (Tt): $Tt = L / V$, where L is the reach length between the upstream and downstream locations.

Area-averaged velocity (V): $V = Q / A$.

After calculating these parameters, you can use them to compute the Muskingum coefficients:

- Coefficient a: $a = [(Q_{down} - Q_{up}) / (Q_{down} + Q_{up})] / 2$.
- Coefficient b: $b = [(S_{down} - S_{up}) / (S_{down} + S_{up})] / 2$.
- Coefficient c: $c = [(Tt_{down} - Tt_{up}) / (Tt_{down} + Tt_{up})] / 2$.

This method can help estimate discharge at locations where flow measurements are unavailable or difficult to obtain. However, the accuracy of the estimates depends on the accuracy and representativeness of the upstream and downstream data and the Muskingum coefficients.

Muskingum-Cunge

Over decades of hydrological practice, the basic Muskingum method has been enhanced numerous times. The Muskingum-Cunge method introduced a physical basis for parameter estimation based on channel properties, allowing its application to ungauged basins. Nonlinear

Muskingum methods addressed the limitation of constant parameters, recognising that storage-discharge relationships often vary with flow magnitude.

Using the Muskingum-Cunge method requires data for the upstream and downstream locations and coefficients a, b, and c. You can then estimate the discharge at an ungauged location using the following equation:

Formula 8: Muskingum-Cunge Discharge

$$Q_{\text{est}} = Q_{\text{up}} * (1 + a * (1 - \exp(-b * (h_{\text{down}} - h_{\text{up}}) / T))) * (1 - c * \exp(-h_{\text{down}} - h_{\text{up}}) / T))$$

Where:

- Q_{est} : Estimated discharge at the ungauged location
 - Q_{up} : Discharge at the upstream location
 - h_{down} : Water surface elevation at the downstream location
 - h_{up} : Water surface elevation at the upstream location
 - T: Time constant, typically related to the travel time
 - a, b, c: Muskingum coefficients calculated from observed data
 - a: Coefficient related to the relative change in discharge
 - b: Coefficient related to the relative change in storage
 - c: Coefficient related to the relative change in travel time
-

The outputs of hydrologic routing provide essential information for risk assessment applications. Water surface profiles enable inundation mapping, while calculated discharges support infrastructure design and floodplain management. The temporal evolution of these parameters offers crucial insights into flood wave propagation and its implications for infrastructure and property exposure.

Modern developments in hydrologic routing have embraced technological advances in computational capabilities and data availability:

- Enhanced computational methods allow for more detailed simulations.
- Improved remote sensing technologies provide better geometric data.
- Integration of real-time monitoring systems supports operational flood forecasting.

- Machine learning approaches complement traditional routing methods.

The intersection of physical process understanding and artificial intelligence presents the most exciting frontier in hydrological routing. Machine learning shows particular promise in addressing long-standing challenges:

- Parameter estimation through machine learning can establish relationships between observable landscape features and difficult-to-measure parameters.
- Model error correction using neural networks can identify and compensate for systematic biases.
- Computational efficiency improvements through surrogate modelling can accelerate simulations dramatically.

Balancing Theory and Practice

The balance between theoretical rigour and practical utility remains a central consideration in hydrologic routing applications. Computational demands increase significantly with model complexity, necessitating pragmatic decisions about spatial resolution and temporal discretisation. Data requirements similarly scale with model sophistication, often necessitating simplifications where detailed information is unavailable.

This theoretical-practical tension manifests in selecting model simplifications appropriate to specific applications:

- More straightforward methods may offer sufficient accuracy for many flood forecasting applications.
- Detailed floodplain mapping may demand more sophisticated approaches.
- Infrastructure design applications may require detailed local modelling and simplified routing for broader system representation.

Uncertainty Considerations

Despite these advances, uncertainty remains an inherent aspect of all routing applications. This uncertainty stems from multiple sources:

- Roughness coefficients and geometry data uncertainty.
- Model structure uncertainty.

- Future land use changes.

Communicating these uncertainties to stakeholders remains crucial for informed decision-making. Modern routing applications increasingly embrace ensemble approaches, sensitivity analysis, and formal uncertainty quantification to provide decision-makers with a complete picture of what is known, what is assumed, and what remains uncertain.

Despite being computationally intensive in some implementations, this advanced framework is an excellent candidate for AI development. It offers the detailed analysis needed for complex systems where more straightforward methods fall short. The ongoing challenge of balancing physical accuracy with computational efficiency continues to propel advancements in the field, highlighting its increasing importance in physical risk assessment applications.

Hydraulic Modelling

Hydraulic modelling is a fundamentally different approach to water flow analysis than hydrologic routing. The distinction is rooted in their mathematical foundations.

While hydrologic routing relies solely on the continuity equation and empirical storage-discharge relationships, hydraulic modelling is firmly grounded in the Saint-Venant equations, which incorporate mass conservation and momentum principles.

“This critical difference means that hydraulic models solve the complete physics of water movement, accounting for acceleration terms, pressure forces, and complex flow interactions that hydrologic methods cannot capture.”-

David Kelly, MKM

Consequently, hydraulic modelling demands significantly more detailed inputs—including precise cross-sectional geometry, roughness coefficients, and channel characteristics—but provides superior resolution of flow dynamics, water surface profiles, and velocity distributions.

Saint-Venant Equations

The Saint-Venant equations, developed by Adhémar J. C. Barré de Saint-Venant in 1871, describe one-dimensional, unsteady flow in open channels through a system of partial differential equations. These equations emerge from the fundamental principles of mass and momentum conservation, providing a mathematical framework that captures the essential physics of

open channel flow while maintaining computational tractability through carefully considered assumptions.

Formula 9: Saint-Venant Equations

The equations consist of:

Continuity equation (Conservation of Mass):

$$\partial h / \partial t + \partial(uh) / \partial x + \partial(vh) / \partial y = 0$$

Where:

h: Water depth

u, v: Velocity components in x and y directions

Momentum equations in the x-direction and y-direction:

$$\partial(uh) / \partial t + \partial(u^2h + \frac{1}{2}gh^2) / \partial x + \partial(uvh) / \partial y = gh(S_{0x} - S_{fx})$$

$$\partial(vh) / \partial t + \partial(uvh) / \partial x + \partial(v^2h + \frac{1}{2}gh^2) / \partial y = gh(S_{0y} - S_{fy})$$

Where:

g: Gravitational acceleration

S_{0x} , S_{0y} : Bed slopes in x and y directions

S_{fx} , S_{fy} : Friction slopes in x and y directions

These equations form the mathematical foundation for hydraulic modelling and are widely used in engineering applications for water resource management, flood prediction, and environmental studies.

The derivation of the Saint-Venant equations is built on a foundation of key assumptions that balance physical accuracy and practical applicability. The framework assumes a one-dimensional flow with a hydrostatic pressure distribution, relatively small channel slopes, and uniform velocity distribution within each cross-section. Regular channel geometry is presumed, while wind and turbulence effects are negligible.

While simplifying the mathematical treatment, these assumptions preserve the essential channel flow characteristics necessary for most practical applications, instilling confidence in the model's reliability.

Different simplifications of these equations yield various practical modelling approaches:

- The **dynamic wave approach**, employing the full Saint-Venant equations, includes all terms and provides the most complete representation of flow dynamics, though at the cost of computational intensity.
- The **diffusion wave approximation** neglects the acceleration terms while retaining the pressure term, making it suitable for gradually varied flows where inertial effects are less significant. The kinematic wave represents the simplest form, neglecting both acceleration and pressure terms, applicable primarily to steep slopes where gravity and friction forces dominate the flow behaviour.
- The **numerical solution** of these equations typically employs finite difference, finite element, or finite volume methods. Each offers distinct advantages regarding stability, accuracy, and computational efficiency. The choice of numerical scheme often depends on specific application requirements and computational constraints.

One and Two-Dimensional Approaches

Hydraulic models are classified based on their spatial dimensionality, with important implications for applications in flood risk assessment. One-dimensional (1D) models represent flow along a single spatial dimension, typically the channel centerline. These models use cross-sections perpendicular to the flow direction to capture channel geometry and solve the 1D Saint-Venant equations for flow and water surface elevation.

They are most appropriate for channels with uniform floodplains and predominantly longitudinal flow. However, they are limited in representing complex lateral flows that often characterise extensive floodplains, which is an essential consideration in their application.

Two-dimensional (2D) models account for flow variations in longitudinal and lateral directions by solving the 2D shallow water equations, extending Saint-Venant principles to multiple dimensions. These models better represent complex terrain, meandering rivers, and irregular flows over floodplains, providing more detailed spatial information about flood characteristics. However, they require significantly more data and computational resources, which can constrain their application in resource-limited situations.

“Coupled 1D-2D models combine the strengths of both approaches by using 1D representation for channel flow and 2D for floodplain flow.” - Johnny Mattimore, MKM.

This hybrid approach allows computational efficiency in well-defined channels while capturing complex floodplain dynamics. Such models require careful interface treatment between 1D and 2D domains to ensure proper mass and momentum transfer.

As outlined in the door-stopper and seminal Flood Handbook Analysis and Modelling, edited by Saeid Eslamian and Faezeh Eslamian, recent advancements in hydraulic modelling have seen coupled 1D-2D models emerge as the preferred approach for flood risk assessment due to their ability to balance computational efficiency with spatial accuracy.

These hybrid systems integrate the strengths of both modelling approaches—employing 1D solutions for efficient channel routing while leveraging 2D frameworks for detailed floodplain dynamics. Research from the handbook demonstrates that coupled models can reduce simulation times by 30–50% compared to complete 2D models while improving flood extent accuracy by 15–25% over standalone 1D systems.

Technological progress has driven this evolution, including improved coupling algorithms that ensure momentum conservation and high-performance computing capabilities that enable basin-scale applications.

The increasing preference for coupled 1D-2D approaches is further evidenced by their successful implementation in complex environments, such as Vietnam's Mekong Delta, where the F28 model effectively captured intricate infrastructure-floodplain interactions while maintaining computational feasibility.

Similarly, regulatory bodies have recognised these advantages, with FEMA guidelines now specifically recommending coupled approaches for systems with disconnected floodplain flows. As computational resources expand, these integrated modelling systems have increasingly become the standard methodology for flood risk assessments requiring both hydraulic precision and operational practicality, particularly in lowland deltas and urbanised floodplains where infrastructure significantly alters natural flow patterns.

Model Requirements and Implementation

Successfully applying hydraulic modelling requires careful data inputs. Input quality directly impacts model reliability. Topographic data forms the foundation for channel and floodplain geometry. Digital Terrain Models (DTMs) represent terrain's bare earth surface. They remove all vegetation and human-made structures from the landscape. LiDAR technology generates these models efficiently. Light Detection and Ranging uses laser pulses to measure distances to Earth. It creates precise elevation measurements over large areas. Survey methods can also produce DTMs when LiDAR isn't available.

High-resolution DTMs improve flood predictions dramatically. They can increase accuracy by up to 30% compared to simplified models. DTM generation faces challenges in certain environments. Dense vegetation blocks laser penetration. Water surfaces reflect signals unpredictably. Supplementary data often fills these gaps. Comprehensive data collection ensures robust hydraulic modelling.

DTMs serve as bare-earth elevation models that capture the underlying terrain by filtering out vegetation and structures from raw data, enabling high-precision representation of riverbed morphology and floodplain features with 10-15 cm vertical accuracies.

LiDAR technology has revolutionised topographic data collection by using laser pulses to measure distances to the Earth's surface, generating dense point clouds that can be processed to achieve spatial resolutions of 1m or finer. This level of detail allows hydraulic models to detect crucial microtopographic elements such as riverbanks, levees, and floodplain depressions that significantly influence water movement during flood events.

Channel and floodplain roughness parameters, expressed as Manning's n coefficients, characterise the resistance to flow and significantly affect predicted water levels and velocities. Boundary conditions, including inflow hydrographs and stage-discharge relationships, define how water enters and leaves the model domain.

Detailed representation of infrastructure elements such as bridges, culverts, and levees ensures their hydraulic effects are appropriately captured. For unsteady flow simulations, initial conditions define the system's starting state.

The integration of high-resolution DTMs with these hydraulic parameters has transformed flood prediction capabilities. In 2019, Xia et al.'s National Water Model (NWM) improvement project documented significant gains when transitioning from simplified trapezoidal channels to realistic geometries derived from LiDAR data.

Xia demonstrated that replacing simplified channel assumptions with realistic geometries extracted from LiDAR-derived DTMs can improve flood stage and discharge predictions by up to 30%.

Despite their advantages, DTM generation faces challenges in areas where dense vegetation or water surfaces obscure ground returns. This often requires supplementary data from sonar surveys or ground measurements to fill gaps in riverbed bathymetry, highlighting the importance of comprehensive data collection approaches for robust hydraulic modelling.

Model Parameters:

- **Physical parameters:** roughness coefficients, expansion/contraction coefficients.
- **Numerical parameters:** time step size, space step size, convergence tolerances.
- **Boundary condition parameters:** upstream flow conditions, downstream stage relationships.

The computational implementation involves several important considerations that affect model performance and reliability. Spatial resolution must balance detail with computational efficiency. The finer resolution provides more accurate results but requires more computational resources.

Temporal discretisation significantly impacts unsteady simulation stability, accuracy, and computational efficiency. Time steps influence how flow changes are captured over time. Smaller time steps generally increase stability by better resolving rapid flow transitions. However, they also increase computational demands. Different numerical schemes offer various trade-offs. Explicit schemes solve equations directly for each time step but require smaller steps to maintain stability. Implicit schemes can use more significant time steps but involve more complex calculations at each step. The Courant-Friedrichs-Lowy (CFL) condition, which relates time step size to grid spacing and flow velocity, provides a key stability criterion.

Recent advancements in computing power have enabled more sophisticated approaches. Research by Oleg Zikanov (2010) established practical guidelines stating that time steps should be less than hydrograph rise time divided by 20 for flood modelling. The theta weighting factor balances stability and accuracy in methods used by software like HEC-RAS (more of which are in the next section).

GPU acceleration has transformed simulation capabilities. Studies by Eslamian and Eslamian (2017) demonstrated that parallel CUDA architectures enable million-particle Smoothed Particle Hydrodynamics simulations with 100 \times speedups compared to CPUs. Adaptive time-stepping approaches dynamically adjust step size based on error estimates, maintaining stability while minimising computational cost.

Modern workflows increasingly combine multiple techniques. Bengt Andersson's work (2011) on error transport equations co-solved with primary equations shows how localised discretisation errors can enable

iterative correction for improved accuracy. Knight and Shamseldin's research (2009) on hybrid implicit-explicit schemes leverages GPUs for stiff terms while explicitly handling non-stiff components, optimising stability and speed. These developments allow hydraulic modellers to achieve high-fidelity simulations of complex phenomena like turbulence and multiphase flows within practical timeframes.

HEC-RAS: Industry Standard Implementation

The Hydrologic Engineering Center's River Analysis System (HEC-RAS), developed by the U.S. Army Corps of Engineers, is the global industry standard for implementing hydraulic modelling principles. Its comprehensive capabilities, including one-dimensional steady and unsteady flow modelling, two-dimensional unsteady flow modelling, and combined 1D/2D approaches, have established it as a worldwide cornerstone tool for flood risk assessment. In addition to hydraulic calculations, it provides modules for sediment transport with movable boundary computations and water quality analysis, ensuring a holistic assessment of water resource systems.

The software's advanced 2D modelling features represent a significant evolution in hydraulic modelling capabilities. Subgrid bathymetry allows detailed terrain representation even with relatively coarse computational meshes, improving efficiency without sacrificing accuracy.

This innovative approach, detailed in the HEC-RAS 2D User's Manual version 3 (2024), enables computational models to represent fine-scale underwater terrain features while using coarse computational grids. The system works by preprocessing detailed bathymetric data into hydraulic property tables for each coarse grid cell. These tables store critical relationships between water elevation, wetted area, volume, and hydraulic roughness at subgrid scales.

A key advantage is the handling of partially submerged terrain. Subgrid bathymetry is a computational technique that incorporates high-resolution underwater terrain data (bathymetry) within larger computational grid cells, essentially storing detailed elevation information below the visible resolution of the model's mesh.

Rather than simplifying terrain by averaging elevations within each cell, this approach preserves critical small-scale features like narrow channels, depressions, and ridges by maintaining their hydraulic properties in lookup tables that inform calculations even when these features are smaller than the grid cell.

Cells can track narrow channels within larger grid elements using elevation-dependent volume curves, thus avoiding simplified binary "wet/dry" approximations. This preservation of flow dynamics in small features

maintains model fidelity, even with larger cell sizes. The Global Modelling and Assimilation Office's 2025 study demonstrates that this approach enables flood modelling across entire basins while still resolving meter-scale hydraulic features.

Coarser grids dramatically lower computational demands—for example, using 500-meter cells instead of 5-meter cells reduces total cell count by a factor of 10,000. The HEC-RAS Model Library (2022) documents how these larger cells enable stable simulations with longer timesteps due to relaxed Courant–Friedrichs–Lewy conditions.

Despite this computational efficiency, accuracy is maintained through terrain fidelity. Small channels are less than 100 meters wide are preserved via subgrid elevation-volume curves, preventing artificial widening that plagued earlier modelling approaches. The software even tracks sediment movement at subcell resolution, avoiding homogenisation issues across large computational elements.

Implementation of rainfall-runoff modelling with infiltration enables direct simulation of pluvial flooding processes. Sophisticated representation of hydraulic structures ensures that bridges, culverts, weirs, and similar features are appropriately integrated into the hydraulic calculations.

Levee and floodwall modelling capabilities permit the evaluation of flood defence systems under various loading conditions. Flexible hydraulic connections between 1D and 2D domains allow appropriate representation of complex flow paths. Wetting and drying algorithms accurately simulate the dynamic progression of inundation across initially dry areas, which is crucial for realistic flood mapping.

Through its RAS Mapper interface, HEC-RAS provides an integrated geospatial environment, supporting terrain model development, comprehensive spatial data management, and advanced visualisation of hydraulic results. This integration with GIS platforms facilitates seamless workflows from data preparation through result analysis and presentation. Computational enhancements, including parallel computing capabilities and GPU acceleration options, allow efficient execution of complex simulations. At the same time, adaptive time stepping improves computational efficiency by adjusting temporal resolution based on flow conditions.

The model setup process involves a systematic approach to represent the physical system accurately. Preparing terrain data incorporates LiDAR, survey data, and other sources to describe the model domain accurately. Developing the computational mesh establishes 1D cross-sections and/or

2D computation cells at appropriate locations and densities to capture relevant hydraulic features.

Specifying boundary conditions defines how water enters and exits the system while assigning hydraulic parameters characterising the flow resistance and structural properties. After executing the simulation, analysing results provides insights into flow characteristics relevant to risk assessment.

HEC-RAS provides a robust framework for implementing the Saint-Venant equations and their variations in practical applications, making it an essential tool for detailed flood risk assessment. Its combination of theoretical rigour, practical usability, and comprehensive features has established it as the preferred platform for hydraulic modelling in professional practice and research applications.

Applications and Limitations

Hydraulic modelling, particularly with advanced tools like HEC-RAS, offers powerful capabilities for flood risk assessment but comes with important considerations. The detailed representation of flow physics enables numerous valuable applications in risk assessment and management. Detailed floodplain delineation provides spatial information on flood extent and characteristics, which are crucial for zoning, insurance, and emergency planning.

Infrastructure design and analysis ensures that bridges, culverts, levees, and other structures perform adequately under design flood conditions. Evaluation of flood mitigation alternatives allows comparison of different intervention strategies to identify optimal solutions. When coupled with hydrologic models, real-time flood forecasting warns of flood conditions.

Outputs for Risk Assessment:

- Inundation extent and depth.
- Flow velocity distributions.
- Water surface profiles.
- Arrival time of flood peaks.
- Shear stress for erosion potential.

Despite its capabilities, hydraulic modelling faces several limitations and challenges that influence its application. Extensive data requirements can constrain implementation in areas with limited topographic or hydrometric information. Significant computational demands may necessitate simplifications or reduced resolution in large or complex systems.

The need for skilled practitioner expertise means that model quality depends heavily on the modeller's experience and judgment. Uncertainty in boundary conditions and parameters propagates through the model, affecting the reliability of results. Model simplifications may not capture all relevant processes, particularly in highly complex or unconventional hydraulic situations.

Several promising developments are shaping the future of hydraulic modelling. Enhanced integration with real-time monitoring systems enables continuous model updating and improved forecast accuracy. Improved representation of complex hydraulic structures allows better simulation of engineered systems.

More efficient numerical schemes lower computational demands while preserving accuracy. Improved uncertainty quantification and communication aid decision-makers in grasping model results' limitations and confidence levels. Integration with artificial intelligence presents significant advancements in model calibration and prediction capabilities, potentially lowering the expertise barrier for effective modelling.

Hydraulic modelling provides a detailed analysis of complex systems where more straightforward hydrologic routing methods fall short. The ongoing challenge of balancing physical accuracy with computational efficiency continues to drive advancements in the field, highlighting its increasing importance in physical risk assessment applications. As computational capabilities advance and data availability improves, the scope and detail of hydraulic modelling applications will continue to expand, enhancing our ability to understand and manage flood risk in an increasingly complex and changing environment.

Chapter 5 - Flood Risk Assessment

While hydrological modelling provides the foundation for understanding water movement, assessing flood risk requires a further layer of analysis that bridges the gap between water dynamics and their impacts on the built environment. The preceding chapter established how we model the fundamental physics of water movement; this chapter examines how we translate that understanding into practical flood risk assessments that inform banking, insurance, and property decisions.

"Flood risk assessment is where the abstract physics of water flow meets the concrete reality of human settlement. It translates hydrodynamic principles into the language of vulnerability, exposure, and banking impact."

Regarding property valuation, the eternal adage of location, location, location also applies to flood hazards." - David Kelly, MKM

Overall, flooding is recognised as the most pervasive and financially devastating natural disaster, with its impacts projected to worsen due to climate change, population growth in flood-prone areas, and continued economic development in high-risk zones. This is supported by several studies:

- **Frequency and Impact:** According to the World Health Organization, floods are the most frequent natural disaster, affecting over 2 billion people between 1998 and 2017. They cause widespread devastation, loss of life, and significant economic damage. The International Association of Hydrological Sciences, in cooperation with UNESCO, emphasises that floods represent the most common natural disaster worldwide, with their frequency increasing dramatically over the past two decades. According to Eslamian and Eslamian (2017), flood

events accounted for 43% of all recorded natural disasters between 1995 and 2015, affecting more communities than any other hazard type.

- **Global Costs:** Research published in the "Flood Handbook" by Eslamian and Eslamian (2017) highlights that floods are among the most devastating natural hazards, with global costs significantly increasing due to factors such as climate change and expanded development in floodplains. Their analysis documents that floods regularly claim over 20,000 lives annually, impacting approximately 75 million people worldwide. Economic assessments from Pender and Faulkner's 2011 "Flood Risk Science and Management" indicate that annual global flood damages range from \$50 to \$60 billion, projected to increase substantially by 2050.

- **Economic Losses:** Knight and Shamseldin's comprehensive study (2009) demonstrates that flood-related losses have grown exponentially, with economic damages increasing sevenfold after adjusting for inflation between the 1960s and 2000s. Their research shows that these impacts disproportionately affect developing nations, where flood damages can represent up to 15% of annual GDP in severely impacted regions. Teegavarapu's analysis (2012) further quantifies that flood events account for approximately 40% of all global economic losses from natural disasters.

- **Regional Vulnerability:** Teegavarapu's 2012 assessment identified East Asia and South Asia as particularly vulnerable regions, with approximately 1.36 billion people exposed to flooding in densely populated, high-risk areas. The Federal Emergency Management Agency analysis presented in the 2019 Flood Risk Report shows that floods represent the most common hazard in the U.S., accounting for over 80% of all federal disaster declarations between 1953 and 2018, with average annual losses exceeding \$8 billion.

At its core, flood risk is the product of two essential components: the probability of a flood hazard occurring and the vulnerability or potential consequences stemming from that event. This dual nature of risk—combining likelihood with impact—forms the fundamental principle underlying all assessment methodologies.

The accurate quantification of flood risk requires specialised analytical processes that go beyond simply modelling water movement. A comprehensive flood risk analysis quantifies the probabilities and consequences of potential flood events, serving as a critical subset of the broader flood risk management process. This management approach represents a cyclical, continuous effort to reduce flood impacts through structural and non-structural measures, with risk assessment providing the analytical foundation for decision-making.

Modern flood risk analyses leverage increasingly sophisticated technology, including advanced software tools, high-performance computing, high-resolution datasets, LiDAR elevation models, and detailed building inventories.

These technologies enable unprecedented precision in evaluating flood impacts across large geographic areas and diverse property portfolios. The outputs of these analyses are often visualised through flood risk mapping, which represents hazard and risk information spatially and shows inundation extent, depth, and velocity for different probability scenarios.

The essential first step in this process is classifying flood types. While all floods involve excess water, their dynamics, predictability, and impacts vary dramatically depending on their source and characteristics. Riverine flooding follows different patterns than coastal storm surges, and flash floods present distinct challenges compared to gradual groundwater rise. Each classification entails different modelling approaches, warning timeframes, and mitigation strategies.

Property-specific assessment forms the crucial next layer of analysis. A structure's position within the landscape—its elevation, proximity to water bodies, and relationship to surrounding terrain—fundamentally determines its exposure to flood hazards. Practitioners must analyse this geospatial positioning and broader environmental factors, including watershed characteristics, natural and local flood defences, and weather pattern change projections that may alter historical patterns.

The resilience of individual properties completes this assessment picture. Structural characteristics, adaptation measures, and system redundancies all influence how a property will respond when flood waters arrive. This component of assessment bridges purely physical analysis with the practical concerns of property owners, insurers, and lending institutions.

As we explore these dimensions of flood risk assessment, we extend the hydrologic and hydraulic models for water flow into frameworks that support the initial assessment of property exposure to a hazard such as flood.

While significant progress has been made in flood risk assessment methodologies, standardised practices and agreed indicators for flood risk mapping, especially at local scales, continue to evolve across different regions and countries. This progression—from hydrological modelling to flood risk assessment—ultimately enables the quantification of impacts discussed in the following chapter.

Flood Classification Systems

The systematic categorisation of flood events is the foundation for practical risk assessment. Classification systems provide the necessary framework to differentiate between flood types, each with distinct characteristics influencing their modelling, prediction, and mitigation approaches. This section explores established and emerging flood classification methodologies, their applications, and limitations in the context of physical risk assessment.

Traditional Classification Methods

Historically, floods have been classified primarily by their source mechanism, with each category representing fundamentally different hydrodynamic processes:

- **Fluvial (Riverine) Flooding:** Occurs when water exceeds the capacity of river channels, causing overbank flow. Riverine floods typically develop more gradually than other flood types, with warning times ranging from hours to days, depending on watershed characteristics. The dynamics of these floods are governed by river channel geometry, floodplain topography, and upstream hydrological conditions. Assessment methodologies typically incorporate river gauge data, precipitation records, and watershed models to establish probability distributions of flood magnitudes.
- **Pluvial (Surface Water) Flooding** occurs when rainfall overwhelms drainage systems' capacity or infiltrates the ground, causing water to pool or flow over the land surface. Unlike riverine flooding, pluvial events can happen anywhere—even far from water bodies—making them particularly challenging to predict and map. The risk is especially pronounced in urban environments with extensive impervious surfaces. Assessment requires high-resolution digital elevation models, detailed drainage system mapping, and precipitation statistics.
- **Coastal Flooding:** Primarily driven by storm surges, high tides, or tsunamis, coastal flooding presents distinct challenges due to its interaction with wave dynamics, tidal patterns, and coastal geomorphology. The combined effect of high water levels and wave action creates complex loading scenarios for coastal structures. Assessment methodologies must account for astronomical tides, barometric effects, wind setup, wave runup, and, increasingly, sea level rise projections.
- **Flash Flooding:** Characterised by rapid onset (typically within six hours of the causative event) and high water velocities, flash floods represent hazardous events with limited warning time. They commonly

occur in steep watersheds (a drainage area or catchment where the land slopes sharply, leading to rapid water movement), urban areas with high imperviousness, or regions susceptible to intense rainfall. Assessment approaches must emphasise the temporal dimension of flooding, incorporating rainfall intensity-duration-frequency relationships and time-of-concentration calculations.

- **Groundwater Flooding:** Results from a rise in the water table to the land surface, typically following prolonged periods of high precipitation. This flood type often has an extended duration but relatively gradual onset. Assessment requires understanding hydrogeological conditions, antecedent moisture patterns, and subsurface water movement—factors frequently underrepresented in standard flood models.

While these traditional classifications provide a helpful starting point, they often fail to capture the complexity of real-world flood events, which frequently involve multiple mechanisms operating simultaneously. This limitation has driven the development of more sophisticated classification frameworks.

Regulatory Preamble

There is never a perfect time or place to mention regulations, but this is the best moment since how governments consider flood risk and their chosen responses is central to how insurance and mortgage lending operate.

The European Floods Directive

Before the establishment of comprehensive flood management frameworks in Europe, the continent experienced several devastating flood events that highlighted the need for coordinated action. The early 2000s saw particularly catastrophic flooding, with the August 2002 floods affecting Central Europe, causing over €15 billion in damages and claiming dozens of lives across Germany, Austria, the Czech Republic, and other nations.

These events, along with rising concerns about the effects of climate change on precipitation patterns and sea levels, prompted European policymakers to create a more integrated approach to flood risk management. Between 2000 and 2006, over 175 significant floods were recorded in Europe, emphasising the urgency for systematic assessment and management protocols that crossed national boundaries.

In response to these challenges, the European Parliament and Council adopted the European Floods Directive (2007/60/EC) on October 23, 2007, which came into force on November 26, 2007. The Directive established a

framework requiring Member States to assess and manage flood risks coordinated across transboundary river basins.

Implementation followed a six-year cycle, starting with Preliminary Flood Risk Assessments completed by December 2011, then detailed flood hazard and risk maps developed by December 2013, and comprehensive Flood Risk Management Plans by December 2015. This revolutionary approach transcended traditional flood defence strategies, shifting towards more holistic risk assessment methodologies that acknowledged flood events' complex, multi-faceted nature and their impacts on human health, economic activities, cultural heritage, and the environment.

U.S. National Flood Insurance Program

Before the National Flood Insurance Program (NFIP) was established, the United States encountered a significant gap in disaster management policy, as many flood-prone properties were uninsurable in private markets, leaving residents primarily dependent on federal disaster assistance.

A series of catastrophic flood events in the 1950s and 1960s, including Hurricane Betsy in 1965, which caused over \$1 billion in damages (equivalent to approximately \$8.5 billion today), underscored the need for a systematic approach to flood risk. These events, coupled with rising federal disaster relief costs and increasing development in floodplains, prompted Congress to commission extensive studies on flood hazards and potential insurance solutions, culminating in the landmark 1966 "Insurance and Other Programs for Financial Assistance to Flood Victims" report that laid the groundwork for a national flood insurance mechanism.

In response to these challenges, the United States Congress established the National Flood Insurance Program through the National Flood Insurance Act of 1968, implementing a comprehensive approach that balanced insurance accessibility with floodplain management requirements.

The program became operational in 1969 under the Department of Housing and Urban Development before transferring to the Federal Emergency Management Agency (FEMA) in 1979 following its creation. The NFIP underwent significant expansions and reforms through subsequent legislation, including the Flood Disaster Protection Act of 1973, which introduced mandatory purchase requirements for properties with federally-backed mortgages in Special Flood Hazard Areas, and the National Flood Insurance Reform Act of 1994, which strengthened compliance mechanisms and established the Flood Mitigation Assistance Grant Program.

These developments transformed what began as a modest voluntary program into a cornerstone of American flood risk management, providing over 5 million policies with more than \$1.3 trillion in coverage nationwide by the early 2000s.

UK's Flood RE

In the decades preceding Flood Re, the United Kingdom experienced increasingly frequent and severe flooding events, highlighting significant market failures in flood insurance provision. The Easter floods of 1998, the widespread flooding of autumn 2000, and the catastrophic summer floods of 2007—which affected over 55,000 properties and caused approximately £3.2 billion in damages—demonstrated the urgent need for reform.

Before 2016, flood insurance in the UK was governed by a series of voluntary agreements between the government and insurers, starting with the informal "Gentleman's Agreement" in the 1960s and culminating in the Statement of Principles (2000-2013). However, this system became increasingly unsustainable as climate change heightened flood risks and premium reached prohibitive levels for many homeowners in high-risk areas.

Following extensive consultation and negotiations between 2010 and 2014, the UK Parliament passed enabling legislation through the Water Act 2014, establishing Flood Re as a not-for-profit reinsurance scheme. The program officially launched on April 4, 2016, representing a novel approach to address the growing insurance protection gap.

In 2022, Flood Re enhanced its offerings by introducing the "Build Back Better" scheme. This innovative initiative provides up to £10,000 in additional funding for property-level flood resilience measures when claims are settled. This forward-thinking program enables homeowners to not repair flood damage and enhance their property's resilience against future flooding events through od doors, raised electrical outlets, and water-resistant flooring.

Designed with a planned 25-year lifespan, Flood Re gradually created a transitional mechanism to move from subsidised to risk-reflective pricing by 2039y. Its Build Back Better scheme exemplifies how insurance can be leveraged to drive adaptation and risk reduction rather than simply transferring financial risk, representing a significant evolution in approach compared to traditional flood insurance models.

UK's Hazard Rating System

The development of comprehensive hazard classification systems in the United Kingdom emerged from a growing recognition of the need to

quantify flood risks more precisely for infrastructure planning and public safety purposes. In the early 2000s, following significant flood events across the UK, particularly the Easter floods of 1998 and the widespread flooding of autumn 2000, research institutions and government agencies began developing more sophisticated approaches to assess flood hazards beyond simple inundation mapping.

This period coincided with the implementation of the Housing Act 2004, which introduced the Housing Health and Safety Rating System (HHSRS), signalling a broader shift toward risk-based assessment methodologies across multiple sectors, including flood management. The Flood Risk to People research project, commissioned by the Department for Environment, Food and Rural Affairs (DEFRA) in 2003, played a pivotal role in establishing the scientific basis for what would become the UK's standard hazard rating formula.

Building on this foundation, the UK Environment Agency and DEFRA jointly published "Flood and Coastal Defence R&D Programme: Flood Risks to People" in 2006, which formally introduced the depth-velocity hazard rating formula that would become widely adopted across the nation.

This methodology represented a significant advancement over previous approaches by explicitly incorporating multiple flood characteristics—water depth, velocity, and debris factor—into a single quantifiable metric directly linked to human stability thresholds in floodwaters.

The formula $HR = d \times (v + 0.5) + DF$ was subsequently incorporated into official flood risk assessment guidance documents, including the 2008 "Supplementary Note on Flood Hazard Ratings and Thresholds" and the Flood Risk Assessment Guidance for New Development (FD2320/TR2). By 2010, this approach had become embedded in standard practice for flood risk assessments throughout England and Wales. It provided emergency planners, local authorities, and developers with a consistent framework for evaluating and communicating flood hazards that directly corresponded to potential impacts on human safety and infrastructure vulnerability.

Australian Rainfall and Runoff Guidelines

The Australian Rainfall and Runoff (ARR) Guidelines represent one of the longest-standing and most comprehensive national frameworks for flood estimation and water management in the world. First published in 1958, these guidelines have undergone several major revisions to incorporate advancements in hydrological science, data collection techniques, and computational methods.

The most significant transformation came with the 4th edition, released in 2016. This edition represented a complete overhaul that integrated over 30 years of additional rainfall observations from more than 10,000 gauging stations across Australia's diverse climatic regions. This landmark update moved away from outdated methodologies like the Rational Method toward more sophisticated regional flood frequency estimation techniques that better reflect Australia's highly variable precipitation patterns.

Following the 2016 revision, Engineers Australia transferred ongoing management of ARR to Geoscience Australia, ensuring its continued development as a government-funded public resource.

ARR has maintained a foundational principle that effective flood risk management requires calibration to local conditions—acknowledging that Australia's diverse landscapes, from tropical northern regions to temperate southern zones, necessitate tailored approaches rather than one-size-fits-all solutions. This emphasis on local calibration has become increasingly important as climate change alters historical rainfall patterns, requiring practitioners to continuously adapt their flood classification methodologies to reflect evolving regional circumstances.

Advanced Classification Frameworks

Recognising the limitations of source-based classifications, more nuanced frameworks have emerged that incorporate multiple parameters to characterise flood events:

- **Multi-parameter Classification Systems:** These approaches classify floods based on combinations of factors, including source, temporal characteristics (onset speed, duration, seasonality), spatial extent, and driving meteorological conditions. The European Floods Directive, for example, employs a multi-parameter approach that enables more precise risk assessment and management planning. These systems recognise that a single flood event may encompass multiple source mechanisms, such as combined coastal and fluvial flooding during severe storms.
- **Duration-Depth-Velocity Matrices:** These three-dimensional classification schemes recognise that flood impacts depend on water depth, flow velocity, and event duration. High-velocity shallow flooding, for instance, can cause more damage than deeper, static flooding due to increased hydrodynamic forces on structures. Similarly, extended duration increases damage to building materials unsuitable for prolonged water exposure. These matrices support more accurate vulnerability assessments and mitigation planning.

- **Recurrence Interval Categorisation:** This approach, which prevails in catastrophe modelling, classifies floods by their statistical return period, such as 1-in-100-year events (1% annual exceedance probability) or 1-in-500-year events (0.2% yearly exceedance probability). While widely used in regulatory contexts and infrastructure design, these classifications face increasing challenges due to non-stationarity in weather patterns. While beneficial to note, this is not what we need to solve in a capital markets scenario.
- **Combined Hazard Classification Approaches:** These frameworks integrate multiple flood characteristics to produce comprehensive hazard ratings that directly connect to potential impacts. For instance, the UK's Hazard Rating system uses an innovative methodology that combines water depth, flow velocity, and a debris factor to calculate risk levels. This integrated approach produces a single numerical value directly correlating with floodwaters' human stability thresholds. By accounting for the combined effects of these variables rather than assessing them in isolation, the system provides a more nuanced and impact-focused classification that better reflects real-world flood dangers to both people and infrastructure.
- **International Standards and Regional Variations:** Classification systems vary significantly across jurisdictions, reflecting regional priorities, data availability, and historical flood experience. For example, the Australian Rainfall and Runoff Guidelines emphasise the importance of local calibration in flood classification. In contrast, the U.S. National Flood Insurance Program classifications focus primarily on regulatory floodplain delineation. These variations create challenges for consistent cross-border risk assessment but also reflect the importance of local context in flood risk management.
- **Urban-Specific Flood Typologies:** The unique characteristics of urban flooding— influenced by complex drainage networks, building configurations, and infrastructure dependencies—have prompted specialised classification systems for urban environments. These typologies incorporate drainage capacity exceedance, surface water flow paths, and infrastructure failure modes. They often employ high-resolution modelling to capture micro-topographic features influencing urban flood routing.
- **Infrastructure Failure Flood Types:** In addition to natural flood mechanisms, increased attention has focused on classifying floods resulting from infrastructure failures, including dam breaches, levee failures, and urban drainage system malfunctions. These classifications incorporate cascade effects and system dependencies, recognising that infrastructure failures often produce flood characteristics significantly different from those of naturally occurring events, particularly onset speed and flow velocity.

Compound Flood Event Categorization

Compound flood event categorisation is an emerging field focused on understanding the complexity of flooding mechanisms that occur simultaneously or in close succession. This research addresses the amplified impact of such occurrences due to their multidimensional nature, involving combinations of drivers like precipitation, storm surge, river discharge, high tides, and more.

The Intergovernmental Panel on Climate Change (IPCC) first formally addressed compound events in 2012, noting that they involve multiple drivers or hazards occurring simultaneously or sequentially, often amplifying the resulting damage (Seneviratne et al., 2012). A notable example is Hurricane Harvey (2017), where record-breaking rainfall, river discharge, and storm surge combined to cause catastrophic flooding.

Research into compound events has proposed several classification frameworks. For instance, one framework identifies four interaction types for compound events:

- **Preconditioned:** Saturated soil that amplifies flood impacts.
- **Multivariate:** Co-occurrence of two hazards like storm surge and river discharge.
- **Temporally Compounding:** Sequential hazards in a region.
- **Spatially Compounding:** Hazards in different locations interact to amplify impacts.

These categories help researchers analyse specific conditions and impacts of compound flood events.

In practical modelling, rainfall runoff, storm surge, and other mechanisms are often studied together to enhance flood hazard assessments tailored to specific geographic areas, like coastal and estuarine environments. Simplified models now aim to break down complex interactions to better understand their effects on flooding and provide actionable insights for flood resilience and management.

The categorisation of compound flood events has emerged as a critical area of study, with several classification dimensions:

- **Flood-Generating Process Combinations:** Compound events frequently involve interactions between different flood types. These combinations include coastal-fluvial events (where storm surge prevents

river discharge, exacerbating upstream flooding), pluvial-fluvial events (where intense rainfall overwhelms drainage systems while simultaneously causing river flooding), and coastal-pluvial events (where storm surge coincides with heavy rain). Each combination produces distinct flood dynamics that standard single-process models may fail to capture.

- **Temporal Sequencing:** Compound events can be categorised based on their temporal characteristics—single events with multiple mechanisms, multiple events clustered closely in time, or events separated by insufficient recovery periods. This temporal dimension is particularly significant for infrastructure resilience and community recovery capacity. Systems still recovering from initial flooding often display amplified vulnerability to subsequent events, even of lesser magnitude.
- **Characteristic Combinations:** The specific combination of flood characteristics—magnitude, duration, timing, onset speed, and spatial extent—provides another classification dimension. For example, prolonged moderate flooding followed by a brief high-magnitude event creates different impact patterns than consecutive moderate events. Seasonal timing significantly influences impacts, particularly for agricultural systems with varying vulnerability throughout growing cycles.
- **Pathway and Mechanism Interactions:** Compound events often involve complex interactions between flood pathways—defence overtopping or breaching, surface water accumulation, groundwater emergence, and infrastructure failure. Each time series pathway may trigger or exacerbate others, creating cascade effects that traditional siloed assessments fail to capture. Classification frameworks increasingly incorporate these interaction pathways to better represent real-world flood complexity.
- **Vulnerability System Memory:** A critical dimension in compound event classification involves the "flood memory" of affected systems—how quickly natural and built environments recover their baseline resilience. Systems with long recovery periods (extended groundwater saturation, damaged flood defences, or depleted community resources) remain at elevated vulnerability, creating compound impacts even when subsequent events would not usually cause significant damage.

These multi-dimensional classification approaches represent a significant advancement beyond single-mechanism frameworks, better reflecting the complex reality of flood events that rarely occur in isolation. Comprehensive compound event classification requires considering hazard drivers, physical characteristics, failure mechanisms, and receptors' vulnerability and resilience characteristics.

Riverbank Breach Modelling

The modelling of riverbank breaches represents a specialised and critically important subset of flood risk assessment. These breaches—whether occurring in natural riverbanks or engineered levees—often result in rapid inundation of previously protected areas, creating distinct hazard characteristics compared to gradual overbank flooding. This section explores the methodologies to model such breaches and their implications for comprehensive flood risk assessment.

Riverbank breach modelling employs a range of approaches that balance physical accuracy with computational efficiency:

- **One-dimensional (1D) Hydraulic Models:** represent the river system as a series of cross-sections perpendicular to the flow direction, solving the Saint-Venant equations for conservation of mass and momentum. While 1D models significantly simplify the three-dimensional flow processes during breaching events, they remain widely used due to their computational efficiency and ability to model long river reaches. The effectiveness of these models depends critically on accurate stage-discharge relationships, particularly for overbank flows where the curve typically shows a distinct inflection point as flow encounters the floodplain's higher roughness and storage capacity.

River schematisation for 1D models necessitates a "broad brush" approach aligned with the overall river features. Cross-sections are positioned to capture significant changes in channel geometry, roughness, or potential breach locations. In these models, breaches are generally depicted as lateral structures with time-varying geometric properties. The breach initiation, widening rate, and final dimensions are either predefined based on geotechnical analysis or modelled dynamically from hydraulic parameters.

- **Two-Dimensional (2D) Hydraulic Models:** divide the floodplain into a grid or mesh of cells, solving the shallow water equations to determine water depth and velocity vectors throughout the domain. These models offer significant advantages for breach modelling by directly representing the spatial variation of flood propagation across the landscape. The explicit terrain representation allows for more accurate prediction of inundation patterns following a breach, particularly in urban areas where buildings and infrastructure create complex flow paths. Additionally, 2D models provide spatial distributions of flow velocity, which are crucial for assessing erosive forces during breach development and potential damages to structures in the breach flow path.

- **Coupled 1D-2D Approaches:** represent the main river channel in 1D while representing the floodplain in 2D, offering a balance between computational efficiency and physical realism. The treatment of the interface between domains is critical for accurate breach modelling, with various methods including lateral structures, vertical links, or horizontal links. Advanced coupled models dynamically activate the 2D domain when water levels exceed threshold values, efficiently focusing computational resources on areas experiencing inundation following a breach.

- **Computational Fluid Dynamics Applications:** These models provide a detailed representation of three-dimensional flow structures for site-specific analyses. They capture complex flow features around breach openings, including vortex formation, supercritical flow regions, and hydraulic jumps influencing breach development and downstream hazards. The extreme computational demands typically limit their application to relatively small spatial domains and short time intervals.

- **Dam Break Modelling:** a specialised subset of breach modelling, focuses on dam failures that can produce particularly catastrophic flooding. In many jurisdictions, dam owners must conduct dam break analyses to predict potential inundation areas and develop emergency plans. Specialised models have been developed for dam break wave propagation, accounting for the distinctive hydraulic characteristics of these extreme events, including supercritical flow transitions and bore formation. In river systems with multiple dams, models must account for potential cascade failures where the breach of an upstream structure leads to overtopping and failure of downstream structures.

Geotechnical Factors in Breach Formation

The hydraulic modelling of breaches must be informed by a geotechnical understanding of the failure mechanisms that initiate and drive breach development:

- **Soil Composition and Erodibility:** fundamentally control breach susceptibility. Different soil types exhibit varying resistance to erosion, with cohesive clay soils generally more resistant than non-cohesive sandy soils. Direct measurement of soil erodibility through Jet Erosion Tests or Erosion Function Apparatus testing provides quantitative parameters for physically-based breach models. Natural riverbanks and even engineered levees often contain spatial variations in material properties, creating potential weak points that may initiate breach formation.

- **Vegetation Effects:** bank stability operates through both mechanical and hydrological mechanisms. Plant roots provide mechanical reinforcement to the soil, increasing erosion and mass failure resistance.

The degree of reinforcement depends on root density, tensile strength, and architecture. Hydrologically, vegetation influences soil moisture regimes through rainfall interception, transpiration, and preferential flow paths along roots. These effects can either enhance stability (by reducing soil saturation) or reduce stability (by creating preferential seepage paths). The stabilising influence of vegetation often varies seasonally, creating temporal patterns in breach susceptibility.

- **Geomorphological Evolution:** recognises that rivers and their banks represent dynamic systems that evolve as a time series. Natural river meandering processes create zones of erosion on the outside of bends and deposition on the inside, continuously altering bank geometry. Long-term processes of channel incision (downcutting) or aggradation (sediment accumulation) alter the hydraulic loading on banks and levees, potentially creating conditions conducive to breaching even without changes in flow magnitude. Analysis of historical breach locations often reveals patterns related to these geomorphological features.

- **Advanced Soil Mechanics Applications:** Advanced soil mechanics applications increasingly incorporate sophisticated principles into breach modelling. Models that couple hydraulic loading with transient seepage can predict the development of internal erosion pathways and stability reduction due to increased pore pressures. Integrating slope stability analysis with hydraulic models enables the assessment of mass failure mechanisms that may initiate or accelerate breach formation. Advanced models incorporate unsaturated soil mechanics to represent the transition from unsaturated to saturated conditions during flood events, including the associated changes in strength and erodibility.

Real-time Monitoring and Prediction

Beyond modelling for planning and design purposes, increasing emphasis has been placed on real-time monitoring and predicting potential breaches during flood events. Since the early 2010s, significant technological advances have transformed flood forecasting capabilities, with AI and machine learning integration revolutionising prediction accuracy and lead times.

Google's AI-driven Flood Hub exemplifies this evolution, providing forecasts up to seven days in advance by analysing global weather data, satellite imagery, and river gauge measurements to enhance situational awareness during flood emergencies.

Modern integrated systems now combine sophisticated GIS tools with hydrological data to produce detailed flood maps, hazard zones, and impact assessments that enable timely alerts to minimise property damage. Companies like Previsico have pioneered high-resolution

modelling systems capable of predicting flood risks at the property level, utilising networks of sensors and "instacasting" to provide actionable warnings based on live radar data and hydrological measurements—dramatically improving the precision with which authorities can identify potential infrastructure failures and implement targeted emergency responses:

- **Sensor Networks:** deploy multiple monitoring technologies to detect precursors to breach formation. Piezometers (a device used to measure the pressure of a fluid, typically groundwater, in a specific location) embedded within levees measure pore water pressures, allowing the detection of anomalous seepage conditions that may precede internal erosion-driven breaches. Deformation monitoring through exotically named inclinometers (precision instruments used to measure angles of slope, elevation, or inclination relative to gravity's pull), extensometers (measure small changes in the distance between two points, specialising in detecting deformation, displacement, or strain in materials and structures), and remote sensing techniques measures bank or levee movement, providing early warning of potential mass failure mechanisms. These sensors are increasingly connected to integrated monitoring systems with automated alert thresholds and real-time data visualisation.
- **Remote Sensing Technologies:** provide broad spatial coverage to complement point-based sensor networks. Synthetic Aperture Radar systems can detect surface deformation and soil moisture changes associated with potential breach conditions, even during storm events when optical systems are limited by cloud cover. Thermal imaging can identify temperature anomalies related to seepage. In contrast, time-series analysis of remotely sensed data identifies progressive changes in bank or levee conditions that may indicate increasing breach susceptibility.
- **Machine Learning Approaches:** apply advanced analytical techniques to the complex, multi-parameter problem of breach prediction. Algorithms trained on historical breach data and precursor conditions identify patterns that may not be apparent in traditional analyses. These approaches can integrate disparate data streams from multiple sensor types, extracting meaningful signals from complex, noisy datasets. Bayesian methods provide probabilistic breach forecasts that update dynamically as new monitoring data becomes available, supporting risk-based decision-making during flood events.
- **Early Warning System Integration:** ensures monitoring and prediction capabilities translate into effective risk reduction. The definition of appropriate alert thresholds balances false positive risk against the need for sufficient warning time to implement emergency measures or evacuations. Clear communication protocols ensure

monitoring data and breach predictions reach decision-makers and affected populations through multiple, redundant channels. Integrating breach prediction into broader flood early warning systems provides comprehensive situational awareness for emergency management.

The modelling and monitoring of riverbank breaches continues to advance through integrating hydraulic modelling, geotechnical analysis, and real-time monitoring technologies. These developments support a more accurate assessment of breach-related risks and a more effective allocation of mitigation resources.

Property Positioning Analysis

The physical positioning of properties within the landscape fundamentally determines their exposure to flood hazards. While flood classification and breach modelling characterise the hazard, property positioning analysis translates these hazard characteristics into site-specific exposure assessments. This analytical approach draws from methodologies developed in real estate valuation—where the mantra "location, location, location" underscores positioning's critical importance—and adapts them to the specific requirements of flood risk assessment.

Typographical Analysis

Topography represents the most fundamental positioning factor for flood risk assessment, directly influencing flow paths, water accumulation, and inundation potential:

- **Elevation Relative to Flood Sources:** is the primary determinant of flood exposure. The vertical distance between a property and potential water sources—whether rivers, coastlines, or surface water accumulation zones—establishes the basic threshold for flooding. This relationship is complicated because absolute elevation alone is insufficient; the hydraulic connectivity between the property and flood sources must be considered. Properties at relatively high elevations may still experience flooding if hydraulic pathways (natural or artificial channels, drainage systems) connect them to flood sources.
- **Digital Elevation Models:** provide the foundation for topographic analysis in modern flood risk assessment. The resolution and accuracy of these models significantly influence assessment quality. While national or regional DEMs may be sufficient for broad-scale risk screening, property-specific assessment typically requires high-resolution data derived from LiDAR or similar technologies capable of sub-meter vertical accuracy. Such precision is essential for identifying subtle topographic features that may significantly influence flood routing.

- **Microtopography:** is particularly important in urban environments, where insignificant elevation differences can significantly alter flood pathways. Features such as curbs, berms, garden walls, and small depressions may redirect or impound water in ways not captured by coarser elevation models. Advanced assessments incorporate these features through ultra-high-resolution elevation data or specific feature recognition and hydraulic representation.
- **Relative Positioning:** within the landscape context often matters more than absolute elevation. Even if their absolute elevation is relatively high, properties at local low points may be subject to ponding during intense rainfall events as surface water flows converge. Similarly, properties positioned along preferential flow paths—natural drainage lines, valleys, or urban street corridors that channel water—face elevated risk during high-intensity events regardless of their elevation relative to typical flood sources.

Location Analysis

Comprehensive property positioning assessment requires analysis at multiple spatial scales, each revealing different aspects of flood risk exposure:

- **Regional Analysis:** examines the broader hydrological context within a property. This includes watershed characteristics, regional weather patterns, and large-scale flood control systems influencing the flood regime. Properties in regions with flat topography, poor regional drainage, or prone to large-scale weather systems capable of producing widespread precipitation (such as tropical cyclones or atmospheric rivers) face fundamentally different risk profiles than those in regions with efficient drainage networks or less extreme meteorological patterns.
- **Macro-location Analysis:** focuses on the property's position within its city or district and its relationship to significant flood sources. The traditional real estate categorisation of locations as central, semi-central, or peripheral takes on a specific meaning in flood risk assessment. Central urban areas often face different flood mechanisms (predominantly pluvial and drainage-system related) compared to peripheral regions (where riverine flooding may dominate). Historical development patterns frequently placed central districts near water bodies for transportation and industrial purposes, creating legacies of flood exposure that modern developments must address.
- **Micro-location Analysis:** examines the immediate surroundings of a property, typically within a radius of several hundred meters, to identify localised flood risk factors. Critical considerations include:

- The proximity to water bodies and drainage systems includes not only significant rivers but also smaller streams, drainage channels, and stormwater infrastructure.
- Local topographic features that might channel or impede water flow.
- Barriers and obstructions that may protect or, conversely, increase flood risk by redirecting flows.
- Neighbouring structures that may influence local flow dynamics.
- Impervious surfaces in the immediate vicinity that generate runoff.
- **Street-level Positioning:** represents the finest scale of analysis, examining how a property's specific position on its street or block influences its flood exposure. Factors such as setback distance from the street, position relative to the street crown and drainage structures, driveway configuration, and even subtle grading around the building footprint can significantly influence how surface water interacts with the structure during flood events.

Built Environment Context

The built environment surrounding a property creates a complex system that significantly modifies natural flood dynamics:

- **Urban Density:** alters hydrological processes by increasing impervious surface area, modifying drainage patterns, and creating the channelling effect of street networks. High-density urban environments typically result in more rapid runoff with higher peak flows than natural landscapes, generating complex flow paths through the built environment. Properties in dense urban areas may experience flooding mainly due to these altered urban hydrological processes rather than from natural water bodies.
- **Street Network Hydraulics:** play a critical role in urban flood routing. Streets often function as secondary drainage systems during intense rainfall events when conventional drainage infrastructure reaches capacity. The orientation, slope, width, and connectivity of streets create a network of flow paths that distribute floodwaters throughout urban areas. Properties at street intersections, low points in the street network, or downstream of large contributing street corridors face elevated exposure to these urban flow paths.

- **Infrastructure Systems:** significantly influence flood exposure through both intended and unintended effects. Stormwater drainage systems reduce flood risk when functioning properly by efficiently removing surface water. However, these same systems can become flood sources when they reach capacity or experience backwater effects. Subsurface infrastructure—including utility corridors, subway systems, and underground parking facilities—can create unexpected hydraulic connections that bypass surface topography, potentially bringing floodwaters to otherwise protected areas.
- **Building Arrangement Effects:** occur when nearby structures alter flow paths to increase or decrease risk to a specific property. Depending on their configuration relative to flow direction and adjacent structures, buildings may block, channel, or redirect flows. In dense urban environments, the arrangement of buildings can create corridor effects, focusing flows between structures or shadow effects, where upstream buildings protect downstream properties. These complex interactions typically require detailed two-dimensional hydraulic modelling to capture accurately.

Exposure Assessment Integration

The comprehensive assessment of property positioning requires the integration of these multiple factors into a coherent exposure assessment:

- **Exposure Pathway Identification:** systematically analyses all potential routes by which floodwaters might reach a property. This includes direct inundation from water bodies, surface flow across the landscape, backwater effects through drainage systems, groundwater rise, and infrastructure-mediated pathways. Each pathway requires specific positioning analysis techniques and may dominate under different flood scenarios.
- **Multi-hazard Positioning:** recognises that property location simultaneously influences exposure to multiple flood types. A comprehensive positioning analysis must consider exposure to riverine, coastal, pluvial, and groundwater flooding and infrastructure-related flood risks. Properties often face different positioning-based exposure levels for each flood type, creating complex multi-hazard risk profiles.
- **Temporal Dynamics:** in positioning assessment acknowledge that exposure changes over time through natural processes and human interventions. Coastal erosion, riverbed migration, urban development patterns, and infrastructure modifications all gradually alter a property's effective position relative to flood sources.
- **Quantitative Exposure Metrics:** translate qualitative positioning analysis into numerical inputs for risk assessment models. These

metrics include elevation relative to various flood levels, distance from flood sources adjusted for hydraulic connectivity, flow accumulation indices from hydrological models, or compound indices that integrate multiple positioning factors. These quantitative metrics enable consistent comparison across properties and incorporation into broader risk assessment frameworks.

Property positioning analysis represents the essential bridge between hazard characterisation and vulnerability assessment. By systematically analysing how a property's location influences its exposure to various flood mechanisms, this approach provides the foundation for targeted resilience strategies and accurate risk quantification. As urban development alters flood patterns, sophisticated positioning analysis becomes increasingly critical for effective flood risk management across the built environment.

Environmental Resilience Factors

While property positioning establishes exposure to flood hazards, environmental resilience factors determine how natural systems buffer, absorb, or amplify these hazards before they reach vulnerable assets. Resilience—the ability of a system to absorb disturbances while maintaining its essential functions and structure—represents a critical dimension of comprehensive flood risk assessment. Natural systems exhibit complex resilience characteristics that fundamentally influence the translation of meteorological events into flood impacts.

Natural Flood Defence Systems.

The landscape surrounding developed areas provides the first line of defence against flooding through various natural mechanisms:

- **Wetland Buffer Capacity:** represents one of the most effective natural flood mitigation systems. Wetlands function as natural sponges, temporarily storing flood waters and releasing them gradually, thereby reducing downstream peak flows. The flood attenuation capacity of wetlands depends on their type, size, antecedent conditions, and position within the watershed. Riparian wetlands directly connected to river systems provide immediate storage during high flows, while isolated wetlands may reduce the volume of water reaching rivers through groundwater recharge. Historical wetland loss—exceeding 50% globally and reaching 90% in some developed regions—has significantly reduced this natural resilience mechanism, amplifying flood risks downstream.

- **Forest and Vegetation Systems:** influence flood dynamics through multiple mechanisms operating at different time scales. The canopy

intercepts precipitation, reducing the volume and velocity of water reaching the ground. Root systems enhance soil infiltration capacity, converting potential surface runoff into subsurface flow, moving more slowly toward water bodies. The forest floor, with its organic litter layer, provides additional water storage and slows surface flow. These mechanisms don't eliminate flooding during extreme events but modify the flood hydrograph by delaying and attenuating peak flows. The effectiveness of these systems varies by forest type, age, health, and management practice, with mature, diverse forests generally providing more significant flood mitigation benefits than young, monoculture plantations.

- **Soil Systems:** play a crucial but often overlooked role in flood resilience. Healthy soils with high organic content and well-developed structures can absorb and store significant water volumes, reducing runoff's volume and velocity. Infiltration capacity—the rate at which water can enter the soil profile—represents a critical parameter for flood mitigation. This capacity varies dramatically across soil types and conditions, with well-structured loamy soils potentially infiltrating ten times more water than compacted clay soils or those depleted of organic matter. Land management practices that maintain soil health contribute significantly to flood resilience, while soil degradation through compaction, organic matter depletion, or contamination reduces this natural defence.
- **Natural Channel Morphology:** influences how efficiently river systems convey flood waters. Naturally, meandering rivers with connected floodplains dissipate energy and store water during high flows, reducing downstream flood peaks. The complexity of natural channel forms—including pools, riffles, side channels, and woody debris—creates hydraulic roughness that slows flow velocities during floods. Channel straightening, artificial levee construction, and floodplain disconnection, while potentially protecting specific areas, often increase downstream flood risk by increasing flow velocity and peak discharge. River restoration efforts increasingly recognise the flood mitigation value of returning rivers to more natural configurations, providing ecological benefits and enhanced flood resilience.

Resilience Mechanisms and Processes

Beyond specific natural features, broader ecological and geomorphological processes contribute to environmental resilience against flooding:

- **Self-Organisation Capacity:** distinguishes resilient natural systems from engineered flood defences. While built infrastructure typically provides static protection until failure thresholds are exceeded, natural

systems dynamically adjust to changing conditions. River systems redistribute sediment, vegetation communities adapt to hydrological changes, and coastal systems migrate in response to sea-level fluctuations. This self-organisation allows natural systems to maintain their flood mitigation functions across various conditions, though within certain thresholds. When these thresholds are exceeded—through extreme events or cumulative human alterations—natural systems may shift to alternative states with different flood response characteristics.

- **Disturbance Absorption Thresholds:** represent the magnitude of flood events that natural systems can accommodate while maintaining their essential functions. These thresholds vary across different ecosystem types and conditions. Healthy floodplain forests may withstand extended inundation during seasonal flooding but lose resilience when flooding becomes too frequent or prolonged. Similarly, coastal dune systems absorb the energy of moderate storm surges but may breach during extreme events, requiring time to rebuild naturally. Understanding these thresholds is essential for a realistic assessment of the protection provided by natural systems and for identifying points of failure where natural resilience may suddenly diminish.
- **Adaptive Learning Processes:** enable natural systems to incorporate previous disturbances into their structure and function, potentially increasing resilience to future events. Vegetation communities in frequently flooded areas develop specialised physiological, morphological, and reproductive adaptations that enhance survival during inundation. Geomorphic systems reorganise after significant floods, sometimes creating more stable configurations that better accommodate future events. These natural adaptive processes operate across multiple time scales, from seasonal adjustments to evolutionary adaptations spanning centuries. Unlike engineered systems designed for static conditions, natural systems continuously adjust to changing flood regimes, though human development time frames often fail to recognise these slower processes.

- **System Recovery Trajectories:** describe how natural flood defence systems return to functional states following disturbance. These trajectories rarely represent simple returns to previous conditions; instead, they often involve reorganisation around new equilibrium states that reflect both the disturbance and changing environmental conditions. For example, a coastal marsh may rebuild after a storm surge with altered channel networks that more efficiently dissipate future surge energy. These dynamic recovery processes contribute significantly to long-term resilience but complicate assessment efforts based on static environmental conditions.

Social-Ecological System Interactions

The resilience of natural systems cannot be fully understood in isolation from human systems, as the two interact in complex ways that influence overall flood risk:

- **Coupled System Dynamics:** recognise that human and natural systems interact through multiple feedback mechanisms relevant to flood resilience. Human modifications to natural systems—through development, resource extraction, pollution, or restoration—alter their flood mitigation capacity. Conversely, natural system responses to these modifications and flood events influence human vulnerability and subsequent adaptation decisions. These coupled dynamics create complex feedback loops operating across multiple spatial and temporal scales. For instance, development in previously natural floodplains reduces flood storage capacity, potentially increasing downstream flooding, which may prompt further flood control measures that alter natural systems.
- **Resilience Trade-offs:** occur when enhancing one aspect of system resilience compromises another. Engineered flood control systems often increase short-term protection for specific assets while reducing the long-term resilience provided by natural systems. Flood walls may protect riverside developments but prevent the natural overbank flows that sustain floodplain ecosystems, providing flood mitigation. Similarly, drainage systems that effectively remove local surface water may increase downstream flood peaks. Comprehensive assessment requires recognising these trade-offs and their implications across different spatial scales and time horizons.
- **Knowledge Integration Challenges:** arise from how human and natural system resilience are conceptualised and measured. Engineering approaches typically evaluate performance against defined design standards, while ecological resilience encompasses more complex, sometimes qualitative characteristics like diversity, redundancy, and adaptive capacity. Practical assessment requires bridging these knowledge systems, incorporating quantitative hydraulic analysis and ecological understanding of system dynamics. This integration remains challenging but increasingly necessary as recognition grows of the substantial flood protection value provided by natural systems.
- **Governance System Alignment:** with natural resilience processes is critical in maintaining and enhancing environmental contributions to flood risk reduction. Governance frameworks—including regulations, incentives, planning processes, and management institutions—strongly influence whether natural flood mitigation systems are protected, restored, or degraded. Effective governance recognises the spatial

misalignment often present between jurisdictional boundaries and watershed or coastal system boundaries. Systems that align governance with natural processes can significantly enhance overall flood resilience.

Weather Pattern Change Implications

The natural shift in weather patterns introduces new uncertainties to environmental resilience assessment, potentially altering both the hazards that natural systems must absorb and the system's capacity to provide protection:

- **Shifting Baseline Conditions:** challenge the assumption that natural systems will continue to provide historical levels of flood protection. Changes in precipitation patterns, rising sea levels, and varying temperature regimes alter the environmental conditions to which local ecosystems are adapted. These shifts may enhance or reduce natural flood mitigation capacity, depending on local circumstances and ecosystem types. For instance, increased precipitation may improve wetland function in some areas, while drought conditions in other regions may decrease soil infiltration capacity and vegetation cover, thus diminishing natural flood protection.
- **Ecosystem Transition Thresholds:** may be crossed, leading to rapid changes in natural flood defence systems. Many ecosystems exhibit non-linear responses to changing environmental conditions, maintaining relative stability until critical thresholds are exceeded, then rapidly transitioning to alternative states. Coastal marshes may keep pace with moderate sea level rise but suddenly convert to open water when a critical rate is exceeded. Forested watersheds may maintain flood mitigation functions through moderate drought but lose them rapidly during extreme moisture stress events that trigger widespread mortality. Identifying these potential transition points represents a critical frontier in resilience assessment.
- **Adaptive Management Imperatives:** emerge from uncertainty when projecting future weather pattern intensity. This uncertainty necessitates approaches that monitor natural system conditions, detect early warning signals of resilience loss, and adjust management strategies accordingly. Adaptive management frameworks incorporate structured learning processes that treat interventions as experiments, generating knowledge to refine subsequent actions. For flood risk assessment, this approach recognises that environmental resilience factors are not static features to be measured once but dynamic processes requiring ongoing evaluation and adjustment of protection estimates.

The assessment of environmental resilience factors completes the hazard component of flood risk analysis by addressing how natural systems mediate the translation of meteorological events into flood impacts on human settlements.

This assessment connects traditional hazard analysis with vulnerability assessment by acknowledging that flood characteristics in specific locations arise from the initiating event and the intricate environmental systems that influence the movement of floodwaters. Understanding and enhancing these natural resilience mechanisms becomes increasingly central to effective flood risk management as human development patterns evolve.

Property-Specific Resilience Assessment

The culmination of flood risk assessment occurs at the individual property level, where the analysis of hazard characteristics, positioning factors, and environmental resilience coalesce into property-specific vulnerability and adaptation assessments. This final analytical layer translates broader flood risk understanding into actionable property-level interventions that significantly reduce physical damage and banking impacts. A structured approach to property-specific resilience assessment provides the critical link between technical flood risk analysis and practical risk reduction outcomes.

Structural Vulnerability Analysis

The built characteristics of individual properties fundamentally determine their flood vulnerability through multiple physical mechanisms:

- **Material Vulnerability Differentiation:** reveals significant variation in how building materials respond to flood exposure. Masonry structures generally withstand hydrostatic pressure better than timber-frame constructions but may experience more significant capillary action, drawing water upward beyond the visible flood line. Concrete structures typically provide superior structural resilience but may create significant drying challenges post-flood. These material-specific vulnerabilities extend to finishing elements—plasterboard deteriorates rapidly when inundated, while water-resistant gypsum boards may maintain structural integrity through moderate flooding events. Comprehensive assessment requires evaluating the complete material assembly from structural elements through insulation to interior finishes.

- **Critical System Placements:** significantly influence both damage potential and recovery timelines. Properties with elevated electrical systems—raised sockets, elevated consumer units, and ring mains routed through upper portions of walls rather than near floor level—typically experience less critical system damage and faster recovery. Similarly, positioning HVAC (Heating, Ventilation, and Air Conditioning) equipment, water heaters, and other mechanical systems relative to anticipated flood levels directly correlate with system survival rates. To prioritise protection or relocation interventions, vulnerability assessment must identify these critical system elevations relative to property-specific flood risk characteristics.

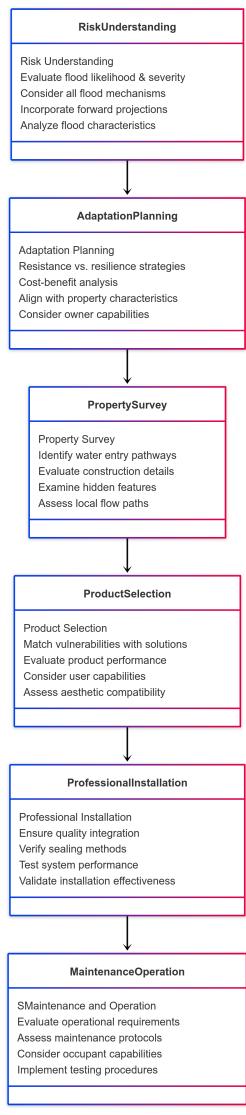
- **Foundation-type assessment:** reveals varying flood response characteristics across foundation systems. Properties with solid floor construction face different vulnerabilities than those with suspended floors—while solid floors may better resist structural damage, they often present more significant drying challenges. Suspended floors allow underfloor drying but may experience greater uplift forces and access points for floodwater. Crawlspaces and basements introduce additional complexities, potentially serving as beneficial buffer zones when properly designed or as significant vulnerabilities when improperly protected. Assessment must consider the foundation type and specific details like damp-proof course positioning and perimeter drainage systems.

- **Dynamic Pressure Considerations:** become particularly important in high-velocity flood scenarios. Properties in breach flow paths or flash flood zones face additional hydrodynamic forces beyond simple inundation. Structural elements require assessment for their capacity to withstand these lateral forces—particularly corner junctions, door and window openings, and non-load-bearing walls that may face disproportionate pressure. This assessment extends to external elements like garden walls and fences, which may provide beneficial shielding or create debris hazards depending on their construction and orientation relative to anticipated flow paths.

Property-Level Adaptation Measures

Practical resilience assessment extends beyond vulnerability identification to evaluation of adaptation options through a structured six-step process:

- **Step 1: Risk Understanding:** provides the foundation for property-specific interventions. Comprehensive assessment requires evaluating the likelihood and potential severity of flooding for the specific property, incorporating all relevant flood mechanisms—fluvial, pluvial, coastal, groundwater, and infrastructure failure. This risk profile should



consider historical events but must also incorporate forward-looking projections. Risk understanding should extend beyond simple inundation mapping to consider anticipated flood characteristics, including depth, duration, velocity, water quality, and seasonal timing—each influencing the selection of appropriate resilience measures.

•**Step 2: Adaptation Planning:** translates risk understanding into strategic decision-making. Property owners must evaluate the full spectrum of available approaches, from resistance strategies (preventing water entry) to resilience measures (minimising damage when water enters). This planning process requires a comprehensive cost-benefit analysis considering initial implementation costs, maintenance requirements, operational reliability, aesthetic implications, and potential insurance benefits. Effective planning aligns selected measures with specific property characteristics and owner capabilities, recognising that even the most sophisticated interventions prove ineffective if improperly deployed or maintained.

•**Step 3: Property Survey:** provides the technical foundation for adaptation implementation. Unlike generalised risk assessments, property-specific surveys identify all potential water entry pathways –doors, windows, air bricks, service penetrations, and construction joints or material interfaces. Effective surveys evaluate construction details that might not appear on standard plans, such as hidden service penetrations, legacy features from previous modifications, or subtle topographic features that influence local flow paths.

Figure 5: Property Adaption Measures

- **Step 4: Product Selection:** matches identified vulnerabilities with appropriate technological solutions. Property-specific assessment must evaluate not only the theoretical performance of these products but also their appropriateness for the specific context, including user capabilities, deployment time requirements, aesthetic considerations, and compatibility with existing building features.

- **Step 5: Professional Installation:** ensures theoretical protection translates to real-world performance. Assessment must consider the quality of installation and selected products, as even superior technologies fail when improperly integrated with existing structures. This evaluation extends to sealing methods, fixing approaches, and interface treatments between different protection system components. Post-installation testing provides critical verification, from simple hose testing of minor interventions to complete deployment exercises for more complex systems.

- **Step 6: Maintenance and Operation:** completes the resilience assessment cycle. The most sophisticated protection systems prove worthless without proper maintenance and timely deployment. Assessment must evaluate operational requirements against occupant capabilities, particularly considering warning time, physical abilities, and the potential need for assistance. Maintenance assessment extends beyond simple physical inspection to consider ageing effects, material degradation patterns, and operational testing protocols, particularly for less frequently used mechanical components that may seize or deteriorate while dormant.

Resilience Scoring Methodologies

Quantitative assessment of property-specific resilience enables more sophisticated risk management and banking protection:

- **Multi-criteria Evaluation Frameworks:** typically incorporate resistance measures (preventing water entry) and recovery-enhancing features (minimising damage when water enters). Effective scoring systems weigh different factors according to their relative importance for specific property types and flooding characteristics. For example, the resistance of building fabric might receive greater weighting in short-duration, high-frequency events, while rapid drying potential might receive greater emphasis in prolonged inundation scenarios. These weighted scores provide a standardised basis for comparing properties, prioritising interventions and demonstrating improvement over time.

- **Resilience Certification Programs:** translate technical assessments into accessible documentation for non-technical stakeholders. These programs, increasingly recognised by insurance markets, provide

standardised verification of implemented measures and their expected performance across defined flood scenarios. Certification typically requires documentation review and physical inspection, often with periodic renewal requirements, to ensure ongoing maintenance. The resulting certifications provide objective evidence of reduced risk that can support insurance negotiations, property valuation, and regulatory compliance demonstrations.

- **Insurance Premium Integration:** represents a critical application of resilience scoring. While historical approaches to flood insurance often relied on simplistic risk categorisation, advanced insurers increasingly incorporate property-specific resilience characteristics into underwriting models. Properties demonstrating robust protection measures—particularly those validated through recognised certification programs—can access more favourable coverage terms, lower premiums, reduced deductibles, or coverage in otherwise uninsurable locations. This insurance recognition creates powerful incentives for resilience investment, creating a virtuous cycle where premium savings help finance further protection measures.
- **Valuation Methodology Alignment:** extends resilience benefits beyond insurance markets. Property valuation increasingly considers flood risk, with high-risk properties facing potential value penalties. Property-specific resilience assessment provides a mechanism to differentiate protected properties within high-risk areas, potentially preserving significant value through demonstrated risk reduction. Advanced valuation methodologies incorporate resilience scoring alongside traditional factors, recognising that properties with equivalent flood zone designations may face dramatically different actual risk profiles based on implemented protection measures.
- **Cost-benefit optimisation:** allows targeting of limited resilience investment for maximum return. Comprehensive assessment enables prioritisation of interventions based on their risk reduction potential relative to implementation costs. This optimisation process typically reveals that specific low-cost measures—such as raising electrical systems or installing removable flood barriers—deliver disproportionate benefits compared to more extensive structural modifications. For properties with recurring flood exposure, these cost-benefit calculations often demonstrate remarkably short payback periods for resilience investments, notably when insurance savings and avoided disruption are fully valued.

The assessment of property-specific resilience completes the risk assessment chain, translating broader flood hazard understanding into practical, location-specific protection strategies. As development pressures continue to alter flood risk profiles across the built environment,

this property-level focus provides the most immediate and effective path to risk reduction.

While community-scale defences and watershed management remain essential components of comprehensive flood risk management, property-specific interventions offer immediate protection benefits regardless of broader defence implementation timelines or funding constraints. Integrating these property-level assessments with risk transfer mechanisms creates powerful banking incentives that accelerate adoption without regulatory mandates, creating a market-driven approach to enhanced resilience.

Chapter 6 - Impact Quantification

The quantification of physical risk impacts represents the culmination of our progression from fundamental principles to practical applications. Having established how water moves through the environment and how we assess the resulting flood risks, we now turn to converting these assessments into measurable impacts on property, infrastructure, and economic systems.

"The value of risk assessment ultimately lies in our ability to quantify potential impacts. Without this quantification, we cannot properly price risk, allocate resources for mitigation, or make informed decisions about development and adaptation." - David Kelly, MKM.

Impact quantification translates the physical reality of hazards into the banking language that drives decision-making. It bridges the gap between scientific understanding and economic consequences, enabling stakeholders to evaluate trade-offs, justify investments, and manage exposures. This translation requires sophisticated methodologies that connect the physical characteristics of hazard events to their banking, social, and environmental outcomes.

Hazard curves are the foundation of this quantification process, characterising the probability distribution of events across different magnitudes and timeframes. These curves encapsulate both historical patterns and future projections. Developing robust hazard curves requires statistical rigour, appropriate distribution selection, and careful treatment of uncertainty—particularly when extending beyond the historical record.

Vulnerability functions then connect these hazard characteristics to expected damages. These functions—often expressed as relationships between flood depth and percentage damage—vary significantly across building types, content categories, and occupancy classes. The development of these functional curves draws upon empirical data from

past events, engineering principles, and, increasingly, advanced simulation techniques.

Applying these functions yields damage estimates across multiple dimensions: direct structural damage, content losses, business interruption, and broader economic impacts. Each category requires specialised methodologies that account for repair costs, replacement values, temporal disruption, and complex economic interdependencies. These estimates must further consider insurance coverage, including policy terms, exclusions, and risk transfer mechanisms that distribute impacts across stakeholders.

The ultimate expression of these impacts often appears in property valuations, where market participants price risk into their investment decisions. Understanding how physical risks translate into changes in market value requires consideration of capitalisation rates, market perception, regulatory influences, and adaptation potential.

As we explore these quantification methodologies, we complete the analytical pathway that began with fundamental physical principles. This progression—from Newtonian physics to banking economics—enables the effective management of physical risks in an increasingly uncertain world.

Hazard Curve Development

Hazard curves form the foundational component of physical risk assessment, providing a quantitative representation of the relationship between hazard intensity and probability of occurrence. This section details the methodological approach to constructing hazard curves, emphasising flood hazards, as exemplified in the framework of the FINOS physical risk project under OS-Climate.

A hazard curve establishes the relationship between the severity of a physical phenomenon and its probability of occurrence. This typically forms a return period curve or exceedance probability curve for chronic hazards like flooding. The OS-Climate methodology adopts principles from natural catastrophe modelling, where:

- Return period (τ) refers to the average time between events with intensity higher than a specified threshold.
 - Exceedance probability represents the probability that in a given period (typically one year), an event of intensity more significant than a specified threshold will occur.
-

Formula 10: Hazard Curve Development

The mathematical relationship between these concepts is expressed as the probability of at least one occurrence in time t:-

$$P(X^h_t \geq 1) = 1 - e^{(-t/\tau)}$$

Where:-

- $P(X^h_t \geq 1)$ is the probability of at least one hazard event occurring within time period t.
- X^h_t represents the number of hazard events of type h occurring in time period t.
- t is the time period of interest (typically 1 year).
- τ is the return period in years (average recurrence interval).

Convert to probability bins. Using the relationship

$$F'_S(h) = F'_S(s_q^{\text{lower}}) - F'_S(s_q^{\text{upper}})$$

Where:

- $F'_S(h)$ is the probability of the hazard intensity falling within a specific bin.
- s_q^{lower} is the lower bound of the hazard intensity bin.
- s_q^{upper} is the upper bound of the hazard intensity bin.
- F'_S represents the probability density function of the hazard intensity

In everyday parlance, a 100-year flood has approximately a 1% probability of occurrence in any given year.

Constructing Hazard Curves

The OS-Climate framework constructs hazard curves through the following process:

- **Identify hazard indicators:** The primary hazard indicator is flood depth (measured in meters) for floods.
- **Obtain spatial hazard datasets:** OS-Climate leverages multiple data sources, including:
 - World Resources Institute (WRI) Aqueduct flood model for global coverage of coastal and riverine inundation
 - Higher resolution regional datasets (e.g., TU Delft pan-European dataset with 100m resolution)
 - Weather-conditioned future projections based on different emissions scenarios

Return Period (years)	Flood Depth (m)
2	0.06
5	0.33
10	0.51
25	0.72
50	0.86
100	1
250	1.15
500	1.16
1000	1.17

Table 4: Flood Depth by Return Period

- **Extract location-specific hazard data:** Extract the relationship between hazard intensity and return period for a given location (latitude, longitude). For inundation, this might look like:
 - Construct probability bins for different hazard intensities. For example, the probability of occurrence of a flood with depth in the range (0.86m, 1.00m] would be $0.02 - 0.01 = 0.01$.
 - Interpolate between return periods. For comprehensive risk assessment, interpolation methods are applied to estimate hazard intensities between the defined return periods.

- Account for weather pattern uncertainty. Compare hazard curves under baseline historical conditions with weather-conditioned curves for future periods to quantify hazard frequency and severity changes.

Connecting Hazard and Vulnerability

The hazard curve and vulnerability curve are connected in the following way:

- The **hazard curve** provides the probability of experiencing different hazard intensities at a location.
- The **vulnerability curve** takes those hazard intensities as inputs and translates them into expected damage/loss values for the exposed assets.

In risk analysis, the hazard curve is first used to determine the range of hazard intensities that must be considered based on their likelihood at that location. Then, the vulnerability curve maps each hazard intensity to a damage ratio or loss estimate for the asset being analysed.

By combining the hazard exceedance probabilities from the hazard curve with the expected damage/losses from the vulnerability curve, catastrophe models can calculate risk metrics like the probable maximum loss or average annual loss for the exposed portfolio of assets.

In summary, the hazard curve characterises the frequency of the peril. In contrast, the vulnerability curve characterises the fragility of the asset to that peril's intensity - both critical components in quantifying potential losses from any flood event.

Vulnerability Curve Modelling

Vulnerability curves represent the relationship between hazard intensity and an asset's resulting damage or disruption. While hazard curves capture the probability of physical phenomena, vulnerability curves translate these phenomena into actual impacts on assets.

The OS-Climate methodology implements vulnerability models through conditional probability distributions that capture the uncertainty in asset response to hazard events.

Formula 11: Vulnerability Curve

For a given hazard intensity s , the vulnerability curve provides the probability distribution of damage/impact d :

$f_{D|S}(d, s)$ = probability that given hazard intensity s , impact d occurs

Where:

- $f_{D|S}(d, s)$ is the conditional probability density function of damage/impact.
- d represents the level of damage or impact to the asset (often expressed as a percentage of total value).
- s represents the hazard intensity (e.g., flood depth in meters, wind speed in km/h)

The effective impact distribution is then derived through the convolution of the hazard and vulnerability distributions:

$$f_D(d) = \int_{-\infty}^{\infty} f_S(s)f_{D|S}(d, s)ds$$

Where:

- $f_D(d)$ is the overall probability density function of damage level d .
- $f_S(s)$ is the probability density function of hazard intensity s .

The integration combines all possible hazard intensities to determine the total damage probability distribution

The vulnerability curve shows the range of possible damage outcomes for a specific hazard level. Rather than predicting a single damage amount, it gives us the probability of different damage levels occurring when faced with a particular flood level.

When we talk about the effective impact distribution, we're combining two critical pieces of information: how likely different hazard intensities are to occur in your area, and what damage each intensity might cause. This gives us the complete picture of possible damages by considering

every potential hazard scenario, weighing each by its likelihood, and adding up all the resulting damage probabilities.

“Vulnerability curves account for how likely different flood depths are in your location and how the building’s resilience responds to each depth, producing your overall damage risk profile.” - David Kelly, MKM

Types of Vulnerability Models

The OS-Climate framework supports multiple vulnerability modelling approaches:

- **Expert-derived vulnerability** curves are used in engineering analysis and judgment.
- **Statistical vulnerability models** derived from empirical loss data.
- **Parametric vulnerability functions** using mathematical distributions to represent damage uncertainty.

Expert-Derived Flood Vulnerability Curves

Expert-derived flood vulnerability curves represent a critical methodological component in physical risk assessment for regions without a sufficiently rich dataset. These curves are developed through a structured analytical process integrating engineering principles with quantitative expert judgment.

The development methodology typically proceeds as follows:

- Structural and hydrological specialists systematically analyse how varying inundation depths affect building typologies and construction materials.
- Quantitative "what-if" analyses determine damage thresholds at specified flood depth intervals, identifying critical points where damage functions exhibit significant gradient changes.
- Mathematical relationships between inundation depth and structural response are established and calibrated to engineering principles and material behaviour characteristics.
- The resulting functions express damage as normalised ratios (0 to 1) or percentage values of total replacement cost across the continuum of possible flood depths.

These expert-derived functions offer several methodological advantages within the risk assessment framework. They can be constructed without extensive empirical flood damage datasets, enabling vulnerability quantification in regions lacking recent flood event data. They incorporate building characteristics and flood parameters that may be absent from historical records while providing a technical foundation for forward-looking assessments under projected climate conditions.

However, significant methodological limitations must be acknowledged when applying expert-derived vulnerability functions. They rely on expert judgment rather than observed damage data and may not capture all multivariate factors influencing flood vulnerability. These functions introduce subjectivity that may vary between expert groups, necessitating careful documentation of assumptions and methodological choices. When integrated with the hazard curves, such functions form the technical basis for quantitative risk assessment, informing infrastructure planning, insurance pricing models, and resilience investment decisions.

Statistical Vulnerability Models

Statistical vulnerability models provide a data-driven alternative to expert-derived curves. These models leverage historical damage data from past flood events to establish empirical relationships between hazard intensity variables and observed damages. Through regression analysis or machine learning techniques, these models can identify key vulnerability factors and provide more objective measures based on empirical evidence.

The OS-Climate framework implements several statistical approaches for flood vulnerability modelling. Random Forest models utilise environmental factors, building characteristics, and historical damage data to predict potential losses at different flood depths. Regression-based approaches establish mathematical relationships between flood parameters (depth, duration, velocity) and observed damage ratios. Empirical damage functions are derived directly from historical loss data across different asset classes and regions, offering direct translation from hazard intensity to expected damage.

Statistical vulnerability models offer significant methodological advantages in risk assessment applications. They are based on actual observed damage rather than theoretical estimates, enabling them to capture complex, non-linear relationships between hazard intensity and damage that might not be apparent through expert assessment alone. They frequently reveal unexpected vulnerability factors not initially considered in expert-driven approaches, providing insights that can refine overall risk understanding.

However, these models face essential limitations in their application. They require sufficient historical damage data, which may not be available

in all regions or asset classes. They may not adequately represent vulnerability under novel weather conditions outside historical experience. Additionally, they can be biased by the specific characteristics of the historical events depicted in the dataset, potentially limiting their generalizability to future conditions.

Parametric Vulnerability Functions

Parametric vulnerability functions represent a third approach in the OS-Climate framework, using standard mathematical distributions to model the relationship between hazard intensity and damage. These functions employ well-defined probability distributions such as Beta, Gamma, or Lognormal to express the damage response of assets to hazard intensities. The parameters of these distributions undergo calibration processes based on available empirical data or structured expert judgment. These functions systematically express damage as a percentage of asset value or absolute monetary loss, providing a standardised format for integration into risk assessment frameworks.

Beta distribution has emerged as a preferred mathematical form for modelling vulnerability in the OS-Climate methodology due to its specific properties and analytical advantages. The distribution is constrained to the interval [0,1], which naturally corresponds to the range of possible damage ratios from no damage to total loss. Its shape parameters can be adjusted to represent different vulnerability profiles across diverse asset types and construction categories. The Beta distribution effectively models the uncertainty in damage estimation at specific hazard intensities, capturing both the expected value and the variance of potential outcomes.

The development process for parametric functions follows a structured methodology that involves categorising assets into similar vulnerability groups based on construction characteristics, occupancy types, or land use categories. For each asset category, the parameters of the chosen distribution undergo estimation procedures using available data or expert assessment. These functions then undergo validation against historical damage data where such information is available, allowing for refinement of the parametric representations.

Parametric functions offer distinct methodological advantages in risk modelling applications. They provide mathematical elegance and computational efficiency, allowing for rapid calculation across large portfolios of assets. Their probabilistic nature enables the representation of uncertainty through complete probability distributions rather than point estimates, reflecting the inherent variability in damage outcomes. These functions integrate seamlessly with probabilistic risk assessment frameworks, supporting advanced risk metrics and uncertainty quantification. This makes parametric functions highly desirable for global

banking systems needing to crunch vast numbers of assets at any defined frequency.

However, several limitations must be considered in their application. Standard mathematical distributions may oversimplify complex damage mechanisms that exhibit threshold effects or non-standard response patterns. Parameter estimation presents significant challenges when working with limited empirical data, potentially introducing additional uncertainty. Furthermore, standard distributions may not accurately capture the actual shape of the damage-hazard relationship, especially for specialised assets or under hazard conditions. Within the OS-Climate framework, these parametric functions frequently complement other vulnerability modelling approaches, particularly in contexts where uncertainty quantification is crucial in comprehensive risk assessment.

The OS-Climate methodology explicitly accounts for two types of uncertainty in vulnerability:

Aleatory Uncertainty

Aleatory uncertainty represents the inherent randomness or natural variability in how assets respond to hazards, even under seemingly identical conditions. This type of uncertainty is irreducible and cannot be eliminated through additional information or improved models.

In the context of flood vulnerability, aleatory uncertainty manifests in several ways. Two identical buildings subjected to the same flood depth may experience different levels of damage due to random factors such as water flow dynamics, debris impact, or subtle differences in construction quality. This natural variability remains even with perfect building characteristics and flood parameters knowledge.

The OS-Climate framework quantifies aleatory uncertainty using probability distributions that capture the range of possible damage outcomes for a given hazard intensity. For example, rather than predicting that a 1-meter flood will cause precisely 30% damage to a specific building type, the model might represent this as a probability distribution with a mean of 30% and a standard deviation that reflects the observed variability in historical flood damage data.

Epistemic Uncertainty

Epistemic uncertainty stems from incomplete knowledge about asset characteristics and their vulnerability to hazards. Unlike aleatory uncertainty, epistemic uncertainty can be reduced through additional data collection or improved modelling.

In the context of physical risk assessment, epistemic uncertainty arises from:

- Limited information about building characteristics (e.g., foundation type, materials, age).
- Incomplete understanding of damage mechanisms.
- Scarce historical damage data for calibration.
- Limitations in modelling approaches.

These uncertainties are modelled through probability distributions rather than deterministic damage functions, allowing robust risk quantification.

Property Clustering Magnifies Flood Risk

When we assess flood risk for residential properties, looking at individual properties in isolation tells only part of the story. The spatial clustering of properties creates complex, interconnected risk dynamics that can dramatically amplify the impact of flood events beyond what individual property assessments might suggest.

Properties that cluster in flood-prone areas create a risk multiplier effect. This occurs through several interconnected mechanisms:

- **Correlated Physical Damage:** When properties are tightly clustered, they share similar flood exposure characteristics. Properties within 1000m of each other (the correlation distance parameter) experience strongly correlated impacts. This means that rather than a random distribution of damages across a portfolio, damages tend to occur in concentrated clusters—creating more significant aggregate impacts than predicted by treating properties as independent risks. A neighbourhood with 50 homes nearby will often see nearly all properties affected simultaneously by a single flood event rather than a random subset. This correlation effect is most substantial in densely developed floodplains, where property values tend to rise and fall in unison based on flood events.
- **Infrastructure Vulnerability Amplification:** Clustered properties typically share critical infrastructure, creating systemic vulnerabilities that amplify flood impacts throughout the neighbourhood. When flood waters overwhelm stormwater systems in a densely populated area, the consequences ripple through all properties connected to that system. The drainage network designed to protect homes becomes a shared

point of failure, causing water to back up into multiple properties simultaneously.

Access becomes another critical issue during flood events in clustered developments. As waters rise, road networks serving entire neighbourhoods can become impassable, cutting off not just individual homes but whole communities from emergency services. This isolation effect extends the impact of flooding well beyond direct water damage, affecting even properties that might remain dry but unreachable.

Perhaps most consequential are the cascading utility disruptions that spread through densely populated areas. A single flooded electrical substation can plunge hundreds of clustered homes into darkness. Water treatment facilities overwhelmed by flood waters can render tap water unsafe across entire subdivisions. These shared service disruptions create prolonged recovery challenges that isolated properties typically don't face.

The failure of shared protective infrastructure represents the most dramatic expression of this vulnerability. When a levee breach or a communal drainage system fails, it doesn't affect just one property but simultaneously unleashes water across countless homes. This synchronised flooding creates resource competition during the emergency response and recovery phases, extending the severity and duration of impacts for all affected properties.

- **Market Value Correlation Effects:** Property clustering creates a robust market value contagion during and after flood events, where banking impacts spread through neighbourhoods much like the flood waters themselves. When a significant portion of homes in a neighbourhood experience flooding, the market reacts not just to the directly damaged properties but to the entire area. Prospective buyers begin to view the whole neighbourhood through the lens of flood risk, creating a collective devaluation effect that touches even homes that remain dry.

Flooding and market stigma

Belanger & Bourdeau-Brien's groundbreaking 2018 study offers significant insights into how flood risk manifests as market stigma in property valuations. By analysing over 600,000 residential properties across England, their research revealed that flood risk creates a complex economic geography where property values respond to actual and perceived hazards.

Properties within designated 100-year floodplains experienced notable price discounts compared to similar properties in safer areas, with the most pronounced effects on waterfront properties. However, proximity to

water bodies alone provides an incomplete picture of how markets process flood risk. Factors like elevation differences and localised topographical features create substantial variations in risk perception and resulting price impacts even within the same floodplain.

The study's examination of stigma's temporal nature revealed that post-flood price discounts typically ranged from 10% to 50%, but these effects were not permanent fixtures in the market. Instead, they gradually diminished as the collective memory of specific flood events faded, with most markets showing substantial price recovery within 4-7 years after major floods. This recovery pattern exhibited significant socioeconomic variation, as lower-priced properties demonstrated greater sensitivity and extended recovery periods than higher-priced homes, which frequently benefited from offsetting amenities such as desirable views or waterfront access.

This finding highlights how market responses to flood risk reflect objective hazard assessments and complex interplays among perceived risk, property characteristics, and location-specific amenities that can mitigate negative risk perceptions in specific contexts. Markets typically respond more strongly to recent visible events than abstract probability calculations or hazard maps, disconnecting technical risk assessments and market perceptions.

The physical recovery process influences these market effects, particularly in clustered developments. When entire neighbourhoods require simultaneous repairs, the demand for contractors, materials, and inspections creates bottlenecks that significantly extend reconstruction timelines. Homeowners compete for limited recovery resources, expanding the visible signs of damage throughout the community and prolonging the period of market impact.

Belanger & Bourdeau-Brien also documented a "halo effect" of property devaluation that spreads through clustered developments after flood events. Their analysis reveals that properties located within the same neighbourhood as flooded homes but themselves untouched by waters still experienced statistically significant devaluations—demonstrating that in clustered developments, flood risk impacts extend beyond physical water damage to affect entire market ecosystems.

The study's findings carry significant policy implications for flood risk management. Traditional floodplain designations frequently overlook the nuanced ways in which buyers evaluate and price flood risk. This disconnect highlights the need for more detailed, property-specific risk assessment tools to align market perceptions with hazard profiles better.

Furthermore, the documented pattern of temporary stigma periods reveals opportunities for targeted recovery interventions that could

accelerate market normalisation following flood events, potentially reducing economic disruption in affected communities. By emphasising these dynamics, Belanger and Bourdeau-Brien's work provides essential guidance for policymakers aiming to address the financial dimensions of flood vulnerability in an era of extensive new property development.

Quantifying the Clustering Effect

The portfolio impact model captures this clustering effect through the correlation matrix Σ , combined with the shock factor σ :

Formula 12: Belanger and Bourdeau-Brien

$$f_D(d) = \int_{-\infty}^{\infty} f_S(s) f_{\{D|S\}}(d, s) ds$$

Where:

- $f_D(d)$ represents the probability density function of the damage level d across the portfolio
 - $f_S(s)$ is the probability density function of the shock intensity s (flood severity)
 - $f_{\{D|S\}}(d, s)$ is the conditional probability density function of damage d given shock s
 - The integration from $-\infty$ to ∞ accounts for all possible shock scenarios
-

The portfolio impact model measures how flood risks spread among closely situated properties by analysing their interconnectedness. This approach goes beyond treating each property as an isolated risk. It recognises that when properties are clustered together, the impact of flooding on one affects the others in predictable ways.

The correlation matrix illustrates the strength of the connections between neighbouring properties—when one property suffers from flooding damage, others nearby encounter similar risks due to their proximity and shared environmental characteristics. The shock factor subsequently measures the extent of this ripple effect across the cluster.

This mathematical framework enables us to calculate the overall probability distribution of damage across an entire neighbourhood or development. It combines the individual probabilities of specific shock events with the conditional probabilities of how those shocks translate into actual damage patterns across properties.

The clustering effect can better predict how flooding impacts cascade through developments where homes share similar construction methods, elevation profiles, and proximity to water sources. The model demonstrates that flooding damage doesn't occur randomly across properties but follows patterns that reflect geography and built environment characteristics.

Understanding these clustering dynamics allows insurers, developers, and policymakers to more accurately assess the actual financial vulnerability of residential areas to flood events rather than simply calculating risks for individual properties in isolation:

- **Higher tail risk:** The 95% Value at Risk (VaR) and Expected Shortfall (ES) metrics increase dramatically compared to diverse portfolios of the same size.
- **Greater volatility:** Portfolio impact variance increases with clustering
- **Reduced diversification benefit:** The risk-reducing effect of portfolio diversification diminishes as spatial correlation increases

A portfolio with properties distributed across different watersheds might see a 95% VaR 2-3 times lower than a portfolio of equal value concentrated in a single flood-prone area.

Clustering Patterns in Real Estate Portfolios

Residential property clustering follows systematic patterns that inadvertently maximise flood risk exposure across investment portfolios. The price premium commanded by waterfront views creates high-value property concentrations precisely in the highest-risk areas, generating natural risk accumulations.

Developmental history contributes significantly to risk concentration, as many neighbourhoods were constructed during single development phases, resulting in uniform risk profiles across entire communities sharing identical construction standards and elevations. Infrastructure determinism further reinforces these patterns, as road networks and utility systems naturally follow topographical features like river valleys, creating property density along the same hydrological pathways that channel floodwaters.

These patterns establish natural risk concentration hotspots where property density and flood susceptibility align—precisely the conditions where correlation effects manifest most strongly. Understanding these clustering dynamics informs several key risk management approaches.

Geographic diversification across multiple watersheds and flood zones becomes essential for reducing portfolio-level exposure.

Correlation-based pricing mechanisms adjust mortgage rates and insurance premiums based on neighbourhood-level clustering characteristics rather than solely on individual property risk metrics. Targeted mitigation strategies focus infrastructure improvements on protecting entire clusters rather than attempting less cost-effective individual property-level interventions. Advanced stress testing methodologies incorporate portfolio simulations with varying correlation parameters to identify vulnerabilities to specific clustering effects.

As flood patterns shift due to additional property development that converts permeable surfaces to impervious concrete, previously diverse portfolios may become inadvertently clustered regarding risk exposure. Areas once considered to have independent flood risks may begin experiencing correlated flooding as weather patterns and hydrological systems change. This dynamic risk landscape necessitates continuous reassessment of portfolio correlation structures rather than relying on static risk assessments based on historical patterns.

The mathematical analysis of spatial correlation confirms what many risk managers have intuitively understood. When evaluating flood risk, the relative geographic positioning of properties within a portfolio is comparable to their risk profiles. By explicitly modelling these clustering effects, analysts develop a more accurate understanding of portfolio-level flood risk. They can formulate more effective strategies for managing that risk in an increasingly uncertain climate future.

These territorial rating methodologies now feed into sophisticated insurance premium calculations incorporating several advanced components. Peril-specific territories establish separate geographic segmentation frameworks for flood, wind, fire, and other hazards. Proximity factors enable rating adjustments based on distance to risk features such as coastlines, floodplains, or fault lines.

Topographic variables incorporate elevation, slope, aspect, and other terrain characteristics influencing risk profiles. Hydrological modelling directly integrates outcomes from the flood models discussed in previous chapters. The following chapter will examine These methodological components in greater detail, forming the foundation for comprehensive risk pricing frameworks.

Quantifying Resilience Using Standard Operating Procedures

Standard Operating Procedures (SOPs) provide a structured approach to quantifying resilience by formalising how systems prepare for, respond to, and recover from disruptive events. When properly implemented, SOPs transform theoretical resilience concepts into measurable outcomes, particularly for built environments facing hazards like flooding.

SOP Resilience Calculation

Formula 13: SOP Fundamental

The fundamental calculation for SOP resilience coverage can be expressed as:

$$S = (\text{Implemented SOPs} + \text{Excluded SOPs}) / (\text{Total Recommended SOPs}) \times 100$$

Where:

- S represents the SOP coverage score (0-100%).
 - Implemented SOPs are procedures actually in place.
 - Excluded SOPs are procedures deliberately omitted with justification.
 - Total Recommended SOPs encompass all procedures advised for the specific context.
-

This calculation yields a normalised percentage that can be integrated into broader resilience assessments.

Resilience Quantification: Assessment

Begin by establishing an unprotected baseline scenario. For a property in a flood zone, this might involve:

- **Unprotected flood depth:** 1 meter during a 1-in-100-year event.
- **Estimated damage:** \$X for building structure and contents.
- **Recovery time:** Y days/weeks before reoccupation.
- **Operational continuity:** Complete disruption for Z days.

Resilience Quantification: SOP Inventory and Analysis

Document all recommended resilience procedures for the property type, categorising them by resilience phase:

- **Preparedness** SOPs: Early warning monitoring, seasonal maintenance, supplies stockpiling
- **Protection** SOPs: Deployment of barriers, equipment elevation, utility isolation
- **Response** SOPs: Evacuation procedures, emergency communication protocols
- **Recovery** SOPs: Water removal sequence, sanitisation procedures, systems restoration

Each SOP should be assigned a weighted value based on its contribution to overall resilience.

Resilience Quantification: Implementation Calculation

Formula 14: Phase-specific SOP

For each resilience phase, calculate a phase-specific SOP score:

$$S_{\text{phase}} = (\sum(w_i \times I_i)) / (\sum w_i)$$

Where:

- S_{phase} is the score for a specific resilience phase.
- w_i is the weight of the i -th SOP.
- I_i is the implementation status (1 for implemented, 0 for not implemented).
- Σ represents the summation across all SOPs in that phase.

This calculation evaluates each SOP in that phase, determines whether it has been implemented (assigning a value of 1) or not (assigning a value of 0), and multiplies this implementation status by the weight assigned to that specific SOP. These weighted values are then summed and divided by the total of all the weights. This results in a score reflecting the implementation level for that particular resilience phase.

The resulting S_phase score provides a normalised measure (typically between 0 and 1) that indicates how thoroughly weighted SOPs have been implemented within a specific resilience phase. A score closer to 1 signifies that most important SOPs (those with higher weights) have been implemented, while a score nearer to 0 suggests that many significant SOPs remain unimplemented.

This method ensures that more critical procedures (those with higher weights) have a proportionally more significant effect on the overall resilience score for that phase.

Resilience Quantification: Hazard Reduction Modelling

The implementation score translates to hazard reduction through modelling relationships between procedures and outcomes:

Formula 15: Hazard Reduction

$$H_{\text{reduced}} = H_{\text{baseline}} \times (1 - R_{\text{eff}})$$

Where:

- H_{reduced} is the reduced hazard level
- H_{baseline} is the baseline hazard level
- R_{eff} is the effective resilience factor derived from the SOP scores

For our flood example:

- Baseline flood depth = 1m
 - SOP implementation score = 80%
 - Effective resilience factor = 0.8
 - Reduced flood depth = $1m \times (1-0.8) = 0.2m$ (20cm)
-

Resilience Quantification: Return Period Translation

The reduced hazard level can then be translated to an adjusted return period:

Formula 16: Return Period Adjustment

$$T_{\text{adjusted}} = T_{\text{baseline}} \times f(R_{\text{eff}})$$

Where:

- $T_{adjusted}$ is the adjusted return period.
 - $T_{baseline}$ is the baseline return period.
 - $f(R_{eff})$ is a site-specific function relating resilience to return period adjustment.
-

In our example, a properly implemented set of SOPs effectively transforms a 1-in-100-year flooding event into a 1-in-20-year event in terms of impact.

Case Example: Residential Property Flood Resilience

Consider a single-family home in a flood-prone area with the following characteristics:

- It is located in a 1-in-100-year floodplain with a 1m projected flood depth
 - Home replacement value: \$350,000
 - Contents value: \$75,000
 - Displacement cost: \$3,000 per month
 - Recovery time without SOPs: 4 months

The homeowner implements a comprehensive set of SOPs:

- Preparatory SOPs
 1. Installation of flood-resistant materials in lower level
 2. Regular cleaning of gutters and drainage pathways
 3. Elevation of electrical systems and HVAC equipment
 4. Creation of a household emergency plan
- Response SOPs
 1. Deployment of door/window flood barriers
 2. Moving valuable items to the upper floors
 3. Secure storage of essential documents in waterproof containers
- Recovery SOPs
 1. The proper sequence for water removal and drying
 2. Mold prevention procedures

Using the formula and appropriate weights for each SOP category:

$$S_{total} = ((4 \times w_{prep}) + (3 \times w_{resp}) + (2 \times w_{recov})) / ((5 \times w_{prep}) + (4 \times w_{resp}) + (3 \times w_{recov}))$$

Assuming weights of 0.4 for preparatory, 0.4 for response, and 0.2 for recovery SOPs, the SOP implementation score is approximately 75%.

This translates to:

- Reduced flood depth: 25cm (instead of 1m).
- Shortened recovery time: 1.5 months (instead of 4).
- Damage reduction: \$85,000 to structure and contents.
- Displacement cost reduction: \$7,500.
- Return period adjustment: Impacts equivalent to a 1-in-25-year event.

Integration with Broader Resilience Frameworks

SOP resilience calculations should be integrated with other resilience metrics for comprehensive assessment:

Formula 17: Total Property Resilience

$$R_{\text{total}} = w_1 \times S + w_2 \times P + w_3 \times A + w_4 \times T$$

Where:

- R_{total} is the total resilience score.
 - S is the SOP coverage score.
 - P is the physical infrastructure score.
 - A is the adaptive capacity score.
 - T is the technological readiness score.
 - w_1 through w_4 are respective weights.
-

Quantifying resilience through SOP implementation provides a structured approach to understanding and enhancing a property's ability to withstand disruptive events. Proper procedures transform a 1-in-100-year event into an effective 1-in-25-year event, demonstrating the tangible value of operational resilience for residential properties.

This methodology allows homeowners, insurers, and policymakers to make informed decisions about resilience investments based on measurable outcomes rather than theoretical frameworks alone. This could result in reduced insurance premiums, decreased recovery costs, and enhanced property values. The next step is to ensure that these fields are included in a property's CDM design.

Chapter 7 - Insurance Risk Assessment

Inurance calculation represents a critical junction in the physical risk assessment process, where the scientific understanding of hazards developed in previous chapters converges with banking risk transfer mechanisms. The preceding chapters have established a comprehensive framework for understanding flood risk—from the fundamental principles of weather prediction and hydrological modelling, hydrological routing and hydraulic modelling to the detailed assessment of flood hazards and property-specific vulnerabilities. Insurance modelling builds directly upon this foundation, transforming these physical insights into economic terms through structured mathematical frameworks.

"Insurance serves as the essential bridge between physical risk science and banking markets—it translates the complexities of hydrodynamic modelling and property vulnerability into an annual premium reflecting a precise risk quantification. While mortgages extend risk over decades, insurance distils it into a single year, creating both a temporal mismatch and a powerful market signal about the true cost of hazards." –Swiss Re Institute.

This connection is particularly evident in territorial rating methodologies, which practically implement the hydrological and hydraulic modelling discussed in earlier chapters. The sophisticated flood modelling approaches—including riverbank breach modelling, digital elevation analysis, and flow path simulation—directly inform how insurers delineate rating territories and assign relative risk factors. These territories essentially translate complex hydrological realities into financially actionable zones for risk pricing.

The insurance industry is uniquely positioned in the physical risk landscape due to its one-year time horizon. Unlike mortgages, which commit capital over decades (as we will explore in Chapter 8), insurance

contracts typically last just twelve months, allowing for rapid adaptation to emerging risk information.

This fundamental difference creates challenges and opportunities—insurers can quickly incorporate new physical risk data into their pricing models. Still, this same flexibility allows them to withdraw coverage entirely from areas they deem too risky, potentially leaving longer-term mortgage exposures unprotected.

This chapter explores how insurers bridge the gap between scientific risk assessment and banking quantification. We begin with premium calculation approaches that translate territorial risk assessments derived from hydrological and hydraulic models into specific pricing structures. We then examine how recent advancements in territorial rating methodologies have enabled more granular, property-specific risk assessments that more accurately reflect the underlying physical realities modelled in previous chapters.

Premium Calculation

Calculating insurance premiums represents the practical application of risk modelling principles within the constraints of market realities, regulatory requirements, and business objectives. For property insurance, particularly in areas prone to flooding and other natural hazards, advanced premium calculation methods are essential for ensuring insurer solvency and market accessibility.

The fundamental insurance premium equation comprises several distinct components:

$$\text{Premium} = \text{Pure Premium} + \text{Expense Loading} + \text{Risk Loading} + \text{Profit Margin}$$

Where:

- **Pure Premium:** The expected loss cost, as determined by comprehensive risk modelling
- **Expense Loading:** Operational costs, including acquisition, administration, and claims handling
- **Risk Loading:** Additional charge for uncertainty and volatility in the loss distribution
- **Profit Margin:** Required return on capital to support the business

The pure premium component derives directly from the risk modelling process through the application of credibility theory, which balances the predictive value of specific risk data against broader class experience:

Formula 18: Pure Premium

$$Z = n / (n + K)$$

Where:

- Z represents the credibility factor.
- n is the number of exposure units.
- K is a constant reflecting the expected variance of the risk class.

The final pure premium calculation then becomes:

$$PP = Z \times \text{Individual Experience} + (1 - Z) \times \text{Class Experience}$$

This approach ensures that properties with sufficient historical data receive individualised rates, while those with limited experience benefit from the stability of class ratings.

Territorial Rating Refinement

Territorial rating has long been a cornerstone variable for property insurance. It is traditionally defined by administrative building blocks such as postal codes, census blocks, or municipal boundaries.

Building on the flood modelling techniques discussed in previous chapters, modern approaches employ significantly more granular spatial segmentation.

The territorial rating process involves two critical phases:

- Determining the boundaries of each territory
- Establishing the rate relativities for these territories

Traditional territory-based methods faced an inherent challenge: balancing the need for territories large enough to provide statistical credibility while small enough to encompass regions with relatively uniform risk exposure. Key considerations for effective territorial boundaries include:

- **Credibility thresholds:** Ensuring sufficient exposure and claim volume
- **Risk homogeneity:** Minimizing variance within territories
- **Intuitive boundaries:** Creating defensible geographic delineations
- **Regulatory compliance:** Adhering to jurisdictional constraints

Modern flood insurers abandon the conventional zone-based rating approach in favor of coordinate-based systems that represent flood risk as a continuous gradient across landscapes.

This approach uses property-specific latitude and longitude coordinates integrated with high-resolution digital elevation models to assess flood exposure with unprecedented precision.

As Mitchell-Wallace et al. note in their practitioner's guide, this creates "continuous risk surfaces that better reflect localised risk gradients," enabling insurers to differentiate between properties even within the same postal code. As we have covered in the previous chapters, modern approaches to hydrology and hydraulic modelling enable advanced tools that allow the following:

- **Micro grid-based pricing:** using uniform spatial grids down to 5m that transcend administrative boundaries and allow different Insurance for two properties in the same Zipcode (Postcode)
- **Risk-based clustering:** Employing machine learning to identify regions of similar risk profiles that have more measurable historic events
- **Dynamic boundaries:** Adjusting territory definitions based on emerging loss patterns and changing hazard conditions
- **Multi-level territories:** Implementing hierarchical structures that balance granularity with credibility

The evolution toward these micro-territories has particular relevance for flood modelling, where risk can vary dramatically within short distances based on elevation, drainage infrastructure, and proximity to water bodies. These approaches enable the incorporation of spatial correlation in premium calculations, recognising that geographic proximity often entails similar flood risk exposures.

Others have begun modelling territorial signals directly at the individual risk level, creating a continuous risk surface rather than discrete territories. This approach involves smoothing and aggregating nearby risks

until sufficient credibility is reached, creating dynamic micro-territories that adapt to emerging patterns. These sophisticated methodologies typically employ holdout validation techniques to ensure model accuracy, with insurers reserving portions of their data to verify that territorial refinements genuinely have indicative power.

The advantages of these advanced geospatial approaches extend beyond technical sophistication. Insurers implementing these methods report significantly increased accuracy in risk assessment, with greater differentiation and more significant properties even within traditionally defined territories. Perhaps more importantly, these approaches enable insurers to respond to localised changes in risk levels without requiring broad rate changes across entire regions.

This granularity promotes fairness and precision in pricing, ensuring policyholders pay rates more closely aligned with their actual risk exposure. From an operational perspective, these systems typically enable incredible speed and efficiency in risk assessments and policy quotes, as territory assignment becomes an automated, algorithmic process rather than a manual lookup procedure.

Multivariate Class Rating

Modern premium calculation employs sophisticated multivariate rating plans that simultaneously consider numerous risk factors through generalised linear models (GLM):

Formula 19: Multivariate Premium Calculation

$$\log(\text{Premium}) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n$$

Where:

- $\log(\text{Premium})$ is the natural logarithm of the premium amount.
 - β_0 is the intercept or base rate coefficient.
 - $\beta_1, \beta_2, \dots, \beta_n$ are the regression coefficients that represent the effect of each risk factor.
 - X_1, X_2, \dots, X_n are the risk factors or predictor variables (both continuous and categorical).
-

These models incorporate both continuous variables (e.g., distance to coast) and categorical factors (e.g., construction type) through:

- **Main effects:** Direct impact of individual factors.

- **Interaction terms:** Combined effects of factor pairs.
- **Offset variables:** External constraints on the rating structure.
- **Smoothing functions:** Techniques like cubic splines for handling nonlinear relationships.
- **Advanced implementations:** extend GLMs through techniques such as:
 - Generalised additive models (GAMs) for capturing complex nonlinear relationships.
 - Mixed-effects models for addressing hierarchical data structures.
 - Elastic net regularisation for variable selection and coefficient stability.
 -

Cost of Capital and Risk Loading

Property insurance—particularly for catastrophe-exposed regions—requires substantial capital reserves to cover potential losses. The risk-loading component of premiums must reflect this capital requirement:

$$\text{Risk Loading} = \text{CoC} \times \text{Required Capital}.$$

Where CoC represents the cost of capital (typically exceeding risk-free rates by 6-12 percentage points), and Required Capital is determined through:

$$\text{Required Capital} = \text{TVaR}_{99.5} - \text{Expected Loss}$$

Where TVaR_{99.5} (Tail Value at Risk) represents the average loss in the worst 0.5% of scenarios. This approach ensures that the premium adequately compensates for the extreme tail risk characteristic of property catastrophe exposures.

Solvency II regulatory framework

The European Union formally implemented Solvency II on January 1, 2016, following nearly a decade of development and several postponements from its originally planned 2012 launch. All 27 EU member states were required to adopt this regulatory framework for insurance and reinsurance companies. While initially implementing Solvency II as an EU member, the UK has maintained these regulations post-Brexit through the UK Solvency II regime. However, some modifications have been introduced over time.

Under Solvency II, insurers must account for catastrophe risk in their Solvency Capital Requirement (SCR) using either the standard formula or

internal models. The standard formula utilises predefined scenarios and correlation matrices, with catastrophe risk integrated within underwriting modules (non-life, life, and health). For non-life catastrophe risk, insurers apply 1-in-200-year loss scenarios across various perils and regions, reporting capital requirements both before and after implementing risk mitigation strategies like reinsurance. Life catastrophe risk involves more straightforward calculations, such as applying 0.15% to the total capital at risk across policies.

Alternatively, larger insurers may develop internal models that reflect their specific risk profiles more accurately. These regulator-approved models use Monte Carlo simulations to model loss distributions with company-specific exposure data and reinsurance structures. While the standard formula offers simplicity with fixed scenarios and correlations, internal models provide greater precision but require more complex data and validation processes.

Both approaches aim to ensure insurers maintain sufficient capital reserves to withstand extreme events at the 99.5% confidence level (1-in-200-year Value at Risk), balancing policyholder protection with operational feasibility.

This regulatory charge directly influences premium pricing, as insurers must generate sufficient returns to compensate for this committed capital. Thus, the regulatory view of extreme physical risk effectively becomes a direct component of consumer premium costs.

While these advancements mentioned earlier offer improved risk segmentation—with some insurers reporting 15–25% reductions in cross-territory rate error compared to conventional methods—they also face regulatory scrutiny. Insurance commissioners increasingly demand transparency in these sophisticated models, requiring detailed documentation of how geographic risk factors are isolated and validated before approving rate filings based on these new methodologies.

The ongoing evolution of flood premium modelling represents a material leap forward in risk-based pricing accuracy, potentially reducing cross-subsidization while better-aligning incentives for property-level flood mitigation investments.

Experience and Schedule Rating

Beyond model-driven base rates, premium calculation often incorporates more nuanced methodologies to account for individual risk characteristics and historical performance. Two key approaches that enable this refinement are experience rating and schedule rating.

Experience rating adjusts premiums based on a policyholder's actual loss history relative to expected losses for similar risks. This approach recognises that historical performance is a reliable predictor of future claims activity.

Calculation Methodology

The experience modification factor is typically calculated as follows:

$$\text{Experience Modification} = (\text{Actual Losses} + \text{ESF}) / (\text{Expected Losses} + \text{ESF})$$

Where:

- Actual Losses: The policyholder's historical claims over a defined period (typically 3-5 years)
- Expected Losses: Industry average losses for similar risk profiles
- ESF (Experience Stabilization Factor): A parameter that dampens the impact of random fluctuations, crucial for smaller portfolios where a single large claim could create disproportionate premium volatility

The resulting factor directly modifies the base premium:

- **Factor > 1.0:** Premium surcharge for worse-than-expected loss experience
- **Factor = 1.0:** No adjustment (experience aligns with expectations)
- **Factor < 1.0:** Premium credit for better-than-expected loss experience

Credibility Weighting

The ESF functions as a credibility mechanism that balances the statistical significance of individual experience against broader risk class data. As exposure volume increases, the formula gives greater weight to personal experience:

$$Z = n / (n + \text{ESF})$$

Z represents the credibility assigned to individual experience, and n reflects exposure units. Exposure units represent the volume of risk-bearing activity (e.g., payroll, property value, or operational scale) used to quantify an insured entity's exposure to potential losses. When calculating premiums, this parameter determines how much weight an insurer assigns to the entity's loss experience versus industry-wide data.

Example: Workers' Compensation

A construction firm with a \$5M annual payroll ($n = 500$ units at \$10,000/unit) and an ESF of 1,000 would have:

$$Z = 500 / (500 + 1,000) = 0.33$$

Here, 33% of the premium adjustment would derive from the firm's claims, while 67% would use industry data.

If payroll grows to \$15M ($n = 1,500$):

$$Z = 1,500 / (1,500 + 1,000) = 0.6$$

Now, 60% of the weight is on the firm's experience, reflecting greater statistical confidence.

Practical Application

For property insurance, particularly in flood-prone areas, experience rating might consider frequency and severity of previous flood claims, loss mitigation responses at earlier events, time elapsed since the last significant claim, and loss ratio compared to premium history.

Example: A commercial property

A commercial property with \$10M in coverage has experienced two minor flood claims totalling \$125,000 over the past decade, compared to an expected loss of \$200,000 for similar properties. With an ESF of \$75,000, their experience modification would be:

$$(\$125,000 + \$75,000) / (\$200,000 + \$75,000) = \$200,000 / \$275,000 = 0.73$$

This 0.73 modification factor would reduce their base premium by 27%, reflecting their better-than-expected loss history.

Schedule Rating

While experience rating looks backward at actual loss history, schedule rating evaluates current risk characteristics that statistical models or territorial classifications may not fully capture.

Implementation Approach

A schedule rating applies structured credits or debits to the base premium based on specific risk factors. Insurance underwriters evaluate properties against defined criteria, with typical adjustment ranges of ±25% from the base rate.

Key Rating Factors

Schedule rating for flood insurance typically evaluates several critical categories. Loss control measures include evaluating property-specific defences such as elevated critical systems, backflow prevention devices, and temporary flood barrier deployment capabilities. Construction features assessment focuses on the building's inherent resilience, examining foundation integrity, water-resistant materials usage, and structural reinforcements against hydrostatic pressure.

Management practices evaluation encompasses organisational readiness through emergency response planning, staff training protocols, and ground-floor exposure minimisation strategies. Maintenance programs review examines the property owner's commitment to regularly inspecting water intrusion points, maintaining drainage systems, and adhering to manufacturer schedules for flood defence systems.

Practical Application

Example: A manufacturing facility located in a moderate flood zone might receive the following schedule rating net computed as the sum of adjustments:

Factor	Condition	Credit/Debit
Elevation	Critical equipment raised 4ft above BFE	-12%
Drainage	Enhanced stormwater system with 150%	-8%
Construction	Water-resistant materials on first floor	-5%
Maintenance	Quarterly inspection program not	4%
Emergency Plan	No formal flood response protocol	6%
Net Adjustment		-15%

Table 5: Insurance Premium Adjustment

This 15% net credit would apply to the base premium (after any experience rating adjustments), reflecting specific risk mitigation efforts despite some management deficiencies.

Premium Optimization

Contemporary insurance pricing extends beyond pure risk assessment to incorporate strategic business considerations through premium optimisation techniques. This approach balances actuarial indications with market realities to achieve broader organisational objectives.

Market-Responsive Pricing

Premium optimisation leverages demand elasticity modelling to understand how price changes affect purchasing behaviour across different market segments. By analysing price sensitivity, insurers can identify segments where modest premium increases may have minimal impact on retention while applying more competitive pricing in highly elastic segments to drive growth.

The optimisation function typically takes the form:

$$\text{Maximize: } \sum_i P_i(r_i) \times (r_i - LR_i)$$

Where P_i represents the purchase probability at rate r_i , and LR_i the expected loss ratio.

Competitive Positioning

Advanced optimisation models incorporate competitive intelligence to position premiums strategically within the market landscape. This requires sophisticated data collection on competitor pricing and regular market basket analyses to understand relative value propositions.

Insurers typically establish target competitive positions based on product features, service levels, and brand strength, then optimise premiums to maintain these relative positions while achieving profitability targets.

Customer Lifetime Value

Rather than optimising short-term profitability, contemporary approaches increasingly focus on maximising customer lifetime value (CLV).

Formula 20: Customer Lifetime Value

$$CLV = \sum_t (P_t \times (r_t - LR_t - E_t)) / (1 + d)^t$$

Where:

- P_t (Probability of renewal): The likelihood that a customer continues their policy in period t . This reflects customer retention rates.
 - r_t (Premium): The amount the customer pays for insurance coverage in period t .
 - LR_t (Expected loss ratio): The anticipated claims costs as a proportion of premium. This represents the core insurance cost.
 - E_t (Expense loading): All operational costs associated with servicing the customer (administration, customer service, etc.).
 - $(r_t - LR_t - E_t)$: This calculates the profit margin per customer in period t .
 - d (Discount rate): The rate used to convert future profits into present value, accounting for the time value of money.
 - $(1 + d)^t$: The discount factor that decreases the value of future profits based on how far in the future they occur.
 - \sum_t : Summation over all periods of the customer relationship.
-

Customer Lifetime Value (CLV) represents the total financial contribution a customer is expected to make to a company throughout their entire relationship. This concept has become increasingly crucial in insurance companies shifting from short-term profitability models to longer-term value maximisation approaches.

In the insurance context, CLV considers several key elements:

- **Renewal Probability:** The likelihood that a customer will continue their policy over time, reflecting customer loyalty and satisfaction.
- **Premium Payments:** The revenue generated from the customer's insurance payments, which may change over time.
- **Expected Claims Costs:** The anticipated payouts for customer claims, which directly impact profitability.
- **Operational Expenses:** The costs associated with servicing the customer, including administration, customer service, and overhead.

- **Time Value of Money:** Future profits are discounted to present value, acknowledging that money received in the future is worth less than money received today.

This approach allows insurers to make more strategic decisions. For instance, they might offer lower initial premiums to customer segments that data suggests will develop into profitable, long-term relationships. This is especially valuable when cross-selling opportunities exist across multiple insurance products.

Insurers can build more sustainable business models that prioritise customer retention and relationship development by focusing on lifetime value rather than immediate profitability. This might mean accepting lower margins or even initial losses on new customers with the understanding that the long-term relationship will ultimately be profitable.

Constrained Optimization

Premium optimisation operates within multiple constraint systems:

- Regulatory constraints on rate changes and rating factor usage.
- Business constraints on growth and retention targets.
- Operational constraints on implementation capabilities.
- Risk constraints on portfolio diversification requirements.

Sophisticated optimisation models employ techniques like linear and nonlinear programming, multi-objective optimisation, and stochastic dynamic programming to navigate these complex constraint environments.

Experience and Schedule Rating

Modern premium optimisation systems integrate with experience and schedule rating methodologies to create a cohesive pricing framework:

Calibrated Elasticity Models

Advanced systems calibrate elasticity models based on historical experience rating data, recognising that price sensitivity often varies based on loss history. Policyholders with favourable experience modifications typically demonstrate higher retention rates even when facing modest premium increases, while those with poor experience may be more price-sensitive.

Optimised Schedule Rating Factors

Rather than applying uniform schedule rating approaches, optimisation techniques identify the most influential schedule rating factors for different market segments. This allows underwriters to focus on the most impactful risk characteristics during property inspections and provides more precise guidance on credit allocation priorities.

Simulation-Based Optimization

Leading insurers employ Monte Carlo simulations that model the combined impact of experience rating, schedule rating, and market factors on portfolio performance. These simulations allow for robust sensitivity analysis and scenario testing before implementing rating changes.

Regulatory Considerations

The premium calculation must operate within regulatory frameworks that vary substantially by jurisdiction. Key regulatory constraints include:

- **Rate adequacy requirements:** Ensuring premiums are sufficient to cover expected costs.
- **Rate equity standards:** Prohibiting unfair discrimination between similar risks.
- **Rate stability provisions:** Limiting the magnitude of premium changes.
- **Filing requirements:** Prior approval vs. file-and-use vs. use-and-file systems.

In many jurisdictions, insurers must demonstrate the actuarial soundness of their premium calculation methodologies through:

- Statistical justification of rating factors.
- Analysis of indicated vs. selected rate changes.
- Compliance with specific limitations on rating variables.
- Documentation of modelling methodologies.

These regulatory considerations create a complex constraint environment for premium calculation, necessitating sophisticated optimisation approaches that balance actuarial indications with regulatory requirements.

Technology-Enabled Pricing Innovations

The insurance industry is undergoing a significant transformation in calculating premiums for flood risk. Technological advancements that allow for more nuanced risk assessment are driving this transformation. These innovations reshape traditional pricing, creating opportunities and challenges for insurers, policyholders, and regulators.

Real-Time Data Integration: The Dynamic Premium

Historically, flood insurance pricing relied on static risk assessments updated infrequently, sometimes only every few years. Insurers are increasingly incorporating real-time data streams that provide continuously updated risk information. Weather forecasts and alerts feed directly into pricing algorithms, allowing premiums to reflect imminent weather patterns rather than just historical averages.

Property modification data—such as newly installed flood barriers or drainage improvements—can be captured and reflected in premiums more quickly. Even occupancy and utilisation patterns are being monitored to adjust risk profiles based on whether properties are continuously occupied or seasonally vacant.

This dynamic approach offers several advantages: premiums more accurately reflect current risk conditions, policyholders can see more immediate banking benefits from risk mitigation investments, and insurers can manage their exposure more effectively during peak risk seasons. However, this approach also introduces volatility into premiums that may frustrate consumers accustomed to stable pricing. Additionally, the constant influx of data creates significant challenges for actuarial teams to validate and incorporate responsibly, potentially leading to unintended biases or errors if not properly managed.

Parametric Pricing: Objectifying Risk

Parametric components are increasingly supplementing traditional indemnity-based insurance models. These components trigger rate adjustments based on objective, measurable parameters rather than subjective assessments. Distance-based flood premium components, for example, might automatically adjust rates based on precise measurements from flood plains or water bodies rather than relying on broad zone classifications. Wind-speed triggered adjustments can modulate premiums based on recorded wind velocities in a region. Most innovatively, satellite-derived vegetation indices are being used to assess ground cover and permeability, factors that significantly impact flood risk but were previously difficult to quantify at scale.

Introducing parametric elements brings greater transparency to pricing—a clear benefit for consumer understanding and regulatory oversight. These objective measures also reduce disputes over claims and pricing decisions. However, parametric approaches may sometimes oversimplify complex risk factors and create discontinuities in pricing that seem arbitrary to consumers. For instance, properties on opposite sides of a distance threshold might see dramatically different premiums despite minimal practical difference in risk.

“Parametric insurance is the foundation for designing a Physical Risk Swap.”-David Kelly, MKM.

Machine Learning: The Algorithmic Revolution

The most profound change in premium calculation comes from machine learning applications that enhance predictive accuracy beyond what traditional actuarial methods could achieve. Gradient boosting algorithms define more granular rate classes by identifying subtle patterns in risk data that human analysts might miss. Neural networks improve claims cost prediction by recognising complex relationships between seemingly unrelated factors. Reinforcement learning algorithms optimise portfolios, balancing risk exposure across different property types and geographical areas.

These advanced techniques deliver remarkable improvements in pricing accuracy, allowing insurers to match premiums to actual risk better and expand coverage to previously uninsurable properties. They also enable more personalised pricing that rewards risk-reducing behaviours. The downside, however, is the "black box" nature of many machine learning models, which creates significant challenges in explaining rate decisions to consumers and regulators. There are also legitimate concerns about algorithmic bias potentially perpetuating or even amplifying existing inequities in insurance access and pricing.

Evolving Risk Cost Assessment

The integration of evolving risk cost patterns into household insurance models reflects a growing recognition of the compounding financial pressures from natural hazards, even as hazard frequency remains relatively stable. In 2023, insured losses from natural catastrophes reached \$108 billion, as the Swiss Re Institute reported. This marked the fourth consecutive year that such losses exceeded \$100 billion globally.

A record 142 natural catastrophes were reported during the year, with severe convective storms (SCS) representing the largest share of losses at \$64 billion. These storms, which include events such as hailstorms and tornadoes, had a particularly significant impact in the United States,

contributing to 85% of SCS-related losses, while Europe experienced the fastest growth in such losses.

The most significant single event was the earthquake in Turkey and Syria, which caused \$6.2 billion in insured losses. Nevertheless, most events in 2023 resulted in medium-severity losses ranging from \$1 billion to \$5 billion. The increasing frequency and severity of these occurrences align with a long-term trend of rising insured losses driven by urbanisation, economic growth, and climate change. Swiss Re estimates that annual insured losses could double within a decade due to intensifying weather hazards and greater exposure in vulnerable regions.

To address these challenges, Swiss Re emphasises the importance of mitigation and adaptation measures, such as enforcing stricter building codes, constructing flood defences, and fostering collaboration between insurers, governments, and communities to reduce risk exposure.

Exposing property developments into high-risk zones is a significant driver of increased costs. Urbanisation and housing demand have pushed construction into floodplains, wildfire-prone regions, and coastal areas, where over 700 U.S. Superfund sites alone face flood risks. Commercial real estate professionals increasingly prioritise climate resilience (46%), yet many properties in disaster-prone areas still encounter elevated insurance premiums and limited coverage options.

In the multifamily housing sector, insurers specifically cite "weather risk" and aging infrastructure as reasons for premium increases ranging from 14% to 45% annually. These significant price hikes have compelled 93% of housing providers to raise deductibles or reduce coverage merely to manage costs.

Construction Quality Concerns

Cost-cutting measures in construction significantly amplify damage severity when natural hazards occur. Issues include:

- Increased property values and development in high-risk areas.
- Higher concentration of properties.
- Rising construction and repair costs.
- Greater concentration of wealth in vulnerable regions.
- More complex building systems and materials.

For example, hailstorms in Italy (2023) caused \$5.5 billion in insured losses—a European record—partly attributed to inadequate hail-resistant building standards. Similarly, construction projects in catastrophe-exposed regions face 20–30% higher insurance premiums if they lack proper risk mitigation protocols.

Banking Sector Implications

The rising costs of natural hazard insurance directly affect mortgage affordability and credit risk. Financial institutions now confront:

- **Collateral devaluation:** Properties in flood zones or wildfire corridors experience value declines of 3.6-10%.
- **Increased default risks:** Following catastrophic events, mortgage delinquencies increase by 15-30% in households in disaster-stricken areas.
- **Capital constraints:** As insurers retreat from high-risk markets (Florida and Texas spring to mind), banks must either absorb more risk or reduce lending in these areas.

The Cost-Frequency Paradox

Swiss Re projects that insured losses could double within a decade, emphasising that socioeconomic factors—not just hazard frequency—drive cost escalation. This underscores the need for models that prioritise exposure growth and construction quality alongside climatic shifts.

The disproportionate rise in costs relative to event frequency demonstrates that vulnerability and exposure have become the dominant factors in determining loss potential. This requires a fundamental reassessment of how risk is priced in household insurance and mortgage underwriting.

To address these changing cost patterns, insurance models now incorporate:

- Dynamic replacement cost estimation accounting for building material inflation.
- Density-adjusted exposure models reflecting wealth concentration.
- Time-dependent vulnerability functions addressing changing building practices.
- Granular rating territories more accurately reflect local risk variations.

Formula 21: Cost Frequency Bayesian

These adjustments typically employ Bayesian hierarchical models that integrate historical data with socioeconomic factors:

$$p(\theta|y) \propto p(y|\theta)p(\theta)$$

Where

- $p(\theta|y)$ is the posterior probability distribution - the updated probability of model parameters θ given the observed data y .
 - \propto means "proportional to" - indicating the relationship without including the normalising constant.
 - $p(y|\theta)$ represents the likelihood function - the probability of observing the data y given parameters θ .
 - $p(\theta)$ is the prior probability distribution - the initial belief about the parameters before observing new data.
 - θ (theta) represents the model parameters (e.g., vulnerability factors, exposure growth rates).
 - y represents the observed data (e.g., historical losses, claims data)
-

The Bayesian approach is very generic and something anyone looking at market data timeseries would recognise. The approach to evolving insurance premiums begins with prior knowledge, where insurers establish initial risk estimates based on historical patterns and expert judgment. When new evidence emerges captured in the likelihood function from claims data, property changes, or environmental shifts), the model updates beliefs by proportionally weighting this evidence against prior assumptions to produce a posterior probability distribution.

The crucial aspect is that this weighting isn't fixed—more reliable or robust evidence receives more significant influence in reshaping premium calculations, while uncertain data has less impact. This creates a continuously refining system where each update incorporates all previously learned information, allowing premiums to evolve gradually rather than overreacting to short-term fluctuations while still remaining responsive to genuine emerging trends in risk factors.

Broad Territorial Rating

Modern insurance risk assessment has evolved significantly from historical broad territorial rating approaches, now leveraging property-specific geospatial analysis powered by advanced machine learning techniques. This transformation combines high-resolution spatial data with sophisticated computational methods to create granular risk profiles.

Traditional rating systems grouped properties into large geographic zones with uniform risk assumptions. This "one-size-fits-all" approach often resulted in similar properties facing different premiums based solely on arbitrary boundaries like zip codes or counties.

Contemporary models now analyse individual properties using:

- **Sub-metre resolution DEMs for precise floodwater flow modelling:** High-definition topographic maps capturing elevation changes with 10-30cm accuracy, enabling detailed water movement simulation and identifying micro-topography features that affect flood vulnerability.
- **LiDAR-derived building characteristics:** Creates 3D digital representations that automatically extract critical structural features influencing wind, hail, and snow damage potential without requiring physical inspection.
- **Satellite-based vegetation analysis:** - Monitors vegetation density, type, and moisture content; calculates defensible space metrics and fuel load profiles; and tracks seasonal changes in wildfire exposure through regular imaging.
- **Hydrological flow models** Simulates water accumulation and movement during rainfall events, accounting for drainage capacity, soil saturation, and infrastructure while providing time-based flood progression scenarios instead of static zone designations.

The predictive engine combines ensemble machine learning methods that handle complex spatial relationships. The model combines outputs from multiple decision trees (typically 100-1000) into a single prediction.

Each decision tree processes specific spatial features like elevation gradients and building shapes. The final prediction is calculated by adding together all these individual tree outputs, with each tree's contribution weighted based on its predictive accuracy. This approach allows the system to capture complex interactions between different geographical and structural factors, creating a comprehensive risk assessment that's far more nuanced than traditional rating methods.

Formula 22: Predictive Tree

$$\hat{f}(x) = \sum_{m=1}^M \gamma_m h_m(x)$$

Where:

- $h_m(x)$ = Individual decision tree processing spatial features (e.g., elevation gradients, building shapes).
 - γ_m = Weighting factor optimising tree contributions.
 - M = Total number of trees in the ensemble.
-

This architecture is an advancement at modelling non-linear interactions between hundreds of geospatial variables:

- **Slope derivatives from DEMs:** Calculates elevation change rates across distances, identifies flow acceleration zones where water velocity increases, and detects natural drainage channels that may not appear on standard maps.
- **Hydraulic connectivity indices:** Measure water travel efficiency between points, identify properties vulnerable to distant flooding impacts, and account for underground drainage systems and culverts that affect water flow patterns.
- **Urban heat island effects:** Quantifies temperature amplification in densely built environments, correlates building density and surface materials with temperature patterns and identifies areas with increased risk of heat-related infrastructure stress based on spatial arrangement.

Behavioral Risk Components

Beyond physical hazard characteristics, effective household insurance modelling must account for human behavioural factors that influence risk outcomes:

- Risk perception variables derived from demographic data.
- Mitigation implementation rates by region and property type.
- Claim filing propensity based on past insurance history.

These factors significantly impact loss experience beyond what physical models alone would predict, necessitating the integration of socioeconomic data into comprehensive risk models.

Insurance modelling thus represents a multidisciplinary integration of actuarial science, hazard modelling, geospatial analysis, and behavioural economics. The resulting framework provides the foundation for premium calculation, coverage design, and capital allocation across insurance markets.

Common Data Models as Drivers of Resilience and Risk Precision

Developing more granular and frequent data points plus more sophisticated flood and insurance models begins the critical journey toward linking lower premiums or more effective insurance with tangible incentives to improve resilience.

While representing significant progress, the true potential lies in standardising data structures through Common Data Models (CDMs) so that all stakeholders operate from a shared definition of reality. Open source initiatives such as FINOS for banking and OASIS for insurance are ideally positioned to facilitate agreement among members on standardised property definitions, creating unified fields to drive calculations for resilience measurements.

CDMs establish the causal foundations for next-generation insurance modelling by creating directional relationships between physical attributes, hazard exposures, and loss outcomes. Unlike traditional models that rely heavily on correlative patterns in historical data, CDM-driven approaches enable a causal understanding of risk factors. This shift from correlation to causation significantly enhances premium calculations' predictive power and explainability, especially in expanding property urbanisation, where historical patterns may become less reliable indicators of future risk.

The transformative power of CDMs lies in their ability to connect disparate data sources through standardised ontologies. For flood risk specifically, CDMs can integrate building characteristics, terrain data, infrastructure performance metrics, and meteorological measurements into a unified framework that preserves the semantic relationships between these elements. This structured approach enables more sophisticated causal modelling, where the impact of specific interventions (such as improved drainage systems or flood barriers) can be quantified with greater precision than previously possible.

Machine learning algorithms applied to CDM-structured data can identify subtle, causal relationships that traditional actuarial methods might miss. Unlike conventional AI applications' "black box" concerns. This combination of advanced analytics with standardised data models represents a fundamental advancement in risk assessment methodology, allowing insurers to develop more targeted and effective pricing strategies that precisely reflect individual risk profiles.

"CDM-based machine learning maintains explainability by operating on semantically consistent data structures with clearly defined relationships." - Johnny Mattimore, MKM.

The revolutionary potential emerges when envisioning local governments maintaining detailed information on taxpayer properties, infrastructure, and terrain characteristics within CDM frameworks, enabling optimised resilience spending that directly translates to lower insurance premiums. This creates a transformative economic model where reduced insurance costs could offset increases in local taxation, leaving property owners financially advantaged and building community resilience.

"CDMs facilitate real-time data integration in ways that ad hoc approaches cannot match. By establishing standardised interfaces for continuous data streams, insurers can incorporate emerging risk information without the required extensive data transformation." - David Kelly, MKM.

This standardisation dramatically reduces the technical complexity and potential errors associated with dynamic premium adjustments, making real-time risk-based pricing more feasible and reliable.

"Perhaps most importantly, CDMs enable transparent risk pricing by clarifying the causal factors influencing premium calculations." - David Kelly, MKM.

When all stakeholders—insurers, policyholders, regulators, and municipal governments—share the same structured understanding of risk drivers, incentives naturally align toward meaningful risk reduction. Property owners gain clear visibility into how specific resilience improvements translate to premium reductions, while insurers can more

confidently offer discounts for measures with demonstrable causal impact on expected losses.

Standardised data frameworks built on CDMs allow insurers to precisely calibrate premium discounts against specific resilience improvements, creating transparent economic incentives aligned with individual and collective risk reduction. This represents the foundation for a virtuous cycle where data standardisation drives improved risk assessment, enabling more targeted resilience investments and ultimately reducing losses and insurance costs across entire communities.

As insurance models evolve from traditional actuarial approaches toward sophisticated machine learning applications, the quality and structure of underlying data become increasingly critical. CDMs provide the essential foundation for this evolution, ensuring that advances in analytical techniques translate to genuine improvements in risk assessment rather than merely amplifying existing data biases or correlative patterns.

By establishing causal clarity with standardised data structures, CDMs are revolutionising property insurance by shifting their role from passive risk compensation to active risk prevention. By standardising data structures among stakeholders - banks, insurers, regulators, and risk modellers - CDMs create causal clarity that reveals root risk drivers while enabling predictive analytics at scale.

Chapter 8 - Mortgage Risk Assessment

Mortgage valuation stands at the intersection of real estate appraisal and banking risk assessment. While traditional property valuation focuses on determining the market or lending value of the underlying collateral, modern mortgage valuation extends beyond this to assess the banking instrument itself. This comprehensive approach accounts for the complex interplay between borrower characteristics, property values, and macroeconomic conditions, determining the probability and impact of default events.

"Creating a reliable mark-to-market framework for mortgages is not merely an accounting exercise—it's fundamental to banking stability. A robust valuation methodology provides the transparency that allows capital to flow efficiently even during periods of stress." – William C. Dudley, Former President of the Federal Reserve Bank of New York.

Three primary banking modelling frameworks have emerged as the foundation for sophisticated mortgage valuation: discounted cash flow analysis, option-adjusted spread methodology, and stochastic house price models. These approaches share common mathematical underpinnings but differ in conceptualising and quantifying risk. The discounted cash flow method projects expected payment streams adjusted for default probabilities and recovery values. Option-adjusted spread methodology treats default as an embedded option exercised by the borrower under certain conditions. Stochastic models directly simulate the evolution of key variables like house prices and interest rates to derive default probabilities under various scenarios.

Capital markets technology has recently introduced an alternative approach that draws from credit default swap (CDS) pricing methodologies. This framework centres on affordability ratios as primary drivers of default probability, establishing a direct relationship between payment burden and credit spread. By implementing sophisticated piecewise functions that capture the exponential increase in default risk as affordability deteriorates, this approach offers particular advantages when

modelling the impacts of physical hazards such as flooding. Incorporating insurance costs, income disruptions, and property value adjustments following such events becomes more straightforward within this framework.

A comprehensive mortgage valuation synthesises these approaches, recognising that default risk stems from two fundamental sources:

- **Inability to pay:** income-related default. Imagine shocks to someone's income-making cash flow, such as redundancy and divorce.
- **Strategic default:** equity-related. Imagine a mortgage holder leaving the keys on the kitchen table and leaving since the property is now uninsurable.

By modelling both the probability of default (PD) and loss-given default (LGD) components dynamically throughout the mortgage's life, these models provide a robust foundation for pricing, risk management, and capital adequacy assessment in residential mortgage lending.

Discounted Cash Flow with PD/LGD Integration

The discounted cash flow approach begins with the fundamental premise that a mortgage's value equals the present value of its expected future cash flows. However, these cash flows must be adjusted for the possibility of default and the resulting losses.

We first calculate the scheduled payment for each period according to the mortgage terms. We then adjust this expected cash flow by subtracting the expected loss from default, which equals the probability of default in that period multiplied by the loss given default and the outstanding balance.

Formula 23: Discount Cash Flow

The formula for the expected cash flow in period t becomes:

$$ECF(t) = \text{Scheduled Payment}(t) \times (1 - PD(t)) - \text{Outstanding Balance}(t) \times PD(t) \times LGD(t)$$

Where:

- $ECF(t)$ is the expected cash flow in period t.
- $\text{Scheduled Payment}(t)$ is the contractual mortgage payment in period t.

- $PD(t)$ is the probability of default in period t (expressed as a decimal between 0 and 1).
- Outstanding Balance(t) is the remaining principal balance in period t .
- $LGD(t)$ is the loss given default in period t (expressed as a decimal between 0 and 1).

These expected cash flows are then discounted using an appropriate risk-adjusted rate to arrive at the present value. The mortgage value is the sum of these discounted expected cash flows across all periods:

$$\text{Mortgage Value} = \Sigma ECF(t) / (1 + r)^t$$

Where:

- Σ represents the summation across all time periods t .
 - r is the risk-adjusted discount rate (expressed as a decimal).
 - t is the time period (typically in months or years)
-

The PD component evolves dynamically as the loan ages. Initially, it reflects the borrower's credit profile and loan characteristics at origination. Over time, it adjusts based on the changing loan-to-value ratio as property values fluctuate and the loan amortises. Macroeconomic factors like unemployment and interest rate environments further modify the default probability.

The LGD component captures the severity of loss when default occurs. It starts with the expected property value at the time of default but accounts for foreclosure costs, property maintenance, and the time value of money during the recovery process. In high-value markets with quick property turnover, LGD might be relatively low. Conversely, in declining markets with extended foreclosure timelines, LGD can be substantial.

Option-Adjusted Spread (OAS) with PD/LGD Factors

The OAS approach recognises that mortgages contain embedded options—the borrower's option to prepay and, implicitly, the option to default. We simulate numerous possible future interest rates and house price scenarios to value these options properly.

We determine whether default occurs in each simulated path based on the evolving conditions. Default typically happens when the borrower crosses certain thresholds - when home equity becomes sufficiently

negative (strategic default) or payments become unaffordable relative to income (payment default).

Formula 24: Mortgage PD and LGD

For path i at time t , the probability of default is modeled as a function of the simulated loan-to-value ratio and other triggers:

$$PD(i,t) = f(LTV(i,t), Payment_Burden(i,t), Unemployment(i,t))$$

Where:

- $PD(i,t)$ is the probability of default for simulation path i at time t .
- $f()$ represents a functional relationship (often logistic regression).
- $LTV(i,t)$ is the loan-to-value ratio for path i at time t .
- $Payment_Burden(i,t)$ is the ratio of mortgage payment to income for path i at time t .
- $Unemployment(i,t)$ is the simulated unemployment rate for path i at time t .

Similarly, the loss severity depends on the simulated house price and market conditions:

$$LGD(i,t) = g(House_Price(i,t), Market_Liquidity(i,t), Foreclosure_Timeline)$$

Where:

- $LGD(i,t)$ is the loss given default for simulation path i at time t .
- $g()$ represents a functional relationship.
- $House_Price(i,t)$ is the simulated house price for path i at time t .
- $Market_Liquidity(i,t)$ is a measure of real estate market liquidity for path i at time t .
- $Foreclosure_Timeline$ is the expected time to complete foreclosure proceedings.

We calculate the present value of cash flows along each path, adjusting for these path-specific default probabilities and loss severities. The mortgage value equals the average present value across all simulated paths:

$$\text{Mortgage Value} = (1/N) \times \sum_i \sum_t CF(i,t) / (1 + r(i,t))^t$$

Where:

- N is the total number of simulation paths.
 - \sum_i represents summation across all simulation paths i.
 - \sum_t represents summation across all time periods t.
 - $CF(i,t)$ is the cash flow for simulation path i at time t.
 - $r(i,t)$ is the discount rate for simulation path i at time t
-

The "option-adjusted spread" is the additional yield required above risk-free rates to make the present value equal to the market price. This spread captures all risks, including default and prepayment risks and their correlation.

Stochastic House Price and Interest Rate Model

The stochastic model takes a more interest rate swaption approach, directly modelling house prices and interest rates as correlated random processes.

Formula 25: Stochastic House Price

House prices might follow geometric Brownian motion, while interest rates could follow a mean-reverting process:

$$dH/H = \mu_h dt + \sigma_h dW_h \quad dr = \alpha(\theta - r)dt + \sigma_r dW_r$$

Where

- Wiener processes W_h and W_r .
 - correlate ρ .
-

Default occurs when the property value falls below a threshold relative to the loan balance or when other borrower-specific triggers are activated. This framework allows us to calculate the probability that the borrower will cross the default boundary during any given period.

The mortgage value is determined by solving a partial differential equation that accounts for the stochastic processes and boundary conditions created by default possibilities. The partial differential equation above incorporates the probability of default and the loss severity upon default, which vary with the property value at default time.

Rather than generating explicit cash flows, this approach directly models the evolving mortgage value over time. The present value emerges naturally from the solution to the PDE, which accounts for all future possibilities in a continuous-time framework.

The advantage of this approach is its ability to capture complex interactions between interest rates and house prices, especially for high LTV mortgages where default risk is susceptible to property value changes.

Synthesis of Approaches

“The most sophisticated valuation models combine elements of all three approaches. They might use the stochastic model’s combined elements of all three approaches. They might use the stochastic model’s mathematical rigour to generate realistic scenarios, apply option-theoretic principles to model borrower behaviour, and structure the results in an intuitive discounted cash flow framework for banking analysis.”-David Kelly, MKM.

The key insight is that mortgage value depends fundamentally on the timing and magnitude of cash flows, which are significantly shaped by default probabilities and loss severities. These PD and LGD factors are dynamic, responding to evolving loan characteristics, borrower circumstances, property values, and macroeconomic conditions.

Given the timing of cashflows - the interest payment, the pre-payment and the proceeds of sale given default, a Monte Carlo simulator is appropriate to value the mortgage:-

- **Path Generation:** The model simulates hundreds of potential future economic scenarios by generating random interest rate paths, projecting house price movements, and modelling key economic indicators like unemployment rates and market liquidity conditions that affect borrower behaviour.
- **Default Risk Modelling:** For each simulated path and period, the model calculates the probability of default by evaluating the borrower's

incentives and constraints, including the current loan-to-value ratio, payment affordability relative to income, local unemployment conditions, and individual credit quality factors.

- **Loss Severity Assessment:** When defaults occur in the simulation, the model estimates the percentage of principal that would be lost by analysing the projected property value, current market liquidity conditions, foreclosure timelines specific to the location, legal recovery processes, and any mortgage insurance protection.
- **Cash Flow Projection:** The model calculates the expected cash flow at each point by combining scheduled payments with adjustments for both prepayment behaviour (driven by refinancing opportunities) and default outcomes (including recovery amounts), creating a probability-weighted cash flow projection for each path.

- **Discounted Valuation:** The final mortgage valuation is derived by discounting each path's adjusted cash flows using path-specific discount rates (incorporating the option-adjusted spread), then averaging these present values across all simulated paths to arrive at a comprehensive risk-adjusted price.

The Capital Markets Approach Using CDS

This approach borrows from credit default swap (CDS) pricing methodology, where the mortgage is valued based on a credit spread that reflects the borrower's default probability. Unlike traditional mortgage models that rely heavily on LTV ratios and house prices, this model centres on affordability as the primary driver of default risk.

The core of this model is the relationship between the affordability ratio (payment burden relative to income) and credit spread. The code implements a sophisticated piecewise function where:

- The credit spread increases linearly with the ratio at low affordability ratios (below 15%).
- There's a steeper linear relationship in the moderate range (15-30%).
- At high affordability ratios (above 30%), the credit spread grows exponentially, reflecting rapidly increasing default risk when housing costs consume too much income.

The loan pricing function then uses these credit spreads to derive hazard rates (instantaneous default probabilities), determining survival probabilities over time. The loan value is calculated as the present value of expected cash flows (weighted by survival probabilities) minus the present

value of expected losses (outstanding balance minus recovery value, weighted by default probabilities).

Integration with PD/LGD Framework

This model elegantly handles both PD and LGD components:

- **Probability of Default (PD):** The credit spread derived from the affordability ratio, is a proxy for default probability converted to a hazard rate using $h = 1 - \exp(-\text{spread}/2)$ for semi-annual periods. Survival probabilities are calculated recursively, with each period's survival probability depending on previous periods.
- **Loss Given Default (LGD):** The model applies a haircut parameter (h) to the property value to determine recovery. LGD is calculated as $\max(0, \text{outstanding_balance} - h * \text{property_value})$. This dynamically changes as the loan amortises, reflecting the evolving equity position

Advantages of Physical Risk Assessment

The key advantage of this approach, especially when considering physical risks like flooding, is that it can easily incorporate additional costs that affect affordability. When a flood event occurs, it can impact a mortgage in several ways:

- Increased insurance costs directly affect the affordability ratio.
- Temporary income disruption during recovery.
- Property value declines due to realised flood risk.
- Dynamically adjusting the insurance parameter (I) after a flood event.
- Modifying the affordability ratio to reflect temporary income impacts.
- Adjusting the haircut parameter (h) to reflect reduced recovery value in flood-prone areas.

The loan pricing function generates several valuable outputs for risk assessment:

- Credit spreads over time.
- Hazard rates (instantaneous default probabilities).
- Survival probabilities.
- Expected loss given default at each point in time.
- Present value of the loan accounting for default risk.

The scenario analysis capability allows you to assess how different LTV ratios and income levels affect loan value, which is particularly useful for stress testing under various physical risk scenarios.

This capital markets approach offers a more direct link between affordability shocks (which can result from physical risk events) and credit quality than traditional mortgage models. Focusing on affordability rather than just collateral value better captures the dual impact of physical risks on property values and household finances.

This approach could be a valuable bridge between traditional mortgage valuation techniques and forward-looking physical risk assessment.

Chapter 9 - Physical Risk Swaps

In the evolving landscape of climate finance, one critical challenge has remained unaddressed: the fundamental disconnect between long-term lending and short-term insurance coverage. This mismatch creates significant exposure for banking institutions, particularly mortgage lenders, who face escalating physical risks to their collateral assets over periods far beyond standard insurance timeframes.

"Physical Risk Swaps (PRS) represent a pioneering banking innovation designed to bridge this gap. By adapting established capital markets instruments to address climate-related physical risks—particularly flooding—PRS offers a mechanism for banking institutions to effectively transfer and manage these exposures over multi-year horizons that align with their asset portfolios."

Johnny Mattimore, MKM

Inception of Physical Risk Swaps by MKM Research Labs

The MKM Exec have been working on Sustainable Finance and Physical risk since partnering with BNP Paribas to lead the OS-Climate "Physical Risk and Resilience" project.

- In November 2023 the Phrase Physical Risk Swap was coined by the Author and Johnny Mattimore in a LinkedIn Post "[Repricing Mortgages using Flood Insurance Risk Transfer](#)" as a way of describing a capital markets equivalent of a parametric insurance. We call this a "Physical Risk" Swap, starting with the most ubiquitous hazard: flood risk. The natural solution was to create a new swap, building on the history of risk transfer. Eureka moment brought together the following principles:-

- Parametric Insurance as a tradeable security
- Analogy with CDS product design and market development

- Analogy with ISDA legal definition of a CDS
- Back to the origins of Tranches, MSB and CMOs
- In October 2024, Johnny Mattimore and the Author presented [“Physical Risk Swap for Flood Risk”](#) as a concept to the FINOS NY gathering. The driving design was to create an acceptable multi-year flood risk insurance protection in capital markets.
- In February 2025, Johnny Mattimore and the Author presented a [Beta version of the Physical Risk Swap CDM](#) alongside property, mortgage and household insurance to the FINOS CDM Working Group.
- The background to the development of a Physical Risk Swap came from the below diagram designed in a LinkedIn Article, [“How to build the next generation financial platforms to fully integrate sustainable finance data”](#), providing the position in the top right corner of the evolution to a risk transfer mechanism.

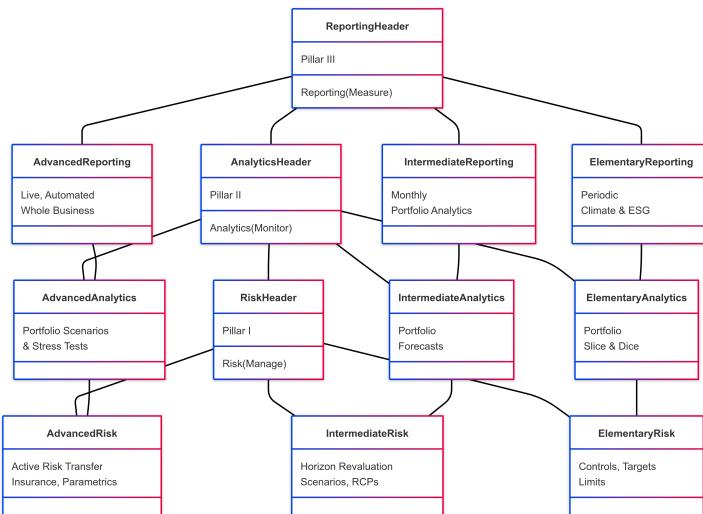


Figure 6: MKM Reporting Framework

Mismatched time horizons between lending and insurance

The banking system faces a critical structural challenge in managing physical risks: a fundamental disconnect between long-term lending and

short-term insurance coverage. Property owners typically borrow for 10-30 years but can only secure property insurance on annual renewal terms. This temporal mismatch creates significant exposure for banking institutions, particularly mortgage lenders, who face escalating physical risks to their collateral assets over periods far beyond standard insurance timeframes.

This disconnect represents more than just an operational inconvenience—it constitutes a growing systemic risk as continued urbanisation intensifies the financial impact of physical hazards. Banking institutions are required to reconcile the long-duration nature of their asset portfolios with the inherently short-term, annually repriced nature of traditional insurance protection. Without addressing this fundamental mismatch, lenders face increasingly uncertain risk profiles that complicate capital planning, threaten portfolio stability, and potentially restrict credit availability in vulnerable regions.

Physical Risk Swaps represent a pioneering banking innovation designed to bridge this gap. By adapting established capital markets instruments to address climate-related physical risks—mainly flooding—PRS offers a mechanism for banking institutions to effectively transfer and manage these exposures over multi-year horizons that align with their asset portfolios. Adapting the proven credit derivatives framework, this approach transforms unpredictable physical hazards into quantifiable, transferable banking risks.

The business problem that Physical Risk Swaps address encompasses three distinct but interconnected risk layers that emerge from the insurance-lending duration mismatch:

- **Property Owners:** Face significant uncertainty about future insurability and premium costs throughout their mortgage term. Annual insurance renewals create vulnerability to sudden premium increases or, more seriously, outright insurance unavailability following major events or shifts in risk assessment. This uncertainty affects property valuations and introduces potential covenant breaches if properties become uninsurable during the loan term.
- **Mortgage Lenders:** Face direct exposure when borrowers cannot secure ongoing insurance coverage. As properties lose insurance protection, lenders bear an increased risk that physical events will impair collateral values without corresponding insurance recovery. This risk compounds at the portfolio level, where geographic concentrations can simultaneously create correlated exposures across multiple properties, facing insurance challenges.
- **Banking System:** Faces systemic correlation risk as weather patterns shift against further urbanisation. Entire geographic regions may

experience concurrent insurance market disruptions, with insurers either raising premiums to unaffordable levels or withdrawing coverage entirely from higher-risk locations.

“This portfolio-wide concentration risk threatens synchronised collateral devaluation across significant portions of mortgage books.” - David Kelly, MKM.

These risks are embedded in existing loan portfolios but remain largely unquantified in traditional risk management systems. Banking institutions must integrate physical risk assessment into their core banking architecture, measuring, monitoring, and managing these exposures with the same rigour applied to market or credit risk. They need mechanisms to translate specialised physical risk data (like flood scores from proprietary vendor models) into quantitative valuations and risk metrics for portfolio management.

For many lenders, flood risk represents the dominant physical hazard threatening collateral impairment. Yet, they lack both methodologies to translate flood risk scores into banking terms and established market mechanisms for transferring this risk. This capability gap prevents active management of embedded physical risks and complicates capital allocation to climate-vulnerable regions, potentially accelerating regional economic disparities through credit availability constraints.

Example Framework for Flood Risk Transfer: Data

The solution framework begins with a sophisticated data transformation pathway that converts environmental monitoring into financially quantifiable events. This process follows a methodical sequence: weather pattern time series generate precipitation measurements, which feed hydrological models producing water flow predictions, ultimately leading to real-time flood level monitoring at strategically positioned gauges. Each gauge is standardised within the Common Domain Model (CDM), creating definitive, reliable trigger points for banking contracts.

The quantification of physical risk begins with comprehensive environmental data integration. Detailed digital terrain models capture the topographical features influencing water movement, while flow distribution models translate precipitation into volumetric water flows across landscapes. These models feed into catchment area response systems that predict how water accumulates throughout watershed regions. The culmination produces flood event distribution sets—statistical representations of flooding scenarios with associated probabilities—transforming unpredictable natural phenomena into quantifiable risk metrics suitable for banking modelling.

Insurance Products: Established Risk Transfer Mechanisms

The architecture of Physical Risk Swaps draws substantial inspiration from parametric insurance contracts rather than traditional catastrophe insurance. Parametric insurance triggers payouts based on objective measurements from specific instruments—such as a water level gauge on the wall of a building—without requiring an assessment of actual damages. These contracts establish direct relationships between physical measurements and banking settlements, creating transparent and rapid trigger mechanisms.

A PRS extends this concept by incorporating multiple gauges throughout a catchment area, typically managed and maintained by the relevant environmental authority, ensuring data integrity and operational reliability. This multi-gauge approach creates a comprehensive risk transfer mechanism covering broad geographic areas while maintaining objective, verifiable and legally enforceable contract triggers.

By establishing tiered trigger levels corresponding to different water depths, parametric structures create granular risk segmentation that allows precise pricing and targeted protection. These established insurance constructs provide proven mechanisms for quantifying and transferring specific segments of the risk distribution without the complications of loss adjustment processes, offering valuable precedent for capital markets approaches that require rapid and unambiguous financial settlement mechanisms. Their performance during flood events demonstrates how objective physical measurements can effectively drive banking settlements in complex environmental scenarios.

Banking Risk Transfer and Physical Risk Swaps

PRS's transformative innovation lies in extending established risk transfer concepts into banking and capital markets infrastructure. While insurance mechanisms have long addressed physical risks, they remain largely separate from banking capital frameworks and lack the standardisation, liquidity, and regulatory integration necessary for effective banking risk management.

PRS bridge this gap by aligning with existing data sets from environmental monitoring and catastrophe modelling while structuring the banking instrument according to ISDA standards—the same legal and operational framework that underpins \$100s of trillions in interest rate and credit derivatives. The valuation models employ methodologies consistent with the Oasis Loss Modelling Framework, creating compatibility with existing insurance analytics while introducing the precision required for mark-to-market valuation.

“Most significantly, PRS development includes regulatory capital methodologies aligned with Basel standards, enabling banks to receive appropriate capital relief for transferred physical risks and creating powerful incentives for market adoption.” - David Kelly, MKM.

Leveraging Existing Monitoring Infrastructure

Implementing PRS benefits from substantial existing infrastructure that can be leveraged without requiring extensive new monitoring systems. For example, the Thames River Basin in the UK represents an ideal development environment with its highly populated, economically critical areas and persistent flood risks affecting assets from residential to industrial properties.

The UK already maintains a robust flood monitoring system managed by government agencies. A network of monitoring stations along the Thames River provides real-time water level data at regular intervals. This existing infrastructure offers a solid foundation for flood risk assessment models and creates reliable triggering mechanisms for banking contracts. Rather than building new monitoring capabilities, PRS implementation can focus on integrating and standardising these data streams for banking applications.

Leverage Experts

Strategic partnerships with leading experts in flood modelling allow the adaptation of sophisticated terrain mapping, flood simulators, and event set development to meet the specific needs of mortgage lending risk assessment. These

collaborations enable access to decades of accumulated expertise in hydrological modelling without requiring banking institutions to develop specialised climate science capabilities internally.

The computational demands of PRS valuation require partnerships providing access to high-performance computing capabilities. While gauge-level monitoring data is available in real-time, the complex task of portfolio revaluation requires substantial computing resources, particularly for stress testing across multiple climate scenarios. Cloud-based computing platforms with dedicated physical risk modelling capabilities allow market participants to access sophisticated analytics without individual infrastructure investments.

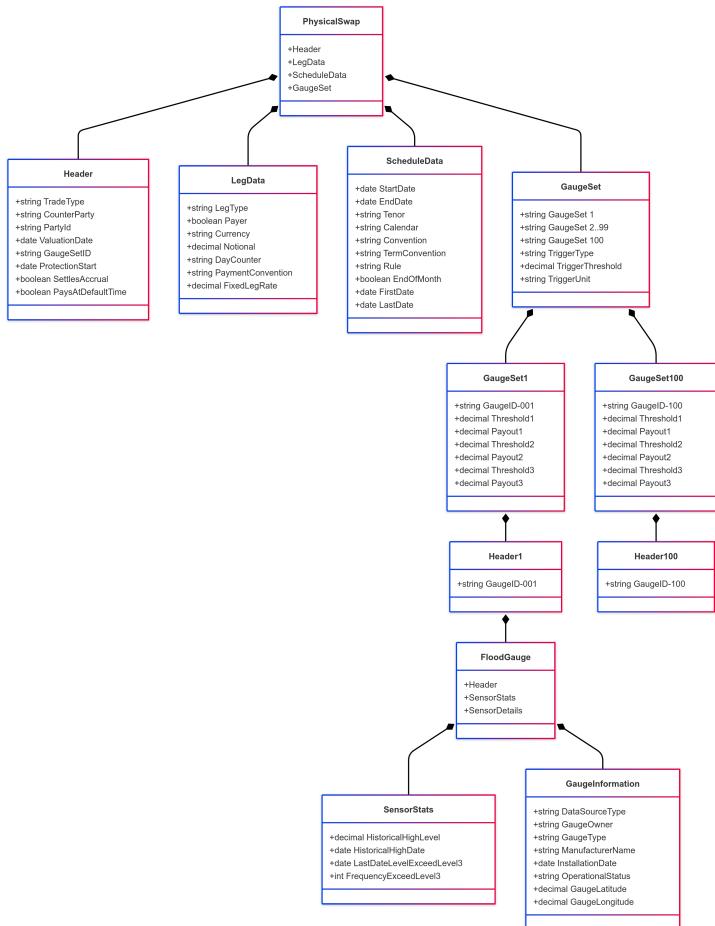


Figure 7: Physical Risk Swap Detail

For effective implementation, governance frameworks must include cross-disciplinary representation where meteorologists, hydrologists, atmospheric, and weather scientists work alongside derivatives structurers and trading strategists to refine risk assessment methodologies.

Insurance underwriters with decades of catastrophe experience advise on threshold calibration and recovery value estimation. This cross-disciplinary integration accelerates product development while strengthening model validation through diverse perspectives.

Structuring a PRS - The Thames River

To illustrate how a Physical Risk Swap functions in practice, we can examine a specific implementation for the UK Thames Flood area. This concrete example demonstrates how the theoretical framework translates into an operational banking instrument.

Without losing generality, we can define a specific swap structure that can be extended to any market segment seeking to manage flood risk. A five-year contract balances significant risk transfer and market liquidity for the Thames catchment area. This instrument operates on a familiar fixed-floating exchange principle: one party pays a fixed rate over the contract period, while the other party makes contingent payments based on realised flood levels (the floating rate component) across monitored points in the Thames basin.

The fixed-rate calculation derives from sophisticated flood modelling, effectively pricing "wholesale" flood risk insurance over the contract term. In a sample implementation, this might be set at £1,150 on a contract notional of £1 million, representing an annualised premium of 11.5 basis points (0.115%). This rate reflects the expected loss based on historical flood data, climate projections, and defined trigger thresholds.

The floating rate reference is determined by measured flood depths at specific gauges throughout the catchment area. These measurements provide transparent, objective triggers for contract settlement. For instance, if five monitored locations experience flooding with depths of 0.1, 0.5, 0.7, 1.0, and 0.2 meters, these might correspond to modelled losses per square meter of £143, £1,371, £1,886, £4,000, and £486 respectively. The floating leg payout would total £7,886 on a £1 million contract.

Settlement mechanics follow established derivative practices, with contingent payments after verified flood events. The contract includes precise definitions of measurement protocols, confirmation procedures for gauge readings, and calculation methodologies for converting flood depths to banking settlements. These standardised terms ensure consistent treatment across market participants and reduce settlement disputes.

This structure provides a precise, objective mechanism for transferring flood risk from mortgage lenders to counterparties better positioned to bear it, with pricing that reflects the underlying physical risk characteristics of the specific catchment area.

Addressing the mortgage lender's dilemma

PRS directly address what can be termed "the mortgage lender's dilemma"—the mismatch between long-duration lending and short-term insurance solutions that creates growing banking vulnerability as climate risks intensify.

From a risk and capital perspective, a PRS provides value through the high correlation between the swap's mark-to-market in the event of a realised flood and the corresponding mortgage impairment (rise in Probability of Default and Loss Given Default) of the portfolio. This correlation creates natural hedge effectiveness for capital relief purposes and portfolio risk management.

However, it's important to note that this correlation is not perfect. Unless gauges are installed on each property—a theoretically possible but practically unrealistic solution—the PRS holder maintains some basis risk due to the difference between gauge locations and actual property locations within the portfolio. While properties and reference gauges share the same catchment area and hydrological system, localised variations in topography, drainage infrastructure, and property-specific flood resilience measures can create a divergence between gauge readings and actual property damage.

This basis risk requires careful calibration of the PRS structure, selecting strategically positioned gauges that maximise correlation with the protected mortgage portfolio while maintaining the objectivity and transparency of environmental authority-managed monitoring stations. Despite this imperfection, the high correlation provides substantial risk mitigation compared to the alternative of unhedged exposure to long-term flood risk.

The primary beneficiaries of instruments like the PRS for UK Thames Flood are mortgage lenders, who face a unique challenge in risk management. These institutions typically hold mortgage assets on their books for up to 20 years, yet the properties securing these loans are often insured annually.

The problem is compounded by correlation risk across mortgage portfolios. Lenders face the continual risk that a concentration of properties will simultaneously fail to secure annual insurance renewals if insurers either raise premiums to unaffordable levels or withdraw from specific geographic areas altogether. This portfolio-wide vulnerability can materialise suddenly, leaving lenders with significant unhedged exposures precisely when those risks are being recognised as most severe.

PRS enable lenders to purchase specific flood protection that aligns with the duration of their mortgage assets. By entering into these contracts, lenders can realise multiple strategic benefits:

- **Reduce Exposure:** To long-term flood risks by transferring this uncertainty to counterparties with appropriate risk appetite and diversification.
- **Improve Risk Accuracy:** Their risk management strategy matches hedge durations to underlying asset exposures.
- **Stabilise Mortgage:** Rates in flood-prone areas by removing the insurance uncertainty premium embedded in current pricing.
- **Enhance portfolio risk profiles:** With associated prudential and regulatory benefits, including potential capital relief.

This approach transforms uncertain, difficult-to-quantify physical risks into defined banking exposures that can be actively managed alongside traditional market and credit risks. Rather than facing unpredictable insurance market disruptions that could simultaneously affect substantial portions of their portfolios, lenders can secure multi-year protection at known costs, significantly improving risk planning and capital allocation.

From CDS to PRS, from credit to hazard curve

Credit Default Swaps (CDS) have established a well-understood pricing framework that can be adapted for Physical Risk Swaps (PRS). To understand the transition from CDS to PRS pricing, we must first examine the fundamental components of CDS valuation.

In a standard CDS contract, pricing involves calculating the present value of two legs:

- **Premium Leg:** The fixed periodic payments made by the protection buyer
- **Protection Leg:** The contingent payment made by the protection seller upon a credit event

The value of a CDS is established when the present values (PV) of these two legs are equal:

$$\text{PV Premium Leg} = \text{PV Protection Leg}$$

The premium leg payments are weighted by the probability that the reference entity survives until each payment date, while the likelihood of default weights the protection leg. This creates a mathematical relationship:

Formula 26: CDS Survival

$$\sum_{(i=1)^n} S \cdot \Delta t_i \cdot DF_i \cdot PS_i = (1 - R) \cdot \sum_{(i=1)^n} (PS_{(i-1)} - PS_i) \cdot DF_i$$

Where:

- S is the CDS spread.
 - Δt_i is the time interval.
 - DF_i is the discount factor.
 - PS_i is the probability of survival to time i.
 - R is the recovery rate.
-

The par spread in credit default swaps (CDS) represents the equilibrium rate where the present value of protection payments (premium leg) equals the present value of potential default payouts (contingent leg). This ensures the CDS contract has zero net value at inception, making it a "fair" pricing benchmark for credit risk.

Event Definition and Triggering Mechanism

Credit Default Swaps (CDS) and Physical Risk Swaps (PRS) share structural similarities but differ fundamentally in their triggering mechanisms and applications. Both instruments allow for customisable risk transfer, but PRS leverage measurable physical parameters—such as flood depths—to define payout conditions, offering granular flexibility akin to credit risk tranching in CDS markets.

CDS: Graded Credit Events and Repricing

CDS contracts are structured to cover credit events across a spectrum of credit grades (AA to B), with payouts triggered by predefined conditions like bankruptcy, failure to pay, or restructuring[1][4][8]. Key features include:

- **Dynamic pricing:** CDS spreads adjust intraday based on market perceptions of creditworthiness, enabling trading at intermediate risk levels (e.g., deteriorating credit from AA to BBB) before a default (D event)

- **Recovery rates:** Post-default payouts account for the residual value of debt, typically determined via cash settlement (e.g., 100% minus post-default bond value)[5][8].

PRS: Flood Thresholds and Adaptive Design

PRS replicate this flexibility using flood-specific triggers:

- **Customizable thresholds:** Contracts can define payout triggers for specific flood depths (e.g., 0.5m-1.0m) or catastrophic levels (e.g., >2.0m), analogous to tranches CDS.
-
- **Objective triggers:** Flood events are determined by remote sensors or gauge data, eliminating subjective judgments required in credit events like restructuring.

The critical adaptation involves replacing default probability curves with physical hazard curves. For a flood PRS, the hazard curve represents the probability of a flood event of a certain magnitude occurring within a specific time frame where meteorological data expressing "1-in-X-year" events are converted to annual exceedance probabilities. Each monitoring gauge has its hazard curve based on historical data and climate projections.

Formula 27: PRS Survival

The hazard rate at time t represents the conditional probability of the physical event occurring in a small interval Δt , given that it hasn't occurred before time t :

$$h(t) = \lim_{\Delta t \rightarrow 0} (P(t < T \leq t + \Delta t | T > t)) / (\Delta t)$$

Where T is the time of the physical event occurrence.

For PRS, survival probability represents the probability that the trigger event (e.g., flood level exceeding the threshold) has not occurred by time t :

$$PS(t) = e^{-\int_0^t h(s) ds}$$

The protection leg is then calculated using:

$$PV \text{ Protection Leg} = (1 - R) \cdot \sum_{i=1}^n (PS_{i-1} - PS_i) \cdot DF_i$$

R represents the "recovery value" in the physical risk context—the residual value after the physical event occurs.

The pricing mechanics have additional considerations regarding the gauge-based triggering system for a physical risk swap. A flood PRS typically references specific water level gauges in a catchment area. Each gauge has:

- A predetermined threshold level that constitutes a trigger event
- Historical data used to develop the hazard curve
- Continuous monitoring and automatic publishing capabilities for event determination

The probability space for pricing is built by:

- Converting traditional flood return periods (e.g., "1-in-100 year flood") into annual exceedance probabilities.
- Developing a continuous hazard rate function based on these probabilities.
- Incorporating weather projections to adjust forward-looking probabilities

For example, a gauge with historical data showing:

- 1-in-10-year flood level = 3.5 meters
- 1-in-50-year flood level = 4.2 meters
- 1-in-100-year flood level = 4.8 meters

These data points fit a curve representing the continuous relationship between flood levels and exceedance probabilities.

Many flood PRS contracts reference multiple gauges within a catchment area, requiring:

- Correlation modelling between different gauge locations.
- Settlement conditions based on one or more gauges exceeding thresholds.
- Weighted payout structures based on the severity and location of flooding

The correlation modelling typically employs copula functions to represent joint probability distributions between gauges, accounting that flood events often affect multiple locations simultaneously.

Form

The spread for a PRS represents the cost of protection and is calculated similarly to a CDS spread:

$$S = ((1 - R) \cdot \sum_{(i=1)^n} (PS_{-i} \cdot DF_i) / (\sum_{(i=1)^n} \Delta t_i \cdot DF_i \cdot PS_i)$$

Where

- S is the PRS spread.
 - Δt_i is the time interval.
 - DF_i is the discount factor.
 - PS_i is the probability of no flood trigger to time i.
 - R is the recovery rate akin to a fixed payout.
-

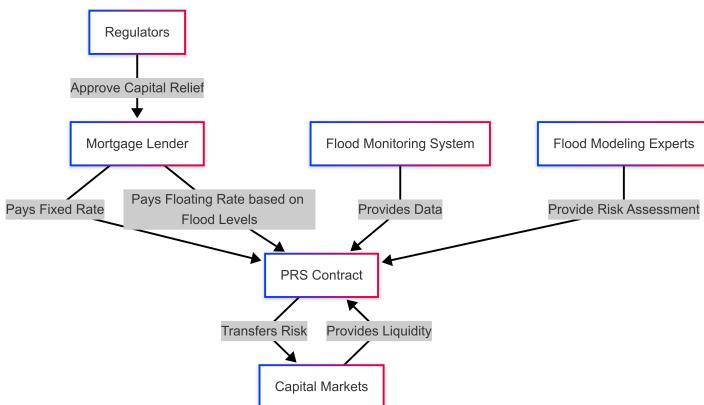


Figure 8: PRS and Flood Participants

According to market conventions for standardised contracts, this spread is converted to an upfront payment plus a fixed coupon.

Implementation Framework

The journey from theoretical pricing models to a functioning Physical Risk Swap market requires careful orchestration across regulatory, technological, and banking dimensions. Rather than developing an entirely new market infrastructure, the PRS ecosystem leverages existing frameworks while introducing targeted innovations to address the unique characteristics of physical risk.

The ISDA Master Agreement and Market Practices

The cornerstone of derivatives market standardisation for decades—provides the contractual foundation for Physical Risk Swaps. This established legal framework reduces documentation uncertainty and creates immediate familiarity for market participants who have executed trillions in interest rate and credit derivative trades under the same structure. The agreement's modular design accommodates new definitions and supplements for physical risk events while maintaining consistency with existing market conventions for trade confirmation, settlement, and dispute resolution.

This legal standardisation is complemented by the FINOS Common Domain Model (CDM), which offers critical data standardisation across the trade lifecycle. The CDM creates a unified representation of trade events, product definitions, and process models, enabling seamless interoperability between counterparties, clearing houses, and regulatory reporting systems.

By incorporating physical risk parameters and gauge-level monitoring data into this existing framework, PRS transactions benefit from day-one compatibility with established market infrastructure. The CDM implementation ensures that gauge readings, threshold triggers, and settlement calculations maintain consistency across diverse market participants and technology platforms.

Most significant for market adoption, Physical Risk Swaps integrate directly with the FINOS Digital Regulatory Reporting (DRR) framework from inception. This pre-established compliance approach eliminates the traditional regulatory uncertainty often accompanying banking innovation.

Because PRS instruments are designed to be DRR-compliant from their first transaction, they avoid the regulatory ambiguity and implementation

lags that have historically slowed the adoption of new derivative products. Banks can confidently deploy these instruments in their regulatory treatment, focusing implementation efforts on risk modelling rather than compliance infrastructure.

Market liquidity development follows a deliberate progression, beginning with bilateral transactions between major banking institutions and strategically expanding to include specialised market makers. Central counterparty clearing adoption occurs early in this evolution, reducing counterparty risk concerns and establishing standardised margin methodologies.

Introducing PRS indices aggregating exposures across multiple catchment areas creates broader hedging vehicles that attract investment funds and insurance-linked securities investors, further deepening liquidity pools. Standard tenors aligned with mortgage and infrastructure financing timeframes (typically 5, 10, and 30 years) create natural hedging opportunities for banks with concentrated real estate exposures.

The computational demands of PRS valuation necessitate shared infrastructure investment across market participants. While gauge-level monitoring data is available in real-time through existing environmental agencies, the complex task of portfolio revaluation requires substantial computing resources, particularly for stress testing across multiple weather pattern timeseries scenarios. Cloud-based computing platforms with dedicated physical risk modelling capabilities allow market participants to access sophisticated analytics without individual infrastructure investments. Significant banking technology providers are extending their derivatives analytics suites to incorporate physical risk modelling components, creating integrated valuation and risk management solutions.

Regulatory capital recognition

The critical incentive for banks to adopt Basel standards develops through targeted implementation rather than new regulatory frameworks. Physical Risk Swaps are structured to qualify under existing capital relief provisions for risk transfer instruments, with careful documentation of how the contracts mitigate specific physical risks within loan portfolios. The transparency of gauge-based triggers, combined with robust modelling of hazard probabilities, creates the verification mechanisms necessary for regulators to acknowledge genuine risk transfer.

“This approach to regulatory acceptance allows banks to realise immediate capital benefits while maintaining prudent risk management standards.” - Johnny Mattimore, MKM

As Physical Risk Swaps move from initial transactions to established market presence, they create dynamic feedback loops that enhance hazard resilience throughout the banking system. Price discovery in the PRS market provides powerful signals about evolving physical risks, informing capital allocation decisions and infrastructure investment priorities. The availability of multi-year hedging instruments encourages lending in flood-vulnerable regions with appropriate risk management rather than wholesale withdrawal of capital.

Instrument pricing

Effective physical risk management through swap instruments requires sophisticated pricing methodologies, regulatory alignment, and substantial computational infrastructure.

Advanced Monte Carlo simulations are employed to price these contracts accurately and assess risk, generating comprehensive sets of potential flood scenarios, including low-probability, high-impact events. Critically, the valuation methodologies use frameworks already approved by regulators, incorporating existing bank model libraries while adding specialised physical risk components.

“This alignment with established regulatory approaches supports the case for capital relief when lenders utilise PRS instruments.” - David Kelly, MKM

The pricing models incorporate several key components, including terrain models that capture topographical features influencing flood behaviours. These flow distribution models translate precipitation scenarios into water movement patterns and historical flood data calibrated with forward-looking weather projections. Integrating physical science with banking modelling requires cross-disciplinary expertise but produces robust pricing that accurately reflects underlying hazard probabilities.

From a regulatory perspective, Physical Risk Swaps benefit from integrating existing frameworks rather than requiring new regulatory structures. The ISDA Master Agreement provides the contractual foundation, while the FINOS Common Domain Model ensures data standardisation across the trade lifecycle. Most significantly, PRS instruments are designed to be Digital Regulatory Reporting (DRR) compliant from inception, eliminating the regulatory uncertainty often accompanying banking innovation.

This pre-established compliance approach allows banks to confidently deploy these instruments in their regulatory treatment, focusing

implementation efforts on risk modelling rather than compliance infrastructure. PRS contracts are structured to qualify under existing Basel capital relief provisions for risk transfer instruments, with careful documentation of how they mitigate specific physical risks within loan portfolios.

The computational requirements for effective risk management are substantial, particularly for portfolio-level analysis across multiple climate scenarios. While gauge-level monitoring data is available in real-time through existing environmental agencies, intraday revaluation demands significant computing resources. Cloud-based shared services offer an efficient solution, allowing market participants to access sophisticated analytics without individual infrastructure investments.

Banking technology providers are extending their derivatives analytics suites to incorporate physical risk modelling components, creating integrated valuation and risk management solutions. These platforms enable stress testing across various weather projections, sensitivity analysis for different threshold configurations, and portfolio optimisation incorporating physical risk dimensions alongside traditional risk factors.

Market Impact - beyond risk transfer

While mortgage lenders are the primary beneficiaries, Physical Risk Swaps have broader implications for property markets, weather resilience, and banking stability.

The introduction of PRS creates significant positive externalities beyond direct risk transfer benefits. By reducing long-term risk exposure, lenders may be more willing to maintain mortgage availability in flood-prone areas, improving property market liquidity in regions that might otherwise face credit contraction. This sustained credit provision supports property values and enables ongoing investment in areas requiring adaptation rather than abandonment.

PRS pricing provides valuable market-based information about perceived flood risks, creating transparent signals that inform urban planning, infrastructure investment, and adaptation priorities. Regional pricing differentials will highlight areas where flood mitigation infrastructure could deliver the most significant economic benefits as the market develops, potentially influencing public investment decisions and encouraging preventative measures rather than post-disaster recovery.

The availability of PRS may incentivise community-wide resilience improvements by creating banking rewards for risk reduction. As lenders better manage their risks, they may offer more favourable terms to borrowers implementing flood resilience measures, creating economic

incentives for adaptation investments. This dynamic could accelerate the adoption of building-level and community-scale flood protection measures through market mechanisms rather than regulatory mandates.

Price discovery in the PRS market will provide ongoing feedback about evolving physical risks, informing capital allocation decisions beyond the mortgage sector. Infrastructure investors, commercial real estate developers, and municipal bond issuers will gain additional market-based insights into physical risk pricing that can guide investment strategies and project designs.

The broader banking system benefits from improved risk transparency and reduced correlation of flood-related losses. By distributing concentrated physical risks more widely through capital markets, PRS minimises the potential for synchronised institutional stress during significant weather events. This distribution mechanism supports banking stability objectives by preventing the concentration of weather risks in institutions with limited capacity to bear them.

Market size estimates and broader implications.

The potential scale of the Physical Risk Swap market is substantial, with estimates suggesting it could develop into a multi-trillion dollar notional market as significant capital markets firms integrate these instruments into their offerings. This projected growth follows patterns observed in other successful derivatives markets, such as interest rate and credit default swaps, which achieved enormous scale by addressing fundamental risk management needs.

The initial focus on flood risk in significant river basins represents only the beginning of potential applications. The same framework can be extended to other physical hazards, including windstorms, wildfires, droughts, and heat stress—each with its observable metrics, historical data sets, and forward-looking projections. Each hazard type requires specific adaptation of triggering mechanisms and loss models, but the fundamental structure of transferring physical risk through structured derivatives remains consistent.

Geographic expansion beyond initial implementations in developed markets with sophisticated monitoring infrastructure presents challenges and opportunities. Emerging markets often face heightened physical hazard risks with more limited historical data and monitoring capabilities. Strategic investment in gauge networks and remote sensing technologies could enable PRS implementation in these regions, potentially supported by development finance institutions seeking to enhance hazard resilience in vulnerable countries.

Product evolution will likely follow patterns seen in other derivatives markets, with initial standardised contracts gradually complemented by more tailored structures addressing specific risk profiles. Index products aggregating exposures across multiple catchment areas will develop to create broader hedging vehicles. At the same time, tranched structures may emerge to segment risk by severity levels, allowing more precise risk transfer matching particular institutional risk appetite requirements.

Integration with climate adaptation finance represents an auspicious direction, where Physical Risk Swaps become components of blended finance solutions. Public sector risk absorption for extreme tail events could be combined with private market capacity for more frequent risks, creating layered protection that maximises the efficiency of limited public resources while maintaining broad market participation.

As physical risk and climate disclosure requirements intensify globally, PRS instruments offer a potential mechanism for organisations to demonstrate active management of identified risks rather than simply reporting exposures.

“This capability could enhance reporting compliance while improving resilience—transforming disclosure from a compliance exercise into strategic risk management.”-

Johnny Mattimore, MKM.

The role of PRS in creating banking resilience

Physical Risk Swaps represent a natural extension of credit derivatives technology to manage natural hazard-related banking risks. The significant advantage of building PRS on established CDS technology, legal frameworks, and data constructs is the reduction in adoption efforts for banks. Banking institutions can leverage their existing systems, knowledge, and operational processes developed for credit derivatives, requiring the integration of physical risk data primarily and modelling rather than wholesale new infrastructure.

By adapting the proven CDS pricing framework to incorporate physical hazard probabilities, these instruments offer a promising approach to bridge the gap between long-term flood risks and existing banking market structures. Implementing the FINOS Common Domain Model provides the critical data standardisation layer that enables interoperability across market participants. At the same time, compliance with digital regulatory reporting ensures smooth regulatory integration from inception.

As we navigate an increasingly uncertain future of urbanisation, instruments like Physical Risk Swaps can play a vital role in creating more

resilient banking systems. By aligning risk management timeframes with the duration of asset exposures, PRS enables more accurate pricing of hazard risks and supports informed decision-making by all market participants. This balanced approach facilitates necessary adaptation measures while maintaining economic vitality in communities navigating flood risk transition challenges.

The successful implementation of PRS requires collaboration between atmospheric and weather scientists, hazard assessment experts, banking engineers, and market participants to develop robust pricing models that accurately capture the complex dynamics of physical risks while maintaining the standardisation and liquidity necessary for market adoption. With the foundation of ISDA documentation and the FINOS CDM, the market can focus on the unique aspects of physical hazard risk modelling rather than reinventing established market infrastructure

Physical Risk Swaps represent not merely a technical, banking innovation but a critical step toward a banking system that properly accounts for and manages physical risks – ultimately supporting both economic stability and climate adaptation in the decades ahead.

Developing Market Liquidity for Physical Risk Swaps

Developing a liquid market for Physical Risk Swaps (PRS) represents a critical innovation in risk management for physical assets. Following the successful evolution of the Credit Default Swap (CDS) market, we can establish a roadmap for creating a robust, efficient marketplace for transferring and managing physical risk exposures. This section explores the key elements necessary to foster liquidity in this evolving market.

Learning from the CDS Market Evolution

The Credit Default Swap market provides valuable lessons for developing PRS liquidity. From its inception in the early 1990s to its peak of \$61.2 trillion in notional value by 2007, the CDS market demonstrated how standardisation, institutional participation, and regulatory frameworks can create a functioning risk transfer mechanism. While the market subsequently contracted to \$9.4 trillion by 2017, the core infrastructure and liquidity mechanisms remained intact, particularly for index products.

The core design of a PRS is a natural evolution of a parametric insurance contract. The only real difference is the legal wrapper, which is

governed by two different regulatory frameworks. The cashflows, however, are identical.

Unlike other innovations in finance, a PRS is incredibly easy to explain. Its simplicity and intuitive nature make it an elegant solution that will drive uptake.

The fundamental design of PRS mirrors the CDS template, replacing credit events with physical risk triggers. This familiar structure allows market participants to leverage existing knowledge and systems:

- **Trigger Mechanism:** Rather than a default event, a PRS is activated when a specified physical measurement (water level, temperature, wind speed, etc.) reaches a predefined alert threshold.
- **Premium Structure:** Regular premium payments from protection buyer to protection seller until maturity or trigger event.
- **Settlement Process:** Predefined payout mechanisms based on the severity of the physical risk event.
- **Term Structure:** Standardized maturities that align with typical investment horizons for physical assets

Unlike the organic and sometimes fragmented development of early CDS documentation, PRS benefits from institutional standardisation from inception that materially reduces documentation risk:

- **ISDA Framework:** Developing standard PRS documentation under the International Swaps and Derivatives Association framework.
- **Common Data Model:** Implementation within the ISDA/FINOS Common Domain Model (CDM) program ensures consistent data structure and event management.
- **Unified Definitions:** Clear, industry-accepted definitions of physical risk events, measurements, and thresholds.
- **Protocol Adoption:** Streamlined adoption through ISDA protocols rather than bilateral negotiation.

Market liquidity requires both willing buyers and sellers with genuine economic interests:

- **Protection Buyers:** Asset owners face increasing physical risks due to climate change and other environmental factors.

- **Protection Sellers:** Institutions with the capacity to diversify and manage physical risks across geographic and categorical dimensions.
- **Insurance Gap:** Insufficient traditional insurance capacity creates a natural demand for alternative risk transfer mechanisms.
- **Portfolio Optimization:** Institutional investors seeking to optimise risk-return profiles of asset portfolios containing physical risk exposure

The substantial and growing gap between physical risk exposure and available insurance capacity creates natural market tension that PRS can address.

Regulatory considerations represent both a challenge and an opportunity for PRS market development:

- **Dodd-Frank Precedent:** Following the implementation pathway created for swaps under the Dodd-Frank Act.
- **Capital Treatment:** Working with regulators to establish appropriate capital requirements that recognise the risk-mitigating nature of properly structured PRS.
- **Reporting Requirements:** Building transparent reporting mechanisms that facilitate regulatory oversight without imposing excessive burden.
- **Cross-Border Considerations:** Developing consistent international regulatory treatment to avoid fragmentation.

Early and constructive engagement with regulators will be essential for market development.

The involvement of several major banking institutions is critical for initial liquidity, requiring a coordinated launch (and ongoing commitment to make markets) by major banks that create an immediate two-way market.

Early central clearing provides significant advantages for market development:

- **Counterparty Risk Reduction:** Mitigating bilateral credit risk concerns through a central counterparty.
- **Netting Efficiency:** Improving capital efficiency through multilateral netting.

- **Standardised Margining:** Creating predictable and efficient collateral requirements.
- **Position Portability:** Facilitating the transfer of positions between counterparties.
- **Risk Aggregation:** Developing methodologies with the banks to aggregate and tier physical risks across diverse geographies and asset classes.
- **Basis Risk Management:** Addressing the inherent basis risk between index products and specific asset exposures.

Liquidity Development Trajectory

Based on the CDS market experience, we can anticipate the following liquidity development pattern for PRS:

- **Initial Phase:** Concentrated bilateral trading among major dealers using standardised documentation.
- **Index Development:** Creation of diversified indices covering major risk categories and geographies.
- **Buy-Side Adoption:** Gradual expansion to asset managers, pension funds, and other institutional investors.
- **Market Bifurcation:** Emergence of highly liquid index products alongside more specialised single-name PRS.
- **Electronic Execution:** Migration from voice trading to electronic platforms as volume increases.
- **Market Maturity:** Development of secondary market trading, curve trading, and basis trading strategies

By following the successful model of the CDS market while implementing early standardisation and central clearing, PRS can evolve into an efficient risk transfer mechanism. The natural economic demand for physical risk protection, combined with careful attention to structural and regulatory considerations, creates a favourable environment for market development.

While challenges remain, particularly regarding scale, basis risk modelling and capital treatment of hedge effectiveness, the pathway to PRS liquidity is clear and achievable.

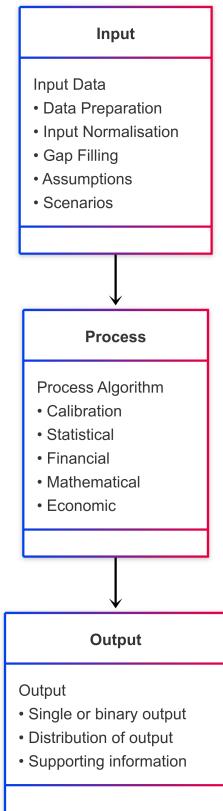
Chapter 10 - Model Risk Governance

As banking institutions increasingly incorporate flood risk into their operations, the governance of flood risk models has emerged as a critical challenge. This chapter explores how traditional model risk governance frameworks must evolve to address the unique complexities of flood risk modelling. We examine the "model stack problem" in flood risk assessment, propose governance solutions, and outline collaborative industry approaches to standardisation.

"Integrating current fragmented vendor offering of flood risk information and analytics into banking operations represents a fundamental challenge for model risk governance." - David Kelly, MKM.

As banks face increasing pressure to incorporate natural hazard considerations into their risk assessments and capital allocations, they must ensure their flood modelling approaches meet rigorous governance standards. Key challenges include:

- **Regulatory evolution:** Regulators are expanding existing prudential frameworks to address climate-related risks, making model governance urgently needed when flood models inform capital requirements or lending criteria.
- **Complex decision-making:** Banks face a web of requirements when implementing flood risk models, particularly for scenarios like assessing flood risk in mortgage portfolios, which can lead to adjusting lending criteria.
- **Societal implications:** Poor model governance creates significant societal risks. For example, inappropriate vendor flood scores used to pre-screen UK mortgage applications profoundly impact property markets and individual homeowners.



Model - Quick Reminder

“Model Risk can lead to financial loss, poor business and strategic decision making, or damage to a bank’s reputation. Model Risk should be managed like other types of risk. Firms should identify the sources of risk and assess the magnitude. A guiding principle for managing model risk is ‘effective challenge’ of models.”

SR 11-7 Guidance on Model Risk Management - Fed / OCC

Figure 9: Components of a Model

A model is thus a quantitative method that includes complex manipulations and expert judgements that applies algorithmic engines to process input data into quantitative estimates for decision-making. A model has to have these three components otherwise it is a process.

Model Governance for Climate Risks

While this book is focused on the present hazard of flood risk, it is worth nothing how the banking industry takes model governance in new areas of risk management of which it is worth adopting as the industry introduces PRS.

The Commodity Futures Trading Commission (CFTC) issued a groundbreaking report in September 2020 titled Managing Climate Risk in the U.S. Financial System, authored by the Climate-Related Market Risk Subcommittee chaired by Bob Litterman. This report marked the first comprehensive effort by a U.S. financial regulator to address climate-related risks in financial markets.

Litterman, a former Goldman Sachs executive known for his work in risk management, led the diverse committee of 34 experts from financial institutions, academia, and non-governmental organizations.

The report emphasizes that effective model governance is essential for managing climate-related financial risks. It acknowledges the unique challenges of climate risk modelling, including data limitations, long time horizons, and deep uncertainty about future scenarios.

The committee recommends that financial institutions establish robust governance frameworks for climate risk models, with clear board-level oversight and integration into existing enterprise risk management systems. According to the report, models must balance standardization with analytical flexibility, while avoiding false precision in their outputs.

These governance principles aim to ensure that climate risk assessment becomes a mainstream component of financial risk management across the U.S. financial system.

Risk Management and Oversight

- **Board Oversight:** Establish clearly defined oversight responsibilities for climate risk at the board of directors level.
- **Integration with Existing Frameworks:** Address climate-related risks through existing risk management frameworks with appropriate governance by corporate management.
- **Enterprise Risk Management:** Integrate climate risks into Enterprise Risk Management (ERM) and Own Risk Solvency Assessments (ORSA) processes.

Climate Risk Data and Modelling Standards

- Standardised Classification Systems: Develop classification systems for physical and transition risks across asset classes and sectors.
- Capacity Building: Implement training and education programs to build climate risk management capabilities.

Scenario Analysis and Stress Testing

- Consistent Scenarios: For assessment, use a consistent and common set of broad climate risk scenarios, guidelines, and assumptions.
- **Analytical Discretion:** Allow firms to decide how they perform scenario analysis, balancing standardisation with flexibility.
- **Continuous Improvement:** Establish mechanisms for ongoing refinement as science, data, tools, and conditions evolve.
- **In-house Capabilities:** Develop internal capabilities to analyse climate scenarios, understand underlying assumptions, and recognise limitations.
- **Avoid False Precision:** Recognize that quantitative results should be treated as illustrative rather than precise in climate scenario analysis.
- **Proportional Analysis:** Ensure the scope, depth, and complexity of analyses are proportionate to the materiality of the impact measured.

Flood Risk Model Stack Problem

At the heart of flood risk modelling lies the "model stack problem." As established in Chapters 3-5, modern flood risk assessment involves the output of one model becoming the input of another, creating a chain of dependencies that amplify uncertainties throughout the modelling process. For flood risk specifically, this stack includes:

- **Weather pattern models:** Generating precipitation forecasts and historical distributions.
- **Hydrological models:** Translating precipitation into river flows and surface water accumulation.
- **Hydraulic models:** Determining how water moves across landscapes and through built environments.

- **Impact models:** Assessing damage to properties based on flood characteristics.
- **Banking impact models:** Translating physical damage into economic and banking consequences

“From a model governance perspective, the practice of hydrologic and hydraulic modelling embodies a certain level of paradox: We use deterministic equations to describe inherently uncertain processes.

We apply simplifications to represent indescribably complex systems. We seek precision in a domain where perfect prediction is impossible. Apart from that it should work perfectly fine so long as we understand and measure uncertainty” - David Kelly, MKM.

Key Characteristics of the Flood Model Stack:

- **Cascading dependencies:** Physical process models feed into statistical models, which feed into banking impact models, creating multiple layers of uncertainty. For example, the previous example of the Thames flood model requires output from weather pattern time series to generate precipitation measurements, which feed hydrological models producing water flow predictions, ultimately leading to real-time flood level monitoring.
- **Simplification of complex processes:** Fundamental physical processes like cloud formation and precipitation patterns directly influencing flooding are often simplified due to computational limitations. As noted earlier, these simplifications significantly impact flood risk assessment.
- **Uncertainty amplification:** Each model in the stack contains its uncertainties, which compound as they move through the chain. Our analysis of flood models showed how errors in flow estimation could be magnified for in-depth calculations, creating significant variations in estimated property damage.
- **Cross-disciplinary integration:** Models originating from different scientific traditions (climate science, hydrology, economics, finance)

with varying validation standards are linked, creating governance challenges across domains.

- **Temporal mismatch:** Models operating at different time scales (from hours for weather to decades for climate) must be integrated coherently, as highlighted in our examination of the mortgage-insurance duration mismatch as described in the previous chapter on PRS.

Model Risk and Complexity

The interconnected nature of model stacks in banking creates unique governance challenges that traditional frameworks struggle to address. Banks operate in increasingly complex regulatory and technical environments where models serve multiple purposes—from credit decisions to capital planning—and inform high-stakes business decisions. These models often form intricate webs of dependencies similar to the flood model stack but with additional banking and regulatory complexities.

Model Dependency Management

Banking institutions typically maintain hundreds or thousands of models across departments, creating a tangled web of dependencies. These models form hierarchical structures, sometimes called a “string-of-pearls”, where there is high dependency. Consider the following:-

- Market data models generate time series of market simulations that collectively create market data.
- Mid-level models link the market data to risk factors used for instrument pricing and hedging.
- Risk models aggregate the risk factors of the originated risk and the hedges, leaving a basis required for P&L attribution and VaR.
- Stress-based models that define tail events are used in the aggregate for capital allocation.

Each layer inherits assumptions and uncertainties from previous layers, often without transparent documentation of these inherited limitations. For example, a mortgage pricing model might depend on interest rate projections, which rely on economic growth forecasts—each introducing its uncertainty profile that compounds through the chain.

Model Boundary Ambiguity

In complex banking environments, the boundaries between models become increasingly blurred. This creates several governance challenges:

- **Unclear ownership:** When multiple teams contribute to different components of an interconnected model system, responsibility for overall performance becomes diffuse.
- **Validation gaps:** Components at the boundaries between models often receive less scrutiny during validation exercises.
- **Inconsistent assumptions:** Adjacent models frequently operate under contradictory assumptions about the same underlying process

One particularly problematic boundary exists where statistical models interface with expert judgment. In flood risk assessment, as in many banking contexts, human experts often adjust model outputs based on experience, creating a "grey area" where governance controls are challenging to implement consistently.

Uncertainty propagation

As highlighted in the flood model stack, uncertainty propagation becomes increasingly important—and challenging—in complex model environments. Banks face particular difficulties with:

- Identifying the sources of model uncertainty.
- Quantifying how uncertainty grows through model chains.
- Attributing performance issues to specific components within the stack.

While techniques like Monte Carlo simulation can help quantify overall uncertainty, attributing this uncertainty to specific model components remains extremely difficult. This creates significant challenges for model improvement efforts as teams struggle to identify which components would benefit most from refinement. Perhaps most concerning from a governance perspective are the emergent properties that arise from complex model interactions.

"Banking history is replete with examples where seemingly well-governed individual models created systemic risks."

David Kelly, MKM.

The following red flags worth calling out in any complex model instances tend to be a combination:

- Models that individually capture risk factors accurately may collectively miss correlation effects during stress periods.
- Feedback loops between market and risk models can amplify volatility during crises.
- Optimisation of individual models may lead to collectively suboptimal outcomes.

These emergent risks often elude traditional governance frameworks focused on individual model performance rather than systemic interactions.

Complex model environment

Banks face significant operational challenges in governing complex model environments such as:

- **Specialised expertise:** Effective governance requires rare combinations of domain expertise (e.g., climate science, banking, risk management).
- **Resource constraints:** Validation resources are typically limited, requiring difficult prioritisation decisions.
- **Technical infrastructure:** Managing model dependencies requires sophisticated technical infrastructure that many banks lack

Uneven Coverage

The result is often uneven governance coverage, with sophisticated models receiving intensive scrutiny while simpler but potentially more consequential models receive less attention.

Comprehensive documentation becomes extraordinarily more challenging as model complexity increases:

- Underlying assumptions may span multiple domains and require specialised knowledge to understand.
- Model interdependencies create complex webs that resist simple documentation.
- Staff turnover leads to knowledge gaps about historical model decisions.

This creates a significant "key person risk," where critical knowledge about model limitations resides with a few individuals rather than in institutional documentation.

Toward Integrated Model Governance

Addressing these governance challenges requires approaches that explicitly recognise model interdependencies and embrace the complexity inherent in modern banking environments. Key principles include:

- **System-level validation:** Complementing component-level validation with holistic assessments of model systems.
- **Consistent metadata:** Developing standardised ways to document model assumptions, limitations, and dependencies.
- **Transparent lineage:** Creating clear documentation of how data and assumptions flow through model chains.
- **Adaptive governance:** Implementing frameworks that can evolve as models and methodologies change.
- **Cross-functional oversight:** Establishing governance bodies with expertise spanning relevant domains.

The following case study on flood risk modelling illustrates how these principles can be applied in practice, highlighting successes and ongoing challenges in governing complex model environments with significant dependencies.

Case Study: Flood Risk Modelling

Our UK flood risk modelling analysis provides an example of a physical risk model stack. The process typically involves the following checklist across the model risk governance process:

- **Measuring typical flow rates:** Officers measure typical flood events in each area, noting how the land responds to rainfall. They manually define scores for pre-defined factors like how quickly water flows, how much is absorbed, and how much runs off.
- **Calibrating flow distributions:** Using the data measured in the first step, statisticians create a distribution of potential flows, including how variable a maximum can be and its frequency.

- **Estimating flood flow frequency:** The next step combines the catchment profiles with the statistical models to assess the relative increase in water flow for less frequent events than typical flow rates.
- **Calculating flood heights by location:** Modelers then take the flow estimates under different probabilities and adapt fluid dynamic simulations to measure how a rainwater burst leads to a rise in water in each location.
- **Inter-model dependency:** The compatibility of underlying assumptions across the different model types (measurement, statistical, fluid dynamics) is often inconsistent. Statistical models are sometimes shoehorned into data, creating human inconsistency.
- **Error amplification:** Final simulations depend on multiple upstream models with distinct error profiles. As shown in our Thames flood analysis, the simulation utilises output from the initial three steps, which consist of perfectly smooth and error-free distributions of future flow severity and frequency. Subsequently, it incorporates its assumptions about landscape response, thereby amplifying errors.
- **Granularity challenges:** Minor differences between two locations can lead to significantly different results, regardless of model sophistication. Chapter 5 revealed how all those affected by floods point out how some properties avoid being affected while those nearby are inundated.
- **Underlying assumptions:** Clear documentation of assumptions and construction of climate variable distributions (e.g., tails of the distribution of precipitation at a catchment area).
- **Data source lineage:** Transparent information about data sources (public and vendor), measurement techniques, curation processes, quality checks, and update frequency. This is particularly important for flood models that rely on high-resolution LiDAR data and real-time gauge measurements.
- **Input data derivation:** Clear explanation of how vendor models derive their inputs, especially for crucial parameters like catchment response rates and infiltration capacities.
- **Portfolio aggregation:** Documentation of how portfolios of assets are aggregated by location, which is essential for understanding concentration risk in flood-prone areas.
- **Model limitations:** Description of known model weaknesses, output uncertainties, and limitations of appropriate use, particularly regarding the granularity issues identified in our flood model analysis.

- **Model use:** Define and document where and how flood model outputs should be used, distinguishing between screening tools and decision-making instruments.
- **Output interpretation guidelines:** Provide clear guidance on how consumers should interpret outputs within their risk-based processes, particularly for flood scores that may appear deceptively precise.
- **Educate end users:** Ensure that risk managers, loan officers, and senior decision-makers fully understand the technical limitations of flood models and the practical constraints on their utility, especially regarding property-specific characteristics that models may not capture.
- **Validation frameworks:** Develop robust validation approaches designed explicitly for flood risk models, incorporating historical flood data and forward-looking climate projections.
- **Uncertainty protocols:** Establish methods for effectively communicating model uncertainty to decision-makers, recognising the compounding uncertainties throughout the model stack.

The Model Risk Cycle

All documents discussing model risk governance have a variation of the following diagram. This book is no exception!

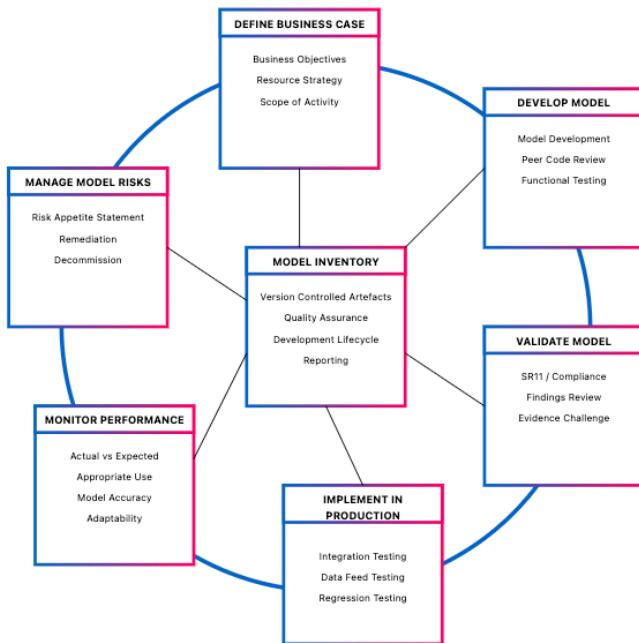


Figure 10: Model Risk Cycle

Chapter 11 - The Cutting-Edge

The fundamental physics governing flood events has been well-understood for centuries. Precipitation falls, water flows downhill, rivers rise, and floodplains inundate according to laws of conservation of mass and momentum that have remained unchanged since they were first mathematically formalised. Yet computational constraints, data availability, and the inherent challenges of modelling complex, multi-scale physical systems have limited our ability to predict specific flood events. Artificial intelligence is now transforming our approach to these challenges, not by changing the underlying physics but by revolutionising how we implement, calibrate, and apply physical models across flood-related processes.

"The future of flood risk management lies not just in better modelling but in better integration of models. When the artificial boundaries between physical science and banking analysis dissolve, we will finally see risk as nature does: as a continuous flow from cloud movements through flood time series to parametric attribution."- Johnny Mattimore, MKM

We stand today at the threshold of a profound transformation. Artificial intelligence is simultaneously revolutionising how we predict weather patterns, model fluid dynamics, and calculate insurance premiums—creating, for the first time, the technological possibility of a truly integrated approach to flood risk. This chapter explores the emerging paradigm of coherent model stacks, where time-series outputs from AI-enhanced precipitation models feed into computational fluid dynamics simulations, generating the probabilistic inputs needed for sophisticated banking and insurance pricing models.

This integration represents a fundamental reconceptualisation of flood risk as a continuum that flows—much like water—from atmospheric conditions through physical infrastructure to banking impacts.

“The coherent model stack connects previously disparate domains: neural networks that identify subtle precursors to extreme precipitation events; reinforcement learning algorithms that optimise the simulation of water movement across complex topographies; and gradient-boosting techniques that translate predicted water levels into location-specific damage estimates and ultimately into dynamic pricing models.” - David Kelly, MKM.

Yet this convergence brings unprecedented challenges. How do we manage uncertainty as it propagates through multiple model layers? What governance structures can ensure both innovation and responsible deployment? How do we maintain transparency when complex AI systems interact across domains? How do we ensure that these powerful new capabilities serve the broader societal goal of creating more resilient communities in the face of increasing flood risk?

This chapter examines the technical, organisational, and ethical aspects of the new integrated modelling paradigm. Drawing on emerging research and early implementation experiences, we outline the vision for the direction this integration is heading, along with the practical steps organisations in the physical sciences, technology, and banking sectors can take to realise its benefits. The way forward requires technical innovation and new forms of collaboration that bridge traditional disciplinary and industry boundaries.

Checklist for the steps way forward

Providing a brief outline of the steps to a new paradigm seems sensible. This is more for reference as each will be discussed in more detail:

- **Standardising weather pattern time series:** Creating consistent formats for representing weather data that can serve as inputs to hydrological models.
- **Developing gauge-level monitoring networks:** Expanding the infrastructure for real-time monitoring of water levels.
- **Building synthetic time series generators:** Creating AI systems that generate plausible future scenarios based on climate projections.

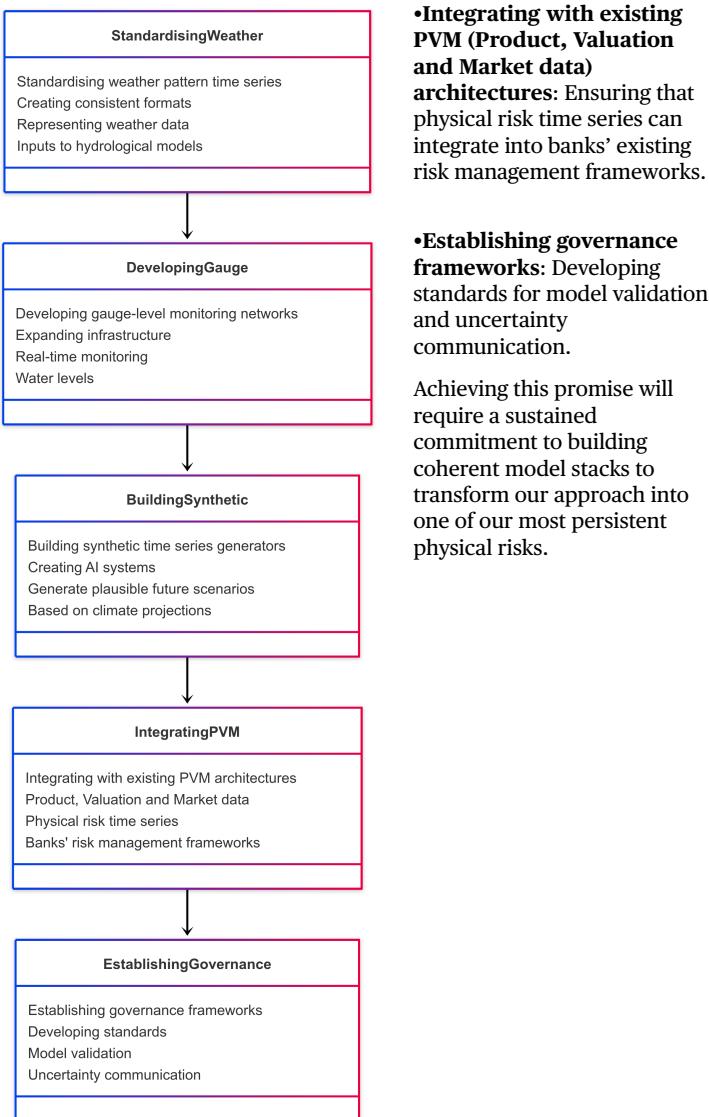


Figure 11: Stages for PRS Development

“The future of physical risk management lies not in isolated models but in integrated time series frameworks that span the entire journey from cloud formation to valuing physical risk swaps and parametric insurance. It also requires a banking ecosystem agreement to do all this consistently and transparently. No bank can do it alone.”

Johnny Mattimore, MKM.

Breaking Down Traditional Barriers

The historical approach to flood risk modelling has been characterised by disciplinary and institutional silos that mirror the organisational structures of academia, government, and industry rather than the physical reality of flooding events. Meteorologists develop precipitation forecasts, hydrologists model river flows, engineers assess infrastructure vulnerabilities, and banking analysts price risk—each group working with different data formats, time horizons, spatial resolutions, and uncertainty frameworks.

Information flows between these domains have typically been manual, episodic, and fragmented, with critical context and uncertainty measures often stripped away at each handoff point. This fragmentation introduces numerous opportunities for misinterpretation, data distortion, and compounding inaccuracies as insights move across disciplinary boundaries.

The translation between specialised technical languages and methodological approaches frequently results in important nuances being overlooked or simplified beyond recognition, leading to analyses that inadequately capture the interconnected nature of flood risks and potentially underestimate systemic vulnerabilities.

This fragmentation has persisted despite widespread recognition of its limitations. Academic papers spanning decades have called for more integrated approaches, and catastrophe modelling firms have made significant progress connecting physical and banking models. Yet the practical challenge of creating genuinely coherent model stacks has remained daunting, particularly for flood risk, which demands high spatial and temporal resolution across multiple physical domains before approaching banking considerations.

Recent advances in computational capacity and machine learning techniques have finally made a new paradigm possible—where time series output from physical models directly feed banking models

through standardised interfaces, with uncertainty appropriately characterised and propagated throughout the system. This shift represents a technical improvement and a fundamental reconceptualisation of understanding and managing flood risk.

Neural Networks and Precipitation Forecasting and NVIDIA

Traditional numerical weather prediction relies on solving partial differential equations representing atmospheric physics across a discretised grid. These models are computationally intensive, necessitating supercomputers to generate forecasts with reasonable lead times. Even with substantial computational resources, spatial resolution remains limited, typically to grid cells of 10-25 kilometres—far too coarse to capture the localised precipitation patterns that often drive flood events.

Neural network approaches like NVIDIA's FourCastNet have transformed this landscape. By learning directly from historical weather data rather than solving equations from first principles, these models can generate global weather forecasts at 25km resolution up to 500 times faster than traditional numerical methods. The speed advantage enables ensemble prediction—running multiple time series with slightly different initial conditions—that better characterises forecast uncertainty. For flood risk modelling, this means more lead time for warnings and the ability to generate many more scenarios for risk assessment.

Perhaps more importantly, machine learning approaches excel at downscaling—generating high-resolution local forecasts from coarse global predictions. NVIDIA's CorrDiff model exemplifies this capability, employing diffusion modelling techniques to refine coarse 25km forecasts to a 2km resolution.

At this scale, orographic effects (changes to air flow when the topography of the land forces air upward), urban heat islands, and other local factors significantly influence precipitation patterns.

“The improvement highlighted by NVIDIA's FourCastNet is not merely cosmetic; it fundamentally alters the information available for subsequent hydrological modelling.” - David Kelly, MKM.

Far from me to mention fractal geometry or the Butterfly Effect, but a summer thunderstorm that appears as a modest area-averaged rainfall in a 25km model cell might manifest as an intense cloudburst over a specific watershed in the downscaled 2km forecast—the difference between a managed event and a flash flood.

These AI weather models learn from observations and the output of traditional physics-based models, effectively distilling decades of meteorological science into neural network weights. The result is not a replacement for physical understanding but rather a computationally efficient implementation that preserves the essential behaviour while enabling previously impossible applications.

Flood Propagation Modelling

Once precipitation reaches the ground, predicting its movement through the landscape requires modelling surface runoff, river channel dynamics, and potentially complex interactions with urban drainage systems or coastal storm surges. Traditional hydrologic and hydraulic models solve simplified versions of the Navier-Stokes equations—the fundamental equations of fluid dynamics—but face significant computational constraints.

Two-dimensional hydraulic models that simulate water movement across a floodplain typically require hours or days of computation for a single scenario covering a modest geographic area. This computational burden has historically forced a tradeoff: either model small areas at high resolution or larger areas at lower resolution, but rarely both simultaneously. AI approaches enable a breakthrough in this domain through several parallel innovations.

Physics-informed neural networks (PINNs) represent one promising direction. These models combine traditional neural network architectures with explicit constraints derived from the laws of physics. For flood modelling, the neural network learns to predict water movement patterns from historical data while being constrained to conserve mass and momentum. The approach combines the computational efficiency of neural networks with the physical realism of traditional models. Early implementations have demonstrated the ability to simulate flood propagation 100 times faster than conventional methods while maintaining comparable accuracy.

Graph neural networks

Graph neural networks offer another powerful approach, particularly well-suited to river networks. These models can learn the complex relationships between upstream and downstream conditions by representing a river system as a graph—with nodes representing gauge locations and edges representing river reaches.

The graph structure naturally captures the connectivity of the river network, allowing the model to propagate information about water movement through the system in a physically consistent manner. These models excel at predicting gauge height time series from precipitation inputs, providing crucial information for flood warning systems.

For urban environments, where complex drainage infrastructure creates additional complications, reinforcement learning algorithms have shown promise in optimising flood simulation. These algorithms focus computational resources on the critical areas for accurate prediction, dynamically adjusting model resolution based on evolving flood conditions.

The result is a more efficient and accurate simulation of urban flooding, where fine-scale features like curbs, drainage inlets, and buildings strongly influence water movement patterns.

Time and Space Together: Multi-Gauge Prediction

The ultimate goal of flood forecasting is not just to predict water levels at a single location but to generate consistent predictions across an entire network of gauges, capturing how flood waves propagate through a river system over time. This multi-gauge, multi-timestep prediction problem represents a particularly challenging test for AI systems, as it requires learning complex spatio-temporal dependencies.

RNNs and LSTMs for Multi-Gauge River Forecasting

Recurrent neural networks (RNNs) and their variants, such as Long Short-Term Memory (LSTM) networks, have demonstrated success in time series prediction tasks like stock price forecasting and hydrological modelling. However, their application to multi-gauge river forecasting introduces unique challenges that demand architectural and methodological adaptations, in particular how to handle non-stationarity and data quality:

- **Multivariate sequential dependencies:** Multi-gauge systems require modelling interactions between spatially distributed gauges (e.g., upstream/downstream relationships). While LSTMs excel at temporal dependencies, they lack innate spatial reasoning. Hybrid

architectures like CNN-LSTM (using convolutional layers for spatial feature extraction) or graph-based LSTMs are often employed to capture catchment topography and flow dynamics.

- **Heterogeneous data fusion:** Integrating gauge data with external variables (e.g., precipitation, soil moisture, satellite-derived flood extents) requires careful feature engineering.
- **Cross-gauge variability:** River networks often exhibit non-uniform hydrological responses due to varying basin sizes, land use, and climate zones.
- **Dynamic normalisation:** Scaling inputs per gauge to address differing magnitudes.
- **Attention mechanisms:** Prioritising relevant gauges or time steps dynamically.
- **Long-term dependencies:** Flood forecasting may require modelling lagged effects (snowmelt-to-discharge delays). LSTMs with stateful configurations or sequence-to-sequence architectures are used to retain memory across extended periods.
- **Missing data imputation:** Sparse or inconsistent gauge readings (common in developing regions) necessitate techniques like bidirectional LSTMs or integration with physics-based models.
- **Transfer learning:** Pre-training on data-rich basins and fine-tuning for data-scarce regions improves generalisation.
- **Interpretability:** LSTMs remain "black-box" models, complicating stakeholder trust.
- **Data hunger:** Effective training typically requires >10 years of daily data, which is unavailable in many regions.

Case Study: Western U.S. River Basins

A 2022 study tested LSTMs across 10 gauges in the western U.S., achieving 20-40% improvement in Nash-Sutcliffe Efficiency (NSE) over traditional models like GloFAS by:

- Training on ERA5 reanalysis data and historical streamflow.
- Using a 7-day forecast horizon with recursive prediction.
- Addressing spatial heterogeneity through basin-specific normalisation.

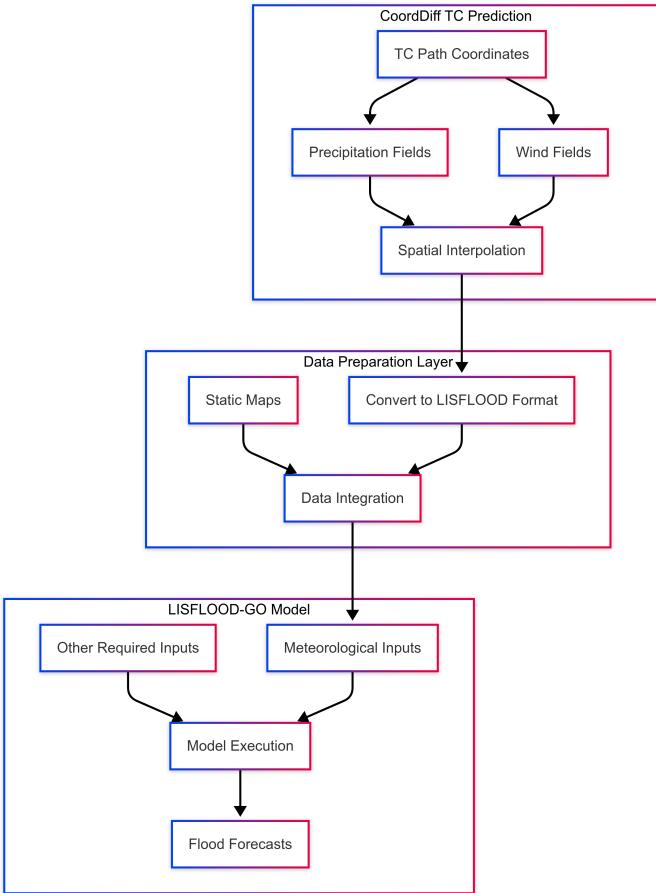


Figure 12: Connection of CorrDiff to LisFlood

For implementation, frameworks like TensorFlow/Keras and PyTorch provide customizable LSTM layers, while tools like NeuralHydrology offer domain-specific optimisations.

Researchers have developed models that can generate coherent predictions across entire river basins by combining LSTM architectures with attention mechanisms that learn which upstream gauges most strongly influence each downstream location.

More recent approaches leverage transformer architectures—initially developed for natural language processing—to capture the sequential

nature of flood wave propagation. Just as a transformer model can understand the relationship between words in a sentence, it can learn how high water levels at one gauge relate to subsequent rises at gauges downstream. These models capture both the time lags between upstream and downstream flood peaks and the attenuation of flood waves as they move through the system.

The most advanced current approaches combine spatial and temporal modelling in unified architectures considering geographic relationships and time evolution into (no surprises here) time series. These spatio-temporal models can ingest data from multiple sources—gauge readings, weather forecasts, soil moisture measurements—and generate consistent flood predictions across regions.

The output is a coordinated set of water level time series for every gauge in the network, providing a comprehensive picture of flood evolution that was previously achievable only with much more computationally intensive physics-based models.

The NVIDIA Earth-2 Vision of an Integrated Digital Twin

NVIDIA's Earth-2 initiative best exemplifies the potential of AI-driven integration across the modelling chain. By combining global weather models (FourCastNet), downscaling techniques (CorrDiff), and increasingly, hydrological and impact models, Earth-2 moves towards a comprehensive digital twin of the earth system. This platform enables rapid exploration of numerous high-resolution scenarios, with potential applications spanning emergency management, infrastructure planning, and banking risk assessment.

The current implementation, while impressive, still requires significant human expertise to bridge between physical and banking domains. As the platform evolves, we can expect increasingly automated translation of physical model outputs into monetary terms, further streamlining the integrated modelling process.

Despite these remarkable advances, AI approaches to physical process modelling face material limitations. The reliance on historical data means these models may struggle with unpreceded events or novel conditions resulting from atmospheric changes. The "black box" nature of many neural network approaches can make diagnosing errors or understanding limitations difficult. This is particularly problematic in high-stakes applications like flood warning systems.

Data quality and availability remain significant constraints, particularly in developing regions with limited historical observations. Models trained

primarily on data from data-rich areas may perform poorly when applied elsewhere, potentially exacerbating inequities in risk management capabilities.

Finally, AI can significantly accelerate computation but not eliminate the fundamental uncertainties in predicting complex natural systems.

Regardless of the modelling approach, chaos dynamics place theoretical limits on forecast accuracy. Well-designed AI systems acknowledge these limitations through appropriate uncertainty quantification rather than promising precision beyond what physics allows.

The most promising direction for AI in physical process modelling is not to replace traditional approaches but to complement them through hybrid systems. Physics-based models provide the foundation of understanding and ensure consistency with natural laws. At the same time, AI components accelerate computation, fill gaps where theoretical understanding is incomplete, and enable new applications not previously possible.

“For flood risk precisely, we can envision a modelling ecosystem where global AI weather models feed downscaled predictions to a mix of traditional and AI-enhanced hydrological models. These physical models generate consistent scenarios that drive banking assessments through machine learning damage functions and business interruption models. The entire system operates coherently, with uncertainty appropriately characterised and propagated throughout.” - David Kelly, MKM.

Achieving this vision requires technical innovation and continued investment in observational networks, data-sharing infrastructure, and cross-disciplinary collaboration. Integrating these advanced AI techniques with banking frameworks represents the frontier of comprehensive flood risk assessment and management.

Integrated Modelling: Green Shoots of Innovation from Across the World

Despite the significant challenges, several pioneering initiatives have demonstrated the value of the integrated modelling paradigm. The Thames Estuary 2100 project in the United Kingdom represents one of the earliest comprehensive efforts to connect climate projections to banking implications, creating a flexible adaptation pathway that has informed billions of pounds of infrastructure investment. While not fully implementing the time series-driven architecture described above, this project established many governance principles and cross-disciplinary collaboration patterns that more technical implementations have followed.

More recently, the Nebraska Department of Natural Resources partnered with the reinsurance industry to develop an integrated model for the Platte River basin that connects weather generators, hydrological models, and insurance loss models. This system produces consistent scenarios across physical and banking domains, allowing policymakers to evaluate the physical effectiveness and the economic efficiency of proposed flood mitigation measures. By maintaining consistency across the modelling chain, the project enables apples-to-apples comparisons of different intervention strategies—something previously impossible when physical and banking assessments used different underlying scenarios.

NVIDIA's Earth-2 initiative represents the most ambitious application of AI to environmental modelling at scale. By combining their FourCastNet global weather forecasting model with CorrDiff's downscaling capabilities, NVIDIA has created a system that can generate high-resolution (2km) weather forecasts up to 500 times faster than traditional numerical weather prediction methods. The platform moves from coarse 25km resolution global forecasts to localised predictions that capture the fine-grained atmospheric dynamics critical for accurate flood prediction.

What makes Earth-2 particularly relevant to integrated flood risk modelling is its end-to-end approach: the same platform that produces weather forecasts can drive hydrological simulations and, increasingly, is being connected to banking risk assessment systems. Early banking service adopters already use Earth-2 outputs to feed their risk models, creating a more seamless connection between cutting-edge climate science and banking decision-making. This integration is particularly valuable for flood modelling, where the spatial resolution of precipitation forecasts directly impacts the accuracy of subsequent hydrological and banking models.

Early implementations have not been without challenges. All required substantial investment in data infrastructure, encountered unexpected

complications in the model coupling, and faced significant change management hurdles in organisations accustomed to more siloed approaches. Yet they also demonstrated tangible benefits: more targeted underwriting, more efficient capital allocation, more effective public investments in resilience, and ultimately, better management of society's exposure to flood risk.

The new modelling paradigm remains in its early stages, with implementations limited to well-resourced organisations and specific geographical regions. Yet the direction is clear: the future of flood risk modelling lies in coherent model stacks that connect physical processes to banking outcomes through calibrated time series, maintaining consistency and appropriately characterising uncertainty throughout the system.

These models can translate predictions of physical disruption—roads flooded, utilities offline, buildings damaged—into estimated banking impacts on the business and local operations that make a property worth living in, providing a much more complete picture of flood risk.

Modelling for Time Series - Role of AI

Artificial intelligence radically transforms our ability to model and predict physical risks, particularly in creating the weather and gauge level time series that form the foundation of our integrated approach. The role of AI in this domain represents a significant evolution from traditional statistical methods, offering enhanced accuracy and computational efficiency.

Weather Time series Generation

The creation of weather time series relies on a hybrid approach that combines the strengths of AI with sophisticated Monte Carlo simulation techniques:

Distribution Calibration

AI models, particularly deep learning architectures, excel at calibrating the distribution of weather patterns based on historical data. These models can identify subtle relationships between atmospheric variables that traditional statistical methods might miss. By analysing decades of meteorological observations, AI systems learn the complex dependencies between temperature, pressure, precipitation, and other factors contributing to flood-generating weather systems.

Pattern Recognition

Neural networks demonstrate remarkable skill in recognising precursors to extreme weather events. By training on historical data that includes normal conditions and those preceding major flood events, these systems can identify early warning signals within emerging weather patterns.

Monte Carlo Path Generation

Once AI has calibrated the underlying distributions, Monte Carlo techniques generate thousands of possible weather scenarios. Each path represents a plausible evolution of weather patterns, with the distribution of these paths reflecting the probabilities informed by the AI's analysis of historical data and climate projections.

This hybrid approach leverages AI's pattern recognition capabilities while maintaining the statistical rigour of Monte Carlo simulations.

Non-linear Gauge Response Modelling

The relationship between precipitation and gauge levels is highly non-linear and influenced by countless local factors. AI models, particularly recurrent neural networks and graph neural networks, are ideal for capturing these complex relationships without requiring explicit parameterisation of every contributing factor.

Real-time Adaptation

Modern AI frameworks can incorporate real-time data to refine predictions as events unfold continuously. This adaptive capability allows for dynamic "intra-day" forecast updating as new information becomes available, improving accuracy throughout the evolution of flood events.

Model Governance Considerations

A significant advantage of running both weather pattern and gauge level predictions through the same AI framework is enhanced model governance. A unified modelling framework ensures consistent assumptions across the entire prediction chain, reducing the risk of incompatible methodologies introducing errors or biases.

An integrated approach allows for greater transparency to track how uncertainties propagate from weather predictions to gauge levels, providing a more comprehensive understanding of confidence intervals in the final forecasts.

A unified AI framework should allow for more efficient allocation of computational resources across different stages of the prediction process, optimising performance where it matters most.

Convergence Challenges

Yet this convergence brings unprecedeted challenges. How do we manage uncertainty as it propagates through multiple model layers? What governance structures can ensure both innovation and responsible deployment? How do we maintain transparency when complex AI systems interact across domains? How do we ensure that these powerful new capabilities serve the broader societal goal of creating more resilient communities in the face of increasing flood risk?

The good news is that the technical challenges follow classic model issues any banking entity has confronted:

- **Uncertainty propagation:** As data flows through the model stack, uncertainties compound at each level.
- **Computational demands:** Running sophisticated AI-driven time series forecasting requires significant computing resources.
- **Data standardisation:** Creating consistent formats for weather, hydrological, and banking data remains challenging unless we stay within the FINOS CDM framework from the outset.
- **Model interoperability:** Ensuring that outputs from physical models can seamlessly serve as inputs to banking models.
- **AI explainability:** Making complex AI-driven predictions interpretable for decision-makers who may lack technical expertise.

The path forward requires technical innovation and new forms of collaboration that bridge traditional disciplinary and industry boundaries. Banking institutions, weather and atmospheric scientists, hydrologists, and technology providers are required to collaborate to build integrated time series frameworks for effective risk management.

The Hidden Architecture of Capital Markets

The advances in AI-driven physical modelling would be of limited value for integrated risk assessment without parallel innovations in translating physical outputs into output that can be ingested directly into banking systems without disrupting existing operating models. Neural network approaches have proven particularly effective at learning the complex, non-linear relationships between physical flood characteristics and resulting damages.

Damage functions—the mathematical relationships that convert flood depths to property damages—traditionally rely on simplified curves derived from limited historical data. Machine learning approaches now enable much more nuanced damage prediction by incorporating additional factors beyond water depth: duration of inundation, flow velocity, water contamination, building materials, occupancy type, and even socioeconomic factors that might influence recovery capacity.

The heart of banking institutions lies what industry insiders, under various guises, call the PVM architecture (Product, Valuation, and Market data).

This classification system, developed over decades, is how banks seamlessly process an astonishing \$500 trillion in instruments daily without significant incidents since the 2008 banking crisis. Each banking instrument is assigned a specific PVM instruction set that:

- Classifies the product type (equity, bond, swap, option, MBS).
- Defines its market data model (what inputs are needed).
- Specifies the analytics model (how to value and risk-manage it).

The standout element frequently underappreciated is the suffix "model" next to market data. While much of the market data is observable, such as a forex rate, a material subset is model-derived.

This model-centric interim step, which creates curves and surfaces and drives higher-order market data such as implied volatility and correlation, is critical as it adds complexity and model risk to the process.

The industry has worked hard since the SABR swaption surface model failure in 2008 and the breakdown of the inflation curve construction in 2022 to establish considerable model governance to ensure their continued performance.

The Time Series Engine for Capital Markets

"Time series—chronological sequences of observations capturing how markets behave over time—of market data, whether observed directly or derived, are the fuel that feeds the capital market's engine." - Johnny Mattimore, MKM.

The importance of time series extends far beyond facilitating market-makers:

- **Algorithmic Trading Strategies:** Quant teams build trading algorithms by detecting pattern anomalies between recent prices and their historical time series.
- **Risk Management Systems:** You cannot discuss time series without acknowledging their crucial role in Value-at-Risk (VaR) engines, which are essentially time series processors.
- **Regulatory Capital Requirements:** Perhaps most critically, the stress tests determining how much capital banks must hold are calibrated using time series data.
- **Research and Origination:** All of those research documents have a chart explaining why their viewpoint now needs attention by the decision-makers and policy deciders.

The Physical Risk Integration Challenge

As banking institutions face mounting pressure to incorporate physical hazard risks into their frameworks, we've hit a fundamental compatibility problem.

Most physical risk across natural hazards into risk information comes in the form of hazard curves—showing, for example, that a 1-meter flood is expected every 20 years, whereas a 2-meter flood is expected every 500 years.

“This works well for the insurance industry, but it does not work for banking. The reason is that any risk information for banking must integrate with banking's PVM architecture. To achieve this, the risk information must be presented as a time series with all of the attributes of market data.” - Johnny Mattimore, MKM.

This integration is far more than a technical exercise. The coherent model stack connects previously disparate domains: neural networks that identify subtle precursors to extreme precipitation events; reinforcement learning algorithms that optimise the simulation of water movement across complex topographies; and gradient-boosting techniques that translate predicted water levels into location-specific damage estimates and ultimately into dynamic pricing models.

A Time Series Approach to Physical Risk

MKM Research Labs, in collaboration with esteemed partners in the FINOS collective, is pioneering an innovative solution to this architectural disconnect. Rather than conforming to the traditional hazard curve paradigm, we're approaching physical risk through the lens of what banks already understand: time series data.

“Rather than reinventing how banks operate, we need to go back to the beginning and build the physical risk equivalent of what powers markets today: robust, reliable, and relevant time series data that can plug directly into existing banking infrastructure.” - David Kelly, MKM

The goal is not to predict the weather but to create comprehensive time-series distributions that adhere to the High-Resolution Rapid Refresh (HRRR) weather data schema. This approach enables us to generate distributions of potential weather patterns, including crucial parameters such as direction, precipitation intensity, and atmospheric conditions.

We employ standard fluid dynamics models combined with precise terrain data from these weather pattern time series to derive second-order time series: water-level measurements at specific river gauges. This methodology creates a continuous chain of causality from weather events to physical impacts that can be expressed in the language of banking markets.

By adopting the local meteorological department's definition of flood alert levels—similar to how a strike price operates in options markets—we establish clear, measurable thresholds that can be used to design parametric insurance products.

FINOS CDM and Physical Risk Swaps

Within the FINOS Common Domain Model, MKM Research Labs has defined a Physical Risk Swap that is remarkably similar to credit default swaps (CDS) as defined by ISDA documentation. While a CDS pays out upon a credit event, these physical risk instruments payout when gauge measurements cross predefined alert levels.

Most importantly, the valuation methodology for these instruments follows established banking market principles:

- Historical time series of gauge measurements inform baseline expectations.
- Synthetic time series generated from climate models provide forward-looking scenarios.
- Standard derivative pricing techniques can be applied to determine fair values.

This approach transforms an insurance-oriented hazard assessment into a market-compatible banking instrument that can be seamlessly integrated into banks' existing PVM architecture.

The introduction of Physical Risk Swaps creates the birth of a brand new asset class that will contribute to the protection—and most importantly, the reduction of insurance cost—of an asset class that, in the US residential market alone, is valued at \$52.5 trillion with \$12.6 trillion of outstanding mortgages. Of these, 12 million properties, in addition to those already in FEMA's Special Flood Hazard Areas, have a significant risk of flooding.

Final Thoughts

The journey through the cutting-edge developments in flood risk modelling reveals a landscape transformed by artificial intelligence and interdisciplinary collaboration. What emerges is not merely a technological evolution but a fundamental reconceptualisation of how we understand and manage flood risk across the physical-financial continuum.

The integration of AI-enhanced weather prediction, sophisticated hydrological modelling, and banking-compatible time series frameworks represents perhaps the most significant advancement in flood risk management since the introduction of computational fluid dynamics. This coherent model stack approach dissolves the artificial boundaries between scientific domains and banking analysis, allowing us to see risk as nature does: as a continuous flow from cloud movements through flood time series to parametric attribution.

While the technical achievements are remarkable, the path forward depends equally on standardisation and governance. The FINOS Common Domain Model (CDM) provides the essential scaffolding for data standardisation, creating a shared language that enables physical risk information to flow seamlessly through banking systems. This standardisation is not merely a technical convenience but a prerequisite for market development, regulatory compliance, and ultimately, societal resilience.

In parallel, legal standardisation through ISDA documentation brings the clarity and certainty that market participants require. By adapting established frameworks for innovative instruments like Physical Risk Swaps, we build on decades of market evolution rather than starting from scratch. This approach speeds up adoption and guarantees compatibility with existing risk management frameworks.

Perhaps most critically, robust model governance must support every aspect of this integrated architecture. As uncertainty propagates through multiple model layers, transparency and validation become not only regulatory requirements but also essential components of a trustworthy system. From neural networks predicting precipitation patterns to the complex valuation models for physical risk instruments, each component must undergo rigorous governance that recognises both its individual characteristics and its role in the broader system.

The convergence of cutting-edge AI techniques with established banking frameworks represents a measured optimism rather than unbridled techno-utopianism. We recognise the challenges: computational demands remain substantial, data quality varies significantly across regions, and model interoperability requires continued attention. Yet these challenges appear increasingly surmountable through collaborative effort and technological innovation.

As we conclude this exploration of future developments, we can observe with quiet confidence that the tools required for comprehensive flood risk management are rapidly materialising. The integration of physical science and banking analysis through standardised time series provides a pathway to more resilient communities, more efficient capital allocation, and ultimately, a more sustainable relationship with our changing environment climate.

I will now end this book with another quote!

“This is not the end of the journey but rather the beginning of a new paradigm—one that promises to transform how we anticipate, mitigate, and respond to one of humanity's oldest and most persistent challenges.”

David Kelly, MKM.

This page has been intentionally left blank.

Glossary of Terms

A

- Chronic Physical Events: Sudden and severe weather occurrences, such as floods, storms, or hurricanes, can immediately damage properties and infrastructure.
- Adaptation Planning: Developing strategies to adjust to actual or anticipated climate change effects is especially important for property-level flood resilience measures.
- Affordability Ratio: The percentage of income allocated to housing costs, including mortgage payments and insurance premiums, utilised in mortgage valuation models to evaluate default risk.
- Aleatory Uncertainty is the inherent randomness in natural processes that cannot be mitigated through additional information or improved models; it contrasts with epistemic uncertainty.
- Annual Exceedance Probability: The likelihood that a specified flood level will be surpassed in any given year, often expressed as a percentage (e.g., 1% annual exceedance probability for a "100-year flood").
- Asset-Level Analysis: A thorough risk assessment carried out at the level of individual properties or assets, contrasting with portfolio-level analysis.
- Atmospheric State Variables: Essential parameters that describe atmospheric conditions, including temperature, pressure, humidity, and wind vectors, which are used in weather prediction models.

B

- Basel III/IV: International regulatory frameworks for bank capital requirements that increasingly include climate and physical risk considerations. These regulations establish standards for how banks must measure, disclose, and hold capital against various risks, with evolving requirements for incorporating physical and transition climate risks into banking operations.
- Bayesian Approaches: Statistical methods that update probability estimates as new data becomes available; widely used in weather

prediction and flood risk assessment to integrate prior knowledge and new observations.

- Bayesian Inference: Statistical method used to update the probability estimate for a hypothesis as more evidence or data becomes available. It is based on Bayes' Theorem, which provides a way to calculate the probability of a hypothesis (or model) given observed data. As the industry gathers more data, we update our belief (posterior probability).
- Beta Distribution: A family of continuous probability distributions defined on the interval [0,1], commonly used to model damage ratios in vulnerability functions.
- Black Box: A system or model whose internal workings are not transparent or easily understood; a concern with specific AI models used in risk assessment.
- Boundary Conditions are the values of variables at the edges of a modelled domain, which are essential for solving hydrological and hydraulic models.
- Boundary layer in weather refers to the thin layer of air at the Earth's surface that is directly influenced by the surface itself (land, water, vegetation, etc.). This layer is typically around 1 to 2 kilometers (km) thick, although it can vary depending on factors like time of day, weather conditions, and geography. It plays a crucial role in weather dynamics because it's where most weather phenomena, like temperature, moisture, and wind, interact with the Earth's surface.
- Building Arrangement Effects: The impact of building configurations in urban environments on flood flow paths and property-specific flood exposure.
- Business Interruption: Banking losses arising from the inability to operate a business during and after a flood event; a critical component of comprehensive flood impact assessment.

C

- Capital markets Approach: A methodology for valuing physical risk that adapts techniques from banking derivatives markets, particularly credit default swaps, to quantify and price flood risk.
- Catastrophe Modelling: Analytical techniques that estimate the physical, social, and economic impacts of natural disasters by combining hazard, exposure, and vulnerability components.

- CDM (Common Domain Model): A standardised data representation framework developed through FINOS that establishes consistent definitions for weather events, flood events, properties, mortgages, and other elements of physical risk assessment.
- Climate Biennial Exploratory Scenario (CBES): A stress-testing exercise conducted by the Bank of England to assess how climate change might impact the UK banking system. It evaluates financial institutions' resilience to transition risks and physical risks across multiple scenarios and time horizons, enabling regulators to understand potential systemic vulnerabilities.
- Coastal Flooding is the inundation of land areas along coastlines, typically caused by storm surges, high tides, or tsunamis. It has distinct characteristics that set it apart from riverine flooding.
- Compound Flood Events: Situations in which multiple flood mechanisms occur simultaneously or in close succession, resulting in more complex and severe impacts than single-mechanism events. Examples include combined coastal and riverine flooding during storms or pluvial flooding coinciding with high groundwater levels, creating amplified impacts through their interaction.
- Convective storm is a type of storm that forms due to the upward movement (convection) of warm, moist air. This process creates instability in the atmosphere, which leads to the development of storms such as thunderstorms, tornadoes, and hailstorms. Convection occurs when warmer air near the Earth's surface rises because it is less dense than the cooler air above it.
- Convolution Integral: A mathematical operation used in hydrograph analysis that combines excess rainfall with the unit hydrograph to calculate direct runoff. It enables hydrologists to predict streamflow response to complex rainfall patterns by treating them as a sequence of simple inputs, each generating its own scaled hydrograph response.
- Convolutional Neural Networks: A class of deep neural networks primarily used for analysing visual data, such as images and videos. They are particularly powerful for tasks like image classification, object detection, and segmentation but have also been adapted to other domains like speech and text processing. CNNs are designed to automatically and adaptively learn spatial hierarchies of features from input data.
- Copula Functions: Mathematical tools that describe the dependence between random variables, beneficial for modeling the correlation between flood-causing factors. Copulas separate the marginal distributions of individual variables from their joint dependency structure, enabling more flexible and accurate representation of

complex relationships like those between precipitation, soil moisture, and river levels.

- Correlation Distance Parameter: The spatial distance (typically around 1000m) within which properties experience intensely correlated flood impacts is crucial for understanding portfolio-level risk.
- Coupled partial differential equations (PDEs) are a system of two or more partial differential equations that are interrelated or "coupled" together. These equations involve multiple unknown functions that depend on the same set of independent variables (such as time and space) and are linked through their terms. The solutions to these equations are dependent on one another.
- Credit Default Swap (CDS): A banking derivative that protects against the risk of default by a particular entity, serving as a model for the development of Physical Risk Swaps.
- Credibility Theory: An actuarial approach that balances individual risk experience with broader class experience when setting insurance premiums.

D

- Damage Functions: Mathematical relationships that translate flood characteristics (depth, duration, velocity) into expected property damage, typically expressed as a percentage of property value.
- Default Probability (PD): The likelihood that a borrower will default on required mortgage payments, potentially influenced by flood events through income disruption or property devaluation.
- Designed neural network architectures refer to the various specific configurations of neural networks that are created and optimised for particular tasks or types of data. These architectures determine how the neurons (or nodes) are connected, how data flows through the network, and how the network learns from the data. Different architectures are suited to various types of problems, and they are typically chosen based on the nature of the input data, the computational resources available, and the desired output.
- Digital Elevation Models (DEMs): Three-dimensional representations of terrain surfaces used in flood risk assessment to determine water flow paths and inundation zones.
- Digital Regulatory Reporting (DRR): A framework within FINOS that standardizes the regulatory reporting process, integrated with Physical Risk Swaps to ensure compliance. DRR transforms

regulatory rules into machine-executable code, streamlining reporting while ensuring consistency across institutions.

- **Discounted Cash Flow Analysis:** A valuation method that estimates the value of an investment based on its expected future cash flows, adjusted for the time value of money and risk.
- **Disturbance Absorption Thresholds:** The magnitude of flood events that natural systems can accommodate while maintaining their essential functions; exceeding this threshold may sharply reduce natural resilience.
- **Distribution Path Modelling:** An approach that characterizes the entire space of potential weather pattern evolutions, especially those that could lead to extreme precipitation events. Rather than producing single forecasts, this method maps the full probability distribution of possible weather trajectories, providing a more complete picture of potential flood-generating conditions.
- **Downscaling:** Producing high-resolution local forecasts from coarse global predictions is essential for capturing localised precipitation patterns that drive flood events.
- **Duration-Depth-Velocity Matrices:** Three-dimensional classification schemes that acknowledge that flood impacts depend on water depth, flow velocity, and event duration.
- **Dynamic Pressure Considerations:** Analysing the hydrodynamic forces exerted by moving floodwaters on structures is critical in high-velocity flood scenarios.
- **Dynamic Wave Approach:** A hydraulic modelling technique that employs the full Saint-Venant equations, including all acceleration and pressure terms, to provide the most complete representation of flow dynamics. While computationally intensive, this approach captures complex hydrodynamic effects crucial for accurate flood modeling in rapidly changing conditions.

E

- **Earth-2 Initiative:** NVIDIA's extensive project aims to create a digital twin of Earth's systems by combining global weather models with hydrological and impact assessments for integrated climate risk modelling.
- **Enterprise Risk Management (ERM) Integration:** The process of incorporating climate-related risks into an organization's comprehensive framework for assessing and addressing all risks that

affect its business objectives. This ensures climate risks are evaluated alongside traditional risks within established governance structures, rather than being managed in isolation.

- Environmental Resilience Factors: These natural systems buffer, absorb, or amplify flood hazards before they impact vulnerable assets, including wetlands, forests, and soil systems.
- Epistemic Uncertainty: This type of uncertainty arises from incomplete knowledge about asset characteristics and their vulnerability to hazards; unlike aleatory uncertainty, it can be mitigated through additional data or improved modelling.
- Event_CDM: A component of the Common Domain Model that standardises the representation of extreme weather event pathways, utilising advanced AI platform components.
- Exceedance Probability: This refers to the likelihood that a specific hazard intensity (such as flood depth) will be met or surpassed within a given time frame, typically one year.
- Expected Cash Flow (ECF): In mortgage valuation, the scheduled payment is adjusted for the probability of default and loss given by default.
- Expected Shortfall (ES), also known as Conditional Value at Risk (CVaR), is the expected loss in the worst $\alpha\%$ of cases. Unlike Value at Risk (VaR), which only measures the minimum loss at a given confidence level, ES calculates the average of all losses that exceed this VaR threshold, providing a more comprehensive picture of tail risk and addressing VaR's failure to capture the severity of extreme losses beyond the threshold.
- Exposure Analysis: This process identifies assets (buildings, infrastructure, etc.) that could be impacted by flooding, including their location, value, and characteristics.
- Exposure Assessment Integration: This involves a comprehensive analysis of all potential paths by which floodwaters could reach a property, incorporating topographic, locational, and built environment factors.
- Extreme Event Prediction: These are specialised approaches for forecasting rare but severe weather events, often necessitating methodologies different from those used for typical conditions.
- Extreme Value Theory (EVT) is a statistical framework for modelling and analysing the stochastic behaviour of extreme events—rare occurrences that lie in the tails of probability distributions. These events, though infrequent, often have significant societal, financial, or environmental impacts, making their quantification critical for risk management across diverse fields

F

- Finite Difference Approximation: A numerical method used in hydrological routing to solve differential equations by replacing continuous derivatives with discrete differences. In flood modelling, this approach enables practical computation of water movement through complex river systems and across floodplains.
- FINOS (Fintech Open Source Foundation): An industry consortium that promotes open collaboration in banking services technology, including developing the Common Domain Model for physical risk.
- Fixed-Floating Exchange: The principle behind Physical Risk Swaps is that one party pays a fixed rate over the contract period while the other makes contingent payments based on realised flood levels.
- Flash Flooding: Rapid-onset flooding characterised by high water velocities, typically occurring within six hours of the initiating event (heavy rainfall or infrastructure failure).
- Flood Classification Systems: Methodologies for categorising flood events based on source mechanism, temporal characteristics, spatial extent, and other relevant factors.
- Flood Depth: The height of floodwater above the ground or floor level is a primary factor in determining damage to structures.
- Flood Gauges CDM: A standardised representation of flood monitoring devices within the Common Domain Model, creating definitive trigger points for banking contracts.
- Flood Insurance Rate Maps (FIRMs): Official maps produced by the Federal Emergency Management Agency delineate flood hazard areas for insurance and regulatory purposes.
- Flood Re: A not-for-profit reinsurance scheme established in the UK to ensure affordable flood insurance remains available to households in high-risk areas. Designed with a planned 25-year lifespan, Flood Re creates a transitional mechanism to move from subsidised to risk-reflective pricing by 2039.
- Flood Risk Swaps: Banking instruments designed to transfer flood risk from one party to another, similar to credit default swaps but triggered by physical events rather than credit events.
- Fluvial (Riverine) Flooding: Inundation is caused when water exceeds the capacity of river channels, resulting in overbank flow across the floodplain.

- FourCastNet: NVIDIA's neural network approach to global weather forecasting can generate predictions up to 500 times faster than traditional numerical methods.

G

- Gauge-Based Triggering System: This is a mechanism used in Physical Risk Swaps, where payouts are determined by water level measurements at specific river or flood gauges.
- Generalised linear models (GLMs) are statistical models utilised in premium calculation that simultaneously consider multiple risk factors.
- Generative Adversarial Networks (GANs) are a class of machine learning models designed for unsupervised learning, enabling synthetic data generation that mimics real-world distributions. GANs employ two neural networks—a “generator” and a “discriminator”—trained adversarially to improve each other iteratively.
- Geographic Diversification: A risk management strategy that spreads property exposures across various watersheds and flood zones to reduce portfolio-level correlation effects.
- Geomorphological Unit Hydrograph (GUH): An approach that connects hydrograph shape to catchment physical characteristics by incorporating stream order statistics and network width functions.
- Gradient Boosting Algorithms: Machine learning techniques that create predictive models by combining multiple weak models, particularly effective for defining detailed insurance rate classes.
- Graph Neural Networks: Specialised deep learning architectures that depict river systems as connected graphs, with nodes representing gauge locations and edges representing river reaches. This approach captures the topological relationships between different parts of a river network, improving the modelling of how flood waves propagate through the system.
- GRIB2 Format: a standardised file format for storing and distributing gridded meteorological data, commonly used for weather model outputs and historical climate data. GRIB2 utilises sophisticated compression techniques to efficiently store large volumes of data but requires specialised software for access and interpretation.
- Groundwater Flooding is inundation caused by a rise in the water table to the ground surface, typically following extended periods of high precipitation.

- Growth-at-Risk (GaR): A framework that associates banking conditions with the distribution of future economic growth, increasingly integrating physical risk factors.

H

- Hazard Curve: A graphical representation of the relationship between hazard intensity (such as flood depth) and its probability of occurrence, which forms the foundation of physical risk assessment.
- Hazard Identification: The first step in risk assessment involves recognising and characterising potential sources and mechanisms of flooding.
- Hazard Rate: In the context of Physical Risk Swaps, the conditional probability of a physical event occurring within a small time interval, given that it has not happened previously.
- HEC-RAS (Hydrologic Engineering Center's River Analysis System): Industry-standard software developed by the U.S. Army Corps of Engineers for hydraulic modelling, capable of simulating one-dimensional and two-dimensional water flow.
- Hidden Markov Models (HMMs): Probabilistic models that treat precipitation as a system transitioning among hidden weather states, each associated with distinct precipitation characteristics.
- High-Resolution Rapid Refresh (HRRR): This is an operational weather prediction model with a 3km spatial resolution that provides frequently updated forecasts, valuable for near-real-time risk assessment.
- Hydraulic Modelling: The simulation of water movement based on fluid dynamics principles, typically solving the Saint-Venant equations to determine water surface elevations and flow velocities.
- Hydrological Modelling: Depicting how precipitation transforms into runoff and river flow, incorporating factors such as infiltration, evaporation, and groundwater movement.
- Hydrostatic Pressure Distribution: An assumption in the Saint-Venant equations that pressure varies linearly with depth at any point in a water column. This simplification enables practical hydraulic modelling while maintaining sufficient accuracy for most flood simulation applications.

I

- Impact Quantification: Converting physical hazard characteristics into measurable effects on property, infrastructure, and economic systems.
- Importance Sampling: A variance reduction technique used in Monte Carlo simulations that focuses computational resources on regions of the state space most relevant to risk assessment. This approach improves efficiency by concentrating sampling in areas of the probability distribution that have the greatest impact on flood risk, particularly in the tails of distributions that represent extreme events.
- Infrastructure Failure Flood Types: Classifications for floods resulting from the failure of human-made structures such as dams, levees, or urban drainage systems.
- Insurance Premium Equation: The mathematical formula that combines pure premium, expense loading, risk loading, and profit margin to determine the cost of insurance coverage.
- Intensity-Duration-Frequency (IDF) Curves: Hydrological modelling uses graphic representations of the relationship between rainfall intensity, duration, and frequency.
- Intensity-Duration-Frequency (IDF) Relationships: Mathematical functions that relate rainfall intensity to storm duration and recurrence frequency. These relationships, typically presented as curves or equations, provide essential inputs for hydrological modelling and are fundamental to flood hazard assessment.
- ISDA Master Agreement: The standardised contract used for over-the-counter derivatives transactions, providing the contractual foundation for Physical Risk Swaps.

J

- Joint Probability Analysis: Statistical techniques considering the simultaneous occurrence of multiple flood-causing factors, such as high river levels coinciding with heavy rainfall.

K

- Kinematic Wave Approximation: A simplified form of the Saint-Venant equations that neglects acceleration and pressure terms, applicable primarily to steep slopes where gravity and friction forces dominate flow behaviour.

L

- LISFLOOD: an open-source distributed hydrological model developed by the European Commission's Joint Research Centre to simulate flood events.
- Loan-to-Value (LTV) Ratio: This ratio of mortgage debt to property value is a key factor in mortgage risk assessment and can be significantly affected by flood events.
- Local Adaptation refers to tailoring global or regional models to capture location-specific weather or flood behaviour in modelling contexts.
- Local Loss Model (LLM): A focused approach to risk assessment that applies detailed vulnerability functions to specific assets or small geographic areas. This methodology provides high-resolution loss estimates by considering individual property characteristics and their particular exposure to flood hazards.
- Loss Given Default (LGD): This is the amount a lender loses when a borrower defaults on a loan, typically expressed as a percentage of the outstanding debt.
- Lumped Models: Hydrological models that treat watersheds as single units with uniform characteristics, sacrificing spatial detail for computational efficiency.

M

- Manning's n Coefficients: Parameters used in hydraulic modelling to represent the roughness or resistance to the flow of a channel or floodplain surface.
- Market Value Correlation: The tendency for property values in a neighbourhood to rise or fall together following flood events, even affecting properties that weren't directly damaged.
- Material Vulnerability Differentiation: Analysis of how different building materials respond to flood exposure, from structural elements to interior finishes.

- Meteorological Models: Computational frameworks that simulate atmospheric conditions to predict weather patterns, particularly precipitation events that may lead to flooding.
- Microtopography: Small-scale terrain features (curbs, berms, garden walls) significantly influence urban flood pathways despite their modest size.
- Model Risk Governance is the framework of policies, controls, and procedures for managing risks that rely on models. It is crucial for flood risk assessment, where multiple models are linked.
- Monte Carlo Simulation is a computational technique that uses repeated random sampling to obtain numerical results. It is handy for exploring the range of possible flood scenarios.
- Mortgage Component: A standardised structure within the Common Domain Model that defines how physical risk affects loan-to-value calculations, default probabilities, and valuation impacts.
- Multi-criteria Evaluation Frameworks: Assessment methodologies that incorporate multiple factors related to property resilience, weighted according to their relative importance.
- Multi-gauge Prediction: The process of generating consistent water level forecasts across an entire network of monitoring points, capturing how flood waves propagate through river systems.
- Multivariate Class Rating: Insurance premium calculation approaches that simultaneously consider numerous risk factors through statistical models like generalised linear models.
- Muskingum Method: A hydrological flow routing model based on weighted storage, used to calculate the outflow of a river reach given an inflow hydrograph. The method conceptualises a river reach as having both prism storage (regular channel volume) and wedge storage (additional volume from flood waves), making it particularly useful for flood wave propagation modeling.

N

- National Flood Insurance Program (NFIP): A U.S. government program established in 1968 to provide flood insurance to property owners in participating communities. It combines insurance availability with floodplain management requirements and employs Flood Insurance Rate Maps (FIRMs) to determine premium rates and regulate development.

- Nash Cascade: A conceptual rainfall-runoff model that represents a catchment as a series of linear reservoirs, each with the same storage coefficient. This approach produces an instantaneous unit hydrograph that captures the essential features of catchment response while maintaining mathematical tractability.
- Navier-Stokes Equations: The fundamental equations of fluid dynamics that describe the motion of viscous fluid substances, forming the theoretical basis for hydraulic modelling.
- Neural Networks: Computational models inspired by the human brain that can learn patterns from data are increasingly applied to weather prediction and flood risk assessment.
- Non-linear Gauge Response Modelling: The use of advanced statistical techniques to capture the complex, non-linear relationship between precipitation and water levels at monitoring points.
- Non-stationarity: The concept that statistical properties of environmental systems change over time is particularly important in climate contexts where historical patterns may not represent future conditions.

O

- One-dimensional (1D) Hydraulic Models: Simulation frameworks that represent water flow along a single spatial dimension, typically the channel centerline, using cross-sections to capture geometry.
- Option-Adjusted Spread (OAS): A methodology that values embedded options in banking instruments by simulating numerous future scenarios for interest rates and housing prices.
- Orographic Effects: The influence of mountains and other terrain features on precipitation patterns, often creating localised heavy rainfall that traditional weather models struggle to account for. OS-Climate Framework: An open-source collaborative initiative for climate risk assessment that standardises approaches to hazard curve development and vulnerability modelling.
- Orographic precipitation is a type of precipitation that occurs when moist air is forced to rise over a mountain range or other elevated terrain. As the air rises, it cools and condenses, leading to the formation of clouds and, eventually, precipitation.
- OS-SFT (Open-Source Sustainable Finance Taxonomy): A project that aims to harmonize regulatory interpretation across jurisdictions

by creating standardized classifications for sustainable finance activities. This taxonomy proposed by Johnny Mattimore supports consistent assessment and reporting of physical risk impacts on financial assets across different regulatory regimes.

- Overbank Flow: When water in a river channel exceeds capacity and spills onto the adjacent floodplain.
- Own Risk Solvency Assessments (ORSA) Integration: The inclusion of hazard-related factors within an insurer's internal process for evaluating its risk management framework and current and future solvency positions. This integration ensures that potential climate impacts on capital adequacy and financial stability are systematically assessed and reported to regulators.

P

- Parametric Insurance is insurance contracts that pay out based on predefined, objective measurements (such as water level at a gauge) rather than assessed damages, providing the foundation for Physical Risk Swaps.
- Parametric Vulnerability Functions are mathematical distributions (like Beta, Gamma, or Lognormal) used to model the relationship between hazard intensity and damage, providing expected values and uncertainty ranges.
- Physical Models: Hydrological approaches based on fundamental water movement and conservation equations demand significant data and computational resources but provide robust theoretical foundations.
- Physical Risk Swap (PRS): A banking derivative instrument that transfers flood risk from mortgage lenders to counterparties better positioned to bear it. It is structured similarly to credit default swaps but triggered by physical events.
- Physics-informed Neural Networks (PINNs): Machine learning models that combine neural network architectures with explicit constraints derived from the laws of physics, preserving physical realism while improving computational efficiency.
- Pluvial (Surface Water) Flooding: Inundation is caused when rainfall overwhelms drainage systems or cannot infiltrate the ground, causing water to pool or flow over the land surface.
- Portfolio Aggregation Techniques are methodologies for combining individual property-level risk assessments into comprehensive portfolio-wide measures. These techniques account for spatial

correlation, exposure concentration, and diversification effects to determine overall risk profiles for mortgage or insurance portfolios.

- Probability of Default (PD): In mortgage assessment, flood events may exacerbate a borrower's likelihood of defaulting on their loan due to property damage or disruption to income.
- Product, Valuation, and Market Data (PVM): A classification system banking institutions use to process banking instruments, offering a framework for integrating physical risk into existing systems.
- Property Clustering: The geographic concentration of properties in flood-prone areas, creating risk multiplier effects through correlated physical damage and market value impacts.
- Property Component: A standardised structure within the Common Domain Model that captures essential physical characteristics, including geolocation, elevation, building specifications, and flood defence relationships.
- Property-Level Flood Resilience (PFR): Specific measures implemented at individual buildings to reduce flood damage, including structural modifications, deployable barriers, and water-resistant materials. PFR approaches can significantly reduce loss potential through either resistance (keeping water out) or resilience (minimising damage when water enters).
- Pure Premium: An insurance premium's expected loss cost component is determined through comprehensive risk modelling.
- PVM (Product, Valuation, and Market Data): A classification system used by banking institutions to process banking instruments, offering a framework for integrating physical risk into existing systems. PVM provides the structural organization for how physical risk factors are incorporated into product definitions, valuation methodologies, and market data flows.

Q

- Quasi-Monte Carlo (QMC) Methods: Computational techniques that systematically explore state spaces through carefully constructed low-discrepancy sequences, providing more uniform coverage than conventional random sampling.
- Queuing Loss Model: A mathematical framework for calculating the probability of system failure when demand exceeds capacity, applied to urban drainage systems during flood events.

R

- Rainfall-Runoff Models: Hydrological frameworks that transform precipitation inputs into water flow outputs across landscapes, incorporating factors like infiltration, evaporation, and groundwater interaction.
- Rating Territories: Geographic zones used by insurers to classify and price flood risk, traditionally defined by administrative boundaries but increasingly based on sophisticated hydrological and topographical characteristics. Modern approaches use high-resolution data to create micro-territories that reflect actual risk gradients more accurately.
- Recurrence Interval: The average time between flood events of a specified magnitude, often expressed in years (e.g., "100-year flood").
- Recurrent Neural Networks (RNNs) are machine learning architectures designed explicitly for sequential data like time series. They can learn temporal dependencies in water level predictions.
- Regularisation: Machine learning techniques prevent overfitting by incorporating a penalty term into the loss function, which is crucial for developing robust flood prediction models.
- Regulatory Capital Methodologies are frameworks that determine how much capital financial institutions must hold against specific risks. These methodologies are evolving for physical risks like flooding to incorporate hazard frequency, severity, and correlation effects within established frameworks like Basel standards.
- Resilience Certification Programs: Standardized verification of implemented flood protection measures and their expected performance, increasingly recognised by insurance markets.
- Return Period: The inverse of the annual exceedance probability; a 100-year return period corresponds to a 1% annual exceedance probability.
- Risk-Reflective Pricing: Insurance premium calculation that accurately reflects a property's actual flood risk based on detailed hazard, exposure, and vulnerability assessment. This approach contrasts with subsidised or community-rated pricing by charging higher premiums for higher-risk properties, creating economic incentives for risk reduction.

- Riverbank Breach Modelling: Specialized simulation of failures in natural riverbanks or engineered levees, often producing rapid inundation with distinct hazard characteristics.
- Risk Loading: A component of an insurance premium that reflects uncertainty and volatility in the loss distribution, crucial for properties exposed to catastrophes.
- Runoff Generation Process: The transformation of precipitation into water movement across landscapes, encompassing surface runoff, subsurface flow, and groundwater contribution.

S

- Saint-Venant Equations: A system of partial differential equations describing one-dimensional unsteady flow in open channels, forming the mathematical foundation for hydraulic modelling.
- Schedule Rating: An insurance pricing approach that applies credits or debits for specific risk characteristics not captured in standard rating factors, such as loss control measures or construction features.
- SCS Dimensionless Unit Hydrograph: A standardized unit hydrograph developed by the Soil Conservation Service (now Natural Resources Conservation Service) that represents the temporal distribution of runoff from a unit of excess rainfall. It provides a template that can be scaled based on watershed characteristics, making it widely applicable across diverse landscapes.
- Self-Organisation Capacity: The ability of natural systems to dynamically adjust to changing conditions, distinguishing resilient natural flood defences from engineered structures.
- Semi-Distributed Models: Hydrological approaches that segment watersheds into sub-basins with unique characteristics, balancing computational efficiency and the representation of spatial variability.
- Sensitivity Analysis: A systematic investigation into how changes in model inputs affect outputs is essential for understanding uncertainty in flood risk assessment.
- Simplifying Complex Processes: The essential reduction of intricate physical processes (such as cloud formation) in models because of computational limitations, leading to downstream effects on flood risk assessment.

- Solvency Capital Requirement (SCR): Under the Solvency II regulatory framework, insurers must hold a specific amount of capital for catastrophe risk, directly influencing premium pricing.
- Spatially Distributed Hydrological Model: A modelling approach that divides watersheds into a grid of cells with unique characteristics to simulate water movement across landscapes with high spatial resolution. These models can capture how terrain, soil, and land use variations influence runoff patterns and flood development.
- Special Flood Hazard Areas (SFHAs): Designated zones on Flood Insurance Rate Maps where the National Flood Insurance Program's floodplain management regulations must be enforced and the mandatory purchase of flood insurance applies. These areas face at least a 1% annual chance of flooding.
- Snyder's Method: An empirical approach to synthetic unit hydrograph development that relates key hydrograph parameters to watershed physical characteristics. Developed in the 1930s, this method uses regional coefficients to estimate lag time and peak discharge based on watershed length, centroid distance, and other measurable features.
- Spatial coherence refers to the consistency or correlation of a physical quantity (such as light, sound, temperature, or other fields) over space. It describes how well a property or signal is correlated across different locations in space, and it's often used to analyse the structure of waves, signals, or phenomena as they propagate or vary over space.
- Spatial Correlation: The tendency for nearby properties to experience similar flood impacts due to shared exposure characteristics, creating risk concentration in mortgage portfolios.
- Stochastic Differential Equations (SDEs): Mathematical equations model random processes that evolve and are applied to weather pattern evolution in flood risk assessment.
- Stochastic House Price Model: A mathematical approach that treats property values as random processes to simulate numerous possible future scenarios for mortgage valuation.
- Strategic Default: A deliberate decision by a mortgage holder to stop making payments despite having the financial ability to continue, typically when a property's value falls significantly below the outstanding loan balance. In flood-prone areas, this can occur when properties become uninsurable or suffer substantial devaluation following flood events.
- Structural Vulnerability Analysis: Assessment of how building materials, critical system placements, foundation types, and other physical characteristics influence flood resilience.

- Survival Probability: In Physical Risk Swaps, the probability that the trigger event (e.g., flood level exceeding threshold) has not occurred by a specific time.
- Synthetic Unit Hydrograph: Techniques for estimating flood response without observed data, including the SCS Dimensionless Unit Hydrograph and Snyder's Method.

T

- Tail Value at Risk (TVaR): A risk measure reflecting the average loss in the worst scenarios (typically the worst 0.5%) utilised to determine the necessary capital for catastrophe exposures.
- Temporal Dynamics: Evaluating how flood exposure shifts over time due to natural processes and human interventions, encompassing coastal erosion, riverbed migration, and infrastructure modifications.
- Territorial Rating: An insurance pricing methodology that employs geographic location as a primary rating factor, increasingly enhanced through high-resolution flood risk data.
- Thames Estuary 2100 Project: A comprehensive initiative linking climate projections to banking consequences, establishing a flexible adaptation pathway for infrastructure investment.
- Time Series Analysis involves analysing data points gathered or recorded at successive intervals, which is crucial for comprehending precipitation patterns and flood frequency.
- Transformer Architectures: Advanced machine learning models initially designed for natural language processing but increasingly employed in sequential flood prediction problems.
- Two-dimensional (2D) Hydraulic Models: Simulation frameworks that partition floodplains into grids or meshes, solving shallow water equations to ascertain water depth and velocity vectors throughout the domain.

U

- Uncertainty Propagation: The process by which errors and uncertainty in input data and model structure are transmitted and potentially amplified through the modelling chain.

- Unit Hydrograph: A mathematical method for understanding how catchments convert excess rainfall into runoff, illustrating the runoff response to a unit of rainfall input.
- Urban Density: The concentration of buildings and impervious surfaces in developed areas, fundamentally altering hydrological processes by increasing runoff volume and peak flows.
- Urban-Specific Flood Typologies: Specialized classification systems for urban flooding that incorporate drainage capacity exceedance, surface water flow paths, and infrastructure failure modes.

V

- Validation Frameworks: Systematic methods for assessing model performance against observed data, particularly crucial for flood risk models that include forward-looking climate projections.
- Value at Risk (VaR): A statistical measure of investment risk, indicating the maximum expected loss over a specified period at a given confidence level.
- Variational Inference: A technique in machine learning that enables networks to learn and update probability distributions over weather patterns rather than producing single-point predictions.
- Vine Copula Structure: A flexible method for modelling high-dimensional dependencies by breaking them down into bivariate relationships, helpful in capturing complex interactions between atmospheric variables.
- Vulnerability Assessment: The process of evaluating how susceptible properties, infrastructure, and communities are to flood damage based on their physical, social, and economic characteristics.
- Vulnerability Curves are mathematical functions that link hazard intensity (such as flood depth) to expected damage, typically represented as a percentage of asset value.
- Vulnerability Function Testing: Empirical validation of damage prediction models compared with actual loss data from historical flood events.

W

- Weather Pattern Distributions: Statistical representations of possible atmospheric conditions, focusing on those that could lead to extreme precipitation events.
- Weather Pattern Time series: Chronological sequences of meteorological data that act as inputs to hydrological models and provide the foundation for integrated flood risk assessment.
- Weighted Payout Structures: Settlement mechanisms in Physical Risk Swaps based on the severity and location of flooding, measured by multiple gauges within a catchment area.
- Wetland Buffer Capacity: Natural wetland systems' ability to temporarily store floodwaters and release them gradually, thus reducing downstream peak flows.

X

- X-Year Flood: A flood with a 1-in-X probability of occurring in any given year (e.g., a 100-year flood has a 1% annual probability of occurrence).

Y

- Year-over-year growth: In mortgage risk assessment, the comparative analysis of property values or default rates across consecutive annual periods is utilised to identify emerging trends in flood-affected areas.
- Yield Curve: In banking contexts, a curve illustrates the relationship between interest rates and various maturity lengths, similar to hazard curves that depict the relationship between flood magnitude and return period.

Z

- Z-Score: A statistical measure that expresses a value's relationship to the mean of a group, used in flood risk analysis to identify statistical anomalies in gauge readings or damage patterns.
- Zarr Format: A cloud-optimized data format for storing and accessing large arrays of meteorological data, enabling efficient processing of specific geographic subsets for regional flood analysis.

- Zone-Based Pricing: Traditional flood insurance rating methods assign premiums based on broad risk zones, contrasted with property-specific assessment approaches that capture individual risk characteristics.

This page has been intentionally left blank.

Property CDM

Level 1	Level 2	Field	Data Type	Data Description	Example Value
Property Header		section			
	Header	section			
		UPRN	text	Unique Property Reference Number	10002333678
		PropertyID	text	Unique identifier for the contract	POL-2024-001
Property Attributes		section			
		Occupancy Type	menu	Primary use classification of the property	Residential owner-occupied
		Property AreaSqm	decimal	Total floor area of property in square meters	140
		Housing Association	boolean	Indicates if property is owned by housing association	FALSE
		Income Generating	menu	Property's income generation status	No
		Paying BusinessRates	boolean	Indicates if property is subject to business rates	FALSE
		Building Residency	menu	Number of separate residences in building	Single
		Property Type	menu	Architectural style/category of property	Semi-detached
		Occupancy Residency	menu	Current occupancy status	Family resident
		Height Meters	decimal	Height of property in meters	2.35
		Number Of Storeys	integer	Number of floors in property	2
		Construction Year	integer	Year property was built	1962
		PropertyPeriod	menu	Construction period category	1950-1975
		CouncilTaxBand	menu	Property's council tax classification	D
		NumberBedrooms	integer	Total number of bedrooms	3

		NumberBathrooms	integer	Total number of bathrooms	2
		TotalRooms	integer	Total number of rooms excluding bathrooms	8
		GardenAreaFront	decimal	Front garden area in square meters	20
		GardenAreaBack	decimal	Back garden area in square meters	120
		ParkingType	menu	Available parking facilities	Driveway and garage
		AccessType	menu	Type of access to property	Public road
		LastMajorWorksDate	date	Date of last significant renovation/works	2018-06-15
	Construction	section			
		ConstructionType	menu	Primary construction method/materials	Brick and block
		FoundationType	menu	Type of building foundations	Strip foundations
		FloorType	menu	Ground floor construction type	Solid concrete
		StiltsHeight	decimal	Height of any stilts/pillars supporting property	0
		PropertyHeight	decimal	Total height of property in metres	7.2
		FloorLevelMeters	decimal	Height of ground floor above ground level in metres	0.6
		BasementPresent	boolean	Indicates presence of basement	FALSE
	Location	section			
		BuildingName	text	Name of building if applicable	Rose Cottage
		BuildingNumber	text	Street number of property	42
		SubBuildingNumber	text	Sub-unit number if applicable	A
		SubBuildingName	text	Name of sub-unit if applicable	Ground Floor Flat
		StreetName	text	Name of street	High Street

		AddressLine2	text	Secondary address line	Millbrook
		TownCity	text	Town or city name	Southampton
		County	text	County name	Hampshire
		Postcode	text	Property postcode	SO16 4AB
		USRN	text	Unique Street Reference Number	8400123
		LocalAuthority	text	Governing local authority name	Southampton City Council
		ElectoralWard	text	Electoral ward name	Millbrook
		ParliamentaryConstituency	text	Parliamentary constituency name	Southampton West
		Country	menu	Country within UK	England
		Region	menu	Geographic region	South East
		UrbanRuralClassification	menu	Urban/rural development classification	Urban
		LatitudeDegrees	decimal	Geographic latitude coordinate	50.9289
		LongitudeDegrees	decimal	Geographic longitude coordinate	-1.4317
		BritishNationalGrid	text	OS National Grid reference	SU 430 110
		What3Words	text	What3Words location identifier	// famout.honest.pizza
	RiskAssessment	section			
		EAFloodZone	menu	Environment Agency flood zone classification	Zone 2
		OverallFloodRisk	menu	Combined flood risk assessment	Medium
		FloodRiskType	menu	Primary source of flood risk	River
		LastFloodDate	date	Date of most recent flood event	15/2/1975
		SoilType	menu	Predominant soil composition	Clay

		GroundLevelMeters	decimal	Height above sea level in meters	45
		RiverDistanceMeters	decimal	Distance to nearest river in meters	250
		LakeDistanceMeters	decimal	Distance to nearest lake in meters	2000
		CoastalDistanceMeters	decimal	Distance to coastline in meters	12000
		CanalDistanceMeters	decimal	Distance to nearest canal in meters	350
		GovernmentalDefenceScheme	boolean	Covered by government flood defence scheme	FALSE
ProtectionMeasures		section			
		InsuranceStatus	menu	Current insurance coverage type	Standard cover
		FloodReEligible	boolean	Eligibility for Flood Re scheme	TRUE
		ClaimsHistory	integer	Number of historical insurance claims	1
		LastClaimDate	date	Date of most recent insurance claim	12/03/2019
		LastClaimType	menu	Nature of most recent claim	Flood damage
Resilience Measures		section			
		WarningSystemStatus	menu	Type of flood warning system in place	Other
		FloodBarriers	boolean	Presence of temporary flood barriers	TRUE
		FloodPanels	boolean	Presence of permanent flood panels	FALSE
		AirbrickCovers	boolean	Presence of airbrick flood covers	TRUE
		WaterproofWalls	boolean	Waterproof wall treatment applied	TRUE
		ResilientDoors	boolean	Flood-resistant door installation	TRUE
		RaisedFoundationsMm	decimal	Height of raised foundations (mm)	450
		ReinforcedWalls	boolean	Presence of reinforced wall construction	TRUE

		PermeablePaving	boolean	Installation of permeable paving	TRUE
		DrainageSystems	boolean	Enhanced drainage systems present	TRUE
		WaterDiversion	boolean	Water diversion measures installed	TRUE
		ElevatedSockets	boolean	Electrical sockets raised above flood level	TRUE
		NonReturnValve	boolean	Non-return valves installed	TRUE
		WaterproofFlooring	boolean	Water-resistant flooring installed	TRUE
		SumpPumpSystem	boolean	Sump pump system installed	FALSE
	NaturalMeasures	section			
		RainGardens	boolean	Presence of rain gardens	FALSE
		DetentionPonds	boolean	Presence of water detention ponds	FALSE
		GreenRoof	boolean	Installation of green roof system	FALSE
		StrategicPlanting	boolean	Strategic flood-resistant landscaping	TRUE
		Bioswales	boolean	Presence of bioswales	FALSE
		WaterStorage	boolean	Rainwater storage systems installed	TRUE
EnergyPerformance		section			
	Ratings	section			
		EPCRating	menu	Energy Performance Certificate rating	C
		CarbonRating	menu	Carbon emissions rating	B
		EmissionsScore	menu	Overall emissions performance rating	Good
	EnergyUsage	section			
		TariffType	menu	Type of energy tariff	Fixed dual fuel

		AnnualCarbonKgCO2e	decimal	Annual carbon emissions in kg CO2e	3000
		HeatingSystem	menu	Primary heating system type	Gas central heating
		RenewableSystem	menu	Type of renewable energy system	Solar PV
		AnnualEnergyKwh	decimal	Total annual energy consumption (kwh_year)	12000
		GridElectricityKwh	decimal	Annual grid electricity usage (kwh_year)	3500
		GasUsageKwh	decimal	Annual gas consumption (kwh_year)	8000
		SolarGenerationKwh	decimal	Annual solar energy generation (kwh_year)	500
		AnnualEnergyBill	decimal	Total annual energy cost (gbp_year)	2950
	BuildingFabric	section			
		WallConstruction	menu	Type of wall construction	Cavity brick
		CavityInsulation	boolean	Presence of cavity wall insulation	TRUE
		ThermalBridgeScore	decimal	Thermal bridging performance score	0.8
		LoftInsulationMm	decimal	Thickness of loft insulation (mm)	270
		RoofType	menu	Type of roof construction	Pitched with tiles
		FloorConstruction	menu	Type of floor construction	Solid concrete
		FloorInsulation	boolean	Presence of floor insulation	TRUE
		HeatingSys	menu	Type of heating system	Combi boiler
		WaterHeating	menu	Type of water heating system	Gas combi
		LightingType	menu	Primary type of lighting	LED
		AirTightnessScore	decimal	Air leakage test score	4.2
		GlazingType	menu	Type of window glazing	Double

		WindowFrameType	menu	Material of window frames	uPVC
		DoorType	menu	Material of external doors	Composite
		SmartMeterType	menu	Type of energy meter installed	Smart meter with export
History		section			
	FloodEvents	section			
		FloodReturnPeriod	integer	Expected flood return period (from modelled output)	100
		FloodDamageSeverity	menu	Severity of flood damage	Minor damage
		LastFloodDate	date	Date of most recent flood event	12/03/2019
	GroundConditions	section			
		SubsidenceStatus	menu	Current subsidence condition	No issues
		ContaminationStatus	menu	Ground contamination status	None detected
		GroundStability	menu	Ground stability assessment	Stable
		LastGroundIssueDate	date	Date of last ground-related issue	22/08/2022
	EnvironmentalIssues	section			
		AirQuality	menu	Local air quality rating	Moderate
		WaterQuality	menu	Local water quality rating	Good
		NoisePollution	menu	Type of noise pollution present	None
		LastEnvironmentalIssueDate	date	Date of last environmental issue	30/11/2018
	FireIncidents	section			
		FireDamageSeverity	menu	Severity of fire damage	None
		LastFireDate	date	Date of most recent fire incident	

TransactionHistory		section			
	Sales	section			
		SalePriceGbp	decimal	Most recent sale price (gbp)	325000
		SaleDate	date	Date of most recent	15/06/2007
		PreviousOwner	text		Smith Family
		MarketingDays	integer	days	45
	Rental	section			
		RentalHistory	menu		Mixed use history
		MonthlyRentGbp	decimal	gbp per month	1200
		VacancyCount	integer		2
		TenancyDuration	menu		12-24 months

The latest version 9 of Property CDM can be found in the public Github:-

https://github.com/MKM-Research-Labs/Phyrsrisk-cdm/blob/main/property/Property_CDM_v9.xlsx

This page has been intentionally left blank.

Bibliography

CLIMATE SCIENCE AND WEATHER MODELLING

1. Ahrens, C.D., & Henson, R. (2024). *Meteorology Today: An Introduction to Weather, Climate and the Environment*, 12th Edition.
2. Allaby, M., & Garratt, R. (2023). *Encyclopedia of Weather and Climate*.
3. Bader, D.C. (2022). *Climate Models: An Assessment of Strengths and Limitations*.
4. Friedunder, S. (Ed.) (2023). *An Introduction to the Mathematical Theory of Geophysical Fluid Dynamics*.
5. Geerts, B., & Linacre, E. (2020). *Climates and Weather Explained*.
6. Kalnay, E. (2003). *Atmospheric Modelling, Data Assimilation and Predictability*. Cambridge University Press.
7. Koonin, S.E. (2023). *Unsettled: What Climate Science Tells Us, What It Doesn't, and Why It Matters*.
8. Lamb, H.H. (2019). *Climate Present, Past and Future Volume 1: Fundamentals and Climate Now*.
9. Lamb, H.H. (2019). *Climate Present, Past and Future Volume 2: Climatic History and the Future*.
10. Neelin, J.D. (2023). *Climate Change and Climate Modelling*.
11. Nebeker, F. (2018). *Calculating the Weather: Meteorology in the 20th Century*.
12. Oliver, J.E. (2024). *Encyclopedia of World Climatology*.
13. Pasini, A. (2020). *From Observations to Simulations: A Conceptual Introduction to Weather and Climate Modelling*.
14. Petersen, J., Sack, D., & Gabler, R.E. (2023). *Fundamentals of Physical Geography*.
15. Potter, T.D., & Colman, B.R. (2023). *Handbook of Weather, Climate and Water Dynamics*.
16. Sanderson, E. (2023). *Six Thousand Years of History*.
17. Singh, S. (2023). *Environmental Geography*.
18. von Storch, H., & Zwiers, F.W. (2019). *Statistical Analysis in Climate Research*.
19. Wallace, J.M., & Hobbs, P.V. (2022). *Atmospheric Science: An Introductory Survey*.

20. Warner, T.T. (2011). Numerical Weather and Climate Prediction. Cambridge University Press.
21. Wilks, D.S. (2019). Statistical Methods in the Atmospheric Sciences, 4th Edition.

HYDROLOGICAL MODELLING

22. Begum, S. (2023). Flood Risk Management in Europe: Innovation in Policy and Practice.
23. Beven, K. (2012). Rainfall-Runoff Modelling: The Primer, 2nd Edition. Wiley-Blackwell.
24. Bombelli, I. (2024). Vulnerability Models and Their Application in Risk Assessment.
25. Bras, R.L. (1990). Hydrology: An Introduction to Hydrologic Science. Addison-Wesley.
26. Caviedes-Voullième, D., & Morales-Hernández, M. (2023). Reconstruction of 2D River Beds by Appropriate Interpolation of 1D Cross-sectional Information for Flood Simulation.
27. Committee on FEMA Flood Maps (2023). Mapping the Zone: Improving Flood Map Accuracy.
28. Dingman, S.L. (2015). Physical Hydrology, 3rd Edition. Waveland Press.
29. Duchon, C., & Hale, R. (2023). Time Series Analysis in Meteorology and Climatology: An Introduction.
30. Eslamian, S., & Eslamian, F. (2023). Flood Handbook: Impacts and Management.
31. Eslamian, S., & Eslamian, F. (Eds.) (2023). Flood Handbook: Analysis and Modelling.
32. Knight, D., & Shamseldin, A. (2023). River Basin Modelling for Flood Risk Mitigation.
33. Maidment, D.R. (1993). Handbook of Hydrology. McGraw-Hill Education.
34. Matczak, P., & Hegger, D. (2023). Flood Risk Governance for More Resilience.
35. Mays, L.W. (2023). Urban Water Supply Handbook.
36. Mazzoleni, M. (2023). Improving Flood Prediction: Assimilating Uncertain Crowdsourced Data into Hydrologic and Hydraulic Models.
37. Meylan, P., Favre, A.C., & Musy, A. (2023). Predictive Hydrology: A Frequency Analysis Approach.
38. Molinari, D., & Menoni, S. (2023). Flood Damage Survey and Assessment: New Insights from Research and Practice.

39. Mukolwe, M.M. (2023). Flood Hazard Mapping Uncertainty and its Value in the Decision-Making Process.
40. National Assessment of Flood and Coastal Erosion Risk in England (2024).
41. National Flood Risk Assessment (NaFRA2). (2025). Annual Report.
42. Pender, G., & Faulkner, H. (Eds.) (2023). Flood Risk Science and Management.
43. Penning-Rowsell, E., et al. (2024). Multi-Coloured Manual: Handbook for Flood and Coastal Erosion Risk Management.
44. Robson, A., & Reed, D. (2023). Flood Estimation Handbook: Statistical Procedures for Flood Frequency Estimation.
45. Rodda, H., & Little, M. (2023). Understanding Mathematical and Statistical Techniques in Hydrology: An Examples-based Approach.
46. Shaw, E.M., Beven, K.J., et al. (2017). Hydrology in Practice, 4th Edition. CRC Press.
47. Teegavarapu, R.S.V. (2023). Floods in a Changing Climate: Extreme Precipitation.
48. Teegavarapu, R.S.V., & Salas, J.D. (2023). Statistical Analysis of Hydrologic Variables: Methods and Applications.
49. Watson, D., & Adams, M. (2023). Design for Flooding: Architecture, Landscape, and Urban Design for Resilience to Flooding and Climate Change.

RISK ASSESSMENT AND FINANCIAL MODELLING

50. Alexander, C. (2008). Market Risk Analysis. Wiley.
51. Allen, S.L. (2024). Financial Risk Management: A Practitioner's Guide to Managing Market and Credit Risk.
52. Andrews, N. (2023). Contract Law.
53. Augustin, P. (2023). Credit Default Swaps: A Survey.
54. Beinstein, E. (2006). Credit Derivatives Handbook. JPMorgan.
55. Bellini, T. (2023). How to Model and Validate Expected Credit Losses for IFRS 9 and CECL: A Practical Guide with Examples Worked in R and SAS.
56. Beltran, A., & Elliott, R. (2018). Environmental Economics and Flood Risk Valuation. Journal of Environmental Economics and Management.
57. Beltran, A., Maddison, D., Elliott, R. (2023). Is Flood Risk Capitalised into Property Values?

58. Blinder, A.S. (2023). After the Music Stopped: The Financial Crisis, the Response, and the Work Ahead.
59. Blyth, S. (2023). An Introduction to Quantitative Finance.
60. Boberski, D. (2009). CDS Delivery Option: Better Pricing of Credit Default Swaps.
61. Bombelli, I. (2023). Vulnerability Models in Risk Assessment.
62. Bruyere, R., Copinot, R., & Fery, L. (2023). Credit Derivatives and Structured Credit: A Guide for Investors.
63. Chang, S.J. (2023). Financial Crisis and Transformation of Korean Business Groups: The Rise and Fall of Chaebols.
64. Chorafas, D.N. (2023). Risk Management Technology in Financial Services: Risk Control, Stress Testing, Models, and IT Systems.
65. Cornuéjols, G., Peña, J., & Tütüncü, R. (2023). Optimization Methods in Finance.
66. Cox, J.C., & Rubinstein, M. (2023). Options Markets.
67. Crouhy, M. et al. (2014). The Essentials of Risk Management, 2nd Edition.
68. Fabozzi, F.J. (2023). The Handbook of Mortgage-Backed Securities.
69. Fabozzi, F.J., & Mann, S.V. (2023). Handbook of Fixed Income Securities.
70. Fabozzi, F.J., & Mann, S.V. (2023). Introduction to Fixed Income Analytics: Relative Value Analysis, Risk Measures and Valuation.
71. Feldstein, S. (2023). Handbook of Municipal Bonds.
72. Financial Crisis Inquiry Commission (2023). The Financial Crisis Inquiry Report: Final Report of the National Commission on the Causes of the Financial Crisis.
73. Gomez-Cunya, L.A., Fardhosseini, M.S. et al. (2024). Analysing Investments in Flood Protection Structures.
74. Gower, L.C.B., & Davies, P. (2023). Principles of Modern Company Law.
75. Hirsa, A., & Neftci, S.N. (2023). An Introduction to the Mathematics of Financial Derivatives.
76. Hull, J.C. (2018). Risk Management and Financial Institutions, 5th Edition.
77. Janssen, J., & Manca, R. (2023). Semi-Markov Risk Models for Finance, Insurance and Reliability.
78. Jarrow, R., & Yildirim, Y. (2023). Theoretical Models for Inflation Risk.
79. Jorion, P. (2016). Financial Risk Manager Handbook, 7th Edition.

80. Klenke, A. (2020). Probability Theory: A Comprehensive Course.
81. Kramer, S.N. (2023). History Begins at Sumer.
82. Kroese, D.P., Taimre, T., & Botev, Z.I. (2023). Handbook of Monte Carlo Methods.
83. Lanczos, C. (2023). Linear Differential Operators.
84. Malz, A.M. (2023). Financial Risk Management: Models, History, and Institutions.
85. McNeil, A.J. et al. (2015). Quantitative Risk Management: Concepts, Techniques and Tools.
86. Meissner, G. (2023). Correlation Risk Modelling and Management: An Applied Guide including the Basel III Correlation Framework.
87. Mitchell-Wallace, K. (2023). Natural Catastrophe Risk Management and Modelling: A Practitioner's Guide.
88. Mounfield, C.C. (2023). Synthetic CDOs: Modelling, Valuation and Risk Management. Mathematics, Finance and Risk Series.
89. Nelsen, R.B. (2023). An Introduction to Copulas.
90. Ramirez, J. (2015). Corporate Equity Derivatives and Equity Capital Markets.
91. Saunders, A., & Allen, L. (2023). Credit Risk Management In and Out of the Financial Crisis: New Approaches to Value at Risk and Other Paradigms.
92. Schultz, G.M., & Fabozzi, F.J. (2024). Investing in Mortgage-Backed and Asset-Backed Securities: Financial Modelling with R and Open Source Analytics.
93. Shreve, S.E. (2022). Stochastic Calculus for Finance I: The Binomial Asset Pricing Model.
94. Tijms, H.C. (2023). A First Course in Stochastic Models.
95. Wagner, N. (2023). Credit Risk Models, Derivatives, and Management.

INSURANCE AND PROPERTY RISK

96. Banks, E. (2004). Alternative Risk Transfer. Wiley Finance.
97. Bayes Business School (2023). Residential Property Flood Risk.
98. Botzen, W.J.W. et al. (2021). Economic Valuation of Climate Risk.
99. Clark, P.K. (2024). Property Risk Assessment in a Changing Climate.
100. Dixon, L. et al. (2023). The Future of Flood Insurance.
101. Elliott, R. (2023). Underwater Loss, Flood Insurance, and the Moral Economy of Climate Change in the United States.
102. Flood RE (2023). Transition Plan Report: Our Call to Action.

103. Grossi, P. and Kunreuther, H. (2005). Catastrophe Modelling.
104. Kousky, C. (2023). Understanding Disaster Insurance.
105. Morri, G., & Benedetto, P. (2023). Commercial Property Valuation Methods and Case Studies. Wiley Finance.
106. Quinto, C. (2023). Insurance Systems in Times of Climate Change: Insurance of Buildings Against Natural Hazards.
107. Wing, O.E.J. et al. (2023). Flood Risk and Financial Stability.
108. Wyatt, P. (2024). Property Valuation: Methods and Applications.

MODEL RISK AND GOVERNANCE

109. Bessis, J. (2015). Risk Management in Banking, 4th Edition.
110. Brandimarte, P. (2024). Handbook in Monte Carlo Simulation.
111. Cherubini, U., & Mulinacci, S. (2023). Dynamic Copula Methods in Finance.
112. Cruz, M.G. et al. (2015). Fundamental Aspects of Operational Risk.
113. Day, D. (2023). The Oxford Handbook of Leadership and Organizations.
114. Franzetti, C. (2023). Operational Risk Modelling and Management.
115. Kemp, M. (2023). Extreme Events: Robust Portfolio Construction.
116. Levi-Faur, D. (2023). The Oxford Handbook of Governance.
117. Panjer, H.H. (2023). Operational Risk Modelling Analytics.
118. Rachev, S.T., & Stoyanov, S.V. (2023). Advanced Stochastic Models.
119. Rösch, D., & Scheule, H. (2023). Model Risk Identification, Measurement and Management.
120. Sayles, J. (2024). Principles of AI Governance and Model Risk Management: Master the Techniques for Ethical and Transparent AI Systems.
121. Tarantino, A. (2023). The Governance, Risk, and Compliance Handbook: Technology, Finance, Environmental, and International Guidance.

AI AND MACHINE LEARNING APPLICATIONS

122. Arlow, J., & Neustadt, I. (2024). Generative Analysis: The Power of Generative AI for Object-Oriented Software Engineering with UML.
123. Bishop, C.M. (2006). Pattern Recognition and Machine Learning.
124. Brownlee, J. (2023). Deep Learning for Time Series Forecasting: Predict the Future with MLPs, CNNs and LSTMs in Python.

125. Camps-Valls, G., Tuia, D., & Zhu, X. (2023). Deep Learning for the Earth Sciences: A Comprehensive Approach to Remote Sensing, Climate Science and Geosciences.
126. Didur, K. (2024). Machine Learning in Finance: Why, What & How.
127. Gelman, A. et al. (2013). Bayesian Data Analysis, 3rd Edition.
128. Goodfellow, I. et al. (2016). Deep Learning. MIT Press.
129. Hastie, T. et al. (2009). The Elements of Statistical Learning.
130. Jansen, S. (2023). Machine Learning for Algorithmic Trading.
131. Jensen, F.B., & Nielsen, T.D. (2023). Bayesian Networks and Decision Graphs.
132. Kaabar, S. (2023). Deep Learning for Finance: Creating Machine Deep Learning Models for Trading in Python.
133. Kirk, A., Timms, S., & Rininsland, A. (2023). Data Visualization: Representing Information on Modern Web.
134. Koski, T., & Noble, J. (2023). Bayesian Networks - An Introduction.
135. McMahon, D. (2023). Quantum Computing Explained.
136. Mishra, D., Yang, X.S., & Unal, A., et al. (2023). Data Science and Big Data Analytics: Proceedings of IDBA 2023.
137. Murphy, K.P. (2012). Machine Learning: A Probabilistic Perspective.
138. Nagaraj, S., & Kumar, K.V.R. (2024). AI-Driven Intelligent Models for Business Excellence.
139. Neapolitan, R.E., & Jiang, X. (2023). Artificial Intelligence With Machine Learning.
140. Peters, J., Janzing, D., et al. (2023). Elements of Causal Inference: Foundations and Learning Algorithms.
141. Pillai, A.S., & Tedesco, R. (2023). Machine Learning and Deep Learning in Natural Language Processing.
142. Raschka, S. (2024). Build a Large Language Model (From Scratch).
143. Satpathy, S., & Mahapatra, S. (2024). AI for Climate Change.
144. Schwabish, J. (2023). Better Data Visualizations: A Guide for Scholars, Researchers, and Wonks.
145. Sharma, H., & Chakravorty, A. (2023). Artificial Intelligence Theory and Applications: Proceedings of AITA 2023.
146. Subramanian, S. (2024). Large Language Model-Based Solutions: How to Deliver Value with Cost-Effective Generative AI Applications.

147. Tabrizi, R. (2023). Behavioral AI: Unleash Decision Making with Data.
148. Thomas, T., Vijayaraghavan, A.P., & Emmanuel, S. (2024). Machine Learning in Cyber Security.

REGULATORY AND CENTRAL BANK PUBLICATIONS

149. Benvenuti, O.P. (2024). Technical Implementation Report: Risk Assessment Framework.
150. Fanzo, J. (2021). Global Food Policy and Climate Change.
151. Litterman, B. (Chair) (2020). Managing Climate Risk in the U.S. Financial System. Report of the Climate-Related Market Risk Subcommittee, CFTC.